

Learning to Select Actions for Complex Hand Tasks: A Comparative Study of Visuomotor Policies and Latent World Models

Graham Pellegrini

Supervisor: Prof. Alexiei Dingli

September 2026

*Submitted in partial fulfilment of the requirements
for the degree of Master of Science in Artificial Intelligence.*



L-Università ta' Malta

Faculty of Information &
Communication Technology

Plagiarism Declaration

I, the undersigned, hereby declare that the work contained in this dissertation is my own original work and that I have not previously in its entirety or in part submitted it at any university for a degree.

I declare that this submission is my own work and that, to the best of my knowledge and belief, it contains no material previously published or written by another person, except where due acknowledgement is made in the text.

I declare that the use of artificial intelligence tools in the preparation of this dissertation was limited to writing assistance, code suggestions, and literature search support. All research design, experimental decisions, implementation choices, result interpretation, and academic claims are my own.

Signature: _____

Name: Graham Pellegrini

Date: _____

Abstract

A latent world model predicts how a scene will evolve in a compact learned representation rather than in raw pixels, judging a proposed action by its predicted consequences before it is executed. Casting an imitation policy as an actor and a world model as the evaluator of its proposed motions, this dissertation asks whether that pairing improves the selection of future bimanual hand motions under a controlled offline benchmark favouring quality of evidence over quantity of tests.

The benchmark evaluates 7,607 held-out windows from the EgoDex egocentric manipulation dataset, each pairing 16 observed frames with a 16-frame prediction target and 18-dimensional bimanual wrist-motion actions, through three branches: the policy alone, world-model scoring with its preferred candidate retained, and scoring with that candidate excluded. Local training exposure was matched across all models, though large pretrained backbones still carry external training this local budget does not offset.

The thesis’s main positive result is that Diffusion Policy candidates scored by V-JEPA2-AC reduce translation error, averaged over the full 16-step predicted chunk and pooled across three training seeds, by 5.98% and rotation error by 0.96 degrees relative to the policy alone, without retraining, since the world model simply re-ranks the five motion proposals the policy already produces at inference time. A perfect selector with ground truth would reduce translation error by at most 16.51% on the same pool, so the world model recovers roughly 36% of the attainable headroom. Other routes, ACT, DexWM, LeWM, DINO-WM, and JEPA-WM, produced negative, mixed, exploratory or bounded outcomes, each linked to a likely explanation or limitation.

Across the model and policy combinations tested, the results suggest that world-model guidance helps when three conditions align: a diverse candidate pool, a scorer receiving meaningful action-conditioned information, and candidate actions compatible with the scorer’s own training support. A secondary reproduction of Chandra et al.’s DiWA on CALVIN supports this interpretation from a different experiment. Weight-level finetuning, which satisfies all three conditions by construction, outperforms this thesis’s own external selection on the same policy and world model by roughly an order of magnitude.

Acknowledgements

I would like to first thank Prof. Alexiei Dingli for his guidance and supervision throughout this thesis and the wider body of work undertaken as an RSO. His knowledge extends well beyond academia, and perhaps his most admirable trait is how effortlessly he interacts with people at every level, making the work itself enjoyable, while taking genuine care not only for the work but for the person behind it.

I am grateful to my parents for their support in all walks of life. Their parental instincts and foresight are always right, even when context is short.

To my girlfriend, with whom I share a very close bond, thank you for your continued belief in me.

Finally, a thanks to my sister for her genuine interest in my achievements, and to her little one, my niece, who is always like a small radar scanning everyone's every move. Between the two of them, they gave me the best possible excuse to step away from the screen.

Contents

Plagiarism Declaration	i
Abstract	ii
Acknowledgements	iii
Contents	vii
List of Figures	x
List of Tables	xv
List of Abbreviations	xvii
List of Notations	xviii
1 Introduction	1
1.1 Motivation	2
1.2 Research Question and Scope	3
1.2.1 Objectives	4
1.3 Contributions	6
1.4 Dissertation Structure	6
2 Background	9
2.1 Learning paradigms	9
2.2 Visuomotor policies	10
2.2.1 Behaviour cloning and compounding error	11
2.2.2 Action chunking	11
2.3 Rigid body pose and the $SE(3)$ representation	12
2.4 Latent representations	13
2.5 Conditional variational autoencoders	14
2.6 Transformer networks and the attention mechanism	16
3 Literature Review	18

3.1	Visuomotor policies	20
3.1.1	Action Chunking with Transformers (ACT)	20
3.1.2	Diffusion Policy	21
3.2	World models	23
3.2.1	World models as policy refiners	25
3.2.2	World models as policy evaluators	26
3.2.3	Latent predictive architectures	27
3.3	A common comparison of the world-model scorers	33
3.4	Research gap and conceptual framework	35
4	Methodology	37
4.1	Dataset and Experimental Setup	37
4.1.1	EgoDex: source and selection rationale	37
4.1.2	Windowing protocol	38
4.1.3	Training splits and test set	39
4.2	Three-Branch Experimental Framework	40
4.3	Policy Variants	41
4.4	World Model Families	42
4.5	Matched Local Training-Exposure Protocol	44
4.5.1	Training exposure budget	44
4.5.2	Evaluation contract	46
4.5.3	Training-seed protocol	46
4.6	Evaluation Metrics	47
4.6.1	Offline trajectory prediction setting	47
4.6.2	Primary metrics	47
4.6.3	Candidate selection comparison	49
4.7	Evaluation Scope and Limitations	51
5	Implementation	53
5.1	Repository and execution model	53
5.2	Data pipeline and windowing	55
5.3	Training configuration	55
5.4	Policy training	58
5.4.1	Action Chunking Transformer	58
5.4.2	Diffusion Policy	63
5.5	World model training	64
5.5.1	V-JEPA2-AC: extraction and transition head	64
5.5.2	LeWM: end-to-end training	65
5.5.3	DexWM: predictor trained from scratch	67

5.5.4	DINO-WM and JEPA-WM exploratory selectors	68
5.5.5	DiWA/CALVIN secondary-study implementation	69
5.5.6	V-JEPA2 (no action conditioning): ablation control	70
5.5.7	PointWorld: action-degradation control	70
5.6	B2 margin threshold calibration	71
5.7	Implementation challenges and resolutions	73
6	Results	75
6.1	World-model selection results	75
6.2	ACT guard results under CEMP	81
6.3	Selection quality across task difficulty	85
6.4	Comparison against DiWA on CALVIN	87
6.4.1	Selection versus finetuning	89
6.4.2	Reproducing DiWA’s own checkpoint	91
6.4.3	C4a and C4b: an inconclusive follow-up	92
7	Discussion and Limitations	95
7.1	What the benchmark shows	95
7.2	Why Branch 2 outperforms Branch 3	97
7.3	Interpreting the model-family differences	98
7.4	Reproducibility and open model releases	101
7.5	Why semantic alignment and object grounding were excluded	101
7.6	Three ways a world model can be used	102
7.7	Main limitations	102
8	Conclusion	106
9	Future Work	109
9.1	Immediate validation work	109
9.2	Mechanism experiments	110
9.3	Longer-term directions	111
	References	114
A	Auxiliary Implemented Modules	122
A.1	Why these modules are excluded, and what future work could do	122
A.2	Semantic alignment pipeline	123
A.3	Object grounding and mask generation	124
A.4	Official-checkpoint world models: DINO-WM and JEPA-WM	125
B	Validity, Scaling, and Sensitivity Studies	127

B.1	Windowing strategy comparison: fixed-stride versus state-change-targeted	127
B.2	Training-split scaling: performance across the three-part corpus	129
B.3	Candidate pool size (K) sweep	131
B.4	Training exposure convergence study	132
B.5	CEMP compute cost	135
B.6	Multi-seed validation	136
B.6.1	Primary EgoDex results	136
B.6.2	DiWA/CALVIN multi-seed validation	139

List of Figures

Figure 2.1	The three classical machine learning paradigms: supervised learning trains on labelled data, unsupervised learning extracts structure from unlabelled data, and reinforcement learning learns from a reward signal over states and actions [13]. Self-Supervised Learning (SSL), central to the world models in this work, is a modern form of learning from unlabelled data in which prediction targets are constructed from the data itself [14].	9
Figure 2.2	A visuomotor policy maps an EgoDex-style camera frame to future bimanual wrist actions in the thesis setting.	10
Figure 2.3	Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [16].	11
Figure 2.4	Translation of the wrist origin along each of the three spatial axes, $\mathbf{t} = [t_x, t_y, t_z]^\top$, following the $SE(3)$ pose convention [17].	12
Figure 2.5	Rotation of the wrist about the y -axis by angle θ ; the same construction applies about x and z , and the wrist's full orientation is the combined rotation about all three, represented as a rotation matrix $R \in SO(3)$ [17]. . .	12
Figure 2.6	Example latent-space plot illustrating the general idea that similar examples can form nearby clusters in a learned representation [19].	14
Figure 2.7	Standard Variational Autoencoder (VAE) (a) and conditional Conditional Variational Autoencoder (CVAE) (b) encoder-decoder pipelines [22]. . .	15
Figure 2.8	VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation) [21].	15
Figure 2.9	Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [23].	16
Figure 3.1	Architecture of the Action Chunking with Transformers (ACT) policy [8].	20
Figure 3.2	Diffusion Policy: general formulation (a), Convolutional Neural Network (CNN) variant (b), transformer variant (c) [9].	22

Figure 3.3	The post-AlexNet split of visual representation learning: a generative branch (top), culminating in diffusion models and the Vision Language Action Models (VLAs) debated above, and a joint-embedding branch (bottom), culminating in the Joint Embedding Predictive Architecture (JEPA)-style world models evaluated in this benchmark. Screenshot from [11].	24
Figure 3.4	The three-stage DiWA recipe: a world model is trained on task-agnostic play data, a diffusion policy is pretrained by imitation in its latent space, and the policy is then fine-tuned with reinforcement learning on rollouts imagined entirely inside the frozen world model [12].	25
Figure 3.5	Classic closed-loop control (a) versus CLOVER’s closed-loop visuomotor control (b) [44].	26
Figure 3.6	Three SSL prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [14].	28
Figure 3.7	Video Joint Embedding Predictive Architecture 2 (V-JEPA2) post-training pathways from internet-scale backbone to specialised applications [24]. . . .	29
Figure 3.8	Standard V-JEPA2 (left) versus action-conditioned Action-Conditioned Video Joint Embedding Predictive Architecture 2 (V-JEPA2-AC) (right) [24]. .	30
Figure 3.9	High-level layout of the DexWM world model pipeline, illustrating how predicted keypoints are incorporated alongside the encoder, Conditional Diffusion Transformer (CDiT) predictor, and decoder stages [25].	31
Figure 3.10	LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [26].	31
Figure 4.1	Fixed-stride windowing scheme (obs16_pred16_s128) used as the canonical training and evaluation index.	39
Figure 4.2	Three-branch evaluation framework (Section 4.2), shown for a policy producing five independent candidates. B2 and B3 reuse the same candidate pool generated for B1, relayed through the branches (dashed lines) rather than regenerated; B3 additionally removes the policy’s preferred candidate before scoring. Both Action Chunking Transformer (ACT) and Diffusion Policy (DP) follow this design; the per-policy candidate-generation mechanics are detailed in Section 4.3.	40
Figure 5.1	Package structure under <code>src/</code> , launched via <code>scripts/</code> and <code>cli/</code> inside Docker.	54
Figure 5.2	Data flow from version-controlled code to the generated, non-version-controlled <code>datasets/</code> and <code>runs/wact/</code> artefacts.	54
Figure 5.3	Shared ResNet-18 visual backbone preceding both policy decoders. Adapted from [60].	58

Figure 5.4	CVAE posterior collapse in the trained ACT checkpoint. Four independently sampled latent vectors produce an identical output $\hat{a}$ (dashed arrows indicate the latent path the decoder effectively ignores).	62
Figure 5.5	Cross-Entropy Method with Policy (CEMP) candidate augmentation, refining SE(3) perturbations of the ACT B1 output over three iterations to form the B2/B3 evaluation pool.	63
Figure 5.6	DDIM subsampling correction. The ten-step inference schedule spans $t \in \{90, 80, \dots, 0\}$. At each reverse step the correct $\bar{\alpha}_{t_{\text{prev}}}$ is the <i>next timestep in the subsampled schedule</i> ; the bug substituted the consecutive training-schedule timestep $\bar{\alpha}_{t-1}$, miscalibrating the $\hat{x}_0$ weighting at every step.	64
Figure 5.7	The two LeWorldModel (LeWM) scoring routes evaluated in this study, both starting from the same candidate action chunk: the rollout scorer (right), which predicts forward through LeWM’s learned dynamics, and the action-prior scorer (left), which does not.	66
Figure 5.8	DexWM pipeline as implemented, comprising a frozen DINOv2 encoder, CDiT predictor, lightweight decoder, and HC loss supervision path. Adapted from [25].	67
Figure 6.1	Reduction in translation error by world model and branch	77
Figure 6.2	Six DiWA conditions, their lineage and module composition	89
Figure B.1	State-change-targeted sampling pipeline yielding one frozen window per episode.	128
Figure B.2	Selection gain against candidate pool size	132
Figure B.3	Normalised validation loss (each family’s value divided by its 1k loss) across six training scales, pooled across 3 seeds except Dexterous Manipulation World Model (DexWM) (single-seed). Solid lines are primary models; dashed lines are controls.	134

List of Tables

Table 1.1	Alignment between the research question, the five objectives above, and the chapters that address each.	5
Table 3.1	Literature roles in the benchmark design.	19
Table 3.2	Average consecutive sub-task completions (<code>avg_len</code>) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [44].	27
Table 3.3	Common-criteria comparison of the world-model families reviewed in this chapter. ”Action-conditioned” asks whether the model’s predictor natively takes a candidate action as input, not whether an adapter can supply one. ”Evaluation evidence” uses the task-success versus offline-prediction-accuracy distinction introduced above.	34
Table 4.1	EgoDex splits under the <code>obs16_pred16_s128</code> windowing contract. Episode and window counts are derived from the canonical index.	39
Table 4.2	Evidence role assigned to each implemented or investigated world-model path.	44
Table 5.1	Training configuration for all model families at $E = 100,000$. Approximate step time is shown for trainable phases only (--- denotes not applicable).	56
Table 5.2	Model input/output contracts used by the implemented routes.	57
Table 5.3	ACT CVAE posterior collapse: four successive retraining attempts, all indistinguishable from zero ($\varepsilon = 10^{-6}$) on every training split, at random seed 42.	60
Table 5.4	ACT CVAE posterior collapse across three random seeds per configuration, <code>train_part1_2_3</code> . <i>collapse</i> marks mean pairwise difference below the same $\varepsilon = 10^{-6}$ threshold used in Table 5.3; <i>ok</i> marks a checkpoint whose five sampled candidates remain meaningfully distinct.	60
Table 5.5	Backbone sourcing and what was trained from scratch for this study, by world model family.	65

Table 5.6	B2 margin threshold calibration on the internal validation split ($K = 5$, V-JEPA2-AC guard, 42,774 windows, disjoint from the test set). Higher δ keeps candidate 0 more often and yields a smaller improvement. Validation-split baseline translation L_2 (candidate 0) is 0.11499, distinct from the pooled test-set baseline of 0.117 reported in Table 6.1, since the two are different windows and this validation-split figure is a single-seed value never re-pooled across seeds.	72
Table 5.7	Material implementation challenges and resolution status.	74
Table 6.1	Branch 1 policy-only baseline results, no world model involved, pooled mean across seeds 42, 43, and 44 [95% CI].	75
Table 6.2	LeWM’s two diagnostic configurations, full 7,607-window test set, neither shown to reliably track which candidate is actually better (Section 5.5.2). Rollout row: pooled mean across 3 seeds [95% CI]. Action-prior row: canonical seed only, a fixed statistical scorer never retrained or retested at the other two seeds (Section 5.5.2); its bracketed figure is a single-seed episode-level CI, not the two-level clustered bootstrap used elsewhere.	78
Table 6.3	V-JEPA2-AC’s per-window Spearman rank correlation between score and true candidate translation error, by candidate pool size K , Branch 2, pooled across 3 seeds [95% CI], on the nested pools behind Appendix B.3. .	79
Table 6.4	Primary DP guard results, pooled mean across 3 seeds [95% CI] (bootstrap method: Section 4.6.3; per-seed detail: Appendix B.6.1). LeWM and DexWM caveats discussed below.	80
Table 6.5	ACT guard results, pooled mean across 3 seeds [95% CI] (selection rule and per-seed detail: Appendix B.6.1), same bootstrap as Table 6.4. Four of five scorers; LeWM held out, reported separately below (same reason as Section 6.1). Negative values are an improvement over the frozen ACT baseline, positive a regression. The DexWM row is single-seed.	81
Table 6.6	LeWM’s two diagnostic configurations on the ACT row, full 7,607-window test set, neither shown to reliably track which candidate is actually better (Section 5.5.2). Rollout row: pooled mean across 3 seeds [95% CI]. Action-prior row: canonical seed only, a fixed statistical scorer never retrained or retested at the other two seeds; its bracketed figure is a single-seed episode-level CI, not the two-level clustered bootstrap used elsewhere. . . .	83
Table 6.7	V-JEPA2-AC capture (percentage of the oracle ceiling recovered, Section 6.1) by baseline-difficulty quintile, $K = 8$, Branch 2, pooled across 3 seeds. A negative value means the selected candidate is worse than candidate 0, not merely less close to the oracle than a perfect selector would be. .	85

Table 6.8	Effect of gap-based abstention on the aggregate translation-error change, pooled across 3 seeds [95% CI].	86
Table 6.9	C1, C2, and C3b on CALVIN: this thesis's own selection harness against DiWA's own imagined-Proximal Policy Optimisation (PPO) finetuning, both read from the same C1 starting point. C3b is DiWA's own trainer-loop evaluation reading; Table 6.11 gives its shared-harness reading. Change is in percentage points (pp), the absolute difference between two percentages, not a further relative change.	90
Table 6.10	C3a and C3b on CALVIN, against DiWA's own reference checkpoint. Both are read on DiWA's own PPO-trainer harness throughout, not comparable to Table 6.9's c1-c2/c2-c3b columns. Confidence intervals are seed-level only, since this harness emits no per-episode trace to cluster on (Section 4.6.3).	91
Table 6.11	C4a: this thesis's selection harness applied to the finetuned C3b checkpoint. The C3b column is measured on C4a's own shared harness, not Table 6.10's trainer-loop reading.	93
Table 6.12	C4b: C4a's harness with the reward classifier recalibrated on C3b's own rollouts, for the 3 tasks it covers. <code>close_drawer</code> excluded, its finetuned rollouts contain no failures to train against. Every seed's classifier, including seed 42's, is collected from its own corrected C3b checkpoint.	94
Table 8.1	Achievement status of each objective from Section 1.2.1.	107
Table A.1	Semantic-pipeline VLM comparison over the shared 64-window sample. SigLIP and OpenCLIP scores are cosine-similarity proxies computed independently within each embedding family and are not directly comparable across the two columns; lower seconds per window is faster.	123
Table A.2	DINO-WM and JEPA-WM validation pilot, 256 windows, <code>train_part1_2_3</code> . Support penalty is the world model's own latent support score [unselected→selected], lower is better fit. The remaining four columns report the selected candidate's change from the same shared unselected (act) baseline, identical for every row since it is the same policy output regardless of which world model scores it: loss 0.0870, translation 0.1053, endpoint translation 0.1126, rotation 30.9°. No B2 reweighted run exists for either model.	126

Table B.1	Candidate-capture rate under fixed-stride (Strategy 1) and state-change-targeted (Strategy 2) windowing, Diffusion Policy row, $K = 8$. Both scorer rows are pooled across 3 seeds. The bottom two rows are a separate, single-seed sub-analysis that splits V-JEPA2-AC's own gain between the state-change test windows themselves and the accompanying state-change support extract, via a two-by-two crossing of window and support-extract strategy; this breakdown alone is not pooled.	128
Table B.2	Translation L_2 by training split. Policy rows are each policy's own unselected (candidate 0) prediction; world-model rows are Branch 2 selected-candidate translation L_2 against DP candidates, not a standalone reconstruction loss. LeWM and DexWM remain "not evaluated" at <code>train_part1</code> and <code>train_part1_2</code> , for the different reasons the paragraph above gives; V-JEPA2-AC's cells at those splits are a fresh rerun, single canonical seed, matching this table's existing single-seed convention.	130
Table B.3	Validation loss across six training scales for all seven model families, pooled across 3 seeds except DexWM (single-seed). Losses are objective-specific and not cross-family comparable. LeWM uses <code>fit/loss</code> for all three seeds, not the <code>validate/loss_epoch</code> metric used elsewhere for this family (the paragraph below explains why), and was trained on CPU rather than GPU throughout this specific study, unlike LeWM's GPU training everywhere else in this thesis.	133
Table B.4	50k-to-100k trend per primary model family, pooled across 3 seeds except DexWM (single-seed). The universal budget is $E = 100,000$ effective items regardless of this trend.	135
Table B.5	Guard runtime, raw pool against CEMP, full 7,607-window test set, Branch 2 only. Branch 3 timings follow the same pattern.	136
Table B.6	Branch 1 policy-only baseline translation error by seed, both policies. Seed 42 is the canonical value reported in Table 6.1.	137
Table B.7	Branch 1 policy-only baseline rotation error by seed, same policies and seeds as Table B.6.	137
Table B.8	DP-row Branch 2 translation change by training seed. Seed 42 is the canonical value reported in Table 6.4.	137
Table B.9	DP-row Branch 2 rotation change by training seed, same scorers and seeds as Table B.8.	138
Table B.10	ACT-row Branch 2 translation change by training seed. Seed 42 uses CEMP (the checkpoint collapses); seeds 43 and 44 use native ACT sampling (neither checkpoint collapses), per the adaptive rule described above. Seed 42 values are the canonical ones reported in Table 6.5 and Table 6.6.	139

Table B.11 ACT-row Branch 2 rotation change by training seed, same scorers, seeds,
and CEMP/native split as Table B.10. 139

Table B.12 close_drawer success rate by seed, all conditions that apply to this
task. C4b is excluded from this task by design. 139

Table B.13 turn_on_led success rate by seed, all conditions. 140

Table B.14 move_slider_left success rate by seed, all conditions. 140

Table B.15 turn_on_lightbulb success rate by seed, all conditions. 140

List of Abbreviations

ACT Action Chunking Transformer.

ALOHA A Low-cost Open-source Hardware System for Bimanual Teleoperation.

BC Behaviour Cloning.

CDiT Conditional Diffusion Transformer.

CEMP Cross-Entropy Method with Policy.

CNN Convolutional Neural Network.

CVAE Conditional Variational Autoencoder.

DDIM Denoising Diffusion Implicit Models.

DDP Distributed Data Parallel.

DDPM Denoising Diffusion Probabilistic Models.

DexWM Dexterous Manipulation World Model.

DiT Diffusion Transformer.

DP Diffusion Policy.

HDF5 Hierarchical Data Format 5.

I-JEPA Image Joint Embedding Predictive Architecture.

JEPA Joint Embedding Predictive Architecture.

KL Kullback-Leibler.

LeWM LeWorldModel.

MSE Mean Squared Error.

PPO Proximal Policy Optimisation.

RMSE Root Mean Square Error.

SSL Self-Supervised Learning.

V-JEPA2 Video Joint Embedding Predictive Architecture 2.

V-JEPA2-AC Action-Conditioned Video Joint Embedding Predictive Architecture 2.

VAE Variational Autoencoder.

ViT Vision Transformer.

VLA Vision Language Action Model.

List of Notations

I Identity matrix; gives the isotropic unit Gaussian covariance.

$\mathcal{N}(0, I)$ Standard isotropic Gaussian prior.

ε Reparameterisation noise, $\varepsilon \sim \mathcal{N}(0, I)$.

μ, σ Encoder-predicted mean and standard deviation.

z CVAE latent variable sampled from the prior $\mathcal{N}(0, I)$.

$\bar{\alpha}_t$ Cumulative noise-schedule product at diffusion timestep t .

ε_θ Noise-prediction network parameterised by θ .

t_{prev} Next step index in a subsampled inference schedule.

$\hat{x}_0$ Predicted clean sample in a DDIM reverse step.

x_T Noise sample at the diffusion schedule endpoint T .

B_{global} Effective global batch size; equals $B \times G \times A$ where B is the per-device batch size, G the gradient-accumulation steps, and A the number of accelerator devices.

E Effective training exposure, measured as total training items consumed, equal to optimisation steps multiplied by effective global batch size.

K Number of sampled candidate action chunks scored for one observation window.

N Number of evaluated samples or windows, depending on the metric context; pairwise comparisons report the number of held-out windows or candidate pairs explicitly.

S Number of optimisation steps.

$\mathbf{a}_{t:t+15}$ Sixteen-step future action chunk predicted from sixteen observed frames; each step has 18 action dimensions.

a_t Action vector at timestep t ; in this thesis each wrist-action step is an 18-dimensional bimanual vector, with left and right hands each represented by 3 translation and 6 rotation coordinates.

z_t Latent embedding of observation frame or video state at timestep t ; dimensionality depends on the encoder family.

1 Introduction

Reinforcement learning and imitation learning provide two important routes to acquiring behaviour in robotics. Reinforcement learning improves a policy through trial and error against a reward signal, while imitation learning acquires behaviour directly from demonstrations [1, 2]. This dissertation works primarily in the imitation setting, which is attractive for manipulation because useful behaviour often depends on subtle motion patterns that are difficult to encode by hand, and demonstrations capture those patterns without requiring a hand-designed reward. Behaviour cloning provides one direct route by learning a mapping from observations to demonstrated actions, while inverse reinforcement learning approaches the same problem by trying to infer the objectives that make a demonstration sensible [2, 3]. Both perspectives frame learning as extraction from observed behaviour, but neither resolves how a system should judge whether a predicted future motion is physically plausible. Reward-driven adaptation returns later in the dissertation, when closely related work that fine-tunes policies inside learned world models is examined.

This issue becomes sharper in dexterous manipulation. Human hand motion is temporally structured and object dependent. This follows from the timing of natural prehension, the visuomotor transformation required for grasping objects, and the postural synergies used to shape the hand for tools and objects [4–6]. In imitation settings, this matters because small prediction errors can compound across a future sequence, so local agreement with a demonstration does not guarantee sequence-level alignment [7]. This thesis is therefore concerned with prediction and evaluation of future hand motion, rather than deployed control. The work remains relevant to robotic control because robotic manipulation systems ultimately require reliable motion selection, but the evidence in this dissertation is restricted to an offline computational benchmark.

Visuomotor policies are a common way to connect visual observations with future motion. This dissertation studies two policy families that span the main design axes of current practice: the ACT, which regresses temporally extended action chunks in a single deterministic pass [8], and DP, which samples action chunks from a learned generative process [9]. Both share the chunked formulation, which suits hand motion because it unfolds over short sequences rather than through single-frame decisions. However, a policy that proposes future chunks still requires a criterion for choosing between candidate futures, and this is the opening for world models. World models learn predictive representations of how a state may evolve, so they can judge a candidate motion by the future it implies. In this benchmark the world model never generates motion of its own. Every candidate in the pool is proposed by the policy, and

the world model acts only at the moment of selection. The comparison then varies how much weight the policy’s own preference carries. One branch keeps the policy’s preferred candidate available while the world model re-ranks the pool, and another removes that preferred candidate so the world model must rank the remaining proposals without candidate 0 being available, though those proposals still originate from the policy. This graded transfer of the final decision from policy to world model motivates the comparative structure of the thesis.

This dissertation studies that structure through a comparative benchmark framework for evaluating whether world-model support changes the quality of future hand-motion prediction. The work began as a broad comparison across several world-model architectures, and during implementation that plan exposed a second research problem, since published world-model systems are not equally reproducible, action-compatible, or fair to compare under one offline benchmark. The study therefore goes beyond training models and reporting scores. Closely related systems are reproduced and examined, each candidate route is instrumented with diagnostic side experiments, and every shortfall is traced to a concrete cause such as an unavailable pretrained checkpoint or an incompatible action space. The final contribution is a controlled benchmark with one clean positive pairing, several boundary cases, and a fully documented account of what world-model adaptation requires in practice before it can support headline claims.

1.1 Motivation

The motivation for this work sits at the intersection of robotics and machine learning. Modern manipulation research increasingly depends on learned models that can process visual observations, predict future motion, and adapt to complex physical interaction. At the same time, the pace of open-source machine learning development can make it difficult to separate robust scientific progress from promising but weakly benchmarked demonstrations. This concern is especially important for industrially relevant settings, where unreliable claims about physical reasoning or manipulation capability can lead to poor engineering decisions.

The direction of the work follows a wider academic debate in the field. LeCun argues that systems which generate actions without predicting the consequences of those actions are limited as planners, a critique he directs at generative VLAs approaches to robotic manipulation [10, 11]. His proposed alternative places a learned world model at the centre of the agent, so that candidate actions are evaluated through their predicted future states before any action is taken. That position gained a concrete instance in DiWA, which fine-tunes a diffusion policy entirely inside a frozen world

model’s imagination using reinforcement learning [12]. Reading DiWA sharpened the question this scholarly debate raises into a testable one: DiWA uses the world model as a training substrate that updates the policy, so a complementary study could ask whether a frozen world model can instead improve the action selected at inference time, without retraining the policy. That question, motivated by this academic disagreement about how a world model should be used rather than by any single deployment need, came to form this work.

Beyond this academic motivation, the work was also shaped by the RSO I activity and the University of Malta M.AI.ESTRO project, where the convergence of robotics, machine learning, and Industry 4.0 raised a practical need for clearer evaluation of emerging model families. This practical context motivated the choice of domain and dataset, not the research question itself, which stands on the scholarly disagreement above independently of any one project’s needs. This dissertation does not claim to deliver a deployment system; the M.AI.ESTRO connection is noted here for context and returns only in the scope discussion of future work.

1.2 Research Question and Scope

The central research question is:

Under a controlled offline benchmark, does a frozen world model, acting only as the final decision stage, improve which candidate bimanual hand-motion chunk is chosen from a pool already proposed by ACT or DP?

This question is answered for the specific policy and world-model combinations this thesis implements and evaluates, not for “world models” as a universal class; Section 1.2.1 lists which combinations those are, and the scope paragraph below states plainly what a positive result on one combination does and does not license.

The phrase *final decision stage* marks the boundary of the claim. The world model is not integrated into the policy and never updates its weights, in contrast to adaptation methods such as DiWA where the world model serves as a training substrate for the policy itself [12]. Here every candidate motion is produced by a frozen policy, and the world model only chooses among those candidates.

This wording is a rescoping of the original proposal question. The proposal asked whether integrating a latent world model into an imitation learning policy would improve physical grounding and robustness, measured through latent agreement error and task success under perturbations. The final implementation does not evaluate deployed task success or perturbation robustness, and latent mean-squared error is used only as a training diagnostic for the world-model transition head rather than as

the main thesis result. The finished thesis therefore frames the question around offline prediction quality, asking whether world model support selects future hand-motion chunks that are closer to demonstrated motion under a shared benchmark.

The benchmark focuses on future bimanual hand motion and transform state. In implementation terms, the main prediction surface is built around future left-hand and right-hand transform chunks, but the detailed representation is defined later in Chapter 4 and Chapter 5.

The scope is deliberately bounded. The thesis does not test closed-loop robotic deployment, real-world task completion, force interaction, contact stability, or general machine common sense. It also does not claim that any one world model family is universally superior, or that a positive result for one tested combination generalises to world models as a class. Instead, it asks whether the specific world-model families implemented here improve an offline prediction benchmark when compared under a shared evaluation surface. The final evidence focuses on DP candidates scored by V-JEPA2-AC as the clean positive pairing. ACT, DexWM, LeWM, DINO-WM, JEPA-WM, PointWorld proxy controls, and a secondary DiWA/CALVIN reproduction and guided-selection study are used to define the boundaries around that result. The benchmark’s 7,607 held-out test cases are windows drawn from 3,229 underlying episodes (Section 4.1.3), several windows per episode on average, so they should not be read as fully independent samples; the statistical treatment in Chapter 4 accounts for this by clustering at the episode level rather than treating each window as an independent draw.

1.2.1 Objectives

The research question is pursued through five concrete objectives:

1. Design and implement a controlled three-branch benchmark that isolates candidate selection from every other source of variance between a policy-only baseline, a hybrid world-model-assisted branch, and a world-model-driven branch.
2. Train two visuomotor policy families (ACT, DP) and five world-model scorer variants (V-JEPA2-AC, V-JEPA2 no-AC, PointWorld, LeWM, DexWM) under a matched local training-exposure budget, so at least one source of variation, how much each family trains locally, is controlled for; families built on externally pretrained backbones still carry an external training advantage this local budget does not equalise (Section 4.5). DINO-WM and JEPA-WM are examined separately, as a smaller-scale pilot comparison outside this matched budget (Appendix A.4).
3. Determine whether world-model-guided selection improves offline trajectory

prediction accuracy relative to the policy-only baseline, and identify the specific conditions under which it does or does not.

4. Reproduce DiWA’s own weight-level world-model finetuning on CALVIN as an independently sourced comparison point against this thesis’s own external, inference-time selection mechanism.
5. For every route that does not yield a positive result, identify the likely reasons for the shortfall rather than reporting only the outcomes that worked.

Table 1.1 maps each objective to the chapters that carry it out.

Table 1.1 Alignment between the research question, the five objectives above, and the chapters that address each.

Obj.	Addresses	Chapters
1	Benchmark design isolating the selection mechanism	Ch. 4
2	Matched local training exposure across model families	Ch. 5
3	Whether and when selection improves prediction	Ch. 6, Ch. 7
4	Independent DiWA/CALVIN comparison	Ch. 6 (§6.4), Ch. 7
5	Diagnosed causes for every non-positive route	Ch. 5, Ch. 7

This scope also explains the structure of the comparison. The benchmark is organised around three branches:

- **Policy-only branch:** the baseline, where the policy’s preferred candidate provides the selected future motion without world-model involvement.
- **Hybrid branch:** the policy proposes a pool of candidate chunks, and the world model re-ranks the pool while the policy’s preferred candidate remains available as a fallback.
- **World-model-driven branch:** the policy’s preferred candidate is excluded from selection, so the world model ranks the remaining policy-generated proposals without candidate 0 being available. Those proposals still originate from the policy.

In every branch the candidate pool itself comes from the policy, and the branches differ only in how much authority the world model holds over the final choice. The purpose of this structure is to separate the value of the world model from the value of the underlying policy as far as possible within an offline benchmark. The description above is a basis framing, stated before any result is available; Section 7.2 revisits what removing candidate 0 turns out to mean in practice once the two branches can be compared.

1.3 Contributions

This thesis makes three distinct kinds of contribution: methodological, empirical, and implementation-level. Each is described in turn below, so that what is a new tool, what is a new result, and what is a new engineering record are not left for the reader to disentangle.

Methodological contribution. A comparative benchmark framework for studying world-model support in bimanual hand-motion prediction. The framework fixes the policy candidates, separates training support from held-out test windows, and records which information a selector is allowed to consume. This provides a controlled basis for asking whether predictive future-state models improve selected hand-motion chunks under the same offline evaluation conditions, and is reusable beyond the specific models evaluated here. Beyond the benchmark itself, the selection mechanism under test is a contribution in its own right, namely a frozen world model used purely as an inference-time re-ranking scorer over a policy’s own candidate proposals, rather than as a training-time substrate the policy is finetuned against, which is the role a world model plays in closely related work such as DiWA [12]. This external-selector formulation is evaluated here as a candidate method, not only as an experimental control against which DiWA’s approach is compared.

Empirical contribution. An implemented evidence hierarchy across the tested combinations. V-JEPA2-AC scoring DP candidates is the main positive semantic world-model result. LeWM action-prior scoring is a secondary action-regularity result. DexWM, DINO-WM, JEPA-WM, ACT, and the PointWorld proxy each return a negative or bounded outcome under the tested configuration, exposing limitations involving checkpoint availability, pretraining scale, action-space compatibility, candidate diversity, and evaluation fairness specific to that configuration, not a claim about the architecture in general.

Implementation and reproducibility contribution. Several model routes were implemented far enough to identify a likely reason they did not support headline claims under the configuration tested. These negative and mixed outcomes are reported as part of the contribution because they show what this thesis’s own evidence indicates is needed before world-model guidance becomes valid scientific evidence in a given configuration, not because they establish a necessary condition for every possible one.

1.4 Dissertation Structure

The remainder of the dissertation is organised as follows:

- Chapter 2 introduces the technical background required to follow the

dissertation, covering transformers, self-supervised learning, variational autoencoders, rigid body pose, visuomotor policies, and world models at a foundational level. Each concept is introduced only to the depth later chapters actually depend on, so readers already familiar with a given area can move quickly through the corresponding section.

- Chapter 3 positions the work within the relevant literature on imitation learning, visuomotor policies, world models, and offline evaluation. It closes by identifying the specific gap this thesis addresses and the conceptual framework, candidate diversity, action conditioning, and action compatibility, that the methodology and later chapters build on directly.
- Chapter 4 defines the benchmark design, scope, branch structure, leakage controls, metrics, and reporting rules. It states precisely what "fair" and "controlled" mean in this study's own terms, so later chapters can be read against an explicit, checkable contract rather than an implicit standard.
- Chapter 5 describes the implemented system, the model-specific routes, and the reproduction constraints encountered during the study. It also documents the concrete engineering problems found and fixed along the way, since several early results looked valid before a defect was traced and corrected.
- Chapter 6 presents the benchmark results for every tested policy and world-model combination, from the clean positive pairing through to the negative and mixed outcomes. Each result is reported with its confidence interval; where relevant, it also includes a ranking-validity check on whether the scorer's ranking is genuinely predictive rather than only correlated with an improved outcome.
- Chapter 7 interprets the findings against the conceptual framework introduced in Chapter 3 and states the main limitations. It separates what the experiments directly show from the interpretation offered for why, and states explicitly which parts of that interpretation remain untested.
- Chapter 8 closes the dissertation by answering the research question under the tested conditions and reporting, for each objective in Section 1.2.1, whether it was fully, partly, or only in a limited sense achieved.
- Chapter 9 identifies future work, grouped by how directly each direction follows from this thesis's own open questions: work needed to validate the current findings further, experiments that would isolate the mechanism behind a given result more precisely, and longer-term directions the current evidence motivates but does not itself test.

- Appendix A documents auxiliary implemented modules excluded from the final benchmark, and labels them explicitly as exploratory rather than validated components. Appendix B reports the validity, scaling, and sensitivity studies that justify choices fixed in Chapter 4, including which of those choices were fixed before the corresponding sensitivity analysis was run and which were examined afterward.

2 Background

This chapter assumes familiarity with the basic mechanics of a neural network and introduces the concepts central to this work: the main learning paradigms, visuomotor policies, rigid-body pose in $SE(3)$, latent representations, conditional VAEs, transformer attention, and self-supervised video prediction. These foundations support understanding of the specific architectures in Chapter 3 and underpin the candidate selection mechanism and world model evaluator design examined in Chapter 4. Together, they explain the thesis question, namely whether a latent world model, acting as the final decision stage, can improve the selection of bimanual action chunks proposed by frozen visuomotor policies before any action is executed.

2.1 Learning paradigms

Machine learning methods can be grouped by how they receive feedback during training, as illustrated in Figure 2.1. Supervised learning trains on data with explicit labels, so the correct answer is provided for every input and the model learns to reproduce it. Unsupervised learning receives no labels at all and instead recovers structure, such as clusters or compact descriptions, from the data alone. Reinforcement learning replaces labels with a reward signal, so the model acts in an environment and learns which state-action sequences lead to good outcomes.




Figure 2.1 The three classical machine learning paradigms: supervised learning trains on labelled data, unsupervised learning extracts structure from unlabelled data, and reinforcement learning learns from a reward signal over states and actions [13]. SSL, central to the world models in this work, is a modern form of learning from unlabelled data in which prediction targets are constructed from the data itself [14].

Reinforcement learning is defined by the interaction loop between an agent and an environment. At each step the agent observes a state, selects an action, and receives a scalar reward, and the objective is a policy that maximises the expected cumulative reward over time [1]. Learning therefore proceeds by trial and error rather than by copying labelled examples. This allows an agent to improve beyond the

behaviour it started with, but it requires an environment, real or simulated, in which the consequences of actions can be observed. The requirement matters later in the dissertation, since the offline benchmark studied here provides no such environment, while closely related work discussed in Section 3.2.1 fine-tunes a policy with reinforcement learning inside a learned world model instead of the real world.

SSL sits between the supervised and unsupervised settings. It trains on unlabelled data by constructing prediction tasks from the data itself, so part of the input is hidden and the model learns to predict it from the remainder [14]. Behaviour Cloning (BC), introduced in the next section, is a direct example of the supervised setting in which the expert demonstration acts as the label, while the world models studied in this dissertation are trained in the self-supervised setting. The close of this chapter develops that idea for video.

2.2 Visuomotor policies

A visuomotor policy is a learned function that takes a visual observation as input and produces an action or sequence of future actions as output. In a manipulation setting, the visual observation is typically a single camera frame showing the hand and the object being manipulated, and the action describes how the hand should move next. The policy is called visuomotor because it connects vision directly to motor output, without requiring a hand-crafted intermediate representation of what the scene contains [15].




Figure 2.2 A visuomotor policy maps an EgoDex-style camera frame to future bimanual wrist actions in the thesis setting.

Learning this mapping from data is appealing because useful hand behaviour depends on subtle visual cues that are difficult to specify by hand. The timing of a reaching movement, the grip shape adopted before contact, and the postural adjustments made in response to object geometry all depend on information visible in the camera image but difficult to encode as explicit rules [4–6]. A visuomotor policy learns these associations directly from recorded demonstrations, capturing the relationship between what the camera sees and what the hand does without requiring the programmer to describe it explicitly.

2.2.1 Behaviour cloning and compounding error

The most direct way to train a visuomotor policy from demonstrations is BC [2]. Here an observed frame predicts the action taken by the expert at that moment, and the policy is trained to minimise the difference between its predicted action and the recorded expert action across all frames in the dataset [7]. The problem arises when a small prediction error causes the hand to move to a state that is slightly different from any state in the training data. This results in a compounding error, since the policy encounters increasingly unfamiliar states and its prediction errors grow larger over time.




Figure 2.3 Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [16].

2.2.2 Action chunking

Action chunking addresses the compounding error problem by changing what the policy predicts. Instead of predicting a single next action at each step, the policy predicts a short sequence of future actions in one forward pass. This sequence is called an action chunk [8, 9]. In this benchmark, the policy observes 16 frames and predicts a 16-frame future action chunk. Each predicted step is an 18-dimensional vector, made up of two hands, each with a six-dimensional rotation representation and a three-dimensional translation. This is why Figure 2.2 shows the action sequence as a stack of multiple action blocks. By committing to a chunk of k future steps, where k is a hyperparameter set per task, the policy makes one prediction decision every k frames rather than one per frame. With fewer decision points, there are fewer opportunities for errors to accumulate, and the hand stays closer to the training distribution for longer [8].

2.3 Rigid body pose and the $SE(3)$ representation

To predict how a hand will move, a system needs a precise way to describe where the hand is and which way it is pointing. These two quantities together are called the pose of a rigid body. For the human wrist, pose is fully described by its position in space and its orientation relative to a fixed reference frame [17].

Translation. Position is represented as a vector of three numbers, one for each spatial dimension: how far the wrist is to the right or left, how far forward or backward, and how far up or down [17]. This vector $\mathbf{t} = [t_x, t_y, t_z]^\top$ is straightforward to predict with no special constraints on the output.




Figure 2.4 Translation of the wrist origin along each of the three spatial axes, $\mathbf{t} = [t_x, t_y, t_z]^\top$, following the $SE(3)$ pose convention [17].

Rotation. Orientation requires more care. A common representation uses three angles known as Euler angles, one per axis. These are intuitive but problematic for learning. At certain configurations the axes align and the system loses a degree of freedom, a singularity called gimbal lock. A subtler problem is discontinuity, since a wrist rotated 359 degrees and one rotated 1 degree are almost physically identical, yet their angle values are numerically far apart. A network minimising prediction error will produce large errors near these boundaries even when the true motion is smooth [18].




Figure 2.5 Rotation of the wrist about the y -axis by angle θ ; the same construction applies about x and z , and the wrist's full orientation is the combined rotation about all three, represented as a rotation matrix $R \in SO(3)$ [17].

A rotation matrix R avoids both issues, since small physical rotations produce small numerical changes in the matrix entries, and there are no singularities.

Homogeneous transform. Translation and rotation are combined into a single 4×4 matrix, the standard representation of a pose in the special Euclidean group $SE(3)$:

$$T = \begin{bmatrix} R & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix} \quad (2.1)$$

The benchmark does not predict the full human body. It predicts left- and right-wrist motion. For each future step, the benchmark stores one rotation and one translation per wrist, giving the 18-dimensional bimanual action vector used by the policy and world-model scorers.

2.4 Latent representations

A latent representation is a learned internal description of an observation. An image starts as a grid of pixel values, where each pixel stores three or four channel intensities (typically red, green, and blue, with an optional transparency value). A model does not usually reason with these raw numbers directly. Instead, it transforms them through several layers into a new set of values that no longer correspond to individual pixel coordinates or colours. These new values form a compact encoding that captures patterns in the image such as shape, position, motion, and spatial structure. The collection of these encodings is often called the latent space. Figure 2.6 shows an example in which similar items cluster together in the learned representation.




Figure 2.6 Example latent-space plot illustrating the general idea that similar examples can form nearby clusters in a learned representation [19].

The purpose of this transformation is to keep task-relevant structure while discarding unnecessary visual detail. For a manipulation scene, a useful latent state might encode hand position, object pose, contact progression, or motion trend without preserving every texture or lighting variation. A latent state is therefore not directly interpretable in the way an image is, but it can still preserve the information needed to reason about what happens next.

This is why latent prediction is often a better fit than pixel-level prediction for high-dimensional video observations. Predicting the next frame at the pixel level requires a model to account for every visual detail of the scene, including lighting changes, texture variation, and background movement that carry no information about task-relevant dynamics. A latent predictor instead operates on compact embeddings and concentrates predictive capacity on the structure that determines what actually happens next [20]. This design trade-off motivates the world model architectures reviewed in Chapter 3.

2.5 Conditional variational autoencoders

A VAE maps an input to a distribution of latent vectors described by a mean μ and standard deviation σ , rather than a single point. Sampling different latent vectors lets the decoder produce different but plausible outputs rather than one fixed answer. Figure 2.7a shows this standard encoder-latent-decoder flow [21].




Figure 2.7 Standard VAE (a) and conditional CVAE (b) encoder-decoder pipelines [22].

For the policy to be trainable end-to-end, gradients must flow from the action prediction loss all the way back to the encoder. Sampling z directly from a distribution blocks this path. The reparameterisation trick resolves this by writing $z = \mu + \sigma\epsilon$, where ϵ is standard normal noise drawn independently of the learned parameters. Figure 2.8 shows the full encoder-latent-decoder pipeline alongside its training loss: a reconstruction term that rewards accurate output and a regularisation term that keeps the latent distribution close to a standard Gaussian, together forming the evidence lower bound [21]:

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}[\log p_{\theta}(\mathbf{x} | \mathbf{z})] - D_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z})). \quad (2.2)$$




Figure 2.8 VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation) [21].

A CVAE adds an extra conditioning input to both encoder and decoder, shown in Figure 2.7b. In a visuomotor policy, this condition is the current observation frame. The latent variable therefore only needs to capture how the future action varies within that scene [22]. ACT’s CVAE trains this way, but at inference the original architecture sets the latent variable to the prior’s mean and decodes a single deterministic action chunk [8]. This thesis instead samples multiple latent vectors from the same observation and decodes each into a separate candidate, which is how ACT’s candidate pool is supplied to the world-model scorers evaluated later (Section 3.1.1), and it is also why ACT’s own candidate pool can fail in the specific way examined in Section 5.4.1. If training collapses the decoder’s reliance on this latent variable, the sampled candidates stop being meaningfully different from one another.

2.6 Transformer networks and the attention mechanism

A transformer processes a sequence as a set of tokens. A token can represent a word, an image patch, a timestep in an action sequence, or any vector produced by an earlier model component. The architectural idea needed here is simple: each token is updated by looking across the other tokens and deciding which of them are most relevant. This relevance-weighted exchange of information is called attention [23].




Figure 2.9 Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [23].

Attention forms three learned views of every token:

- **Query** (Q) – what this token is looking for.
- **Key** (K) – what another token contains.
- **Value** (V) – what another token contributes if attended to.

Figure 2.9 shows how each token (highlighted) sends its query outward to every other token, receives weighted values in return, and produces an updated output. The bidirectional arrows indicate that every token attends to every other token simultaneously, and the crossing arrow beneath represents the weighted sum that combines those values into the final output. For each token, the transformer compares its query with the keys of all tokens. The resulting scores become weights, and those weights are used to combine the corresponding values. This operation is usually written as [23]:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right) V. \quad (2.3)$$

Transformers usually repeat this operation in several heads and add positional information so the model knows where each token sits in the sequence. Multiple heads allow different relationships to be represented at the same time, while positional information preserves order [23]. Attention over tokens is the mechanism DP’s transformer denoising backbone and ACT’s transformer decoder both build on (Section 3.1), and it is also the operation the latent world-model predictors reviewed in Chapter 3 use to relate a candidate action to the visual tokens it is meant to explain.

These concepts meet in self-supervised video prediction, the training signal behind the world models in this work. Given a sequence of observed frames, one or more future frames are hidden and the model learns to predict them from the frames that came before; the recording itself provides the target, so no annotation is needed. Predicting raw pixels runs into the ambiguity problem described in Section 2.4, where the model averages over possible futures and produces a blurred image. Predicting in latent space avoids this, since the model predicts the compact embedding of the future frame, discarding background motion and fine texture so that capacity is spent on the structure that determines what happens next [20].

When the predictor is also given the action the agent intends to take, the result is a world model: a learned function that predicts how the scene embedding will change in response to a proposed action. If a policy proposes five wrist-motion chunks, a latent world model can predict the future embedding of each and favour the candidate whose predicted future is most plausible under its learned dynamics. This ties the chapter together: visuomotor policies produce action chunks, $\text{SE}(3)$ describes the wrist motion being predicted, latent representations provide the scoring space, and self-supervised prediction supplies the world-model training signal. The architectures built on these foundations are reviewed in Chapter 3.

3 Literature Review

This chapter reviews the literature that leads directly to the central research question set out in Section 1.2, namely whether a world-model evaluator can improve action-chunk selection for bimanual EgoDex wrist-action prediction, and whether any such improvement holds once the underlying policy family changes. The reviewed work therefore has two roles: first, to justify ACT and Diffusion Policy as the candidate policy families; and second, to motivate latent, action-conditioned world models as candidate selectors. Whether an offline benchmark of this kind can be trusted at all is a question the surveyed literature also bears on directly, and it is picked up again once results are presented.

The review is also selective about evidential role. Some papers motivate the final positive test, especially Diffusion Policy and V-JEPA2-AC; others explain why negative or exploratory routes are still scientifically useful, especially DexWM, LeWM, DINO-WM, and JEPA-WM. DiWA is treated as the closest adjacent policy-improvement study because it uses a world model to improve Diffusion Policy performance, but its training-time adaptation setting differs from this thesis’s inference-time candidate selection setting. Table 3.1 summarises how each literature group maps to the benchmark design.

The chapter assumes familiarity with the technical foundations in Chapter 2; world models are introduced here.

Table 3.1 Literature roles in the benchmark design.

Work	Relevant strength	Relevance and limitation for this thesis
ACT [8]	Efficient bimanual action chunking with a structured latent variable.	Provides the transformer policy baseline, but its sampled candidates collapsed in the final implementation, limiting its use as a selector stress test.
Diffusion Policy [9]	Strong multimodal visuomotor policy with diverse sampled action chunks.	Main policy family for the head-line result, with higher inference cost than ACT.
V-JEPA2 and V-JEPA2-AC [24]	Internet-scale video representation and action-conditioned latent prediction.	Its action-conditioned head gives the most direct architectural match to the candidate-scoring task of any evaluator considered, motivating its inclusion; whether that match translates into a positive result is a question for Chapter 6, not assumed here.
DexWM [25]	Dexterous manipulation world model with hand-consistency supervision.	Important comparator, but no public pretrained checkpoint was available and full upstream-scale pretraining fell outside the fair local budget.
LeWM [26]	Lightweight end-to-end latent dynamics model.	Useful for action-prior evidence, but the visual rollout route did not support the main semantic-selection claim.
DINO-WM and JEPA-WM [27]	Official checkpoint route for image/JEP-style latent prediction.	Exploratory official-checkpoint comparators; action-interface mismatch required an adapter and results remain appendix-level.
DiWA [12]	Uses imagined experience to improve Diffusion Policy.	Closest related study; it adapts the policy at training time rather than selecting candidates at test time. Its released policies are examined in a secondary reproduction and guided-selection study.
SimplerEnv and WorldEval [28, 29]	Evidence that offline/simulated rankings can correlate with real-world outcomes.	Supports the offline-evaluation premise, while requiring careful claims about calibration and deployment validity.

3.1 Visuomotor policies

This thesis focuses on two prominent policy families in imitation learning for manipulation, which attack the same pair of failure modes, compounding error and multimodal demonstrations, from opposite representational directions. This section reviews each in turn and then examines what the comparative literature actually establishes about their relative strengths, which is less than is commonly assumed.

3.1.1 Action Chunking with Transformers (ACT)

ACT was introduced by Zhao et al. to address the compounding error and multimodal distribution challenges that make bimanual manipulation demanding for imitation learning, achieving success rates of up to 90% on fine-grained tasks with the A Low-cost Open-source Hardware System for Bimanual Teleoperation (ALOHA) platform [8]. In SE(3) wrist-pose prediction specifically, this thesis treats such pose deviations as important because they can alter the geometry of contact during grasping, and a policy trained with a naive regression loss produces averaged predictions that may not represent any plausible trajectory when the demonstrated distribution is multimodal [9]. ACT's CVAE architecture addresses both of these failure modes directly.




Figure 3.1 Architecture of the Action Chunking with Transformers (ACT) policy [8].

ACT does so through the CVAE design of Section 2.5, depicted in Figure 3.1. A CVAE encoder compresses the demonstrated action sequence and joint observation into a style variable z that captures the multimodal distribution of plausible futures, and a transformer decoder predicts the future action chunk from the visual observations, joint positions, and z . The encoder is discarded at test time. In the original formulation, z is set to the mean of the prior (zero) to produce a single deterministic action chunk per observation. At inference, Zhao et al. apply temporal ensembling, averaging overlapping chunk predictions with recency weights to smooth the executed motion [8]. When a pool of candidate chunks is required instead, as in this thesis,

multiple values of z can be sampled independently from the prior and each decoded into a separate candidate, all conditioned on the same observation (Section 4.3).

ACT carries known limitations alongside these strengths. Its candidate diversity depends on the decoder making meaningful use of z , and nothing in the standard CVAE objective guarantees this. The decoder conditions on both the observations and the latent variable but is never forced to draw on z when the observations alone suffice to reconstruct the demonstrated actions. If the decoder learns to ignore the latent variable, sampling different values of z no longer changes the predicted chunk and the candidate pool degenerates into copies of a single output; this failure mode is known in the VAE literature as posterior collapse [30].

The VAE literature also offers mitigations, including annealing the Kullback-Leibler (KL) weight during training [30] and reserving a minimum information budget for the latent through free bits [31], yet the standard ACT recipe trains with a fixed KL weight and adopts neither. Zhao et al.’s own ablation sharpens the picture of when the latent matters, since removing the CVAE objective makes almost no difference on scripted, deterministic demonstrations, but on human demonstrations success collapses from 35.3% to 2% [8]. The latent variable is therefore load-bearing exactly when the data is multimodal, which is also the regime in which its training is most fragile. Beyond the latent, chunk length is a fixed hyperparameter that must be tuned per task, and ACT additionally uses recency-weighted temporal ensembling to smooth overlapping predictions [8]; such averaging may introduce lag under abrupt environmental changes.

3.1.2 Diffusion Policy

Diffusion Policy approaches the same compounding-error and multimodality challenges from the opposite representational direction. Instead of a structured latent variable, it starts each prediction from Gaussian noise and iteratively denoises it into an action chunk, conditioned on the recent observations [9]. The authors attribute part of Diffusion Policy’s advantage to its ability to represent multimodal action distributions without collapsing onto a single deterministic regression target, and independent samples from the same observation form a naturally diverse candidate pool. Figure 3.2 shows the general formulation and its two published heads, a CNN variant with FiLM observation conditioning and a transformer variant with causal cross-attention; their architectural details are given in the original papers [9, 32, 33]. The transformer variant is the form used in this benchmark.

Chi et al. benchmark Diffusion Policy across twelve tasks from four robot manipulation benchmarks and report an average improvement of 46.9% over the prior state of the art, including implicit and mixture-density policy baselines [9]. This gain is




Figure 3.2 Diffusion Policy: general formulation (a), CNN variant (b), transformer variant (c) [9].

consistent with the multimodal-representation advantage described above. ACT is not among the compared methods, since the two papers appeared within weeks of each other and neither evaluates the other, a point taken up below.

The main limitation of Diffusion Policy is inference cost. Generating a single action chunk requires multiple sequential network evaluations, since standard Denoising Diffusion Probabilistic Models (DDPM) variants use on the order of 100 denoising steps compared with the single forward pass of ACT’s decoder, which substantially reduces the achievable control frequency [9]. The scale of the follow-up effort is itself evidence of how seriously the field treats this defect. Efficient transformer variants reduce diffusion-policy cost through mechanisms such as sparse denoising experts, noise-conditioned routing, and expert caching [33]. Consistency Policy distills the denoising chain into a one- or few-step generator [34]. The π_0 generalist policy goes further still, abandoning discrete-step diffusion for flow matching to reach real-time control rates [35]. Iterative denoising in its original form is thus increasingly treated as a starting point to be distilled away rather than a final design.

In practice, the two families suit different pressures. ACT’s single decoder pass makes it the cheaper option and well suited to structured, repeatable bimanual tasks, but the CVAE latent it depends on for diversity is exactly what degrades on highly multimodal human demonstrations, as Zhao et al.’s own ablation shows [8]. Diffusion Policy pays for its iterative denoising with inference cost, but that same process is what lets it explore multimodal action distributions without collapsing to an average [9]. Both are used extensively for hand and bimanual manipulation. Consistent with this trade-off, the ACT authors’ later ALOHA Unleashed system uses a diffusion policy for challenging bimanual tasks involving deformable objects and complex contact-rich dynamics [36]. What the main literature surveyed here does not offer, however, is a direct comparison of the two at the scale and setting this thesis studies. Diffusion Policy’s twelve-task study predates ACT and does not include it; ACT is benchmarked only against BC-ConvMLP, BeT, RT-1, and VINN; and later multi-policy studies such as Lee et al.’s behaviour-generation comparison include Diffusion Policy but still omit

ACT [8, 9, 37]. One study does compare the two directly on shared bimanual data, Drolet et al.’s small-data comparison of imitation-learning algorithms on a simulated dual-arm peg-insertion task [38], but at a much smaller scale, 200 demonstrations on a single contact-rich insertion task, than the human-video hand-motion setting studied here, and without a world-model-assisted branch. The natural three-way comparison a reader might expect at that scale, one policy against the other against a world-model-assisted version of either, is one I found no study in the literature reviewed here reporting.

3.2 World models

A world model is a learned predictor of how a system evolves over time. Given the current state and a proposed action, it estimates the state that is likely to follow. In a manipulation setting, the state may be a raw image, a latent representation of the scene (Section 2.4), or a structured description of hand and object configuration. Regardless of the representation chosen, the underlying principle is the same, since the model internalises the dynamics of the environment and uses them to project forward from any given state [39].

The case for world models is itself contested, and the disagreement has deep architectural roots, traced from a common ancestor in Figure 3.3. The generative camp treats scale as the path to competence, so VLAs such as RT-2, OpenVLA, and π_0 couple web-scale vision-language backbones to autoregressive or flow-based action decoding, and their authors report broad generalisation to novel objects and instructions as evidence that this route is working [35, 40, 41]. LeCun has argued the opposite, that architectures which generate actions without predicting their consequences cannot plan reliably, and that scaling autoregressive models toward physical intelligence is a dead end [10]. The joint-embedding branch, and the latent world models built on it, is the constructive form of that critique. This dissertation does not adjudicate the debate; it tests one specific claim from the world-model side, that predicted latent futures carry enough signal to rank candidate actions.




Figure 3.3 The post-AlexNet split of visual representation learning: a generative branch (top), culminating in diffusion models and the VLAs debated above, and a joint-embedding branch (bottom), culminating in the JEPA-style world models evaluated in this benchmark. Screenshot from [11].

World models become especially useful when a pool of candidate actions must be evaluated before any one is committed to. Each candidate is rolled forward through the model, producing a predicted future state; those predictions are ranked by predicted cost or consistency. The quality of this ranking depends heavily on the chosen representation, since pixel-space world models must reconstruct every visual detail and suffer from mode-averaging at multimodal transitions, while models that operate in the compact learned latent space introduced in Section 2.4 can concentrate predictive capacity on the dynamics that matter for task outcomes [20]. The following subsections survey how world models have been used to improve policies, the literature supporting the evaluator design, and the architectures implemented in the benchmark.

Two distinct kinds of evidence recur throughout this survey, and they are not interchangeable. Some results report physical or simulated task-success evidence, an executed policy actually completing or failing a manipulation task. Others report offline prediction-accuracy evidence, a world model’s predicted future matching a recorded or held-out trajectory more or less closely, with no execution involved. A result of the first kind speaks directly to deployment; a result of the second kind speaks only to representational fidelity, and the literature on objective mismatch reviewed below is precisely about how loosely the two can be connected. Each specific claim in this section names which kind of evidence it is.

3.2.1 World models as policy refiners

The most direct way to use a world model to improve a policy is to adapt the policy inside it. DiWA develops this idea on CALVIN, a long-horizon manipulation benchmark where a Franka arm operates a tabletop drawer, sliding door, lightbulb, LED, and coloured blocks, shipped with hours of unstructured teleoperated play alongside expert demonstrations [42]. DiWA treats the world model as a training substrate, fine-tuning a diffusion policy with reinforcement learning (Section 2.1) entirely inside a frozen world model’s imagination, following the recipe in Figure 3.4. That world model is a DreamerV2-style recurrent state-space model trained on roughly six hours of task-agnostic play, about 500,000 transitions, with reward supplied by a binary success classifier trained on latent embeddings of the expert demonstrations. The denoising chain of the diffusion policy is framed as a Markov decision process, letting proximal policy optimisation run on imagined rollouts while a behaviour-cloning regulariser keeps the policy close to its pretrained initialisation [12].




Figure 3.4 The three-stage DiWA recipe: a world model is trained on task-agnostic play data, a diffusion policy is pretrained by imitation in its latent space, and the policy is then fine-tuned with reinforcement learning on rollouts imagined entirely inside the frozen world model [12].

Across all eight CALVIN skills tested, this purely imagined fine-tuning raises task success substantially. Baseline success ranges from 35.63% to 62.55%, and DiWA fine-tuning lifts every skill above 74%, up to 91.95% on close-drawer, averaged over three seeds [12]. A model-free baseline that fine-tunes the same policy in the real environment needs hundreds of thousands to millions of physical steps to reach comparable performance, where DiWA uses none. The authors also report the first fine-tuning of diffusion policies for real-world robotic skills through an offline world model, improving three physical manipulation skills after imagined fine-tuning on a world model trained from roughly four hours of teleoperated play [12]. They are

explicit about the limits of this evidence. Because the world model is frozen, its errors persist through fine-tuning and can be exploited by the policy, and gains measured in imagination do not always survive execution on the physical robot [12].

3.2.2 World models as policy evaluators

WorldGym demonstrates that a learned world model used as a simulated environment for policy evaluation produces rankings that align with real-world task success with meaningful fidelity [43]. Candidate policies are evaluated not by deploying them physically but by rolling them out inside the model’s predicted environment, reducing reliance on repeated physical deployment.




Figure 3.5 Classic closed-loop control (a) versus CLOVER’s closed-loop visuomotor control (b) [44].

CLOVER extends this to closed-loop control, as shown in Figure 3.5. A visual planner generates a sequence of sub-goal frames; the executor measures the embedding-space error between the predicted sub-goal and the current observation. When the error exceeds a threshold, the system triggers an adaptive transition and replans rather than committing to the original trajectory. The feedback loop, absent from open-loop imitation policies, is the mechanism that translates world model prediction accuracy into task completion gain [44]. Together, WorldGym and CLOVER establish that world model rankings are informative both before and during execution.

The practical value of this evaluator loop is measurable. Table 3.2 reproduces CLOVER’s real-world evaluation on a multi-step robot manipulation suite, reporting the average number of consecutive sub-tasks completed (`avg_len`) for representative baselines and CLOVER [44].

Table 3.2 Average consecutive sub-task completions (`avg_len`) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [44].

Method	avg_len
ACT [8]	0.6
R3M (pretrained visual repr.) [45]	0.7
RT-1 [46]	1.1
CLOVER (closed-loop)	2.1

Alongside the task completion gain, CLOVER also reduces embedding-space prediction Root Mean Square Error (RMSE) relative to open-loop baselines, suggesting prediction accuracy and task success move together under the evaluator framing. This correlation is not guaranteed by construction. However, Lambert et al. document an objective mismatch at the heart of model-based reinforcement learning, since world models are trained to minimise prediction error, yet low prediction error does not by itself produce good control, and the correlation between the two can be weak [47]. A world model’s rankings therefore inherit no guarantee from its training objective, and whether they are informative, including whether they transfer to physical outcomes and remain reliable out of distribution, must be established empirically in each setting, a question revisited when results are presented.

3.2.3 Latent predictive architectures

The energy-based framework of LeCun and Dawid establishes why prediction should operate in latent rather than pixel space [20]. The same representational argument introduced in Section 2.4 applies directly to the prediction target, favouring compact embeddings over pixel reconstructions that force the model to account for irrelevant visual detail and average over all plausible continuations at ambiguous transitions.

The pixel-space route is nonetheless pursued at scale, and its supporters reach the opposite conclusion. GAIA-1 predicts future driving video directly [48], and Genie learns action-controllable interactive environments from unlabelled internet video [49], treating faithful visual generation itself as evidence that a useful world model has been learned. The joint-embedding position is that this fidelity is bought with capacity spent on task-irrelevant detail. The architectures reviewed below are built on that counter-position, as latent predictors, and this benchmark evaluates them on whether their predictions rank candidate actions well, not on whether they can generate a convincing video.

Image Joint Embedding Predictive Architecture (I-JEPA) establishes the




Figure 3.6 Three SSL prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [14].

joint-embedding predictive architecture as a viable approach to learning rich visual representations without any pixel reconstruction objective [14]. As Figure 3.3 traces, this predictive objective is the direct ancestor of most world models surveyed in this section. Panel (c) of Figure 3.6 shows the data flow directly. A context block x from the image is passed through the context encoder to produce latent tokens z_x . A lightweight ViT predictor then takes z_x together with positional tokens indicating where the masked target regions y are located, and outputs predicted latent representations $\hat{s}_y$ for those regions. A separate target encoder, whose weights are a slow exponential moving average of the context encoder, independently encodes the actual target pixels into ground-truth latents s_y . The loss $\text{MSE}(\hat{s}_y, s_y)$ is computed entirely in latent space: no decoder, no pixel generation, no reconstruction of lighting or texture. The outcome is a context encoder whose representations capture semantically meaningful structure (object shape, spatial layout, and scene dynamics) because that is all the predictor ever needs to get the target embeddings right. Training is approximately 2.5 times faster than pixel-reconstruction methods as a result. V-JEPA2 extends this same joint-embedding predictive architecture from images to video sequences. The following subsections cover the fundamental literature of the architectures explored throughout this benchmark.

V-JEPA2

V-JEPA2 extends the I-JEPA framework from static images to video, pretraining on over one million hours of internet video [24]. This base pretraining is action-free. The context encoder receives masked spatio-temporal patches of video, and the predictor recovers the masked patches' representations, regularised against a momentum (EMA) target encoder in the same way as I-JEPA. The result is a frozen ViT-Large encoder, approximately 300 million parameters, that encodes physical priors about scene and object dynamics acquired from internet-scale pretraining. As shown in Figure 3.7, the pretrained backbone supports three post-training pathways: Language Alignment, which grounds visual embeddings in natural language; Attentive Probe, which adds a lightweight classification head; and Action Conditioning, which forms the basis of V-JEPA2-AC.




Figure 3.7 V-JEPA2 post-training pathways from internet-scale backbone to specialised applications [24].

V-JEPA2-AC

V-JEPA2-AC keeps the frozen V-JEPA2 ViT-Large encoder as its backbone but is a genuinely separate post-training stage, not a lightweight add-on to the base pretraining. The published system post-trains an autoregressive, approximately 300-million-parameter transformer predictor on roughly 62 hours of unlabelled robot interaction video from the Droid dataset [24]. Training combines two regimes: a teacher-forcing loss, where the predictor is always shown the true preceding frame’s embedding at each step, and a multi-step rollout loss, where it instead conditions on its own earlier predictions. The former keeps short-horizon predictions accurate, and the latter keeps longer-horizon trajectory representations consistent once prediction errors are allowed to compound. At planning time, the original work scores candidate actions by the distance between their predicted future embeddings and the embedding of a given goal image [24].

An important qualification is raised by Tsagkas et al., who show that pretrained visual representations can encode spatial plausibility well while providing a weaker signal for fine-grained temporal dynamics [50]. A frozen encoder trained on internet video may represent hand shape and position accurately at each individual frame while providing a less reliable signal about whether the predicted trajectory is kinematically consistent across the full chunk horizon. This temporal limitation provides context for the path shape results in Chapter 7.

DexWM

DexWM [25] is a dexterous manipulation world model from Meta FAIR, shown in Figure 3.9. Each input frame is encoded by a frozen DINOv2 image encoder into a grid




Figure 3.8 Standard V-JEPA2 (left) versus action-conditioned V-JEPA2-AC (right) [24].

of spatial patch tokens, a CDiT predictor advances the latent state under action and camera-motion conditioning, and a lightweight decoder maps the result back to image and keypoint space. Its distinguishing feature is the Hand Consistency loss, an auxiliary head that predicts fingertip and wrist heatmaps from the latent state, imposing a hand-geometry prior that is absent from both V-JEPA2-AC and LeWM. Goswami et al. report two task-success figures from physical execution: an average improvement of over 50% over Diffusion Policy across grasping, placing, and reaching tasks, and an 83% zero-shot real-world grasping success rate under latent planning [25]. Both are task-success evidence, not offline prediction-accuracy evidence. These figures come from a single study with no public checkpoint at the time of writing and have not been independently reproduced. Where V-JEPA2-AC transfers a video-pretrained backbone, DexWM transfers one pretrained on static images, so the comparison in this benchmark isolates the contribution of video-specific temporal pretraining.




Figure 3.9 High-level layout of the DexWM world model pipeline, illustrating how predicted keypoints are incorporated alongside the encoder, CDiT predictor, and decoder stages [25].

LeWorldModel

LeWM is a latent world model trained end-to-end from random initialisation on raw pixel observations [26]. As shown in Figure 3.10, a ViT-Tiny encoder of roughly 15 million parameters maps each frame to a single 192-dimensional token, and a dynamics predictor estimates the next latent from the current token and the action. Training combines a next-embedding prediction loss with SIGReg, a regulariser that enforces approximate Gaussianity of the latent distribution along random univariate projections. This single regularisation term stabilises end-to-end training without contrastive pairs or a moving-average target encoder, reducing the tunable loss hyperparameters of prior JEPA-style designs from six to one [26].




Figure 3.10 LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [26].

The design trades spatial resolution for speed. LeWM trains on a single GPU in a few hours; at planning time, its compact latent achieves a $48\times$ speedup over models that predict over frozen DINOv2 patch grids, while still detecting physically implausible transitions and tracking agent and object locations across frames [26]. As with DexWM, these results are self-reported by a single group, and the model is recent enough that independent evaluations do not yet exist.

Read together, the three architectures occupy different points on a cost-versus-prior trade-off, though only LeWM’s speedup claim is directly benchmarked against a comparable model in its source paper. V-JEPA2-AC carries a heavy but frozen video-pretrained Vision Transformer (ViT)-Large backbone with only a lightweight trained transition head. DexWM is the most resource-intensive of the three, since on top of a similarly heavy DINOv2 backbone it adds a second learned CDiT predictor and an auxiliary hand-geometry head, the extra machinery behind its reported Hand Consistency gains. LeWM, unlike DexWM, discards pretraining altogether for a from-scratch ViT-Tiny encoder, which its authors report is far cheaper to train and run, consistent with the $48\times$ planning speedup above. This compact from-scratch route is what the benchmark tests against the pretrained-transfer routes of V-JEPA2-AC and DexWM, a comparison returned to with results in Section 6.1.

DINO-WM and JEPA-WM

Two further latent world models enter the benchmark as exploratory official-checkpoint routes, and each pairs with one of the architectures above. DINO-WM performs latent prediction over a frozen DINOv2 patch grid [27], the same backbone family DexWM builds on, and Figure 3.3 places both on the self-distillation branch rather than the predictive JEPA branch. JEPA-WM denotes the officially released JEPA world-model checkpoints trained on robot manipulation data [51, 52]. Like LeWM, it sits directly on the predictive branch established by I-JEPA above, though as an independently pretrained checkpoint rather than a from-scratch design. Both DINO-WM and JEPA-WM provide pretrained checkpoints, but their expected action interfaces do not match the 18-dimensional bimanual wrist space, so each is evaluated through an adapter and treated as an exploratory comparator paired with its architectural relative. Their architectures and validation status are documented in Appendix A.

PointWorld

PointWorld is a large pretrained 3D world model for in-the-wild robotic manipulation [53]. Unlike the latent-space models described above, it operates directly

in 3D geometry. Scene state is lifted from RGB-D images into a full-scene point cloud, and robot actions are represented as the 3D point flows traced by robot surface geometry through forward kinematics. This embodiment-agnostic representation allows the model to reason over physical interaction geometry rather than learned abstract features.

Such a model is not designed to handle an 18-dimensional bimanual wrist-pose action space directly. PointWorld expects RGB-D observations and represents actions through robot-geometry-conditioned 3D point flow, whereas this benchmark provides RGB observations and 18-dimensional bimanual wrist transforms without a corresponding robot surface model, so a direct application would require an interface adaptation.

3.3 A common comparison of the world-model scorers

The subsections above describe each world-model family on its own terms, architecture by architecture. Read only that way, the differences between them can look like an unordered list of design choices rather than a small number of axes that actually predict how each family will behave as a candidate scorer. Table 3.3 places the five evaluator families reviewed above, plus DiWA’s own world model, on four shared axes: whether the model natively conditions on a candidate action at all, where its representation and any pretraining come from, what its latent representation actually is, and what kind of evidence its source work uses to evaluate it.

Two patterns stand out. First, every family that is genuinely action-conditioned without an adapter, V-JEPA2-AC, DexWM, LeWM, and DiWA’s own world model, was evaluated in its source work with some form of task-relevant evidence, either task success directly or a ranking objective tied to a goal; DINO-WM and JEPA-WM, evaluated here only through an adapter, are the two families whose source-work evidence this thesis cannot independently confirm transfers to the 18-dimensional bimanual interface at all. Second, only DexWM and DiWA’s own world model report physical or simulated task-success evidence directly; V-JEPA2-AC and LeWM’s own source works report offline evidence only, which is exactly the gap this thesis’s own offline benchmark is positioned to speak to for those two families, without extending that speaking-to into a task-success claim itself.

Table 3.3 Common-criteria comparison of the world-model families reviewed in this chapter. "Action-conditioned" asks whether the model's predictor natively takes a candidate action as input, not whether an adapter can supply one. "Evaluation evidence" uses the task-success versus offline-prediction-accuracy distinction introduced above.

Model	Action-cond.	Training data	Latent representation	Evaluation evidence (source work)
V-JEPA2-AC	Yes, native	Internet-scale video pretraining ($>1\text{M h}$) + a locally trained head	ViT-Large patch tokens (video-pretrained)	Offline: candidate ranked by distance to a goal-image embedding
DexWM	Yes, native	DINOv2 static-image pretraining + official EgoDex+DROID predictor (checkpoint unreleased)	DINOv2 patch tokens + CDiT predictor state	Task-success: physical grasping/placing success rates
LeWM	Yes, native	From scratch, no external pretraining	Single 192-d token, ViT-Tiny encoder	Offline: planning speedup and implausible-transition detection, not task success
DINO-WM	Via adapter only (native 7D DROID)	DINOv2 patch grid, official DROID/RoboCasa checkpoint	DINOv2 patch grid	Offline, official benchmark (not independently reproduced here)
JEPA-WM	Yes, native (7D DROID)	Gated DINOv3 backbone + official robot manipulation data	JEPA-style predictive embedding	Offline, official benchmark (not independently reproduced here)
DiWA's world model	Yes, implicit via imagined rollout	$\approx 6\text{h}$ task-agnostic play ($\approx 500\text{k}$ transitions), no internet-scale pretraining	Recurrent state-space latent (DreamerV2-style)	Task-success: CALVIN skill success after PPO finetuning

3.4 Research gap and conceptual framework

Three observations from this chapter, taken together, define the gap this thesis addresses. First, the two prominent visuomotor policy families, ACT and Diffusion Policy, are rarely compared directly on shared bimanual data, and never at the scale and setting this thesis studies (Section 3.1.2); each is benchmarked against different baselines in its own paper, and even the one existing small-scale comparison stops short of the natural three-way comparison a reader might expect, one policy against the other against a world-model-assisted version of either. Among the literature reviewed here, I found no study that reports that comparison. Second, world models are used in the literature in two disconnected ways: as a training-time substrate that updates policy weights (DiWA, Section 3.2.1) or as a task-success evaluator inside a closed-loop control system (WorldGym, CLOVER, Section 3.2.2). Neither published route tests the narrower, cheaper mechanism this thesis asks about, namely a frozen world model re-ranking a fixed pool of candidates a policy has already produced, without retraining the policy and without closed-loop execution. Third, Table 3.3 shows that the world-model families with the strongest published task-success evidence, DexWM and DiWA’s own model, are exactly the ones this thesis cannot reproduce at full published scale or fidelity (an unreleased checkpoint, and a training-time integration this thesis’s offline setting does not support respectively), while the families this thesis can evaluate cleanly at inference time, V-JEPA2-AC and LeWM, have only offline evidence behind them in the literature. No existing study closes this particular loop, comparing ACT and Diffusion Policy directly, under a frozen, inference-time, world-model selection mechanism, for the specific families where the published evidence is representational rather than task-success.

Closing that gap requires a way to predict, before running the benchmark, which world-model and policy combinations should be expected to work at all. Three properties recur across the discussion above as the ones that actually distinguish a world model that can plausibly help candidate selection from one that cannot.

- **Candidate diversity.** A scorer can only be as useful as the pool it selects from, and Section 3.1.1’s account of ACT’s narrow-multimodality risk versus Diffusion Policy’s sampling diversity (Section 3.1.2) shows this is already a live concern for the policies themselves, before any scorer is involved.
- **Action conditioning.** Table 3.3 shows this is not universal even among the families reviewed here, and a scorer with no native way to consume a candidate action, or only an adapted one, has no obvious mechanism to prefer one candidate over another based on what that candidate actually does.
- **Action compatibility.** Even a natively action-conditioned scorer was trained on some action distribution, and Lambert et al.’s objective-mismatch finding

(Section 3.2.2) is precisely a warning that a scorer’s own training objective carries no guarantee its ranking stays informative once the candidates being scored fall outside what it was trained to judge.

These three properties are the conceptual framework the methodology in Chapter 4 is built around and the results in Chapter 6 are read against; whether they turn out to be necessary, sufficient, or merely correlated with a positive result in the specific combinations this thesis tests is an empirical question this review does not presume to answer in advance.

This framework, together with the gap identified above, motivates the research question and objectives set out in Section 1.2, namely whether a frozen world model, scored against these three properties, improves candidate selection for ACT and Diffusion Policy specifically, under the controlled offline benchmark Chapter 4 defines next.

4 Methodology

This chapter defines how the central research question (Section 1.2) was translated into a controlled, testable benchmark. Answering whether a frozen world model can improve candidate selection, and whether that holds across policy families, requires isolating the selection mechanism from every other source of variance. Section 4.1 sets out the dataset and action space; Section 4.2 the three branch definitions; Section 4.3 the policy and candidate-generation routes; Section 4.4 the evidence roles assigned to each world-model family; Section 4.5 the fairness and leakage rules that separate held-out evidence from exploratory implementation work; and Sections 4.6 and 4.7 the evaluation metrics and scope limitations.

4.1 Dataset and Experimental Setup

4.1.1 EgoDex: source and selection rationale

The benchmark is built on EgoDex [54], a large-scale egocentric dataset of bimanual dexterous manipulation collected using Meta Quest Pro head-mounted cameras. It spans more than 100 task categories across its full release, including folding, slotting, pouring, and kneading, covering a wide range of fine-grained hand motions. The full release comprises five training parts and supplementary additional-data splits; this study uses the first three training parts (Section 4.1.3). Each episode provides synchronised RGB video and full six-degree-of-freedom $SE(3)$ wrist annotations for both hands at approximately 30 frames per second.

The choice of dexterous manipulation as the evaluation domain is not incidental. This dissertation forms part of the University of Malta’s ongoing M.AI.ESTRO research project [55], which investigates the use of AI for scene and task prediction on dexterous manipulation tasks to support industrial transfer learning and human skill digitisation. M.AI.ESTRO’s immediate pilot targets human hand pose prediction and guidance from egocentric demonstration; a second, later pilot addresses robotic hand manipulation, and the intended pathway between them translates poses learned from human demonstration into low-level actuator commands for the robot controller, rather than having the system generate robotic actions directly. The research domain for this dissertation was therefore set by the scope of the first pilot. The introduction situates this motivation more broadly, and the present chapter focuses on the methodological consequences of the domain choice.

Within the dexterous manipulation setting, EgoDex was selected on three properties:

- The benchmark required a fine-grained bimanual hand pose action space, provided through an eighteen-dimensional bimanual wrist pose representation that large-scale corpora built around rigid grippers or arm contact do not provide.
- The first-person perspective aligns with the intended downstream deployment setting of a wearable assistive device (a VR guidance application).
- The corpus is large enough to support simultaneous training of five world model families and two policy architectures on declared non-overlapping splits, while retaining a held-out test set of meaningful size.

EgoDex’s demonstrations are human rather than robotic, and it carries no contact or force signal, only kinematic $SE(3)$ pose. Neither is a limitation as scoped here; both are discussed further in Section 4.7. This study collects no new human-subject data and uses EgoDex as a pre-existing public dataset; no additional participant recruitment or data collection was undertaken.

4.1.2 Windowing protocol

The dataset is stored as paired `.mp4` and `.hdf5` episode files. A deterministic index stage discovers all valid paired episodes, infers their lengths from HDF5 metadata, and writes canonical observation and prediction windows for each split. This shared index is the single source of truth for all downstream modules: policy training, world model extraction, and evaluation all consume the same window file for their respective split, ensuring that output differences across the experimental branches can be attributed to branch design rather than to differences in the data presented to each component. The benchmark uses a fixed-stride windowing scheme as the canonical contract. A supplementary state-change-targeted scheme was also developed during pipeline iteration, to correct the canonical scheme’s under-sampling of high-motion segments, but was not adopted for the canonical benchmark. Its design and a standalone comparison against the canonical scheme, run to check whether it would have been worth adopting, are both reported in full in Appendix B.1.

Strategy 1: Fixed-stride sampling. Each window spans 16 observation frames followed by 16 prediction frames, with consecutive window starts separated by a stride of 128 frames. At 30 frames per second this places starts approximately 4.3 seconds apart, corresponding to an observation context and prediction horizon of approximately 0.53 seconds each. This contract is encoded as `obs16_pred16_s128` and forms the foundation for all policy training and world model training (see Figure 4.1).




Figure 4.1 Fixed-stride windowing scheme (obs16_pred16_s128) used as the canonical training and evaluation index.

4.1.3 Training splits and test set

Three cumulative training splits were assembled from the first three parts of the EgoDex release, each a strict superset of the previous; Table 4.1 summarises their composition. This three-part subset is used by design rather than the full five-part release, chosen for the broader task-category coverage each successive part adds (26, then 39, then 63 task categories), not because translation accuracy was found to improve monotonically with corpus size under the fixed training-exposure budget available on the shared cluster GPUs. Appendix B.2 reports this directly. The two policies trained across all three splits show a flat or worsening trend in translation error as the corpus grows, and the three world-model scorers were only ever evaluated at the final split, so this appendix cannot speak to their own scaling behaviour. How the three splits relate to each other during training is addressed in Section 4.5.1. The held-out test split is the official EgoDex test set and is never used during training at any stage. The B2 selection-margin δ (Section 5.6) is calibrated on a separate, disjoint internal validation split carved from `train_part1_2_3` itself, not on this held-out test split, so the reported headline results are evaluated only once, on data no calibration decision has seen.

Table 4.1 EgoDex splits under the obs16_pred16_s128 windowing contract. Episode and window counts are derived from the canonical index.

Split	Tasks	Episodes	Windows	Role
<code>train_part1</code>	26	46,091	140,239	Stage 1 training
<code>train_part1_2</code>	39	140,654	291,369	Stage 1+2 training
<code>train_part1_2_3</code>	63	194,333	430,962	Canonical training
<code>test</code>	111	3,229	7,607	Frozen evaluation

The canonical results in Chapter 6 use `train_part1_2_3`. The test split spans 111 task categories, including all 63 present in `train_part1_2_3` plus 48 drawn from the excluded parts; this exposes the benchmark to some task-level distribution shift, though it is not a controlled generalisation test, since the excluded categories were never held out as an experimental variable. Its window index is frozen throughout.

4.2 Three-Branch Experimental Framework

Answering the central research question (Section 1.2) cleanly requires holding as many other experimental factors constant as the design allows. The benchmark therefore organises evaluation around three branches that share the same policy weights and the same evaluation windows; only the candidate pool composition and selection rule differ across branches by design (Figure 4.2), so a measured performance difference between them is attributable to the selection mechanism rather than to policy weights or evaluation data, though it cannot rule out every other possible source of variance between models, such as architectural differences between world-model scorers themselves (Section 4.4).

The candidate pool size is fixed at $K = 5$ throughout this study. No literature precedent dictates this figure specifically. A common-random-number sweep over $K \in \{2, 3, 4, 5, 6, 8, 10, 12, 16, 20\}$, detailed in Appendix B.3, confirms that the ranking of world-model scorers is identical at every pool size tested, so no comparative conclusion drawn in this work is an artefact of fixing $K = 5$. Given this, K is treated as a compute-budget variable rather than one derived from theory or tuned for effect. The limitation is discussed further in Section 4.7.




Figure 4.2 Three-branch evaluation framework (Section 4.2), shown for a policy producing five independent candidates. B2 and B3 reuse the same candidate pool generated for B1, relayed through the branches (dashed lines) rather than regenerated; B3 additionally removes the policy’s preferred candidate before scoring. Both ACT and DP follow this design; the per-policy candidate-generation mechanics are detailed in Section 4.3.

Branch 1 (B1): Policy-only baseline. The policy generates its full candidate pool and selects by its own internal preference; no world model is queried. The ideal design

produces several structurally independent candidates so the downstream branches can compare meaningfully different trajectories. Neither policy family provides a calibrated preference score among independently sampled candidates by design, so “internal preference” is operationalised per policy (Section 4.3). For Diffusion Policy specifically, the published architecture does not itself rank independent samples, so this thesis defines the heuristic as the candidate with the lowest final denoising residual [9]. B1 establishes the trajectory prediction quality each policy achieves acting alone and is the baseline against which B2 and B3 are measured.

Branch 2 (B2): World-model-inclusive selection. The world model scores all candidates in the pool and selects the highest-scoring one unconditionally; no margin threshold is applied. The candidate pool is the same as B1. Scoring proceeds by predicting forward in the world model’s latent space from the current observation and measuring prediction fidelity against each candidate action. The policy’s own output is not privileged; it is selected only if the world model assigns it the highest score. For V-JEPA2 without action conditioning, the transition head required to produce candidate-specific latent predictions is absent; if the model assigns indistinguishable scores to all candidates, B2 degenerates to B1, directly testing whether action conditioning is the operative mechanism.

Branch 3 (B3): World-model-exclusive selection. The policy’s preferred output is excluded from the evaluation pool; the world model selects the highest-scoring remaining candidate unconditionally. The ideal design ensures the remaining candidates are structurally independent so that exclusion meaningfully changes the selection outcome. B3 tests whether the world model maintains selection quality when the policy’s top choice is unavailable. All five world model families apply this protocol.

The three branches select by *argmax* at three different points in the pipeline, and the distinction matters: B1’s *policy-preferred* *argmax* is made without reference to any external, world-model-derived score, using the per-policy internal-preference rule defined above; B2’s *unconditional* *argmax* is the world model’s choice over that same pool; and B3’s *independent-pool* *argmax* is the world model’s choice once the preferred candidate is removed, over the remaining structurally independent candidates only. Both policy families evaluated, ACT and DP (Figure 4.2), follow this design, each generating its own five-candidate pool for B1 and B2 and scoring the remaining four for B3, using the different sampling procedures detailed in Section 4.3.

4.3 Policy Variants

The benchmark evaluates two imitation-learning policy families. Their architectures and training objectives are reviewed in Chapter 3; the properties relevant to this study

are how each generates its candidate pool and what that implies for world-model scoring.

- **Action Chunking Transformer** (Section 3.1.1). ACT’s trained CVAE [8] is used here to produce diverse candidates, by sampling latent vectors $z \sim \mathcal{N}(0, I)$ from its prior and decoding each into a 16-step bimanual $\text{SE}(3)$ action chunk, in place of the single deterministic chunk the original inference procedure decodes from the prior’s mean. Its single-shot decoding makes it the cheaper of the two policies to run at candidate-generation time. The trained architecture itself follows the original directly, unmodified, and is described in Chapter 5, including a checkpoint-level CVAE posterior-collapse issue and the inference-time candidate-augmentation mitigation it required (Section 5.4.1), neither of which changes the architecture itself.
- **Diffusion Policy** (Section 3.1.2). DP [9] generates each candidate through an independent reverse-diffusion chain, using a Denoising Diffusion Implicit Models (DDIM) schedule [56] and the transformer (Diffusion Transformer (DiT)) variant. Its iterative denoising is designed to capture multimodal demonstrations without averaging across modes, at a higher inference cost than ACT. The implementation follows the original formulation directly and is described in Chapter 5.

Evaluating both families under the same three-branch protocol allows the study to ask whether world-model augmentation improves candidate selection in general or only under a particular generation mechanism. A result that holds across both constitutes stronger evidence for the generality of the approach.

4.4 World Model Families

The model families have different evidence roles in the final thesis. They are not treated as equal candidates for the headline claim. A result is promoted to main evidence only when it uses frozen candidates, training-only support, held-out test windows, and a scorer whose action representation is compatible with the candidate pool.

- **V-JEPA2-AC** (Section 3.2.3), the headline scorer. A frozen V-JEPA2 [24] ViT-Large backbone provides a visual embedding, and a lightweight transition head predicts a future embedding from the current embedding and the candidate action chunk; candidates are scored by distance to the reference latent. It combines a large pretrained video representation with explicit action conditioning, the most direct match to the candidate-scoring task.

- **Mechanistic controls:** V-JEPA2 without action conditioning and PointWorld. Both reuse the V-JEPA2-AC visual pathway; the former removes candidate-specific action information, the latter degrades the bimanual action chunk into an incompatible robot-arm-style representation. Together they test whether a positive V-JEPA2-AC result depends on meaningful action conditioning rather than the visual backbone alone.
- **LeWM** (Section 3.2.3) [26], a secondary action-regularity scorer. It is implemented first as a visual rollout and latent-support world model, evaluated against the same causal-validity gates as V-JEPA2-AC before being promoted to semantic evidence. A narrower action-prior variant, using training-only action statistics to prefer candidates closer to the training distribution, is retained as a secondary result, a valid action-regularity signal, distinct from a visual-rollout imagination claim.
- **DexWM** (Section 3.2.3) [25], a fixed-budget negative comparator. It is the closest architecture-level comparator for dexterous world modelling. Public pretrained weights for the upstream EgoDex+DROID recipe are unavailable, so the thesis evaluates a reproducible fixed-budget implementation, treating its result as evidence only for that implementation; full official pretraining is excluded because it needs DROID raw-data onboarding and a training scale outside the fair-budget contract.
- **DINO-WM and JEPA-WM** (Section 3.2.3) [27, 51], exploratory official-checkpoint comparators. Both require an adapter from this study’s 18-dimensional bimanual wrist chunks to the 7-dimensional DROID-style action interface the checkpoints expect. Because of this interface mismatch, they are treated as validation-only routes and reported in the implementation appendix rather than promoted to held-out evidence.
- **DiWA/CALVIN** [12, 42], a secondary reproduction and comparison. DiWA studies world-model-guided Diffusion Policy improvement through imagined training, whereas this study evaluates test-time candidate selection without retraining the policy; the CALVIN study compares the two improvement operators under shared components. The methodology’s plan is to first fully reproduce DiWA’s published results under its original frozen world model, reward classifier, and policy checkpoints, then compare that unmodified upstream baseline against this study’s own test-time selector applied to it, against the paper’s own imagined-PPO fine-tuning, and against this study’s selector applied to the finetuned checkpoint in turn, under identical evaluation episodes. This is a secondary study, kept separate from the primary EgoDex branch definitions

rather than a substitute for it. Implementation status is reported in Chapter 5, and the six resulting conditions and their results are reported in Chapter 6.

Table 4.2 Evidence role assigned to each implemented or investigated world-model path.

Path	Evidence role	Main claim status
V-JEPA2-AC	Headline semantic scorer	Held-out test claim
V-JEPA2 no-AC	Action-conditioning ablation	Held-out control
PointWorld proxy	Action-degradation control	Held-out mixed control
LeWM visual rollout	Diagnostic scorer	No-go for semantic rollout claim
LeWM action-prior	Training-action regularity scorer	Secondary held-out claim
DexWM fixed-budget	Domain-specific comparator	Held-out negative/null claim
DINO-WM / JEPA-WM	Official-checkpoint exploratory routes	Validation-only, appendix
DiWA/CALVIN	Secondary reproduction and comparison	Reported separately in Chapter 6

4.5 Matched Local Training-Exposure Protocol

Comparing the five world model families and two policy architectures against each other requires holding as many non-design variables constant as this study’s local training controls this. Two aspects are addressed: training exposure and evaluation contract. Both are matched at the level of what this thesis itself trains locally. Several families, V-JEPA2-AC and DexWM foremost among them, build on backbones pretrained externally at a scale far beyond what any local budget could reproduce or equalise (Section 4.5.1); matching local exposure controls one source of unfairness between families, not every source, and this distinction is carried through the rest of this section rather than described as full fairness.

4.5.1 Training exposure budget

Different model families have different computational costs per training step and different sensitivities to data volume. Without a principled exposure budget, a comparison could conflate the effect of architecture with the effect of receiving more or fewer training samples. To determine an appropriate shared budget, a convergence study was conducted prior to full training, training each model family independently across six data scales and examining the resulting loss curves for plateau behaviour. The full study and its per-model results are reported in Appendix B.4; the finding relevant to the experimental design is summarised here.

The effective training exposure E is defined using the following symbols (see the List of Notations):

E	effective training exposure (items consumed)
S	number of optimisation steps
B	per-device batch size
G	gradient-accumulation steps
A	number of accelerator devices
B_{global}	effective global batch size

$$E = S \times B_{\text{global}} \quad (4.1)$$

where:

$$B_{\text{global}} = B \times G \times A \quad (4.2)$$

This formulation counts the total training items consumed regardless of how steps are distributed across hardware, allowing models with different batch sizes and GPU counts to be normalised to a common scale.

The convergence study established a universal budget of $E = 100,000$ effective items, chosen as follows. Pooled across the 3 training seeds used throughout this thesis, none of the five primary families shows a clear plateau within the six-scale sweep (Appendix B.4 gives the full pooled trend for all seven, including why ACT’s own seed-42 checkpoint looks flat in isolation): ACT changes by -2.8% from 50,000 to 100,000 items, LeWM by $+7.1\%$, Diffusion Policy by -14.9% , V-JEPA2-AC by -2.0% , and DexWM by -19.9% (single-seed) over the same doubling. $E = 100,000$ is therefore not a saturation point for any family. It is the largest exposure this study’s compute budget could apply uniformly across all seven families without giving any one family a disproportionate share of the available cluster time, accepting that every family’s reported result is a compute-constrained lower bound on what it could achieve rather than its ceiling (Section 4.7 returns to this point). This figure is applied uniformly across all model families and all three training splits, controlling for one specific confound, namely that a family trained on more local data or more local optimisation steps might simply have seen more, rather than learned better. It does not equalise every architecture’s total exposure to training data. V-JEPA2-AC and DexWM each build on a frozen backbone pretrained externally, on internet-scale video and static-image corpora respectively, well before this thesis trains anything on top of them (Section 4.4); LeWM, by contrast, trains entirely from this same $E = 100,000$ budget with no external pretraining at all. A measured difference between these families can therefore reflect architecture, or it can reflect which family arrived at the local budget already carrying an external, uncounted training advantage. The local budget is fair in the narrower sense that every family receives the same opportunity to adapt within this thesis’s own compute envelope; it is not a claim that every family started that adaptation from the same place. Each split is trained independently from a fresh local initialisation at this same budget; later splits do not continue training from an earlier

split’s checkpoint. Any difference between the `train_part1`, `train_part1_2`, and `train_part1_2_3` results therefore reflects the data available at that scale, not a warm start carried over from a smaller split. The full data-construction and training curriculum is described in Chapter 5.

4.5.2 Evaluation contract

All models receive identical inputs and evaluation conditions. At evaluation time, every family sees the same RGB observation frames from the canonical test window, without additional depth, point cloud, or semantic metadata; PointWorld’s distinctive characteristic is its degraded action conditioning rather than a different observation modality. The evaluation windows are fixed at indexing time, shared across all branches and families, and never selected, filtered, or re-sampled after training. Policy and world model weights are frozen before evaluation and are not updated. The only variable at evaluation time is the selection rule applied to the candidate pool.

The causal validator checks four leakage boundaries before a run can be treated as evidence:

1. Support episodes must not overlap with evaluation episodes.
2. Support windows must not temporally overlap the held-out test window.
3. Task labels may not be used to retrieve test support unless the same restriction is declared for every compared family; the final canonical runs use global retrieval rather than test-task-scoped retrieval.
4. Repeated-subject metadata are preserved when available but not used as a selection feature.
5. Threshold and hyperparameter choices, in particular the Branch 2 margin threshold (Section 5.6), must be fixed on a disjoint internal validation split rather than the frozen test set. Section 5.6 records one point in this thesis’s own development history where an earlier inspection did not fully meet this boundary, and how the final threshold was chosen.

If a model route cannot provide these guarantees, it remains exploratory or appendix material rather than a headline result. Every evaluation run also emits a frozen manifest recording its selection mode, candidate source, and aggregate score statistics, and writes per-window results as it goes so an interrupted run resumes from its last completed window rather than restarting.

4.5.3 Training-seed protocol

A single trained checkpoint cannot distinguish a genuine architectural result from an artefact of that one initialisation. To guard against this, every trainable family this study

can afford to retrain, ACT, Diffusion Policy, V-JEPA2-AC’s transition head, and LeWM, is trained at three random seeds (42, the canonical seed used throughout this thesis’s other design choices, plus 43 and 44) rather than one, and every headline result in Chapter 6 is reported as the mean across these three seeds rather than a single run. DexWM is the one exception, excluded from this protocol on compute-cost grounds detailed in Section 4.7 and Section 5.5.3; results tables involving DexWM state this exclusion explicitly rather than silently reporting a single-seed number alongside three-seed means. Three seeds is not a large cohort, and does not guarantee every possible training-time anomaly is caught; it is a compute-constrained choice, the same kind of choice already made for the candidate pool size (Section 4.2) and the training-exposure budget above. It does, however, provide a basic robustness check, and in this thesis that check surfaced a real seed-dependent effect for one family; the full per-seed picture, not only the resulting mean, is reported alongside the headline results in Chapter 6.

4.6 Evaluation Metrics

4.6.1 Offline trajectory prediction setting

EgoDex does not define an official offline evaluation metric. The original benchmark measures policy effectiveness via online task success rate on a physical robot deployment [54], which is outside the scope of this study. The evaluation setting here differs fundamentally. Rather than measuring task execution, the study evaluates offline visuomotor trajectory prediction, measuring how closely each configuration’s predicted wrist pose sequences align with ground-truth demonstrated trajectories in the held-out test set. Results should therefore be interpreted as trajectory prediction quality under each branch and world model configuration, not as claims about downstream task success. The relationship between offline trajectory prediction quality and online task performance is addressed in Section 4.7 and revisited when results are presented.

4.6.2 Primary metrics

The action representation throughout this study is a 16-step sequence of bimanual wrist actions, each wrist pose the 6-DOF position-plus-orientation representation introduced in Section 2.3. At step t , hand $h \in \{L, R\}$ has predicted translation $\hat{\mathbf{p}}_{t,h}$, target translation $\mathbf{p}_{t,h}$, predicted rotation $\hat{R}_{t,h}$, and target rotation $R_{t,h}$. A prediction can fail in two physically distinct ways. The hand can end up in the wrong place, or it can end up in the right place while pointing the wrong way. A manipulation policy can fail

at either independently of the other. The primary metrics below therefore report translation and rotation error separately, rather than folding them into one combined score.

Translation error, denoted e_{trans} or `translation_12` in tables throughout this thesis, is the mean per-axis root-mean-square wrist-position error, averaged over both hands and *all sixteen* predicted steps, never only the first:

$$e_{\text{trans}} = \frac{1}{2T} \sum_{t=1}^T \sum_{h \in \{L, R\}} \sqrt{\frac{1}{3} \sum_{i \in \{x, y, z\}} (\hat{p}_{t,h,i} - p_{t,h,i})^2}, \quad T = 16. \quad (4.3)$$

It is reported in metres. This is a per-axis RMSE convention, not the true Euclidean wrist-position distance; the latter is exactly $\sqrt{3}$ times larger for the same prediction, since the per-axis form divides the squared-error sum by 3 before taking the square root where the Euclidean norm does not. On average, e_{trans} is how far the predicted wrist position is from where the demonstration actually placed it, over the full 16-step horizon and both hands, expressed as a per-axis rather than a straight-line quantity. Task success itself is not measured in this thesis (Section 4.6.1), so no claim is made here about which of the two error types matters more for it; translation and rotation are reported separately because they capture physically distinct failure modes, not because one is asserted to dominate task outcomes.

Rotation error is the mean geodesic distance on $\text{SO}(3)$ [17], reported in degrees:

$$e_{\text{rot}} = \frac{1}{2T} \sum_{t=1}^T \sum_{h \in \{L, R\}} \cos^{-1} \left(\frac{\text{tr}(\hat{R}_{t,h}^\top R_{t,h}) - 1}{2} \right) \frac{180}{\pi}. \quad (4.4)$$

This measures, in degrees, how far the predicted wrist orientation is from the demonstrated one on average, using the geodesic (shortest-path) angle between two rotations rather than a per-axis angle difference. Per-axis Euler differences are exactly the representation flagged in Section 2.3 as discontinuous and prone to gimbal lock [18]; the geodesic distance on $\text{SO}(3)$ avoids both problems and gives a single physically meaningful angle regardless of how the underlying rotation is parameterised. A correctly positioned hand with the wrong orientation is a real manipulation failure in its own right; a wrist rotated the wrong way can miss a grip, collide with an object, or approach a handle from an angle that cannot close around it, even while sitting at the exact target position.

Root-mean-square error is used for diagnostics where a squared aggregate is more appropriate than a mean absolute distance, penalising a few large deviations

more heavily than many small ones:

$$\text{RMSE} = \sqrt{\frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2}. \quad (4.5)$$

Here M denotes the number of scalar components being compared.

Endpoint translation uses Equation 4.3 over the final four prediction steps only, capturing how accurately the trajectory arrives at its destination specifically, distinct from its accuracy earlier in the chunk. Velocity and acceleration diagnostics apply the same paired-error logic after first and second finite differences of the wrist position sequence, capturing whether the predicted motion moves at a demonstration-consistent speed and smoothness rather than only whether its sampled positions happen to be close.

4.6.3 Candidate selection comparison

For each window w , let $m(c, w)$ be the trajectory error of candidate c and let $c_0(w)$ be the policy-only candidate. A branch selecting $c_s(w)$ records a pairwise win when:

$$\mathbb{K}_{\text{win}}(w) = \begin{cases} 1, & m(c_s, w) < m(c_0, w), \\ 0, & \text{otherwise.} \end{cases} \quad (4.6)$$

The win rate is:

$$\hat{p}_{\text{win}} = \frac{1}{W} \sum_{w=1}^W \mathbb{K}_{\text{win}}(w), \quad W = 7607. \quad (4.7)$$

Mean deltas are reported as selected candidate minus candidate 0, so negative values indicate improvement. Every delta in Chapter 6 is reported alongside a 95% confidence interval on that estimate. Informally, under repeated sampling and the assumptions of the resampling procedure, 95% of intervals constructed this way would be expected to contain the target parameter. A narrower interval means the effect is more precisely estimated, and an interval that excludes zero provides evidence that the estimated effect is consistently different from zero under this resampling procedure. Confidence intervals are computed by clustered bootstrap over evaluation episodes [57]: 2,000 resamples, each drawing whole episode clusters with replacement, with the interval taken as the 2.5th and 97.5th percentiles of the resulting resampled means (a percentile bootstrap, not a bias-corrected or normal-approximation interval), using a fixed random seed of 42 for the resampling itself so every reported interval in this thesis is exactly reproducible from the same underlying per-window results. Task-clustered intervals [58] are used as a robustness check where task-level

dependence is the relevant uncertainty source. Small win-rate differences are not interpreted without their clustered interval or a corresponding paired delta analysis.

Every headline result pooled across the three training seeds (Section 4.5.3) uses a two-level extension of this same bootstrap rather than the single-level version above. Treating all three seeds' windows as one undifferentiated pool would overstate precision, since there are really only three independent draws (the seeds), not thousands (the windows); resampling episodes alone would silently assume the three seeds themselves carry no uncertainty. Each of the 2,000 resamples therefore proceeds in two steps: first, which of the three seeds contribute to this resample is itself drawn with replacement (so a resample can include a seed twice, or not at all); second, within each drawn seed's own run, episodes are resampled with replacement exactly as in the single-level version above. The reported interval is again the 2.5th and 97.5th percentiles of the resulting 2,000 resampled means. For the pooled three-seed analysis, this thesis extends the episode-cluster bootstrap hierarchically, resampling seeds first and episodes within each sampled seed. This follows the general rationale for clustered bootstraps with nested dependence structures [59]. This two-level resample is what "pooled across three seeds" means everywhere in Chapter 6 unless stated otherwise, not an average of three independently computed point estimates. Because only three training seeds are available, the top level of this hierarchy resamples from just three clusters; these intervals are best read as a descriptive uncertainty summary rather than a precise population-level inference over random initialisation.

A scorer that prefers one candidate over another is not automatically a scorer that ranks well. A preference can look decisive on paper while still failing to track which candidate is actually the better one, if what the scorer favours does not systematically correspond to lower trajectory error. Win rate and mean delta, defined above, describe the outcome of a single selection decision. They do not on their own say whether a scorer's ordering of candidates agrees with the candidates' true quality ordering across windows. A separate measure is needed for that question.

Spearman rank correlation provides it. For a set of n candidates, let $R_1(i)$ be the rank assigned to candidate i by the scorer's own score, and $R_2(i)$ the rank assigned by the candidate's true trajectory error. The correlation is:

$$\rho_s = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2 - 1)}, \quad d_i = R_1(i) - R_2(i). \quad (4.8)$$

ρ_s ranges from -1 (the scorer's ordering is exactly reversed) through 0 (no monotonic rank association between the scorer's ordering and true quality) to $+1$ (the scorer's ordering exactly matches the true one). Confidence intervals on ρ_s follow the same clustered-bootstrap procedure defined above, resampled over evaluation episodes [57, 58].

This metric matters wherever a scorer’s aggregate error reduction alone cannot settle whether the world model is doing genuine work. Chapter 6 applies it as a ranking-validity check throughout, most importantly for validating LeWM and the other world models on whether their candidate rankings are actually predictive rather than merely correlated with an improved outcome by some other route.

4.7 Evaluation Scope and Limitations

The following scope constraints and known limitations apply to the experimental design.

Offline evaluation only. No physical robot is used and no task-completion signal is available; all metrics are computed against held-out demonstrated trajectories (Section 4.6.1). Results should be read as evidence about candidate selection quality in the offline trajectory prediction setting, not as claims about downstream robot task success.

Kinematic action space and embodiment. The action representation is the bimanual SE(3) wrist space described in Section 4.4, with two properties worth noting rather than treating as open scope gaps:

- *No finger, contact, or force signals.* Finger-level articulation, contact forces, and haptic feedback are absent. Given how accurately EgoDex already fits hand pose, an additional contact or force channel would have been a valuable enrichment for finer-grained manipulation analysis, though the kinematic evaluation performed here does not require it.
- *Human-to-robot kinematic gap.* Demonstrations are produced by human subjects, whose wrist kinematics differ from robot manipulators in range of motion and joint compliance. This is not a limitation of the present study, which makes no robotic-deployment claim; it becomes relevant only once M.AI.ESTRO’s later robotic-translation pilot (Section 4.1.1) is pursued.

Training budget as a compute-constrained lower bound. The universal 100k-item budget (Section 4.5.1) is not a saturation point for any family, pooled across the 3 training seeds used throughout this thesis. Doubling the budget from 50k to 100k changes pooled validation loss by -2.8% for ACT, $+7.1\%$ (an increase) for LeWM, -14.9% for Diffusion Policy, -2.0% for V-JEPA2-AC, and -19.9% for DexWM (single-seed; Appendix B.4). Seed 42’s own ACT checkpoint looks flat in isolation because it has collapsed (Section 5.4.1), not because the family has converged; LeWM’s own trend is a small increase at 100k present across all three seeds, not a plateau. DexWM’s final doubling is larger than its preceding one (10.8% from 25k to 50k),

showing no sign of decelerating within the swept range. For every family the shared budget is therefore a genuine compute-constrained lower bound, not an approximate saturation point, so reported results likely understate what a longer training run would achieve, potentially substantially for DexWM. Performance differences may partly reflect differing distances from respective convergence points rather than intrinsic architectural differences.

Candidate pool size as a compute-constrained choice. The five-candidate pool size (Section 4.2) is a compute-constrained choice rather than a literature-derived one; the full sweep and its findings are detailed in Appendix B.3. Absolute effect sizes throughout should be read as a property of this budget, since margins grow monotonically, without a knee, as pool size increases. Comparative conclusions are unaffected, since scorer ranking is invariant at every pool size tested, but that invariance was only established on the Diffusion Policy row (Section 4.2), not the ACT row.

DexWM pretraining abstention. Published DexWM weights are unavailable [25]. A fixed-budget implementation was evaluated, but the full upstream EgoDex+DROID pretraining recipe was abstained from. The reasons are practical and methodological: DROID raw onboarding is storage- and preprocessing-heavy, DexWM inference and training are substantially more expensive than the other routes, the fixed-budget result did not suggest a likely headline gain, and full upstream-scale pretraining would move outside the common fairness contract used for the main benchmark.

Training-seed cohort size. The three-seed protocol declared in Section 4.5.3 is itself a compute-constrained cohort size, the same kind of constraint already applied to the candidate pool (Section 4.2) and the training-exposure budget above; three seeds is not large enough to guarantee that every possible training-time anomaly is caught, only enough to provide a basic robustness check. Whether three seeds is in fact enough to catch what needs catching is an empirical question this design choice cannot settle on its own. Chapter 5 reports which family the seed axis exposed a real training-time anomaly in, and Chapter 6 reports how that anomaly was handled and whether the resulting guard results are themselves stable once it is. DexWM is excluded from the seed protocol entirely, on compute-cost grounds detailed in Section 5.5.3, and every table involving it states this exclusion explicitly.

5 Implementation

This chapter describes what was actually implemented and what each route can support as evidence. The implementation does not end with a uniform comparison where every model family becomes a headline result. Instead, it produces an evidence hierarchy: V-JEPA2-AC is the main semantic world-model scorer, LeWM's rollout scorer is a genuine but negative visual-dynamics result, its action-prior scorer is an illustrative action-regularity reference rather than a comparable result, DexWM is a fixed-budget negative comparator, DINO-WM and JEPA-WM are exploratory official-checkpoint routes, and DiWA/CALVIN is a secondary reproduction and comparison study, reported separately from this evidence hierarchy.

The chapter is organised around six concerns: repository ownership and reproducibility (Section 5.1); data and windowing (Section 5.2); training and hardware configuration (Section 5.3); policy candidate generation (Section 5.4); model-specific implementation paths (Section 5.5); and implementation challenges that changed the scientific interpretation of the project (Section 5.7).

5.1 Repository and execution model

The research stack is organised as a single Python package under `src/`, with one top-level module per concern: `data/` handles indexing and windowing, `policy/` implements the two policy architectures and the world-model-aware guard, `world_model/` implements each world model family, and a `cli/` layer exposes every stage as a runnable command (Figure 5.1). A sibling `scripts/` directory collects the Docker flags, GPU assignment, and mount paths for a given model or operation into one launch script (e.g. `dexwm.sh`, `lewmm.sh`), backed by a small set of shared bash libraries and a generic, config-driven `run_stage.sh` driver for ad hoc runs (a one-off invocation with a specific flag combination that does not warrant its own dedicated launch script).

Docker is the canonical runtime surface for training and evaluation, built from a pinned `Dockerfile` on PyTorch with CUDA 12.4 support. Reproducibility depends on keeping code and generated output on separate, auditable layers. The version-controlled `src/`, `cli/`, and `scripts/` read and write `datasets/` (EgoDex episodes, the window index, cached backbone checkpoints), which in turn produce `runs/wact/` (checkpoints, manifests, metrics, and guard results, the generated evidence read by the thesis). Only the code is tracked by git; both data directories are regenerated by the pipeline rather than version controlled (Figure 5.2), and this separation holds regardless of which host machine executes a given stage. Pretrained backbones (V-JEPA2, DINOv2, DINOv3) are pulled once from PyTorch Hub or

HuggingFace and cached locally rather than re-downloaded per run. Individual routes use isolated dependency caches where an upstream project's requirements would otherwise conflict with the main environment, notably the JEPA-WMS and DiWA checkouts used in Sections 5.5.4 and 5.5.5.




Figure 5.1 Package structure under `src/`, launched via `scripts/` and `cli/` inside Docker.




Figure 5.2 Data flow from version-controlled code to the generated, non-version-controlled `datasets/` and `runs/wact/` artefacts.

The run output directory follows a split-first, then model layout. Each training job writes a `result.json`, model checkpoints, metrics, and a manifest. Canonical evidence additionally records candidate hashes, support-role metadata, window counts, and fair-evaluation receipts. These manifests are treated as part of the result, not as optional logs.

5.2 Data pipeline and windowing

The index stage discovers valid paired EgoDex episodes (a .mp4 video plus a matching .hdf5 pose file), infers episode lengths from the Hierarchical Data Format 5 (HDF5) metadata, and writes the canonical observation/prediction windows defined by `obs16_pred16_s128` (Section 4.1.2). It is resumable and excludes unpaired files up front, so a missing video or annotation surfaces here rather than as a downstream training error. The canonical index contains 194,333 training episodes and 430,962 training windows for `train_part1_2_3`, plus 3,229 held-out test episodes and 7,607 held-out test windows.

Every downstream component reads from this same index rather than re-deriving windows, which keeps the branch comparison of Section 4.2 attributable to branch design rather than to accidental differences in the data each component received. The two policy architectures consume the RGB frame and 18D action target directly; world model extraction additionally consumes the action-delta encoding used to build transition-head training pairs; guard evaluation consumes the frozen index to generate and score candidate pools.

Three cumulative training splits were assembled from the first three parts of the EgoDex release (Table 4.1 of the methodology). Each split is a strict superset of the previous. The larger splits were assembled with relative symbolic links rather than by copying data, avoiding duplication of approximately 1 TB of earlier data on disk while allowing the data-scaling axis of the benchmark to be evaluated directly, since all model families are trained independently on each split at the same universal budget (Section 4.5.1).

5.3 Training configuration

Table 5.1 consolidates the hardware and batch configuration for all model families at the universal budget $E = 100,000$. Settings marked (*shared*) are kept identical across all trainable models. The effective-item budget is verified for each trainable family via $E = S \times B_{\text{global}}$, where S is the number of optimisation steps and B_{global} is the effective batch per step. V-JEPA2 (no AC) and PointWorld require no model training and reuse the Phase 1 extraction shared with V-JEPA2-AC.

The canonical checkpoint for every family is trained at random seed 42; seed-robustness coverage beyond this canonical run is reported in Section 4.7. Training manifests record optimiser steps, effective exposure, device count, batch size, and checkpoint paths. V-JEPA2-AC uses AdamW for the transition head; the JEPA-WMS adapter calibration also uses AdamW with learning rate 10^{-3} . Every

Table 5.1 Training configuration for all model families at $E = 100,000$. Approximate step time is shown for trainable phases only (--- denotes not applicable).

Family	GPUs	Batch/device	Steps S	B_{global}	s/step
ACT	4×L40S	32	781	128	≈0.5
Diffusion Policy	4×L40S	32	781	128	≈0.8
V-JEPA2-AC Ph. 1 (extraction) [†]	2×L40S	—	—	—	—
V-JEPA2-AC Ph. 2 (AC head) [†]	1×L40S	32	3,124	32	<0.01
LeWM	4×L40S	4	6,250	16	≈0.047
DexWM [‡]	4→3×L40S	4	7,234	≈14	≈28
V-JEPA2 (no AC)	—	—	—	—	—
PointWorld	—	—	—	—	—

[†] Phase 1 is a one-time offline extraction (2×L40S, sharded, 27.4 windows/s, no backward pass). Measured wall times: `train_part1` ≈0.75 h, `train_part1_2` ≈1.4 h, `train_part1_2_3` ≈2.2 h. Phase 2 head training completes in under ten seconds per split (≈0.01 s/step on cached embeddings) on one GPU, using a deterministic budgeted-exposure sampler that oversamples with a fixed permutation-based schedule whenever raw training pairs fall short of the target, so its effective exposure matches the universal budget exactly (99,968 items), the same as every other family.

[‡] DexWM trained on 4×L40S for `train_part1` and `train_part1_2` (6,250 steps each, global batch 16); the split reported here, `train_part1_2_3`, resumed on 3×L40S after step 3,299 following a GPU-availability change, ending at 7,234 steps with the effective-item budget preserved. DexWM total wall time across all three training splits was approximately 150 h (≈28 s/step throughout, 6,250 + 6,250 + 7,234 steps); LeWM total wall time was approximately 15 minutes (≈0.047 s/step), substantially faster than DexWM despite the identical effective budget.

trainable family, ACT, Diffusion Policy, V-JEPA2-AC’s transition head, LeWM, and DexWM’s predictor, monitors a held-out validation loss during training and selects its best .pt checkpoint by the lowest value reached, exactly as described for ACT above; V-JEPA2 without action conditioning and PointWorld train nothing of their own and simply reuse V-JEPA2-AC’s already checkpoint-selected weights, so no separate selection applies to them. Where this per-model validation surface is unavailable for a specific diagnostic run, the frozen run manifest and hash instead identify the exact artefact used for evaluation.

None of the implemented routes natively consume EgoDex’s RGB and bimanual wrist-state observation, or emit an 18-dimensional, 16-step action chunk. Every family required a purpose-built input or output adapter. ACT and Diffusion Policy are architecture specifications rather than pretrained models with a fixed interface, and every world model family was pretrained (where a pretrained checkpoint exists at all) on a different observation or action convention. Table 5.2 records, per route, the interface actually consumed and produced and whether that interface required a custom adapter (Section 5.5 details each adapter).

Table 5.2 Model input/output contracts used by the implemented routes.

Route	Input	Output	Loss / score	Adapter
ACT / Diffusion Policy	RGB frame + wrist state	16-step, 18D action chunk	ACT: reconstruction + KL DP: diffusion denoising	18D bimanual action head
V-JEPA2-AC	V-JEPA2 embedding + 18D action chunk	Future latent embedding	Latent prediction distance	Action-delta encoder into frozen latent space
LeWM rollout	RGB/action sequence	Predicted compact latent	Latent-support score	—
LeWM action-prior	Candidate 18D action chunk	Scalar regularity cost	Training-action z-score prior	—
DexWM fixed-budget	DINOv2 patch tokens + pose actions	Future patch/key- point latent	CDiT + HC training losses	—
DINO-WM / JEPA-WM	Official visual latent + adapted 7D action	Future official latent	Support/action- policy score	18D→7D linear adapter

5.4 Policy training

Both policy families share the same observation encoder: a ResNet-18 visual backbone [60] processes the single observation RGB frame (96×96 pixels) through four residual stages (64, 128, 256, and 512 channels), producing a 512-dimensional feature vector that is concatenated with the bimanual wrist state before being passed to the policy-specific decoder (Figure 5.3). Image resizing to 96×96 is applied at loading time and is consistent across all branches. Wrist pose sequences are normalised to zero mean and unit variance using statistics computed from the training split; these statistics are frozen after training and consumed by the world model scoring stages to ensure consistency.




Figure 5.3 Shared ResNet-18 visual backbone preceding both policy decoders. Adapted from [60].

5.4.1 Action Chunking Transformer

The ACT architecture is reviewed in Section 3.1.1 and illustrated in Figure 3.1. The implementation follows the architecture and training procedure of Zhao et al. [8] with no dependency on an external codebase. Custom components include the EgoDex HDF5 window loader, the 18-dimensional bimanual wrist action space adapter, and a checkpoint serialisation layer that freezes the CVAE configuration and normalisation statistics alongside the model weights. The CVAE encoder-decoder and transformer decoder are realised directly from the paper specification. At training time, the CVAE encoder compresses a ground-truth action chunk together with the current observation into a style variable $z \sim \mathcal{N}(\mu, \sigma^2)$; the transformer decoder reconstructs

the chunk from z and the observation. The encoder is discarded at evaluation time, and z is drawn from the standard Gaussian prior $\mathcal{N}(0, I)$. The design intent is to generate up to N diverse candidates by sampling independent z vectors and decoding each under the same observation. Training uses four L40S GPUs (Distributed Data Parallel (DDP)) with a per-device batch of 32 (global batch 128): $E = 781 \times 128 \approx 100,000$ (Table 5.1). Each run writes a `best.pt` checkpoint selected by lowest validation loss, a `last.pt` checkpoint, per-epoch metrics in `metrics.jsonl`, and a frozen `manifest.json` recording the configuration and window file consumed.

CVAE posterior collapse: investigation and retraining. This is primarily a model and training limitation of ACT as trained here, not a property of the downstream candidates the CEMP augmentation below produces. The trained ACT checkpoint exhibits CVAE posterior collapse, so at inference time the decoder is latent-insensitive, and sampling five independent z vectors from $\mathcal{N}(0, I)$ produces five identical action chunks. A diagnostic confirms the extent of the collapse. Injecting $z = 10\sigma$ (ten standard deviations from the prior) via a forward pre-hook on the latent projection layer shifts the output by approximately 3×10^{-7} in absolute value, which is effectively zero. The mean pairwise difference across five independently drawn candidates is below 10^{-6} on all three training splits, i.e. within floating-point noise of exactly zero; this threshold is denoted ε throughout the remainder of this section.

This is a well-documented failure mode of conditional VAE models rather than a mis-implementation [30]. When the conditioning signal (the visual observation) is sufficiently informative to reconstruct the target without consulting z , the decoder learns to ignore z entirely. The KL divergence term in the training objective regularises the *encoder* distribution toward the prior but cannot compel the *decoder* to use the latent; the collapse is therefore not resolved by increasing the KL weight. Four retraining runs were conducted to test this:

- v1: `kl_weight=10` (default): collapse unchanged.
- v2: `kl_weight=100`: weighted KL contribution ($100 \times 0.25 \approx 25$ nats) exceeded reconstruction loss (~ 15), yet collapse persisted.
- v3: cyclical KL annealing [61] (`kl_cycle_batches=200`, `kl_cycle_ramp_frac=0.5`) and 30% observation dropout to force decoder reliance on z : collapse persisted.
- v4: cyclical KL annealing combined with free bits [31] (`free_bits=0.5` nats/dim, no observation dropout), run for completeness after v3: collapse persisted.

All four runs produced mean pairwise differences indistinguishable from zero on every split (Table 5.3). The EgoDex visual observations are sufficiently informative that the decoder finds a loss minimum that ignores z regardless of training pressure, and correcting the collapse without an architectural change (discrete latents, vector-quantisation, or an explicit diversity head) is outside the scope of this thesis.

Table 5.3 ACT CVAE posterior collapse: four successive retraining attempts, all indistinguishable from zero ($\varepsilon = 10^{-6}$) on every training split, at random seed 42.

Run	kl_weight	Cyclical KL	Obs. dropout	Free bits	Mean pair- wise diff
v1	10 (default)	no	0 %	no	$< \varepsilon$
v2	100	no	0 %	no	$< \varepsilon$
v3	100	yes (200-batch 50 % ramp)	cycle, 30 %	no	$< \varepsilon$
v4	100	yes (200-batch 50 % ramp)	cycle, 0 %	0.5 nats/dim	$< \varepsilon$

Is collapse deterministic, or seed-dependent? Table 5.3 tests all four configurations at a single random seed (42) each. This leaves open whether collapse is a deterministic property of a given configuration, in which case one seed is enough to characterise it, or whether it is itself seed-dependent, in which case a single seed under-samples how often a configuration actually collapses. To distinguish these, each of the four configurations was retrained at two further seeds (43 and 44) on the canonical `train_part1_2_3` split, and the same collapse diagnostic applied to all twelve resulting checkpoints. Table 5.4 reports the result.

Table 5.4 ACT CVAE posterior collapse across three random seeds per configuration, `train_part1_2_3`. *collapse* marks mean pairwise difference below the same $\varepsilon = 10^{-6}$ threshold used in Table 5.3; *ok* marks a checkpoint whose five sampled candidates remain meaningfully distinct.

Configuration	Seed 42	Seed 43	Seed 44	Collapse rate
v1 (default)	collapse	ok	ok	1/3
v2	collapse	ok	collapse	2/3
v3	collapse	ok	collapse	2/3
v4	collapse	collapse	ok	2/3

Collapse is seed-dependent, not deterministic, for every configuration tested. This matters directly for how Table 5.3 should be read. At a single seed each, none of the four configurations looks better than the others, all collapse, but that comparison itself only had one draw per configuration to work with. Across three seeds, v1, the default configuration already used for this thesis’s canonical checkpoint, has the lowest collapse rate of the four (1/3, against 2/3 for v2, v3, and v4). This does not mean v1 has solved collapse; one of its three seeds still exhibits it. It does mean the added training pressure in v2 through v4 buys no seed-robustness advantage over the simpler default, which is a more complete justification for retaining v1 as canonical than the single-seed comparison alone could offer.

Why seed identity matters more than configuration. Table 5.4 also shows a pattern worth reading by seed rather than by configuration. Seed 42 collapses under all four configurations tested (4/4). Seed 44 collapses under two (2/4, under v2 and v3). Seed 43 collapses under only one, and only the most heavily regularised of the four (1/4, under v4 alone). Read this way, which random seed a checkpoint happens to be trained with predicts whether it collapses better than which of these four regularisation strategies is applied. This is consistent with the mechanism behind CVAE posterior collapse [30]. Collapse results from a race, early in training, between the decoder learning to reconstruct the target from the conditioning signal alone and the encoder learning an informative latent posterior. If the decoder’s randomly initialised weights already reconstruct reasonably well from the observation before the encoder has learned anything useful to place in the latent, the KL divergence term drives the encoder’s posterior toward the prior with no reconstruction cost, and once the posterior matches the prior the decoder has no further incentive to consult the latent at all; this is self-reinforcing for the remainder of training. Both the initial network weights and the order training examples are seen in are set by the random seed, so which side of this race wins in the first few hundred batches, and therefore whether the run collapses, can differ by seed even under an identical configuration. This thesis’s own diagnostics do not extend to per-step encoder and decoder gradient tracking during training, so a precise mechanistic account of why seed 42 specifically loses this race under every configuration tested while seed 43 mostly wins it is not established here. What Table 5.4 does establish is that seed identity, not configuration choice, is the larger factor governing whether a given ACT checkpoint collapses.

Section 6.2 reports how this seed-dependence is handled at evaluation time: CEMP is applied only to checkpoints the collapse diagnostic actually flags, not uniformly across every seed, so a checkpoint whose native sampling already produces a diverse pool is evaluated on that native pool rather than having CEMP imposed regardless.




Figure 5.4 CVAE posterior collapse in the trained ACT checkpoint. Four independently sampled latent vectors produce an identical output $\hat{a}$ (dashed arrows indicate the latent path the decoder effectively ignores).

The collapse is treated as an empirical finding, illustrated schematically in Figure 5.4. It motivates the use of an external candidate generator, described below, and is reported in Section 6.2 as a characteristic of the ACT policy on this dataset.

CEMP candidate augmentation. To restore the five-candidate evaluation pool, ACT-based guards use a CEMP planner. This is an adaptation of the Cross-Entropy Method [62] initialised from the policy’s output rather than a uniform prior, inspired by the policy-guided model-based planning strategy used in TD-MPC2 [63]. A perturbation is a small SE(3) displacement applied to the base action chunk, shifting each wrist pose by a sampled translation and rotation offset while preserving temporal coherence across the 16-step chunk. Starting from the collapsed ACT B1 output, three iterations each draw 64 such perturbations around the current mean, score them with the world model’s latent prediction function, and retain the top 25 % (16 elites); the standard deviation decays by $0.7\times$ per iteration. The four highest-scoring candidates across all iterations are returned (Figure 5.5).

For Branch 2 the pool contains these four alternatives plus the ACT B1 output ($N=5$); for Branch 3 the B1 output is excluded, leaving the four alternatives only. Appendix B.5 reports CEMP’s compute cost against the raw pool it replaces.

Circularity caveat. The same world model that scores candidates inside CEMP’s three refinement iterations (Figure 5.5, from Iteration 1 onward) also performs the final B2/B3 selection over the resulting pool. This is a deliberate design choice, not an oversight. The ACT policy itself cannot generate diverse candidates (the posterior collapse above), so some external proposal mechanism is required; using the world model under test to shape that proposal is the only way to obtain a five-candidate pool at all for this policy family. The consequence is that ACT B2/B3 results are not a clean measurement of world-model scoring ability in isolation; they measure the combined world-model-plus-planner system, where the planner has already biased the pool toward the scorer’s own preferences before selection ever runs. This caveat does not

apply to Diffusion Policy, whose five candidates come from independently seeded DDIM chains with no world-model involvement in their generation. Any ACT B2/B3 advantage over B1 should accordingly be read as evidence for the world-model-plus-CEMP pipeline as a unit, not as an isolated measurement of scorer quality comparable to the Diffusion Policy branch.




Figure 5.5 CEMP candidate augmentation, refining SE(3) perturbations of the ACT B1 output over three iterations to form the B2/B3 evaluation pool.

5.4.2 Diffusion Policy

The Diffusion Policy implementation follows the transformer-backbone variant (DP-T) described in Section 3.1.2 and illustrated in Figure 3.2. The implementation follows Chi et al. [9] with no external codebase dependency. Custom components include the 18-dimensional bimanual wrist action space adapter, the EgoDex HDF5 window loader, and the five-candidate DDIM evaluation protocol (five independent noise seeds decoded under the shared conditioning signal). The same ResNet-18 observation encoder produces a conditioning vector `obs_cond` passed to a DiT denoising backbone via cross-attention. At evaluation time, five independent DDIM denoising chains are run on the same observation to produce five candidate action chunks, each starting from independently sampled Gaussian noise and converging to a distinct trajectory under the shared conditioning signal. Branch 1 selects among them by taking the chain with the lowest final denoising residual. Training uses four L40S GPUs (DDP) with a per-device batch of 32 (global batch 128): $E = 781 \times 128 \approx 100,000$ (Table 5.1). The same checkpoint, metrics, and manifest artefacts as ACT are written at the end of each run.

DDIM subsampling correction.

$\bar{\alpha}_t$	cumulative noise schedule product at step t (List of Notations)
t_{prev}	next step in the subsampled inference schedule (List of Notations)
$\hat{x}_0$	predicted clean sample (List of Notations)
x_T	Gaussian noise at the schedule endpoint (List of Notations)

An implementation audit identified an error in the reverse-diffusion sampler [56]. The original implementation passed $\bar{\alpha}_{t-1}$ (the consecutive index in the training schedule) instead of $\bar{\alpha}_{t_{\text{prev}}}$:

$$\text{Bug: } \bar{\alpha}_{t-1} \longrightarrow \text{Fix: } \bar{\alpha}_{t_{\text{prev}}} \quad (5.1)$$

For a ten-step schedule $t \in \{90, 80, \dots, 0\}$ this substitutes $\bar{\alpha}_{89}$ for the correct $\bar{\alpha}_{80}$ at the first denoising step, miscalibrating the $\hat{x}_0$ weighting throughout. The diversity guarantee is unaffected, since each candidate begins from an independent $x_T \sim \mathcal{N}(0, I)$, so the collapsed-latent failure mode present in ACT cannot arise. The fix supplies the explicit next-in-schedule index to `ddim_step`; all three evaluation splits are re-run under the corrected sampler and the corrected results supersede the initial runs. The correction did not change the qualitative conclusion that Diffusion Policy supplied a diverse candidate pool, but it improved the integrity of the candidate source used by all downstream world-model selectors. The before/after audit is retained in the implementation record rather than used as a separate thesis result.




Figure 5.6 DDIM subsampling correction. The ten-step inference schedule spans $t \in \{90, 80, \dots, 0\}$. At each reverse step the correct $\bar{\alpha}_{t_{\text{prev}}}$ is the *next timestep in the subsampled schedule*; the bug substituted the consecutive training-schedule timestep $\bar{\alpha}_{t-1}$, miscalibrating the $\hat{x}_0$ weighting at every step.

5.5 World model training

Table 5.5 summarises what each family reused from an external checkpoint versus trained for this study; Table 5.2 already covers the interface adapters this required. The remainder of this section focuses on the design decisions and defects specific to each family, rather than repeating the loading mechanics that are common across all of them.

5.5.1 V-JEPA2-AC: extraction and transition head

The V-JEPA2-AC architecture is reviewed in Section 3.2.3 and illustrated in Figure 3.8. This implementation keeps the same frozen ViT-Large backbone and scoring principle as the published system. It replaces the published system’s autoregressive, approximately 300-million-parameter predictor with a substantially smaller, single-step MLP transition head. This head is trained on cached embeddings from the EgoDex demonstrations used throughout this study, rather than on Droid robot interaction video, and without the published teacher-forcing/rollout training regime. It is therefore a V-JEPA2-AC-inspired scorer built on the same frozen backbone and

Table 5.5 Backbone sourcing and what was trained from scratch for this study, by world model family.

Family	Backbone / checkpoint source	Trained for this study
V-JEPA2-AC	Frozen ViT-Large, official checkpoint (PyTorch Hub)	Action-conditioned transition head
LeWM	None; no pretrained backbone released	Encoder, predictor, action embedder, and projector (fully, from scratch)
DexWM	Frozen DINOv2 ViT-Large, official checkpoint (HuggingFace)	CDiT predictor, hand-consistency head
DINO-WM / JEPA-WM	Official DROID/RoboCasa checkpoints (Meta)	None; frozen official weights
DiWA/CALVIN	Official world model and reward classifier (checkpoints obtained from the authors)	None; C3 reproduces the paper's own PPO fine-tuning
V-JEPA2 (no AC)	Reuses V-JEPA2-AC Phase 1 embeddings	None
PointWorld	Reuses V-JEPA2-AC backbone and transition head	None

action-conditioning principle, not a reproduction of the published predictor at its original scale or training regime. Training separates into two phases. Phase 1 is a one-time offline extraction. The frozen encoder processes all indexed windows in parallel across two L40S GPUs, writing one embedding vector per window with no backward pass or optimiser. A critical bug in an early run applied a 2,000-window default cap, producing only around 1,000 training pairs from a 140,000-window split (Section 5.7). Phase 2 trains a small MLP transition head on the stored embeddings to predict the next embedding from the current embedding and an action-conditioning vector derived from the bimanual wrist-state delta (Table 5.1).

5.5.2 LeWM: end-to-end training

LeWM is reviewed in Section 3.2.3 and illustrated in Figure 3.10. It is trained entirely from scratch, with no external weights. The encoder, autoregressive predictor, action embedder, and projector are jointly optimised together. The training objective combines a next-embedding prediction loss with the SIGReg regulariser, which enforces a Gaussian prior on the latent distribution [26] (Table 5.1).

Two export-pipeline defects affected early guard runs. A mismatched frame cap between the training and evaluation exports silently excluded 97.9 per cent of test windows. A later normalisation-statistics defect then introduced a $1.74\times$ scale mismatch in the LeWM cost function, caused by loading test-set action statistics instead of training-set ones. Both defects are resolved and detailed in Section 5.7.

LeWM exposes two distinct scoring routes, evaluated separately in this study and shown side by side in Figure 5.7. The rollout scorer predicts a future state through LeWM’s learned visual dynamics, in the same role as V-JEPA2-AC. The action-prior scorer instead scores a candidate directly against a training-only action distribution, without predicting anything.




Figure 5.7 The two LeWM scoring routes evaluated in this study, both starting from the same candidate action chunk: the rollout scorer (right), which predicts forward through LeWM’s learned dynamics, and the action-prior scorer (left), which does not.

The *rollout scorer* takes the candidate action chunk and the observation history into the LeWM predictor, one rollout step at a time. Each predicted latent is appended back to the observation history before the next step, so the predictor consumes its own prior output. After all steps, only the final predicted latent is looked up against a training-only latent support bank, via nearest-neighbour search, to produce a support distance. This is the mechanism LeWM was built to provide: a genuine forward prediction through its learned visual dynamics.

The *action-prior scorer* takes the same candidate action chunk, but scores it directly against a training-only action set’s per-dimension mean and standard deviation, producing a per-dimension z-score from which an action regularity score is derived. It does not consult the observation history, the latent space, or the LeWM predictor at all. It never rolls the candidate forward, and the predictor plays no role in the score it produces. It measures whether a candidate’s actions are typical of LeWM’s training distribution, not whether LeWM predicts that candidate will produce a better outcome.

These two routes are not interchangeable, and only one of them exercises LeWM’s learned dynamics at all. Section 6.1 reports both, and treats them accordingly: the rollout scorer as LeWM’s substantiated, reportable result, and the action-prior scorer as an illustrative reference only, not a measurement of LeWM’s visual prediction.

5.5.3 DexWM: predictor trained from scratch

DexWM is reviewed in Section 3.2.3. Every conclusion this thesis draws about DexWM applies only to the locally trained, fixed-budget version implemented here, not to the officially reported architecture at its published pretraining scale, since the official EgoDex+DROID-pretrained checkpoint was never released (Section 3.2.3) and is not reproduced by this implementation. Its pipeline pairs a frozen DINOv2 image encoder with a CDiT predictor and incorporates a hand-consistency loss that supervises fingertip and wrist heatmap predictions (Figure 5.8).




Figure 5.8 DexWM pipeline as implemented, comprising a frozen DINOv2 encoder, CDiT predictor, lightweight decoder, and HC loss supervision path. Adapted from [25].

Checkpoint sourcing and abstained full pretraining

The published DexWM architecture [25] exposes architecture code and dataset instructions, but no public pretrained checkpoint could be sourced. The associated HuggingFace repository contains trajectory data rather than weights, and the GitHub repository describes pretraining on EgoDex and DROID without linking a downloadable model checkpoint. This matters because the published result depends on a large upstream training recipe that is not reproduced merely by having the code.

This study’s implementation therefore uses a fixed-budget local route. The frozen DINOv2 encoder provides patch tokens, and the CDiT predictor and hand-consistency head are trained from random initialisation on the declared training splits at the universal budget ($E \approx 100,000$, Table 5.1). Training is a genuine

action-conditioned next-patch-embedding prediction task. The 18D SE(3) wrist-action delta conditions every predictor block via adaLN-Zero modulation [64], and the loss compares the resulting prediction against the true future embedding rather than a reconstruction of the input, so the route is not degraded to unconditional pretraining.

A later preflight tested whether the full official EgoDex+DROID pretraining suite should be reopened instead of this fixed-budget route, and passed the EgoDex-side checks, but the DROID raw data needed for the published recipe was not available locally. Reopening it would also have trained DexWM on a strictly larger corpus than every other family, breaking the fairness contract that keeps the main comparison attributable to model design rather than to training-data scale. Since the fixed-budget route was already valid, full upstream pretraining was abstained from rather than launched as a new open-ended experiment; its performance is expected to fall short of the published recipe as a result. This is a reproducibility limitation of the available release as well as a project-scope decision. Code without a pretrained checkpoint leaves independent researchers to repeat a large-scale pretraining programme before they can evaluate the published model under a new benchmark.

Single-seed evaluation. Unlike every other trainable family in this study, DexWM is not retrained at the additional seeds declared in Section 4.5.3. Two compounding costs drove this decision. First, DexWM is already the most expensive family to train at a fixed exposure budget, ≈ 28 seconds per step against LeWM's ≈ 0.047 (Table 5.1), so two further seeds alone were estimated at roughly 112 GPU-hours. Second, a smoke-test ahead of committing to that cost found DexWM's own training-data export (the frozen-DINOv2 patch-token extraction its predictor trains on) was also missing from disk, adding a further ≈ 9.3 hours of largely GPU-bound prerequisite regeneration before retraining could even begin, so roughly 121 extra GPU-hours in total. Weighed against a family whose result is already negative under this fixed-budget substitute and was never treated as load-bearing, that cost was judged disproportionate to what additional seeds could plausibly change. DexWM therefore remains a single-seed result, and its result carries a higher seed-specific uncertainty than the other trainable families for that reason; every table reporting it states the single-seed status explicitly.

5.5.4 DINO-WM and JEPA-WM exploratory selectors

DINO-WM and JEPA-WM were added after the initial closeout as bounded official checkpoint comparators, and both are kept exploratory throughout this thesis rather than promoted to a held-out claim, precisely because the action-interface reduction described below makes their comparison to the other families less direct than any of the natively 18-dimensional routes. The implementation uses the official JEPA-WMS code path and public Meta checkpoints, downloaded from HuggingFace. DINO-WM

DROID loaded directly. JEPA-WM DROID also required the official gated DINOv3 ViT-L/16 backbone; access to the gated repository was obtained under a university research-program licence, and the backbone was then downloaded and converted into the native format expected by the upstream loader.

The central implementation issue is action compatibility. This study’s candidates are 16-step chunks in an 18-dimensional bimanual wrist action space. The official DROID/RoboCasa JEPA-WMS checkpoints expect a 7-dimensional single-arm robot action. A direct truncation would invent semantics and discard one hand, so the implementation uses a train-only linear adapter from EgoDex action chunks into the expected 7D interface. This keeps the route auditable, but it also weakens the claim. The scorer is an adapted transfer experiment, not a native EgoDex world model.

Both models passed load, extraction, and scorer smoke tests, and were carried through a bounded validation pilot rather than the held-out evaluation used for the main comparison. No held-out test claim is reported from this route; the pilot outcome and its interpretation are reported in Appendix A.4 rather than here, since neither model was run to a comparable, fully-evaluated state.

5.5.5 DiWA/CALVIN secondary-study implementation

DiWA/CALVIN is a closely related study, planned in Section 4.4 of the methodology, and reproduced here for integrity. It lets this study’s findings be checked alongside an independent, world-model-guided policy improvement method under shared components rather than in isolation.

A dedicated DiWA/CALVIN checkout was prepared alongside the official CALVIN benchmark. Public artefacts, including the frozen world model, reward classifier, and environment assets, were downloaded, hashed, and checked for archive safety before use. The world model and reward classifier were then verified loading on CPU and in a bounded GPU timing check.

The official base-policy and fine-tuned policy checkpoints, together with the statistics and sampler metadata needed to load them, were not published with the initial release. Access was requested directly from the DiWA authors, and a corresponding issue was also filed on the official repository. The needed files were obtained through this direct correspondence, covering `close_drawer`, `turn_on_led`, `move_slider_left`, and `turn_on_lightbulb`.

A six-condition experiment contract (`contract.py`) holds the base policy, world model, reward classifier, environment, and initial states fixed across all six conditions, so the improvement operator is the only variable that changes between them: **C1**, DiWA’s own base policy acting alone; **C2**, the same policy with this study’s test-time best-of-five selector applied, using the frozen DiWA world model and success classifier

without retraining the policy; **C3a**, the identical checkpoint to C1 at the moment PPO fine-tuning begins; **C3b**, the same policy after DiWA’s own imagined-rollout PPO fine-tuning, which updates the policy itself inside the frozen world model; **C4a**, this study’s selector applied to C3b’s fine-tuned checkpoint; and **C4b**, C4a with the reward classifier recalibrated on episodes rolled out from C3b’s own policy. Full definitions and results for all six conditions are reported in Chapter 6.

This secondary study departs from the primary EgoDex protocol by design. C2, C4a, and C4b execute the selected policy in the real CALVIN simulator, and C3b updates the policy itself through reinforcement learning. The primary study, by contrast, scores fixed candidate chunks offline against held-out demonstrations without simulation or policy updates. This departure is required to reproduce DiWA under conditions faithful to its own design; the two studies remain separate and are not merged into a single evidence hierarchy. Despite that separation, this secondary study is not a minor add-on. It is what turns this thesis’s own three-condition account of when world-model guidance helps (Section 3.4) from an untested framework into a finding with a second, independent data point behind it, since DiWA’s own design satisfies all three conditions by construction and its result is read against this thesis’s primary EgoDex findings throughout Chapter 6.

5.5.6 V-JEPA2 (no action conditioning): ablation control

The methodological role of this ablation is described in Section 4.4. In implementation, the variant reuses the Phase 1 embeddings extracted for V-JEPA2-AC, requiring no additional training (Table 5.1). At guard time the scorer passes a zero action vector to the transition model for all candidates, producing action-blind scores that reflect visual plausibility alone. A diagnostic persistence loss (Mean Squared Error (MSE) plus $0.1 \times$ cosine distance between consecutive stored frame embeddings, the same weighted objective as the AC head training loss) is computed directly from `embeddings.npy` without invoking the transition model, confirming the scorer makes no action-differentiated prediction under a zero action vector.

5.5.7 PointWorld: action-degradation control

The methodological role of the PointWorld control is described in Section 4.4. The implementation uses the same frozen V-JEPA2-AC backbone and transition head as the V-JEPA2-AC family, requiring no additional model training (Table 5.1). The sole difference is that candidate action chunks are degraded before being passed to the transition head. This is a deliberate experimental degradation, built to create a controlled ablation, distinct from the incidental information loss in the

DINO-WM/JEPA-WM action adapter (Section 5.5.4), where one hand’s pose is discarded only because it is required to match an external checkpoint’s native interface, not to construct a control.

The adapter keeps only the left-wrist rotation and collapses the remaining twelve dimensions (left translation, right rotation, right translation) into a single scalar summary before zero-padding back to 18 dimensions. The right-hand pose is not recoverable from this output; it does not survive as a distinguishable signal in the degraded representation. The transition head therefore receives the correct visual features but systematically corrupted action conditioning.

The convergence study diagnostic additionally computes a proxy loss:

$$\mathcal{L}_{\text{PointWorld-proxy}} = \mathcal{L}_{\text{persist}} + 0.5 \cdot \text{mismatch} \quad (5.2)$$

where $\mathcal{L}_{\text{persist}}$ is the frame-to-frame embedding distance and mismatch is the fraction of action-vector norm discarded by the projection, clipped to $[0, 1]$:

$$\text{mismatch} = 1 - \frac{\|\mathbf{arm}\|}{\|\mathbf{hand}\|} \quad (5.3)$$

This diagnostic confirms that the adapter introduces genuine information loss. The bimanual hand action space does not map faithfully to the 7-DOF arm representation, as expected for a degradation control.

5.6 B2 margin threshold calibration

Branch 2 selects the world model’s top-scoring candidate whenever it beats the policy’s own candidate 0 by at least a margin δ , measured in the guard’s own score units; below that margin, candidate 0 is kept instead. The canonical configuration used throughout this thesis is unconditional selection, equivalent to $\delta = 0$, so the top-scoring candidate is always taken, however small its lead over candidate 0. Raising δ makes the guard progressively more conservative, since a larger lead is required before it will override the policy, so it keeps candidate 0 more often and moves the branch closer to the Branch 1 policy-only baseline. The final threshold choice reported in this thesis was made on a 42,774-window internal validation split drawn from `train_part1_2_3` itself (`split_windows_by_episode`, 10% held-out episodes, seed 42), disjoint from both the training windows and the frozen 7,607-window test set. A calibration run compared seven settings of δ on that split, on the canonical $K = 5$ candidate pool and the V-JEPA2-AC guard (Table 5.6). An earlier, superseded version of the same calibration had been run directly on the test set, before this validation split was introduced; that earlier test-set inspection is a limitation of this thesis’s

development process and is recorded as such in Section 7.7. What the final choice rests on is the validation-split calibration below.

Table 5.6 B2 margin threshold calibration on the internal validation split ($K = 5$, V-JEPA2-AC guard, 42,774 windows, disjoint from the test set). Higher δ keeps candidate 0 more often and yields a smaller improvement. Validation-split baseline translation L_2 (candidate 0) is 0.11499, distinct from the pooled test-set baseline of 0.117 reported in Table 6.1, since the two are different windows and this validation-split figure is a single-seed value never re-pooled across seeds.

Margin δ	Candidate 0 kept	Translation L_2	Translation change
n/a*	22.54%	0.10778	-6.26%
0.05	48.09%	0.10873	-5.44%
0.10	70.55%	0.11069	-3.73%
0.15	84.26%	0.11239	-2.26%
0.25	95.60%	0.11415	-0.72%
0.50	99.78%	0.11493	-0.05%
1.00	100.00%	0.11499	0.00%

*Canonical configuration used throughout this thesis: the guard takes the world model’s top-scoring candidate unconditionally, with no margin gate.

Read down the table, and the pattern is monotonic in one direction only. As δ rises, candidate 0 is kept more often and the improvement shrinks toward zero. The 100% row is not the best-performing setting; it is the opposite. At $\delta = 1.00$ the margin is large enough that the gate never clears on this score scale, so Branch 2 keeps candidate 0 on every window and exactly reproduces the Branch 1 policy-only baseline, zero change by construction, not by merit. This is the same monotonic shape and the same unconditional setting winning that the earlier, test-set version of this calibration found; that the validation-split calibration, on windows no tuning decision has ever seen, reaches the same conclusion is reassuring rather than a full defence, since the earlier test-set inspection cannot be undone. DexWM was excluded from this calibration because its initial guard runs were vacuous due to a candidate-loading defect (Section 5.7); its canonical evaluation instead uses the dedicated `guards_native_dexwm/` pipeline.

Two considerations motivate choosing the unconditional setting over a margin gate:

1. **CEMP already pre-filters toward world-model preferences.** As the circularity caveat of Section 5.4.1 describes, the internal CEMP scoring loop pre-filters the candidate pool toward the world model’s own preferences before the guard’s final selection ever runs. Adding a margin gate on top would reintroduce a hyperparameter that varies across model families and pool types, without addressing that deeper redundancy.

2. **Unconditional argmax isolates discriminative quality directly.** Measuring the world model’s preference without a threshold suppression layer produces a cleaner experimental signal and avoids per-family calibration.

The improvement lives entirely in the windows where the guard overrides the policy, and the unconditional setting overrides most often, which is why it produces the largest reduction in Table 5.6 and is the same configuration reported as the headline result in Table 6.4. This is the chosen configuration for every Branch 2 result in this thesis, since it dominates every tested margin gate on translation error while still keeping candidate 0 a meaningful fraction of the time rather than never or always.

5.7 Implementation challenges and resolutions

The implementation challenges are part of the contribution because several early runs looked plausible but were not valid evidence. Table 5.7 records the material issues, their impact, the corrective action, rerun status, and remaining risk. Longer forensic notes are retained in the project archive and appendices.

The types distinguish four different kinds of problem, which do not all deserve the same interpretation. *Implementation* issues are code-level defects, fixed by correcting the code, and generally the lowest-risk once resolved because the fix is independently verifiable against the source. *Experimental-design* issues are protocol-level choices that needed rethinking, not code defects, so their resolution is a design decision rather than a bug fix. *Resource* issues are external unavailability, a missing checkpoint, that no local fix can fully close, only work around. *Model* issues are architectural incompatibilities inherent to comparing an externally trained checkpoint against this study’s own action space, which an adapter can bridge but not remove.

The remaining risk is concentrated in routes that were not promoted to main claims. The headline V-JEPA2-AC result, LeWM rollout result, DexWM fixed-budget negative result, and ACT collapse diagnostic all use frozen candidate artefacts, held-out windows, and manifest-level validation.

Table 5.7 Material implementation challenges and resolution status.

Issue	Impact	Fix	Rerun status	Type	Risk
DDIM subsampling error	Miscalibrated reverse-diffusion updates for DP candidates.	Passed the explicit previous timestep t_{prev} into <code>ddim_step</code> .	DP evaluations rerun; corrected candidates used downstream.	Impl.	Low.
Oracle goal leakage	Early B2/B3 scoring runs derived the latent goal from each episode's ground-truth future frame, information unavailable at deployment time.	Set <code>goal_source=none</code> in all canonical benchmark configurations.	Rerun with the ground-truth goal removed; every guard result reported in this thesis reflects this leakage-free configuration, not the earlier cheating one.	Exp. design	Low.
Branch 3 selection-mode error	B3 initially duplicated B2 because candidate 0 remained eligible.	Relaunched with <code>world_model_only</code> ; manifest checks require zero policy fallbacks.	Corrected B3 runs replace earlier outputs.	Impl.	Low.
Score-cache memory collision	Tensor memory-address reuse could silently corrupt scale factors after garbage collection.	Re-keyed cache on explicit run-scoped identifiers.	Affected runs excluded or rerun.	Impl.	Low.
Evaluation window shortfall	Default guard limits evaluated 787 rather than 7,607 windows.	Added full-evaluation flags and manifest count checks.	Full 7,607-window runs used.	Impl.	Low.
LeWM wrong normalisation source	Test-set action statistics distorted candidate costs.	Loaded training-only statistics and reran diagnostics.	Rollout scorer still failed; action-prior route separated.	Impl.	Medium for rollout only.
Candidate-source asymmetry	Earlier families used different candidate sources, confounding comparison.	Final causal matrix uses shared frozen candidates and explicit candidate hashes.	Superseded runs archived.	Exp. design	Low.
DexWM missing checkpoint	Published pretrained model unavailable; zero-score stub risk found.	Stub replaced by hard failure; fixed-budget local route evaluated.	Valid negative/null for local route only.	Resource	Medium for general DexWM claims.
DINO-WM/JEPA-WM action mismatch	Official 7D DROID action interface does not match this study's 18D bimanual chunks.	Train-only adapter and validation-only pilots.	Kept exploratory; no held-out claim.	Model	High for headline use.
DiWA missing checkpoints	Reproduction could not proceed without the official base/fine-tuned policy checkpoints and sampler statistics, unpublished with the initial release.	Access requested directly from the authors and via a repository issue; needed files obtained.	C1–C4b secondary study proceeded; results reported separately (Chapter 6).	Resource	Low, resolved.
V-JEPA2-AC Phase-2 exposure shortfall	AC transition head trained on 85–90k items instead of the declared $E = 100,000$ budget (10–15% short) on all three splits, an undisclosed-until-caught deviation from Table 5.1's uniform-budget claim.	Wired the trainer's existing deterministic budgeted-exposure sampler into the launch script (was implemented but never invoked).	All three splits rerun at the corrected budget; old checkpoints archived, val loss improved slightly on all three.	Impl.	Low.

6 Results

This chapter works through the results in the order set by the experimental design in Chapter 4. It starts with the world-model scorers on the DP row, from weakest to strongest, then reports ACT’s CEMP-augmented guard results as secondary corroboration of that ranking, examines how selection quality varies with task difficulty, and closes with a secondary reproduction study on DiWA’s own CALVIN benchmark. The candidate pool size ablation that bounds how these margins scale with K is reported separately in Appendix B.3, since it is a sensitivity study on the chosen operating point rather than a headline result in its own right. Unless stated otherwise, every number below is a single mean pooled across three training seeds, 42, 43, and 44 (Section 4.5.3), not a report of any one seed alone; the full per-seed breakdown behind these pooled means is given separately in Appendix B.6, and this chapter discusses seed-to-seed mechanics only where a result specifically calls for it. All reported numbers come from the same frozen EgoDex `train_part1_2_3` test split of 7,607 windows, 3,229 episodes, and 111 task categories, evaluated once under branch and support rules fixed before interpretation.

The chapter does not measure robot task success, human preference, or closed-loop deployment. A negative result means a model did not improve this offline metric under this frozen protocol. It does not mean the model family is generally unusable.

6.1 World-model selection results

Branch 1 is the policy-only baseline every later branch in this chapter is measured against. The policy generates its candidate pool and selects by its own internal preference, with no world model involved (Section 4.2). Table 6.1 reports this baseline for both policies studied in this thesis. Translation follows the RMS-per-axis convention used throughout this chapter, not a true Euclidean distance; rotation is the mean geodesic angle in degrees. The bracketed figure beside each mean is a 95% confidence interval, not a range of per-seed values, defined in Section 4.6.3.

Table 6.1 Branch 1 policy-only baseline results, no world model involved, pooled mean across seeds 42, 43, and 44 [95% CI].

Policy	Translation error	Rotation error
Diffusion Policy (DP)	0.117 [0.114,0.120]	40.01° [39.61,40.42]
Action Chunking Transformer (ACT)	0.065 [0.024,0.143]	22.35° [12.66,41.35]

DP’s baseline is stable and consistent across the three tested seeds and is the

stronger of the two policies on both axes. ACT’s pooled figure above is not a single representative number, for the same, already-established reason its guard results need the adaptive CEMP treatment described below, seed-dependent CVAE posterior collapse (Section 5.4.1). Appendix B.6.1 gives the full per-seed breakdown behind both baselines, including how far seed 42’s collapsed checkpoint pulls ACT’s pooled figure from either seed’s own regime; this is why ACT’s guard results are kept in a separate section (Section 6.2) rather than pooled alongside DP’s in this one. Every delta reported below in this section is measured against DP’s Branch 1 baseline above.

Figure 6.1 summarises the headline comparison across four of the five scorers on the DP row, as the reduction in translation error under Branch 2 and Branch 3. LeWM is held out of the figure and reported in Table 6.2 instead, for reasons the discussion below explains. The two branches share the same candidate pool and world-model scorer, and differ in one respect (Section 4.2). Branch 2’s pool includes the policy’s own preferred candidate, and the scorer picks freely among all five. Branch 3 removes that preferred candidate first, so the scorer must choose among the four remaining candidates. Reading the two branches together says more than either alone. If Branch 2 improves on the baseline but Branch 3 does not, the policy’s own candidate was already close to the best available choice, and the scorer’s main contribution is avoiding the cases where that candidate was wrong. If Branch 3 improves more than Branch 2, the policy’s preferred candidate was actively misleading, and better options existed elsewhere in the pool. If neither branch improves on the baseline, the scorer is not distinguishing candidate quality at all. The following paragraphs work through all five scorers on this row, from the weakest result to the strongest, reading both branches together throughout.

The zero-action ablation (V-JEPA2 no-AC) returns the baseline unchanged in Branch 2. It assigns tied scores and never overrides candidate 0, exactly by construction, the expected outcome for a scorer with no action information to rank on. Our other proxy control, PointWorld, worsens translation while improving rotation. This is consistent with the degraded-conditioning design described in Section 5.5.7. It was built to reuse V-JEPA2-AC’s visual pathway with weakened action conditioning, specifically so that a mixed, non-improving result would confirm that action information, not visual quality alone, drives selection.

DexWM worsens translation error in both branches, and by a larger margin than the PointWorld negative control. This is a valid negative result for the implementation actually evaluated here, from a single training seed; unlike ACT, DP, and V-JEPA2-AC, DexWM is not retrained at additional seeds in this thesis, for the cost-benefit reasons stated in Section 4.7, so this result carries a seed-specific-bias risk the other rows in this table do not. Section 5.5.3 describes the reproducible fixed-budget training route substituted for DexWM’s original recipe, since public pretrained weights for its




Figure 6.1 Reduction in translation error against the frozen DP candidate-0 baseline, by world model and branch. Positive values are an improvement over the policy-only candidate. Values are pooled across three training seeds except DexWM, which remains single-seed (Section 4.7).

published EgoDex+DROID recipe were not available. This is a good place to state a caveat that applies to every scorer in this chapter, not only DexWM. This study is not designed to find the best possible result each architecture can reach. Every model family here is trained under the same universal exposure budget (Section 4.5.1), deliberately chosen so a comparison reflects architecture and candidate-selection ability, not who happened to receive more training. DexWM’s original recipe trains at a substantially larger scale than this budget allows, so its result here is evidence about DexWM under this study’s fair-budget implementation, not a ceiling on what a fully-scaled DexWM could achieve. The same caveat holds for every other result in this section, and is revisited in Chapter 7.

LeWM is held out of Figure 6.1 because its margin comes with a caveat none of the earlier scorers in this section needed, explained below and reported in its own table rather than plotted alongside a cleanly comparable result. LeWM predicts a future state through its own learned visual dynamics, the same role V-JEPA2-AC plays. LeWM’s diagnostic scorer weighs several terms together, and one of those weights acts like a knob for how much its own predicted future state counts toward the final score. Turned up, the knob is the intended design. It calls LeWM’s image encoder, predicts a future state, and scores each candidate by how well that prediction matches it, described mechanically in Section 5.5.2. Before this is trusted, it has to pass a ranking-validity gate, a check for whether a candidate’s score actually predicts whether

that candidate turns out better.

At the full scale this chapter reports on, 7,607 held-out test windows, LeWM fails that gate. Its Spearman rank correlation (Section 4.6.3) between score margin and actual translation improvement is -0.507 (B2) and -0.518 (B3), essentially no relationship between what the scorer prefers and what actually turns out better. This is not a small-sample artefact. It gets worse, not better, at this larger scale than it did on a smaller diagnostic set used earlier to tune the gate. As a control, turning that knob all the way down disables LeWM’s model entirely. The result is almost identical, and still fails the same check.

The honest summary of LeWM’s result is this. Its scorer does pick candidates with lower translation and rotation error on average, by the largest margin in this chapter. But its own score does not reliably predict which candidate will turn out better, at the scale this chapter reports on. Something other than the scorer’s stated logic is driving the improvement, and disabling the model entirely does not change that outcome. LeWM is not validated as a causal mechanism here.

Table 6.2 reports both configurations’ full-scale numbers together. One is the version that calls LeWM’s model, and the other is the control that disables it.

Table 6.2 LeWM’s two diagnostic configurations, full 7,607-window test set, neither shown to reliably track which candidate is actually better (Section 5.5.2). Rollout row: pooled mean across 3 seeds [95% CI]. Action-prior row: canonical seed only, a fixed statistical scorer never retrained or retested at the other two seeds (Section 5.5.2); its bracketed figure is a single-seed episode-level CI, not the two-level clustered bootstrap used elsewhere.

Configuration	Branch	Translation change [95% CI]	Rotation change [95% CI]
LeWM rollout	B2	-10.094% [-10.338, -9.839]	-0.827° [-0.902, -0.755]
LeWM rollout	B3	-8.873% [-9.152, -8.601]	-0.725° [-0.804, -0.640]
LeWM action-prior	B2	-10.236% [-10.671, -9.816]	-0.849° [-0.972, -0.726]
LeWM action-prior	B3	-9.009% [-9.468, -8.549]	-0.741° [-0.885, -0.603]

Both configurations post the largest percentage improvements in this chapter, larger than V-JEPA2-AC’s. That size is exactly why they cannot be trusted at face value. A more negative number here is a bigger improvement, but neither configuration’s score is shown to predictively track which candidate produces it, so this is the largest result in the chapter, not the best one. That is why LeWM is held out of the headline comparison rather than reported alongside the other four scorers.

The rollout scorer’s own margin also reproduces reliably (Appendix B.6.1), which does not rescue the ranking-validity problem above, since a scorer can reproduce the same large, untracking margin every run and still not be tracking which candidate is actually better, the clearest instance in this thesis of a result reproducing cleanly while

failing to be validated as a predictive ranking mechanism.

V-JEPA2-AC is the strongest validated instance of the mechanism this thesis studies. It is an inference-time scorer that predicts a future outcome for each candidate and selects on that prediction. Its validity rests on the same test that just ruled LeWM out. Its Spearman rank correlation (Section 4.6.3) between score and each candidate’s true translation error, measured directly per window rather than assumed, is real and well above zero at every pool size tested. Table 6.3 reports this correlation at four pool sizes, pooled across all three training seeds, using the same nested-pool sweep described in Appendix B.3. The pooled correlation is small, modest rather than strong, and essentially flat across pool size, 0.162 to 0.164 from $K = 2$ to $K = 20$, but consistently and clearly above zero at every size tested, in sharp contrast to LeWM’s negative correlation. At $K = 5$, the pool size used throughout this study, the pooled correlation is 0.164. This is why V-JEPA2-AC’s percentage gains, unlike LeWM’s, come from a score shown to track genuine forward prediction rather than an unvalidated causal claim.

Table 6.3 V-JEPA2-AC’s per-window Spearman rank correlation between score and true candidate translation error, by candidate pool size K , Branch 2, pooled across 3 seeds [95% CI], on the nested pools behind Appendix B.3.

$K = 2$	$K = 5$	$K = 10$	$K = 20$
0.162 [0.155,0.167]	0.164 [0.155,0.170]	0.163 [0.155,0.169]	0.163 [0.155,0.168]

The translation-change figures in this chapter show the reduction in error the world model’s selection adds on top of what the policy alone would have produced, not the policy’s absolute error level. Both Branch 2’s -5.980% and Branch 3’s -5.368% (also presented in the full DP table of results, Table 6.4) clear that bar. More precisely, Branch 2 wins on translation in 49.88% [49.00, 50.71] of windows, ties in 23.17% [22.48, 23.84], and loses in 26.95% [26.26, 27.62]. The tie rate equals its candidate-0-keep rate by construction, a window can only tie when the guard keeps candidate 0, so the two statistics are the same count under any resampling and this identity holds exactly at every one of the pooled bootstrap’s 2,000 resamples. This is why Branch 2 beats Branch 3 overall (Appendix B.6.1 confirms this holds seed by seed too). It falls back to the policy’s own best sample when the world model is not helpful.

Set against what is achievable, this gain is precisely estimated, its 95% confidence interval does not cross zero, but modest. DP’s Branch 1 baseline translation error is 0.117 (Table 6.1). At the fixed five-candidate pool used throughout this study, the best selection possible from that pool reduces this by 16.508% in Branch 2 and 14.683% in Branch 3. The oracle ceiling is a definition of how much error is fixable by choosing the best candidate among the ones DP already generated, something no real scorer can do without access to ground truth. V-JEPA2-AC reaches 36.22% of that

ceiling in Branch 2 and 36.56% in Branch 3, meaning its selection recovers a little over a third of what a perfect, ground-truth oracle could have recovered from the same five candidates, at both branches essentially equally. The rest of the baseline error, everything past the oracle ceiling, is not a selection problem at all. It is a property of the candidates DP generates, and no amount of world-model selection can reach below it without a better candidate pool to select from. This draws a clear line around what this thesis’s mechanism can and cannot fix. That is, it can only re-rank what the policy already proposes, not raise the ceiling of what the policy is capable of proposing. Five is a compute-driven choice, not a ceiling. Appendix B.3’s full ablation across pool sizes two to twenty shows scorer ranking is stable at every size tested, and the achievable gain grows steadily with pool size rather than plateauing, so the fixed pool of five understates the ceiling reported above rather than overstating it.

Table 6.4 gives the point estimates and episode-clustered 95% confidence intervals for every scorer on this row, including LeWM, which the figure above omits. Each delta is the selected candidate minus candidate 0, so negative values mean the guard selected a better trajectory than the policy-only baseline; against the 0.117 m Branch 1 baseline (Table 6.1), V-JEPA2-AC’s -5.980% in Branch 2 corresponds to roughly $0.117 \rightarrow 0.110$ m.

Table 6.4 Primary DP guard results, pooled mean across 3 seeds [95% CI] (bootstrap method: Section 4.6.3; per-seed detail: Appendix B.6.1). LeWM and DexWM caveats discussed below.

Scorer	Branch	Translation change [95% CI]	Rotation change [95% CI]
V-JEPA2-AC	B2	-5.980% [-6.465, -5.432]	-0.960° [-1.071, -0.830]
V-JEPA2-AC	B3	-5.368% [-5.773, -4.950]	-0.854° [-0.946, -0.758]
V-JEPA2 no-AC	B2	0.000% [0.000, 0.000]	0.000° [0.000, 0.000]
V-JEPA2 no-AC	B3	-0.009% [-0.297, +0.289]	-0.004° [-0.150, +0.125]
PointWorld proxy	B2	+0.265% [-0.039, +0.579]	-0.390° [-0.490, -0.282]
PointWorld proxy	B3	+0.249% [-0.115, +0.629]	-0.354° [-0.468, -0.237]
DexWM	B2	+0.570% [+0.118, +1.020]	-0.051° [-0.187, +0.077]
DexWM	B3	+0.563% [+0.086, +1.037]	-0.084° [-0.230, +0.060]
LeWM	B2	-10.094% [-10.338, -9.839]	-0.827° [-0.902, -0.755]
LeWM	B3	-8.873% [-9.152, -8.601]	-0.725° [-0.804, -0.640]

LeWM’s row shows the largest numerical translation change in this table, but, as established above, is not validated as a predictive ranking mechanism; only V-JEPA2-AC’s margin is.

The three semantic world models, V-JEPA2-AC, DexWM, and LeWM, also contrast fundamentally different representational strategies and backbone-pretraining provenance, so the selection deltas above should be read model by model, not as a

head-to-head architecture ranking. Section 7.3 interprets this difference, and the relative inference cost it implies, in full.

6.2 ACT guard results under CEMP

The two policies are not evaluated on equal footing, and this shapes how every later section should be read. DP produces five truly distinct candidates per window, so a world-model scorer has something real to choose between. ACT’s candidate pool does not always: at the canonical seed its CVAE posterior collapses at inference time, a property of that checkpoint, not of the guard evaluator (Section 5.4.1). This collapse is itself seed-dependent, so this thesis’s evaluation pipeline applies CEMP, a cross-entropy candidate search, only to the checkpoint that needs it, and scores the other two seeds’ own natively sampled candidates unmodified (Appendix B.6.1 gives the full account). The table below reports the single pooled mean this adaptive rule produces, read below as one result, not three separate seed results.

Where CEMP applies, every candidate in that pool other than candidate 0 is produced by CEMP’s own search process, not sampled directly from ACT, so that portion of the pooled result is evidence for the world-model-plus-CEMP system as a unit, not an isolated measurement of scorer quality comparable to the DP row, exactly as Section 5.4.1’s circularity caveat states. Branch 2 keeps the real ACT output as a fifth eligible candidate throughout; Branch 3 excludes it, the same exclude-the-baseline rule the DP row uses for Branch 3 (Section 4.2).

Table 6.5 ACT guard results, pooled mean across 3 seeds [95% CI] (selection rule and per-seed detail: Appendix B.6.1), same bootstrap as Table 6.4. Four of five scorers; LeWM held out, reported separately below (same reason as Section 6.1). Negative values are an improvement over the frozen ACT baseline, positive a regression. The DexWM row is single-seed.

Scorer	Branch	Translation change [95% CI]	Rotation change [95% CI]
V-JEPA2-AC	B2	-0.577% [-0.856, -0.139]	-0.021° [-0.049, -0.003]
V-JEPA2-AC	B3	-0.585% [-0.880, -0.167]	-0.021° [-0.046, -0.004]
V-JEPA2 no-AC	B2	0.000% [0.000, 0.000]	0.000° [0.000, 0.000]
V-JEPA2 no-AC	B3	+0.248% [-0.397, +0.403]	+0.010° [-0.003, +0.032]
PointWorld proxy	B2	+0.342% [-0.226, +0.501]	-0.012° [-0.032, -0.001]
PointWorld proxy	B3	+0.337% [-0.277, +0.504]	-0.013° [-0.033, -0.001]
DexWM	B2	+0.263% [+0.186, +0.343]	+0.007° [-0.001, +0.015]
DexWM	B3	+0.263% [+0.184, +0.348]	+0.007° [-0.001, +0.015]

V-JEPA2 no-AC’s Branch 2 result is exactly 0.000%, with a zero-width interval, the same structural tie already described on the DP row (Section 6.1): with no action

information to rank candidates on, it assigns every candidate an identical score and never overrides candidate 0. This is exact by construction at every seed individually, not an artefact of averaging across seeds.

PointWorld proxy and DexWM both post a positive translation change in this table, meaning their selections are worse than the ACT baseline on average, not better, though neither is distinguishable from zero. PointWorld proxy’s Branch 2 and Branch 3 point estimates are close, +0.342% and +0.337%, because its scorer rarely selects the ACT candidate in Branch 2 either, so excluding that candidate in Branch 3 changes little. As a negative control, its role is to show what a scorer with no real predictive signal produces here. The result is a small, statistically inconclusive regression rather than a confirmed improvement or worsening. Its translation change here is broadly comparable to its DP-row counterpart, +0.265%, both with wide, zero-crossing intervals. Its rotation change remains a small, non-zero improvement on both rows, smaller here, -0.012° to -0.013° , than on the DP row, -0.390° to -0.354° ; CEMP, engaged only where a checkpoint needs it, narrows the candidate pool’s diversity compared to DP’s five independently sampled candidates, which plausibly explains why this row leaves less rotation error on the table for even a null scorer to stumble into. DexWM’s translation change is a similarly small regression, +0.263% in both branches, and its rotation change does not clear the interval away from zero, consistent with the negative result already reported for DexWM on the DP row (Section 6.1), where its rotation change also did not clear the interval away from zero. Both rows agree DexWM moves rotation error in neither a reliably better nor worse direction, even though the point estimates happen to have opposite signs. Unfortunately, the fixed-budget training route substituted for DexWM’s original recipe (Section 5.5.3) does not translate into a working scorer here either.

LeWM’s effect on this row is the largest in the chapter, translation improved by roughly 1.8% to 2.2% depending on configuration and branch, but it is not the strongest evidence of world-model prediction, which is why it is held out of Table 6.5 rather than listed alongside the other four scorers. As Section 6.1 explains in full, LeWM’s diagnostic scorer fails its own ranking-validity gate at the scale this chapter reports results on, even with the weight that lets it invoke its own learned dynamics turned all the way up. The margin is real, but the score behind it is not shown to predictively track which candidate is actually better (Section 5.5.2).

These two configurations suffer the same caveats already described for the DP row’s own held-out table (Table 6.2), and the same trend appears here too. Their translation and rotation changes are broadly comparable and their confidence intervals overlap, and neither configuration’s score is validated as a predictive ranking mechanism. Directly checking the Spearman rank correlation between score margin and translation improvement, the same test applied to the DP row, gives $\rho = -0.455$

Table 6.6 LeWM's two diagnostic configurations on the ACT row, full 7,607-window test set, neither shown to reliably track which candidate is actually better (Section 5.5.2). Rollout row: pooled mean across 3 seeds [95% CI]. Action-prior row: canonical seed only, a fixed statistical scorer never retrained or retested at the other two seeds; its bracketed figure is a single-seed episode-level CI, not the two-level clustered bootstrap used elsewhere.

Configuration	Branch	Translation change [95% CI]	Rotation change [95% CI]
LeWM rollout	B2	-1.844% [-2.151, -0.479]	-0.016° [-0.030, -0.007]
LeWM rollout	B3	-1.841% [-2.154, -0.471]	-0.016° [-0.030, -0.007]
LeWM action-prior	B2	-2.150% [-2.235, -2.068]	-0.031° [-0.044, -0.016]
LeWM action-prior	B3	-2.150% [-2.235, -2.068]	-0.031° [-0.044, -0.016]

for the action-prior configuration and $\rho = -0.466$ for the rollout configuration. Both are strongly negative and close to each other. This is consistent with the DP-row finding that turning LeWM's own model on or off barely changes the outcome. The score never tracked outcome quality in the first place.

Unlike the DP row, this margin is not tight across seeds, following the same seed-42-collapses split established for this row's other scorers rather than a new instability in LeWM itself (Appendix B.6.1). None of these margins is shown to predictively track which candidate is better regardless, so the spread is orthogonal to the causal-validity problem above, not a resolution of it.

Given that LeWM's margin cannot be read as a validated guard pass, V-JEPA2-AC is this row's actual headline result. Its translation change is smaller, **-0.577%** in Branch 2 and **-0.585%** in Branch 3, and both intervals distinguish from zero, and it is the only scorer on this row whose score is validated to track real forward prediction rather than an action-distribution regularity. Rotation improves too, by a small but consistent margin in both branches. This is the one result on this row that can be stood behind as a working world-model guard. It follows the same analysis given for the DP row above (Section 6.1). The direction of the effect matches exactly, an improvement in both translation and rotation, in both branches. The magnitude is far smaller here, -0.58% against -5.98% in Branch 2, consistent with this row's secondary, corroborative status rather than a differing finding.

V-JEPA2 no-AC serves as a second control. Its scorer cannot distinguish candidates by action content, so every candidate in Branch 2's pool scores identically, and ties resolve toward the ACT baseline, kept 100% of the time. Branch 3 forces a choice among scores that are still effectively tied, and the result comes out marginally worse than the baseline on both translation and rotation, though this small regression does not clear the interval away from zero. This is the expected null result for a scorer with no action information to rank on. It confirms that action conditioning, not CEMP

search alone, is doing the work in the other rows, and it makes V-JEPA2 no-AC the only scorer on this row that keeps the ACT candidate anywhere close to 100% of the time rather than close to never. Both this row’s Branch 3 translation change for this scorer, +0.248% [-0.397, +0.403], and the DP row’s, -0.009% [-0.297, +0.289], now have intervals that cross zero, consistent with neither branch’s small point estimate for this null-signal control being a robust finding. CEMP’s perturbations around the ACT baseline are not the exchangeable, independently sampled candidates DP provides, so a scorer with no real signal is not guaranteed to land on the same null result in both rows, only on the same underlying lack of signal.

Therefore, every other scorer keeps the ACT output close to never, roughly 0% of windows in Branch 2. This is not comparable to the DP row’s candidate-kept rates, and should not be read as the world model strongly disagreeing with the policy in the same sense. CEMP explicitly searches for the candidate that scores highest under the same objective the scorer then uses to select, so it is expected to beat a fixed, unoptimised ACT baseline on that same metric most of the time. The low kept rate across this row is a mechanical consequence of pairing a search procedure against a fixed baseline, not a finding about how much the world model and the policy disagree.

CEMP’s compute cost against the raw pool it replaces is reported in Appendix B.5. It costs more for every scorer, most severely for DexWM, whose full-transformer scoring makes the 40-fold increase in candidate count a genuine 40-fold increase in forward-pass compute, a cost accepted rather than reducing CEMP’s search budget for DexWM alone.

Every result in this section is secondary to the DP row reported above, but the reason is not uniform across the pooled mean. Two of the three seeds behind it, 43 and 44, are natively sampled exactly as DP’s candidates are, with no CEMP involved (Appendix B.6.1); only seed 42’s contribution carries the circularity caveat that follows from ACT’s CVAE posterior collapse there (Section 5.4.1), since CEMP search stands in for the candidate diversity that checkpoint alone lacks. The scorers rank the same way here as they did in Section 6.1, and that agreement is worth having. The pooled row as a whole still does not carry the same evidential weight as DP’s, since a third of it reoptimises the candidate pool against a fixed baseline that a search procedure is expected to beat almost by construction, not against a fixed pool the policy never saw. None of the numbers in Table 6.5 should be read as validating or invalidating a world model on their own. Their role is to reinforce the DP row’s ranking of scorers, not to stand as a main result in their own right.

6.3 Selection quality across task difficulty

Section 6.1 reports capture as flat across pool size when aggregated over the full test set, but a single flat number averaged over every window can hide a serious local failure, and this section finds exactly one. On the fifth of windows where the policy’s own unselected output was already most accurate, the selector makes the outcome worse on average, not merely less improved. This section checks for that directly. The canonical DP Branch 2 test set is sorted by baseline (candidate-0) error, the error the policy already makes before any world-model selection, and split into five equal-sized quintiles. Q1 is the fifth of windows where the policy’s own uncorrected output was already most accurate, and Q5 is the fifth where it was least accurate. Capture, the percentage of the oracle ceiling recovered (Section 6.1), is then computed separately within each quintile rather than pooled across all of them, so a failure confined to one difficulty band is visible instead of being averaged away. This is done at $K = 8$ rather than the canonical $K = 5$, and for both windowing strategies compared in Appendix B.1: S1 (Strategy 1, the fixed-stride sampling used everywhere else in this chapter) and S2 (Strategy 2, a state-change-targeted alternative). Both columns are pooled across all 3 seeds, the same standard applied throughout this chapter.

Table 6.7 V-JEPA2-AC capture (percentage of the oracle ceiling recovered, Section 6.1) by baseline-difficulty quintile, $K = 8$, Branch 2, pooled across 3 seeds. A negative value means the selected candidate is worse than candidate 0, not merely less close to the oracle than a perfect selector would be.

Quintile	S1 capture	S2 capture
Q1 (easiest)	-33.09%	+2.90%
Q2	24.90%	43.14%
Q3	41.10%	57.39%
Q4	49.65%	65.77%
Q5 (hardest)	49.84%	64.40%
All	37.26%	53.94%

On the easiest fifth of windows under S1, selection reverses direction (Table 6.7). Translation error rises by +5.81% [+4.49, +7.29] relative to keeping candidate 0, a capture of -33.09% , rather than falling as it does on every other quintile, and the reversal holds individually at all 3 seeds, not only in aggregate (per-seed margins +5.38%, +4.85%, +7.21%, all worsening). This asymmetry is broadly consistent with how much headroom each quintile leaves for a re-ranking step. Q1 is defined as the fifth with the lowest baseline error, so candidate 0 is already close to the best the policy can produce there; there is little for selection to gain by overriding it, and a comparatively larger relative cost when it overrides it wrongly. On the harder quintiles

the baseline is further from optimal, so an incorrect override costs proportionally less and a correct one recovers substantially more error, which is why capture climbs from Q2 through Q4 and holds roughly flat into Q5. Only Q1 inverts under S1. The aggregate -5.980% reported in Table 6.4 therefore conceals a regime in which the selection mechanism performs worse than not selecting at all.

The Q1 inversion is largely specific to S1, not an intrinsic property of the scorer. Under S2, Q1’s margin is -0.54% $[-1.26, +0.12]$, a confidence interval that crosses zero, against S1’s robust and consistently positive (worsening) $+5.81\%$ $[+4.49, +7.29]$. Every other quintile keeps the same sign and a comparable or larger capture under S2 as under S1 (Table 6.7), so this is not S2 simply inflating every number; it specifically removes the one regime where selection was actively harmful. This suggests the failure mode identified above is tied to how the easiest windows get selected under standard, fixed-stride sampling specifically, rather than a fixed weakness of V-JEPA2-AC’s scoring itself. Section 7.7 returns to this as a second possible route to closing the gap, alongside the score-based gating explored below, which is evaluated only under S1.

A natural response to a scorer that fails specifically on its easiest cases is to make it more cautious. Only trust its selection when it is confident, and fall back to the policy’s own candidate otherwise, an approach analogous to selective prediction [65]. Confidence here is operationalised as the score gap, how far apart the top-ranked and second-ranked candidate’s scores are. A large gap means the scorer strongly prefers one candidate over the rest; a small gap means it is nearly indifferent between its top choices. Abstaining on the fraction of windows with the smallest gaps, meaning falling back to the policy-only candidate on those windows instead of trusting the scorer, does not recover the Q1 loss at any threshold or pool size tested. It erodes the aggregate gain instead.

Table 6.8 Effect of gap-based abstention on the aggregate translation-error change, pooled across 3 seeds [95% CI].

Abstain % (lowest gap)	$K = 5$	$K = 8$	$K = 20$
0% (as run)	-6.001% [-6.504,-5.093]	-7.236% [-7.845,-6.279]	-9.137% [-9.769,-8.023]
10%	-5.630% [-6.129,-4.772]	-6.715% [-7.275,-5.857]	-8.445% [-9.051,-7.434]
30%	-4.701% [-5.132,-4.149]	-5.563% [-5.976,-4.927]	-6.850% [-7.333,-6.078]

Table 6.8 gives the aggregate effect at three abstention thresholds and three pool sizes. The aggregate loss shrinks monotonically as the abstained fraction grows, for example from -7.236% at 0% abstention to -5.563% at 30% for $K = 8$, but this is

not evidence that the gate is finding and removing the Q1-style regressions specifically. Abstaining on a larger fraction of windows simply falls back to candidate 0 more often across the whole test set, diluting both the gains and the losses uniformly rather than targeting the losses alone. Abstention loses aggregate gain steadily across the whole test set rather than selectively removing the Q1-style regressions it was meant to catch. The reason is that the score gap carries essentially no information about whether a selection will be correct. The pooled Spearman correlation (Section 4.6.3) between the gap and the resulting baseline error is -0.0018 $[-0.0159, +0.0147]$, indistinguishable from zero. A wide gap does not mean the scorer’s confident pick is more likely to be right, and a narrow gap does not mean it is more likely to be wrong. Gating on the maximum score itself performs no better. The selector therefore has no selection-time signal, available from the candidate scores alone, that identifies when it is in the loss-making Q1 regime. The Q1 failure is not a confidence-calibration problem a wrapper can fix from the candidate scores alone. It is reported here as a limitation of the selection mechanism itself, not a solved problem (Section 7.7).

6.4 Comparison against DiWA on CALVIN

The primary results above are all measured on a frozen offline metric on EgoDex, so this section asks a different question. It checks whether the same underlying idea, choosing better actions with the help of a world model, reproduces on the benchmark and protocol that motivated it in the first place. Chandra et al.’s DiWA [12] pairs a diffusion policy with a frozen world model and reward classifier on CALVIN, and uses PPO finetuning against imagined rollouts rather than the inference-time selection used everywhere else in this chapter. This section reproduces DiWA’s own protocol on four gated tasks, using the authors’ own checkpoints throughout, as a secondary study rather than a like-for-like extension of the primary results. It reports five conditions, labelled C1 through C4b, each combining DiWA’s own policy and world model differently. Despite its secondary status, this comparison is where the three-conditions account of when world-model guidance helps (Section 3.4) stops being a framework this thesis’s own EgoDex results alone motivate and becomes a finding independently corroborated on a second benchmark: DiWA’s own design satisfies all three conditions by construction, and §6.4.1 below reads its result against exactly that account. Every condition in this section, C1 through C4b, is pooled across the same three training seeds used throughout this thesis, 42, 43, and 44 (Section 4.5.3), the same seed-robustness standard applied to the primary EgoDex results elsewhere in this chapter; this is stated once here rather than repeated at each table below. Appendix B.6.2 gives the full per-seed picture behind these pooled means, mirroring

Appendix B.6.1 for the primary results.

- **C1:** DiWA's own base policy acting alone, DiWA's own checkpoint, no world model, no finetuning. The same role Branch 1 plays throughout this thesis's own EgoDex results.
- **C2:** the same policy and DiWA's own frozen world model and reward classifier, unchanged, with this thesis's own test-time selection harness added on top. The world model sits outside the policy here. It only re-ranks action chunks the policy has already generated, the same role every world-model scorer plays elsewhere in this chapter. No finetuning.
- **C3a:** identical to C1's own checkpoint. No PPO has been applied to either. C3a exists only as a named reference point, the policy evaluated at the exact moment DiWA's own imagined-PPO finetuning is about to begin, so that C3b's own effect can be measured against it. Finetuning itself happens between C3a and C3b, not within C3a.
- **C3b:** the same policy once DiWA's own imagined-PPO finetuning completes, DiWA's own recipe run in full. The world model now sits inside the policy's own training loop, generating imagined rollouts in place of real environment interaction, and PPO updates the policy from those rollouts. Gradients flow only through the policy's own parameters, not through the world model itself, so the result is written into the policy's weights, not just into which action gets picked at inference time.
- **C4a:** this thesis's own selection harness, as in C2, applied on top of C3b's already-finetuned checkpoint. Both conditions use the identical, hyperparameter-corrected finetuned weights; C4a asks whether adding outside-the-policy selection on top of that same checkpoint changes anything.
- **C4b:** the same setup as C4a, but with the reward classifier itself swapped out. DiWA's original classifier was trained before C3b's finetuning happened, so it may be a poor judge of candidates generated by the now-different, finetuned policy. C4b retrains that classifier on samples from C3b's own policy, then reruns C4a's selection with this policy-matched classifier in place of the original, testing whether an outdated classifier, rather than the selection mechanism itself, explains what C4a shows.

Figure 6.2 summarises how the six conditions relate: which checkpoint each one starts from, and which of DiWA's three modules (policy weights, world model, or reward classifier) each one changes.




Figure 6.2 How the six conditions relate. Arrows show which checkpoint each condition starts from and what changes on the way to it. DiWA’s policy, world model, and reward classifier are the same three modules used throughout this section; only their role changes between conditions.

6.4.1 Selection versus finetuning

The first comparison worth making is C2 against C1. Does this thesis’s own selection harness (Figure 6.2), applied to DiWA’s own native policy and world model, produce any real gain on its own? Pooled across the three tested seeds, it does not, reliably. Three of the four tasks move backward slightly, by 2.30 to 3.12 points, and `turn_on_lightbulb` moves forward by 6.06 points, the largest single-task change in either direction, though C1’s and C2’s own confidence intervals overlap substantially on every task, so none of these four changes should be read as a confirmed effect on its own. This is not because C2 is a weakened test. It draws the same $K = 5$ candidate pool used throughout the primary EgoDex results (Section 6.1), and DiWA’s diffusion policy gives it something substantive to choose between, the reward classifier picks a candidate other than the first one on 79 to 81% of decisions, unlike ACT’s collapsed CVAE posterior (Section 5.4.1), which offers no real choice at all. CALVIN’s task success is also binary, so a candidate only registers if it flips an entire episode from failure to success, unlike EgoDex’s continuous translation and rotation errors, and each task here is pooled over roughly a hundred episodes across the three seeds, still one to two orders of magnitude fewer than EgoDex’s 7,607-window test set. C2 is a fair test. Its result is still small and mixed. Table 6.9 reports this alongside C3b.

Placed next to C3b in Table 6.9, the honest reading is that this thesis’s own method looks weak. C3b (Figure 6.2), DiWA’s own complete imagined-PPO recipe, is

Table 6.9 C1, C2, and C3b on CALVIN: this thesis’s own selection harness against DiWA’s own imagined-PPO finetuning, both read from the same C1 starting point. C3b is DiWA’s own trainer-loop evaluation reading; Table 6.11 gives its shared-harness reading. Change is in percentage points (pp), the absolute difference between two percentages, not a further relative change.

Task	C1	C2	C3b	c1-c2	c2-c3b
close_drawer	59.77% [48.28,70.11]	57.47% [47.13,67.82]	93.10% [93.10,93.10]	-2.30pp	+35.63pp
turn_on_led	52.87% [40.23,64.37]	50.57% [39.08,62.07]	90.81% [86.21,96.55]	-2.30pp	+40.24pp
move_slider_left	57.29% [47.92,67.71]	54.17% [43.75,65.62]	82.29% [75.00,87.50]	-3.12pp	+28.12pp
turn_on_lightbulb	67.68% [58.59,76.77]	73.74% [63.64,82.83]	87.88% [81.82,96.97]	+6.06pp	+14.14pp

the right comparison here, not C3a. C3a is only the pre-finetuning reference used to isolate finetuning’s own direct effect, reported separately below; C3b is what DiWA’s method actually is once run in full, and reproducing it faithfully is this section’s purpose. Against that fair comparison, finetuning moves the same starting point by 14 to 40 points on every task, consistently, roughly an order of magnitude larger than selection’s own effect. This is a result to accept plainly, not soften. Section 7.1 states three conditions world-model-guided selection needs aligned before it helps at all: a diverse candidate pool, a scorer receiving meaningful action-conditioned information, and candidate actions compatible with the scorer’s own training support. DiWA’s imagined-PPO finetuning satisfies these by construction rather than by chance. Its reward classifier is not an external judge trained once and left to match whatever the policy later becomes. It is the training signal PPO optimises against directly, so the policy is pushed to become exactly what the classifier already rewards. External selection carries no equivalent guarantee, since its scorer and the policy it ranks were shaped independently and can drift out of alignment, the same distribution-shift concern C4b later in this section exists to test directly.

DiWA’s own authors make a closely related observation. Chandra et al. report that an offline-RL baseline built on a separately trained critic diverged for exactly this reason, sparse coverage of successful behaviour prevented the critic from generalising to the states the policy actually visited, and describe their own on-policy imagined-rollout design as avoiding this failure mode specifically because the reward signal and the policy update share the same rollout distribution throughout training (Chandra et al., Appendix S.4.5). That is the same compatible-training-support condition stated above, observed independently by the authors whose method this section reproduces, not only inferred from this thesis’s own results.

Does this mean deeper integration is simply the better approach, and this thesis’s own selection-only mechanism the wrong design choice throughout? The answer depends on what is available to integrate against. DiWA’s finetuning trains the policy on rollouts imagined inside its own world model, in place of real environment

interaction, over 800 PPO iterations per task, checked against the real CALVIN simulator along the way. This thesis’s own EgoDex setting has no equivalent environment to imagine rollouts against in the first place. The demonstrations are pre-recorded human video, not a controllable simulator a policy can be finetuned against. Test-time selection is not a weaker choice made in place of a better one. On EgoDex, it is the only mechanism this thesis’s data supports. Where an interactive environment exists, as it does for DiWA’s own CALVIN setting, this result suggests deeper integration is worth its cost. Where it does not, inference-time selection is what remains, and the primary results earlier in this chapter show it can still deliver a real, if smaller, gain under that constraint. This is the same boundary Chapter 7’s limitations (Section 7.7) draw around generalising beyond EgoDex. What a mechanism can achieve depends on what kind of environment its evidence was gathered in, and the two settings compared in this section offer genuinely different ones.

6.4.2 Reproducing DiWA’s own checkpoint

C3a and C3b are measured on a different harness from C1 and C2, DiWA’s own PPO-trainer evaluation loop rather than the shared harness C1, C2, and C4a use, which is a narrower technical point worth flagging before reading their own table. It is why C3a’s own number does not match C1’s, most visibly on `turn_on_led` (62.07% against 52.87%), even though both describe the same starting checkpoint before any finetuning. Table 6.10 gives C3a and C3b together, with the direct effect of finetuning between them, alongside a zero-shot evaluation of DiWA’s own already-finetuned checkpoints as a reference point for what the recipe should achieve.

Table 6.10 C3a and C3b on CALVIN, against DiWA’s own reference checkpoint. Both are read on DiWA’s own PPO-trainer harness throughout, not comparable to Table 6.9’s `c1-c2/c2-c3b` columns. Confidence intervals are seed-level only, since this harness emits no per-episode trace to cluster on (Section 4.6.3).

Task	C3a	C3b	c3a-c3b	DiWA checkpoint reference
<code>close_drawer</code>	64.37% [55.17,72.41]	93.10% [93.10,93.10]	+28.73pp	93.1%
<code>turn_on_led</code>	62.07% [58.62,65.52]	90.81% [86.21,96.55]	+28.74pp	82.8%
<code>move_slider_left</code>	54.17% [50.00,56.25]	82.29% [75.00,87.50]	+28.12pp	84.4%
<code>turn_on_lightbulb</code>	60.61% [54.55,66.67]	87.88% [81.82,96.97]	+27.27pp	90.9%

In Table 6.10, pooled across the three tested seeds, `close_drawer` lands exactly on DiWA’s reference checkpoint, `turn_on_led` lands 8.01 points above it, and `move_slider_left` and `turn_on_lightbulb` land 2.11 and 3.02 points below. An exact match was never the expectation. This reproduction reruns DiWA’s training and evaluation code end to end on this thesis’s own hardware and seeds, so stochastic differences between runs are normal and do not need to vanish for the reproduction to

count as faithful. Pooling across seeds also surfaces actual seed-to-seed spread that a single run cannot, most visibly on `turn_on_lightbulb`, whose own per-seed rate ranges from 81.82% to 96.97% (Appendix B.6.2), so its closeness to the reference at the canonical seed alone was partly a seed-specific artefact rather than a fully representative reading. What the pooled closeness across four independent tasks still establishes is stronger than any single exact match would be on its own. It would not happen by chance if the reproduction’s policy, world model wrapper, PPO loop, and evaluation harness did not faithfully match DiWA’s own implementation. This validates the reproduction pipeline itself, which is what gives every comparison in this section its meaning. It does not validate this thesis’s own five-world-model pipeline evaluated elsewhere in this chapter, since every condition here uses DiWA’s own policy and world model throughout, not one of the five studied on EgoDex.

Reaching C3b required one correction worth recording. An initial run left the behaviour-cloning regularisation coefficient α_{BC} at zero for every task. `move_slider_left` and `turn_on_lightbulb` collapsed under this setting. DiWA’s published hyperparameter table (Chandra et al., Table S.5, Appendix S.1.4) sets $\alpha_{BC} = 0.05$ by default and $\alpha_{BC} = 0.025$ for `close_drawer` and `turn_on_led`. Rerunning the two collapsed tasks at the published default fixed both. `close_drawer` and `turn_on_led` converged at $\alpha_{BC} = 0$ initially, below the published value for those two tasks, but were subsequently rerun at the paper’s own recommended $\alpha_{BC} = 0.025$ as well, for consistency with the published recipe rather than because the zero-coefficient run had failed; the C3b, C4a, and C4b numbers reported throughout this section use this paper-aligned rerun for all four tasks. The coefficient mattered exactly where the paper’s default applies, and was not strictly necessary where the paper recommends a smaller value, a reminder that DiWA’s own reported robustness depends on a hyperparameter not shared across tasks and easy to miss without checking the source paper’s appendix.

6.4.3 C4a and C4b: an inconclusive follow-up

With C3b’s own reproduction established, C4a tests the natural next question it raises: can C2 be improved by adding the same finetuning C3b underwent? It starts from the identical checkpoint C3b reports (Figure 6.2), the same hyperparameter-corrected weights, not a separate or untuned run, and adds only this thesis’s own selection harness at inference time. Its own baseline column is C3b’s performance measured under C4a’s harness rather than Table 6.10’s cross-harness number, since C1, C2, and C4a share one evaluation code path and Table 6.10 uses DiWA’s separate PPO-trainer evaluation loop, exactly the same reason C1 rather than C3a is C2’s point of comparison above. Table 6.11 gives the result, alongside the total change from C1, the

original baseline.

Table 6.11 C4a: this thesis’s selection harness applied to the finetuned C3b checkpoint. The C3b column is measured on C4a’s own shared harness, not Table 6.10’s trainer-loop reading.

Task	C3b (same harness)	C4a	c3b-c4a	c1-c4a
close_drawer	93.10% [87.36,97.70]	93.10% [87.36,97.70]	0.00pp	+33.33pp
turn_on_led	93.10% [86.21,97.73]	88.51% [77.01,97.70]	-4.59pp	+35.64pp
move_slider_left	84.38% [77.08,91.67]	80.21% [70.83,88.54]	-4.17pp	+22.92pp
turn_on_lightbulb	87.88% [79.80,93.94]	85.86% [76.77,92.95]	-2.02pp	+18.18pp

C4a is flat-to-slightly-negative on every task. `close_drawer` shows exactly zero net effect; the other three move backward by 2.02 to 4.59 points. Every one of these differences sits well inside C3b’s and C4a’s own overlapping confidence intervals, so this is read as consistently, marginally negative, but not distinguishable from no effect at all, rather than either a confirmed harm or a real improvement. With four tasks and one checkpoint each pooled over three seeds, this is still not enough evidence to conclude that selection helps, hurts, or does nothing once a policy is already finetuned, only that none of the four changes is large relative to its own uncertainty. Unlike C1 against C2 against C3b above, this comparison does not support a conclusion either way. The `c1-c4a` column shows the two corrections stack, most of the gain over the original baseline still comes from finetuning, and stacking selection on top neither adds to nor meaningfully erodes it.

Given how strong `close_drawer`’s own finetuned result already is, its policy succeeds on every episode in its own rollout data, so a recalibrated classifier for that task would have no failure case to train against. C4b is therefore a direct descendant of C3b rather than of C4a (Figure 6.2), restricted to the other three tasks. It retrains DiWA’s own reward classifier on episodes rolled out from C3b’s own policy, then reruns C4a’s harness with the recalibrated classifier in place of the original. This follow-up was motivated by an earlier, buggy version of C4a that appeared to show selection actively harming the finetuned policy, traced at the time to a plausible distribution shift between the classifier’s training distribution and the finetuned policy’s own candidate distribution. Once the bug, a denoising-schedule setting that silently routed every step through the frozen base actor, was fixed, the harm disappeared, but the distribution-shift explanation remained worth testing directly. Table 6.12 gives the result, again alongside the total change from C1.

C4b is close to C4a. Two of the three tasks land on exactly the same pooled point estimate as C4a, with the recalibrated classifier selecting the identical candidate at every one of the three seeds individually, not only in aggregate; `turn_on_led` moves by 1.15 points, smaller than C4a’s own confidence interval. This is read as confirming that recalibration is not doing detectable work here, rather than as confirming or

Table 6.12 C4b: C4a’s harness with the reward classifier recalibrated on C3b’s own rollouts, for the 3 tasks it covers. `close_drawer` excluded, its finetuned rollouts contain no failures to train against. Every seed’s classifier, including seed 42’s, is collected from its own corrected C3b checkpoint.

Task	C4a	C4b	c4a-c4b	c1-c4b
turn_on_led	88.51% [77.01,97.70]	89.66% [81.61,96.55]	+1.15pp	+36.79pp
move_slider_left	80.21% [70.83,88.54]	80.21% [70.83,88.54]	0.00pp	+22.92pp
turn_on_lightbulb	85.86% [76.77,92.95]	85.86% [75.76,94.95]	0.00pp	+18.18pp

refuting the distribution-shift explanation in general. This is the same underlying question Section 5.6 raises about DP’s own Branch 2 margin threshold, whether calibrating a selection mechanism to a specific distribution earns its added complexity, and here, in the one place this thesis tests recalibration directly, the answer is no clearer than the unconditional choice already made there. One caveat still bounds how much weight C4b can carry: each recalibrated classifier is trained on only 150 to 420 rows from a single 29-to-33-episode rollout per seed, with no held-out validation split, and reaches perfect training precision and recall, a pattern more consistent with memorising a small dataset than with a genuinely recalibrated scorer. C4b is reported as a directional probe throughout, not a validated improvement.

This section is not merged into the primary comparison in Section 6.1. CALVIN is a simulated, binary-success benchmark using DiWA’s own policy, world model, task set, and success metric, none shared with the primary EgoDex evaluation, and improvement is measured by finetuning weights rather than by inference-time selection, so the two studies share no comparable axis. It stands as secondary supporting evidence that *supports* the interpretation developed on the primary evaluation, not as proof that this interpretation generalises beyond it, and not as a second data point in the same comparison.

Taken together, the two studies in this chapter point the same way. On the primary EgoDex benchmark, external world-model selection helps only where its candidate pool, action-conditioning, and training support all align, and DP scored by V-JEPA2-AC is the one pairing that clears all three. On CALVIN, the same three conditions explain why DiWA’s own weight-level finetuning outperforms this thesis’s own external selection applied to the identical policy and world model by roughly an order of magnitude. Chapter 7 interprets both results against that shared framework.

7 Discussion and Limitations

7.1 What the benchmark shows

The final benchmark gives a narrower and more defensible conclusion than the earlier project plan. World-model guidance can improve offline trajectory selection, but the positive semantic result is concentrated in Diffusion Policy candidates scored by V-JEPA2-AC. Branch 2 reduces translation error, averaged over the full 16-step predicted chunk and pooled across three training seeds, by 5.98% and rotation error by 0.96 degrees relative to the policy-only candidate. Branch 3 also improves both metrics, but less strongly. This supports the policy-first version of the guard framework this thesis proposes. Keep the policy candidate available, then use the world model to choose among diverse policy-generated alternatives.

Section 1.2's research question asks whether a frozen world model, acting only as the final decision stage over candidates a policy already proposes, improves the selected motion chunk. The evidence above answers this directly. Under the tested pairing of DP candidates and V-JEPA2-AC, it does, by a margin that survives every causal check applied in this chapter. That question was deliberately asked wide before any evidence came in. Five world-model families, V-JEPA2-AC, LeWM, DexWM, DINO-WM, JEPA-WM, and the PointWorld proxy control, were evaluated against two structurally different policies, ACT and DP, not because every family was expected to succeed, but because answering "do latent world models improve selection" honestly means testing more than one candidate world model and letting the evidence decide which pairing works, rather than committing in advance to a single architecture. One family succeeding cleanly and the other four each failing for a different, diagnosed reason is the answer this thesis set out to find, not a shortfall against a target of five positive results.

The other four families each carried a different, diagnosed shortcoming rather than a uniform failure to work. Section 7.3 works through each in turn, architecture, result, and why, rather than repeating that detail here.

This creates a useful scientific result. The strongest positive evidence behind it comes from one tested combination, DP candidates scored by V-JEPA2-AC; the other four tested world-model families did not produce the same result under this benchmark (Section 7.3). Read across all five, the tested results suggest that world-model-guided selection performs better when three conditions align:

- (i) the candidate pool is diverse,
- (ii) the scorer receives meaningful action-conditioned information, and
- (iii) the candidate actions remain compatible with the scorer's training support.

This is offered as a pattern this thesis’s own experiments support, not a general law established across models or datasets beyond the ones tested here. When one of these conditions fails, the world-model label alone is not enough.

This thesis’s approach to a non-working case is diagnostic rather than dismissive, and that approach is itself part of what this chapter contributes. ACT’s own candidate pool is degenerate at inference time (Section 5.4.1), which would ordinarily leave no basis for a world-model comparison on that policy at all. Rather than exclude ACT from the benchmark, this thesis built CEMP specifically to recover a usable, searched candidate pool around its frozen output, then reran the same five world-model scorers against it (Section 6.2). Those results are reported as evidence for the world-model-plus-CEMP system as a unit, not promoted to the same footing as the DP row’s primary claim, but they still let the same five-family comparison run on a second, structurally different policy rather than being abandoned when the first policy’s own candidates turned out unusable. LeWM received the same treatment on its own terms, tested through two distinct configurations rather than a single reported number, detailed in Section 7.3. Testing a further configuration or a workaround when the first result is null, rather than writing the model off, runs through every non-positive result in this chapter.

A secondary reproduction study on DiWA’s own CALVIN benchmark (Section 6.4) sharpens this same point using a different policy, world model, and benchmark entirely. Applying this thesis’s own test-time selection harness to DiWA’s own frozen policy and world model produces a small, mixed effect, no better than the baseline in any reliable sense. Reproducing DiWA’s own imagined-PPO finetuning on the same starting policy instead moves task success by 14 to 40 percentage points on every task, roughly an order of magnitude larger, and not close. The three conditions above explain why. DiWA’s reward classifier is not an external judge that happens to match a policy it played no part in shaping. It is the training signal PPO optimises against directly, so the policy is pushed to become exactly what the classifier already rewards, an alignment guaranteed by construction. This thesis’s own external selection, on CALVIN and throughout the primary EgoDex results alike, carries no equivalent guarantee, since its scorer and the policy it ranks are shaped independently. Where an interactive environment exists to finetune against, as it does for DiWA’s own CALVIN setting, deeper integration is worth its cost. This thesis’s own EgoDex setting, pre-recorded human video rather than a controllable simulator, has no such environment, which is why external selection is what this thesis evaluates throughout, not evidence that it is the stronger mechanism in general.

7.2 Why Branch 2 outperforms Branch 3

Section 1.2.1 introduced Branch 3 in basis terms, as the world model ranking the remaining candidates with candidate 0 removed. That framing holds structurally, candidate 0 genuinely is unavailable to Branch 3, but it is not the same claim as Branch 3 producing better or more neutral judgement than Branch 2. This section refines that basis description now that results are available.

Branch 2 and Branch 3 share the same candidate pool and world-model scorer, and differ only in whether the policy's own preferred candidate remains available to select (Section 6.1 sets out the three-way interpretive logic this comparison relies on in full). What the experiments directly show: in the final DP plus V-JEPA2-AC result, Branch 2 improves translation by 5.98%, Branch 3 by 5.37%, both pooled across three training seeds; pooled the same way, Branch 2 keeps candidate 0, rather than a world-model-selected alternative, on 23.17% [22.48, 23.84] of windows; and Branch 3 has no such fallback and must choose from candidates 1 to 4 on every window. Everything beyond these three observations below is this thesis's own interpretation of why the pattern looks this way, not a further measurement.

The interpretation offered here is that Branch 3 is best read as a check on Branch 2's judgement, not a separate, higher-authority condition competing with it, since removing candidate 0 does not give the world model anything extra to work with, it only removes the option of falling back to the policy's own choice. Branch 2's edge is read as practical rather than as a sign Branch 3's judgement is worse, since the windows where Branch 2 keeps candidate 0 are read as cases where the policy's own candidate is already competitive and the world-model score does not justify changing it, not as failures of the selection mechanism, though the experiment design does not directly distinguish this reading from the alternative that Branch 2 is simply defaulting to a safe choice in truly ambiguous cases.

Read this way, Branch 3 is doing real work rather than merely losing to Branch 2. Forced to choose without candidate 0 available at all, it still improves on the baseline by 5.37%, which this thesis reads as the world model's judgement being substantive rather than dependent on always having the policy's own preferred output within reach, though this interpretation rests on the observed delta alone and is not independently verified by a separate test. The result favours a fallback deployment shape over an unrestricted one, not evidence that the world model's independent judgement is unreliable when tested without a safety net.

7.3 Interpreting the model-family differences

The model-family comparison should be read as an evidence hierarchy, not as a simple leaderboard. Each family is summarised here in outline, architecture, result, and why, referencing Chapter 3 for full architectural description and Chapter 6 for the underlying evidence rather than repeating either in full.

V-JEPA2-AC (Section 3.2.3) is the only family that gives a clear positive semantic world-model result on the headline DP evaluation (Table 6.4), and its own architecture is the most direct explanation why. Its frozen ViT-Large encoder is pretrained on over one million hours of internet video, so it carries real, video-specific temporal priors before this study trains anything on top of it. The only component trained within this thesis’s fair exposure budget (Section 4.5.1) is a lightweight action-conditioned head. The zero-action ablation, V-JEPA2 no-AC, removes only the action-conditioning input from the same pretrained backbone and collapses to the baseline in Branch 2 (Table 6.4), confirming the action-conditioned head, not the pretrained visual features alone, is doing the work that produces the positive result.

DexWM (Section 3.2.3) offered the clearest opportunity for a genuine comparison between two pretrained backbones, and its own shortfall is a caveat about checkpoint availability rather than architecture. Like V-JEPA2-AC, it is built on a frozen, externally pretrained backbone, DINOv2 rather than a video encoder, feeding a CDiT predictor. Unlike V-JEPA2-AC, that predictor is the substantial, published component of the architecture, and the officially reported EgoDex+DROID-pretrained predictor checkpoint was never released. This thesis’s evaluated DexWM therefore substitutes a predictor trained from scratch under the same fair exposure budget as every other family, unlike V-JEPA2-AC’s head, which was always a small, from-scratch component by original design. Reproducing DexWM’s own full-scale pretraining recipe independently of the missing checkpoint was also not a reasonable option within this thesis’s budget. Even at the same fixed exposure every family trains under, DexWM is already by far the most computationally expensive family to train, ≈ 28 seconds per step against LeWM’s ≈ 0.047 (Table 5.1), so attempting the officially reported EgoDex+DROID pretraining scale on top of that was a compute cost this thesis deliberately did not take on, not only a checkpoint it could not obtain. Under this fixed-budget route, DexWM does not improve translation error in either branch (Table 6.4), a valid negative result for the implementation actually evaluated, not for the architecture at its intended pretraining scale. What was observed is the negative delta under the fixed-budget substitute; the suspected cause is that a from-scratch CDiT predictor at this budget cannot match a predictor pretrained on the official EgoDex+DROID recipe; the experiment that would test this directly is training DexWM at that published scale, once a checkpoint or the raw DROID data becomes available,

and comparing the resulting delta against this fixed-budget result under otherwise identical evaluation.

LeWM (Section 3.2.3) is investigated through two distinct configurations rather than a single reported number, specifically to isolate which part of its mechanism, if any, is doing real work. The rollout configuration actually invokes LeWM’s own learned visual dynamics; the action-prior configuration turns that mechanism off and scores only action-distribution regularity. Table 6.2 reports both. Neither passes the ranking-validity gate at the full 7,607-window scale (Section 6.1), so LeWM’s visual-rollout claim is not validated as a semantic world-model result. The narrower, valid result is that training-only action regularity, on its own, still improves frozen DP candidate selection, useful information but a smaller and different claim from a working learned-dynamics scorer. What was observed is a rollout scorer whose margin is real but whose ranking does not track true candidate quality under a direct Spearman check; the suspected cause is that LeWM’s compact, from-scratch single-token latent does not carry enough task-relevant structure for its rollout predictions to discriminate between candidates reliably; the experiment that would test this is retraining a rollout scorer with a validation-only split reserved specifically for tuning the ranking-validity gate itself, isolated from the split used to develop the scorer, so a genuine rollout improvement could be distinguished from one that only looks better on the same data it was tuned against (Chapter 9).

DINO-WM and JEPA-WM (Section 3.2.3) are official-checkpoint routes rather than models trained within this study, one paired with DexWM’s DINOv2 backbone family and the other with V-JEPA2-AC’s predictive branch (Figure 3.3). Both expect a 7-dimensional DROID-style action interface incompatible with this study’s 18-dimensional bimanual wrist chunks, so each is evaluated only through an adapter, in a validation-only pilot rather than a held-out claim (Section 5.5.4, Appendix A.4). The pilots show a model can improve its own latent support score while failing to improve the trajectory metric this thesis actually measures, a further, concrete instance of the action-compatibility boundary this section describes throughout. What was observed is an improved latent support score alongside a flat or worse trajectory metric; the suspected cause is information lost by the train-only linear adapter compressing this study’s 18-dimensional bimanual actions into the checkpoints’ native 7-dimensional interface; the experiment that would test this is scoring the same candidates through a version of either model finetuned natively on 18-dimensional actions rather than through the adapter, which this thesis’s exploratory, validation-only budget for these two routes did not extend to.

The PointWorld proxy control is not an evaluation of Huang et al.’s own point-cloud architecture (Section 3.2.3), which reasons over rigid-arm forward kinematics structurally incompatible with continuous $SE(3)$ wrist transforms from

human hand demonstration. Instead it reuses V-JEPA2-AC’s own visual pathway with deliberately weakened action conditioning (Section 5.5.7), a control built specifically to separate the contribution of visual quality from the contribution of action information. Its mixed result, a translation point estimate no longer distinguishable from zero once pooled across seeds but a small, consistent rotation improvement (Table 6.4), is consistent with that design. Visual quality alone does not reproduce V-JEPA2-AC’s positive translation result, though it still carries some useful rotation signal.

Read together, V-JEPA2-AC, DexWM, and LeWM also contrast three fundamentally different representational strategies, video-pretrained transfer, image-pretrained transfer with geometric supervision, and self-supervised task-data learning from scratch, respectively (Section 3.2.3), a difference in backbone provenance rather than a departure from the fair-budget protocol governing training (Section 4.5.1). All three receive the same 100,000-item exposure budget within this study, but differ in what each backbone already carries in before that budget starts. V-JEPA2-AC’s ViT-Large backbone and DexWM’s DINOv2 encoder are pretrained externally, on video and image corpora far larger than this study’s own budget; LeWM’s ViT-Tiny backbone has no such external pretraining and is trained entirely from scratch inside the same budget that governs every other model here. The selection deltas above should therefore be read model by model, not as a head-to-head architecture ranking, since a from-scratch model and two externally-pretrained ones start this study’s own fair-budget training from very different places, even though that training itself is held equal. Inference cost is the one dimension directly comparable regardless of this asymmetry:

- DexWM’s DINOv2 encoder and CDiT predictor make it the heaviest of the three at planning time, consistent with its own measured raw-pool runtime (Appendix B.5, reported for the ACT row rather than this one, the most detailed runtime comparison in this study);
- V-JEPA2-AC’s frozen backbone with a lightweight transition head sits in between;
- LeWM’s from-scratch ViT-Tiny is the cheapest by design, consistent with its reported $48\times$ planning speedup, though this study did not separately measure its raw-pool runtime to confirm that figure directly.

This reinforces the central point running through every family above. World models must be action-compatible and metric-aligned before they can become useful evaluators.

7.4 Reproducibility and open model releases

The DexWM investigation raises a broader reproducibility issue. Open-source code is valuable. It lets independent researchers inspect architectures and understand how a result was constructed. But code alone is not a reproducible model release when the published claim depends on large-scale pretraining. DexWM’s architecture and data recipe are public, but the pretrained checkpoint for the reported model was not, and reproducing that recipe independently would need DROID raw-data onboarding, preprocessing, storage, and compute outside this thesis’s fairness contract.

A similar gap appeared while reproducing DiWA (Section 5.5.5), which also did not publish the base and fine-tuned policy checkpoints its own reported numbers depend on. There, direct correspondence with the DiWA authors, including a repository issue, produced the missing files quickly enough for the reproduction in Section 6.4 to proceed. This responsiveness is worth naming as a standard other released benchmarks should aim for, not as a criticism of any project that could not offer it. A reproducible release is not only a matter of what is published on day one. It is also a matter of how quickly a gap can be closed once one is found.

The lesson is not that incomplete releases have no value. They do. It is that open code, open data, open checkpoints, and fully reproducible evidence are different levels of openness with different claim boundaries, and that active maintainer engagement can matter as much as what shipped on day one.

7.5 Why semantic alignment and object grounding were excluded

Semantic alignment and object grounding were part of the original implementation plan, and both modules were fully built (Appendix A), but neither was carried into the final benchmark claims. Adding them only to V-JEPA2-AC would give one model family privileged inputs not available to LeWM, DexWM, or the controls, changing the question from whether the world model improves selection under a shared protocol to whether extra semantic information helps one specific model. The exclusion applies symmetrically. Neither module conditions any family’s result in the final benchmark.

Later exploration of other model families suggests the original intuition was not misplaced. PointWorld’s own architecture lifts each scene directly into a 3D point cloud (Section 3.2.3), an object-and-scene-level representation built into the model itself rather than added afterward, and DexWM’s Hand Consistency loss imposes a related structural prior over hand geometry (Section 3.2.3). That independently designed architectures build in richer scene or geometry information than the plain

visual and action inputs used here is a reasonable signal that enriching a world model’s own representation this way, not only its selection procedure, is a promising direction, not merely a design choice this thesis happened to try and set aside. Testing it properly needs a symmetric design, where every compared family receives equivalent enrichment rather than one alone, and is proposed as future work in Section 9.

7.6 Three ways a world model can be used

World models are used three distinct ways across this thesis and DiWA, worth naming explicitly since CEMP (Section 6.2) sits between the other two. DiWA’s world model is training-time and internal. PPO trains on rollouts imagined inside it, updating the policy’s weights directly, so the world model’s judgement is baked into the policy itself even though gradients never pass through the world model (Section 6.4). The DP row’s world-model scorers (Section 6.1) are the opposite extreme, inference-time and fully external. They only rank candidates the policy has already generated, with no perturbation and no gradient fed back into the policy at all. CEMP sits in between, inference-time and external like the DP row, but actively searching rather than only ranking. It perturbs ACT’s frozen output and asks the world model to rank the perturbations, all after training has finished. This is why CEMP works around ACT’s CVAE posterior collapse (Section 5.4.1) rather than repairing its root cause. It compensates for a broken sampler at decision time instead of fixing the policy that produced it. DiWA’s approach would arguably be the more direct fix, retraining ACT under a world-model-augmented objective to correct the collapse itself, future work not attempted here.

7.7 Main limitations

- **Generalization beyond EgoDex.** Every result in this benchmark is validated on EgoDex alone. This is not simply an evaluation gap to be filled by testing on more datasets later. V-JEPA2-AC’s result depends on an action-conditioned head trained to relate actions to outcomes on this thesis’s own EgoDex data, so there is no reason to expect that relationship to transfer automatically to a different dataset’s action distribution, the same way a policy trained on one demonstration set does not automatically act well on another. Testing generalization properly means action-conditioning the same architecture on new data, not only evaluating the existing checkpoint against a different test set, and is a natural extension for future work rather than something this benchmark’s current evidence could already show.

- **Scope.** The benchmark is offline and trajectory-level, measuring how closely selected wrist-motion chunks match expert demonstrations, not task success or physical execution. DiWA's own C1 through C4b reproduction (Section 6.4) is the one exception, executing live against the real CALVIN simulator; extending the rest of this thesis's own selection harness to live inference the same way is proposed as future work (Section 9).

This is a different caveat from the one DiWA's own authors state about their frozen-world-model evidence, that its errors persist through fine-tuning and gains in imagination do not always survive execution (Section 3.2.1). This thesis's own limitation is a modality mismatch rather than a training-substrate risk. Several families evaluated here, DexWM especially, were built and validated as task-success evidence under closed-loop, physical execution (Table 3.3), a different modality from the offline trajectory-matching this thesis applies to them, not a claim that offline error is what those architectures were built to be judged on.

- **Candidate diversity and policy coverage.** The ACT results show that a valid evaluation can still be uninformative if the candidate generator does not produce meaningfully different actions. This mechanism needs both a scorer and a useful candidate set. This is also why the primary positive claim rests on one policy family, DP, rather than a meaningfully diverse set of policy architectures. ACT could only be evaluated with CEMP's searched pool substituted for its own degenerate one. A secondary reproduction (Section 6.4) does validate this thesis's selection harness against a second, independently-trained diffusion policy under DiWA's own world model, evidence the mechanism is not an artefact of this thesis's own DP training run specifically, though it does not diversify beyond the diffusion-policy family itself.
- **A shared world model in C4a and C4b.** Both conditions (Section 6.4) apply this thesis's own external selection on top of a policy already finetuned by DiWA's imagined-PPO recipe, using the same frozen world model in both roles, inside the PPO loop that shaped the policy's weights, and outside it as C4a's and C4b's selector. Little new information should be expected from a second consultation of a model whose opinion the policy has already been trained against directly; C4a's and C4b's flat, statistically inconclusive results are consistent with that, not only with selection being unhelpful once a policy is finetuned in general. A more informative test of that broader question would pair the finetuned policy with a second, independently trained world model as the external selector, optimised separately from the one used during finetuning, so any residual gain would reflect genuinely new information rather than a repeated opinion. That test was not run here, and is left as future work (Section 9).

- **Model coverage.** LeWM visual rollout, DexWM full upstream-scale pretraining, DINO-WM, JEPA-WM, and DiWA/CALVIN all reached bounded implementation states but not headline evidence states. These routes are valuable because they define limits, but they do not become additional positive claims.
- **Metric coverage.** Translation and rotation errors, each averaged over the full 16-step predicted chunk, are meaningful for offline imitation fidelity, but they are incomplete. They do not measure grasp stability, finger articulation, object interaction, or long-horizon recovery after an error. These other properties matter in their own right, but none of them is a standardised, off-the-shelf metric the way translation and rotation error are. Each typically needs its own tuning and validation before it can be trusted, which is why translation and rotation, limited to two metrics as they are, remain this thesis’s primary and most directly interpretable measures.
- **Checkpoint availability.** DexWM’s and DiWA’s own base and finetuned policies are the two clearest instances of a broader reproducibility-level limitation, discussed in full in Section 7.4. Both are judged on the implementations actually available, not as a verdict on the official architectures at their intended scale.
- **Selection quality is not uniform across the test distribution.** Windows are grouped into five bins by how accurate the policy’s own unselected output already is, from least to most accurate. On the most accurate fifth, the cases where the policy needed help the least, the selector’s own choices average out slightly worse than simply keeping the policy’s original output, and this specific regime is not currently detectable from the candidate scores alone at selection time (Section 6.3). This does not undo the overall positive result. The aggregate gain reported for DP Branch 2 already nets this loss-making fifth against the gains on the other four fifths and still comes out positive. It does mean the gain is not spread evenly across every input, and that a small, identifiable subset of easy cases would be better served by not applying selection at all, a gap the gating direction proposed in Section 9 could close. This regime is also largely specific to the fixed-stride windowing (S1) used throughout the primary evaluation; under the state-change-targeted alternative (S2), the same easiest-fifth loss nearly vanishes (Section 6.3), so a gate is not the only route to closing this gap. A windowing choice that already avoids selecting the easiest cases the way S1 does is a second, untested alternative to a scorer-side fix.
- **Training-seed coverage.** The three-seed protocol is a basic robustness check, not a proof of full stability. The headline DP + V-JEPA2-AC result showed consistent results across the three tested seeds (Section 4.7), as did LeWM’s rollout scorer on both rows. ACT is diagnostic in a different way, since whether its CVAE posterior collapses is itself seed-dependent, not deterministic (Section 5.4.1), and

this thesis’s evaluation pipeline is built to detect that per checkpoint and apply CEMP only where collapse actually fires, rather than imposing it uniformly across every seed regardless of whether the checkpoint needs it. Once that adaptive rule is applied, ACT’s own guard results are consistently signed across the three tested seeds. DexWM remains a single-seed result by deliberate decision, not an oversight. Two further seeds were estimated at roughly 112 GPU-hours of training, and its own training-data export was also found missing, adding a further ≈ 9.3 hours of largely GPU-bound regeneration before retraining could even start, so roughly 121 extra GPU-hours in total, a disproportionate cost for a family whose result is already a negative, heavily-caveated one that this thesis does not treat as load-bearing (Section 5.5.3). DexWM’s result therefore carries a higher seed-specific uncertainty than the other trainable families.

- **Threshold calibration history.** The final Branch 2 margin threshold was selected on a disjoint internal validation split (Section 5.6). Earlier in this thesis’s development, before that validation split was introduced, a superseded version of the same calibration was run directly on the frozen test set. The validation-split calibration reaches the same conclusion, the unconditional setting, on windows no tuning decision has seen, which is reassuring, but the earlier test-set inspection cannot be undone and is recorded here as a limitation of the development process rather than smoothed over.

Despite these limitations, the way this benchmark was constructed is itself scientifically defensible. Several earlier results looked stronger but were not clean enough for final claims. Some used test-derived support, some allowed future-goal access, some used inconsistent candidate generation. The final result set is weaker in headline size but stronger academically because it removes these shortcuts, and because the thesis reports negative outcomes when the evidence does not support a claim rather than only the outcomes that worked. That process, not only the headline result, is part of this dissertation’s contribution.

8 Conclusion

This chapter closes the dissertation: it answers the research question under the specific conditions this thesis tested, reports how far each objective from Section 1.2.1 was met, and states this thesis’s main contribution in one sentence before handing off to future work.

Section 1.2 asked whether a frozen world model, acting only as the final decision stage over candidates a policy already proposes, improves the selected motion chunk. Chapter 7 answers this directly, partially and under the tested conditions rather than as a general yes. Under the tested pairing of DP candidates and V-JEPA2-AC, it does. Branch 2 reduces translation error, averaged over the full 16-step predicted chunk and pooled across three training seeds to confirm the result is not a single-run artefact (Appendix B.6.1), by 5.98% and rotation error by 0.96 degrees on the held-out EgoDex test set; Branch 3 also improves both metrics, though less strongly. The other four tested world-model families did not reproduce this result under the same benchmark (Section 7.3), so the positive answer above is specific to this one combination, not a claim that holds across every combination this thesis tested. That 5.98% is more meaningful read against what the same five-candidate pool could achieve at best. Section 6.1’s oracle ceiling, the error reduction a perfect, ground-truth-informed choice among the same candidates would reach, pooled the same way, is 16.51% for Branch 2 and 14.68% for Branch 3. V-JEPA2-AC’s selection recovers 36.2% and 36.6% of that respectively, a little over a third of what this candidate pool makes achievable, not a small residual improvement against an unbounded target.

The result supports the policy-first design this thesis adopts throughout. The best-performing branch is not full world-model authority. It is the branch that lets the world model rank policy-generated candidates while still retaining the policy’s own preferred candidate when that candidate is already best.

Table 8.1 reports, for each of the five objectives set out in Section 1.2.1, whether it was fully, partly, or only in a limited sense achieved.

The broader result is not that every world model helps, and this thesis treats that outcome as informative rather than as a shortfall (Section 7.3 works through why each of DexWM, LeWM, DINO-WM, JEPA-WM, and PointWorld carried a different, diagnosed limitation). Beyond the individual results, the evaluation protocol itself is a contribution. Fixing test windows, using frozen candidate pools, blocking future-goal access, restricting support to training data, and keeping branch authority explicit is what let earlier, historically stronger-looking results be identified as relying on shortcuts, transductive support, or oracle access, and be corrected rather than reported. The final evidence is smaller in headline magnitude than those earlier results

Table 8.1 Achievement status of each objective from Section 1.2.1.

Obj.	Status	Basis
1	Fully	The three-branch benchmark was implemented as designed and used for every family reported in Chapter 6.
2	Partly	Local training exposure was matched, but families built on externally pre-trained backbones still carry an external training advantage this local budget does not equalise (Section 4.5).
3	Fully within the tested configurations	Whether and when selection helps was determined for every tested combination, with a positive result for one and a diagnosed reason for each of the other four; the conclusion applies to the implemented model combinations, not to world models in general.
4	Fully	DiWA's own weight-level finetuning was reproduced and compared against this thesis's own external selection under shared components (Section 6.4).
5	Partly	A likely reason is identified for every non-positive route, but "likely reason" is this thesis's own diagnosed interpretation in most cases, not an independently proven cause; Section 7.3 states, for each family, what further experiment would test that interpretation directly.

but stronger as academic evidence, because it survives the same scrutiny it was built to apply to every other claim in this thesis.

A related lesson, discussed in full in Section 7.4, concerns reproducibility. Open code is not the same guarantee as a reproducible result when a published claim depends on large-scale pretraining that cannot itself be repeated, as DexWM's missing checkpoint illustrates.

The conclusions drawn here concern offline wrist-trajectory prediction on EgoDex. They do not claim closed-loop robot success, physical task completion, or solved dexterous manipulation. Within that scope, this thesis's own experiments across five world-model families and a secondary CALVIN reproduction consistently point to the same finding, namely that world-model scoring improved inference-time action selection specifically where the three conditions identified in Section 7.1 aligned. This is reported as a pattern found across the combinations this thesis tested, not as a general law about world models established beyond them. It did not help merely by being added to a system. The secondary DiWA/CALVIN reproduction (Section 6.4) supports this same pattern from a different experiment. DiWA's own imagined-PPO finetuning satisfies all three conditions by construction, since its reward classifier is literally the training signal the policy is optimised against, and it outperformed this thesis's own external selection applied to the same policy by roughly an order of magnitude as a result. External selection carried no equivalent guarantee in the cases tested, since its scorer and the policy it ranks are shaped independently rather than jointly. The negative and mixed results reported throughout this thesis are as much a part of this finding as the positive one.

This dissertation’s main contribution, in one sentence, is a controlled offline benchmark and reproducible protocol showing that a frozen world model can improve which of a policy’s own already-proposed motion chunks is selected, for at least one tested policy and world-model combination, without retraining the policy and without any claim yet made about real-robot deployment or physical task success. Chapter 9 sets out where this evidence points next.

9 Future Work

Several of the following directions are motivated by discussion already given in Chapter 7. Each entry below points back to that motivation rather than repeating it, and adds what is new here: how the work would proceed and what a positive or negative result would mean. Directions are grouped into three kinds, ordered by how directly each follows from this thesis’s own open questions: work needed to validate the current findings further before anything else is built on them, experiments that would isolate the mechanism behind a given result more precisely, and longer-term directions the current evidence motivates but does not itself test.

9.1 Immediate validation work

- **Online evaluation, highest priority.** Motivated as this thesis’s first limitation (Chapter 7). The benchmark is offline, and this is the single most important open question this thesis leaves unanswered. Every other direction below presumes the offline gain reported here is worth building on; this is the direction that actually tests that presumption against real task performance. DiWA’s own C1 through C4b reproduction (Section 6.4) already shows a route to closed-loop, simulator-executed evaluation. A concrete next step is running this thesis’s own DP + V-JEPA2-AC branch structure through the same live CALVIN protocol DiWA uses, keeping the branch definitions identical so the offline gain already measured (Section 6.1) can be checked against real task success rather than assumed to transfer. A positive outcome would show Branch 2’s offline translation gain corresponds to a measurable rise in task success. A null outcome would not necessarily contradict the offline result. The Q1 quintile analysis (Section 6.3) already shows this gain concentrates unevenly across inputs, so task success may not be linearly sensitive to trajectory-level improvement at this scale.
- **Cross-dataset replication, and closing the DexWM seed gap.** Motivated by the single-dataset limitation noted in Section 4.7. This thesis retrains ACT, Diffusion Policy, V-JEPA2-AC, and LeWM at additional seeds to check the headline result’s stability across training runs (Section 4.7), but no family here is retested on a dataset other than EgoDex, and DexWM remains single-seed by deliberate cost-benefit decision rather than technical necessity. The direct next steps are replicating the DP + V-JEPA2-AC pairing specifically on a second bimanual manipulation dataset, and revisiting DexWM’s own seed coverage should a cheaper route to its training-data export become available, so its single negative result is eventually checked against the same standard applied to every other

family here. A positive cross-dataset replication would show the headline finding is not an artefact of EgoDex’s own action distribution or recording setup; a failure to replicate would bound how far the three-conditions framework (Section 7.1) generalises beyond the dataset it was found on.

- **Q1 inversion gating.** Motivated by the loss-making regime identified in Section 6.3, where the current score-gap abstention mechanism cannot detect the case because the gap carries no information about baseline quality. The direct next step is a gate conditioned on an estimate of how good candidate 0 already is, rather than on scorer confidence, since it is baseline quality that this thesis’s analysis identifies as the actual variable separating the regime where selection helps from the regime where it hurts. Such a gate needs its own held-out validation split, developed and calibrated entirely apart from the test set used throughout this thesis, matching the validation-split correction already applied to this thesis’s own B2 margin calibration (Section 5.6). Resolving it would remove the one remaining reason to withhold world-model guidance on any subset of inputs, letting a deployment use the selector’s aggregate gain across the full input distribution rather than only on the harder cases this thesis’s evidence currently supports.

9.2 Mechanism experiments

- **Candidate generation.** Motivated by ACT’s degenerate native candidate pool (Section 7.3), which needed CEMP’s search before a world-model comparison was possible at all. A direct next step is applying the same kind of measurable-diversity criterion CEMP was built around directly to policy sampling itself, so diversity is designed in rather than repaired afterward. A useful diagnostic before running any new candidate generator is the same pairwise-distance check that first revealed ACT’s collapse (Section 5.4.1), applied as a gate before a scorer is ever trained against the resulting pool.
- **Stronger causal LeWM development.** Motivated by LeWM’s split outcome across its rollout and action-prior configurations (Section 7.3). Future work should separate visual-rollout development from action-prior development from the start, each with its own validation-only split, so a rollout scorer is never tuned or selected using the same test resource its own ranking-validity gate is later checked against.
- **Selection design that targets the same alignment PPO gets by construction.** Section 7.1 identifies three conditions world-model-guided selection needs aligned before it helps at all. DiWA’s imagined-PPO finetuning (Section 6.4)

satisfies all three by construction, since its classifier is literally the training signal the policy optimises against, which is the main reason it outperforms this thesis's own external selection by roughly an order of magnitude on CALVIN. A direction that keeps selection external, avoiding weight-level retraining and the compute cost that comes with it, is to design the selection harness itself to actively satisfy these same three conditions rather than relying on them holding by coincidence:

- (i) measuring and enforcing candidate diversity before scoring,
- (ii) calibrating the scorer specifically against the deployed policy's own action distribution rather than a fixed offline dataset, and
- (iii) re-validating that compatibility as the policy or its outputs drift.

Closing this gap without touching policy weights would keep test-time selection's practical advantages, no retraining and no interactive environment required, while capturing more of the benefit weight-level integration currently holds almost exclusively.

9.3 Longer-term directions

- **Semantic and object-grounding enrichment.** Motivated in Section 7.5. Both modules are already implemented, and PointWorld's and DexWM's own architectures independently build in related scene or geometry information. The direct next step is a symmetric study, the same semantic and grounding inputs offered to every compared family rather than one, so any gain can be attributed to the information itself rather than to which model happened to receive it. A positive result would mean a world model's own representation, not only its selection procedure, is worth enriching. A negative result would mean the pretrained visual and action inputs already used here capture what these modules would have added.
- **Selection-aware finetuning.** C4a and C4b (Section 6.4) stack this thesis's own selection harness onto a policy finetuned in total ignorance that selection would ever be applied to its outputs, and show no consistent benefit from doing so. A genuinely deeper combination would make the finetuning objective itself aware of the selector, for example sampling several candidates per PPO step and crediting the update to the selected candidate's outcome rather than a single sampled rollout's, so training and test-time selection co-adapt instead of composing after the fact. This is a different training algorithm, not a configuration change, and would cost roughly an order of magnitude more compute than the finetuning runs already reported in this thesis to evaluate at the same standard. It is left as future work for that reason, not because the idea is unpromising.

- **An independent world model for C4a/C4b selection.** A related but distinct gap (Section 7.7): C4a and C4b both select using the same frozen world model that DiWA’s imagined-PPO recipe already finetuned the policy against, so any residual selection gain is asked of a model whose opinion the policy has already absorbed directly into its weights. Pairing the finetuned policy with a second, independently trained world model as the external selector, one that never participated in finetuning, would test whether external selection adds genuinely new information after finetuning, rather than repeating information the policy already has. This requires training or sourcing a second CALVIN-compatible world model, which was outside this thesis’s own scope.
- **A more efficient video-pretraining backbone.** V-JEPA2-AC’s advantage over every other family in this thesis traces to its backbone’s video-scale pretraining (Section 7.3), not to any property of the selection mechanism itself. Kuhn et al.’s LeVJEPA [66], released after this thesis’s own results were finalised, reports an efficient, heuristic-free video-pretraining recipe in the same self-supervised joint-embedding family V-JEPA2 belongs to. Whether a backbone pretrained this way would change the margin over the other four world-model families studied here, and whether its efficiency gains would ease the pretraining-scale barrier that ruled out a fully-scaled DexWM in this thesis (Section 7.3), are open questions this thesis’s own protocol is already set up to answer directly, by swapping in a LeVJEPA-pretrained encoder under the same fair evaluation surface.
- **Policy distillation.** World-model guidance could be distilled back into a policy. If the world model reliably identifies better candidates, accepted candidates could become a teacher signal for policy retraining, moving this thesis’s approach from test-time selection toward training-time policy improvement, the direction DiWA itself takes (Section 6.4).

Finally, this thesis’s findings feed directly into the M.AI.ESTRO research project it was conducted within. The project’s VR pilot already plans to drive real-time ghost-hand guidance from a V-JEPA2-AC checkpoint fine-tuned from this thesis’s own EgoDex training, and is expected to extend into transfer-learning and robotic-arm application phases of the wider project. Whichever configuration such a deployment ultimately uses, the one evaluated in this thesis or one closer to DiWA’s training-time integration, the argument this thesis makes for keeping a world model’s authority bounded by a policy’s own proposed candidates, rather than replacing the policy with an unconstrained generative model, remains the same. A policy-constrained selector can only choose among motions the policy already considers plausible. A fully generative predictor has no equivalent constraint, and no built-in way to distinguish a physically safe prediction from an unsafe one. For a deployed system guiding a novice

user's hand or a robot arm, that bound is not an incidental design choice. It is what keeps world-model guidance a step toward movement prediction that is powerful and plausible, rather than powerful and ungrounded.

References

- [1] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. Cambridge, MA: MIT Press, 2018, ISBN: 978-0-262-03924-6.
- [2] B. D. Argall, S. Chernova, M. Veloso, and B. Browning, "A survey of robot learning from demonstration," *Robotics and Autonomous Systems*, vol. 57, no. 5, pp. 469–483, 2009. DOI: 10.1016/j.robot.2008.10.024
- [3] A. Y. Ng and S. J. Russell, "Algorithms for inverse reinforcement learning," in *Proceedings of the Seventeenth International Conference on Machine Learning*, Morgan Kaufmann, 2000, pp. 663–670.
- [4] M. Jeannerod, "The timing of natural prehension movements," *Journal of Motor Behavior*, vol. 16, no. 3, pp. 235–254, 1984. DOI: 10.1080/00222895.1984.10735319
- [5] M. Jeannerod, M. A. Arbib, G. Rizzolatti, and H. Sakata, "Grasping objects: The cortical mechanisms of visuomotor transformation," *Trends in Neurosciences*, vol. 18, no. 7, pp. 314–320, 1995. DOI: 10.1016/0166-2236(95)93921-J
- [6] M. Santello, M. Flanders, and J. F. Soechting, "Postural hand synergies for tool use," *Journal of Neuroscience*, vol. 18, no. 23, pp. 10 105–10 115, 1998. DOI: 10.1523/JNEUROSCI.18-23-10105.1998
- [7] S. Ross, G. Gordon, and D. Bagnell, "A reduction of imitation learning and structured prediction to no-regret online learning," in *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 15, 2011, pp. 627–635.
- [8] T. Z. Zhao, V. Kumar, S. Levine, and C. Finn, "Learning fine-grained bimanual manipulation with low-cost hardware," in *Proceedings of Robotics: Science and Systems XIX*, 2023. DOI: 10.15607/RSS.2023.XIX.016 arXiv: 2304.13705 [cs].
- [9] C. Chi et al., "Diffusion policy: Visuomotor policy learning via action diffusion," in *Proceedings of Robotics: Science and Systems XIX*, 2023. DOI: 10.15607/RSS.2023.XIX.026 arXiv: 2303.04137 [cs].
- [10] Y. LeCun, *A path towards autonomous machine intelligence, version 0.9.2*, OpenReview preprint, 2022. [Online]. Available: <https://openreview.net/forum?id=BZ5a1r-kVsf>
- [11] Welch Labs, *Yann LeCun's \$1B bet against LLMs*, <https://www.youtube.com/watch?v=kYkIdXwW2AE>, Video essay; accessed July 2026, 2026.

- [12] A. L. Chandra, I. Nematollahi, C. Huang, T. Welschehold, W. Burgard, and A. Valada, “DiWA: Diffusion policy adaptation with world models,” in *Proceedings of the 9th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 305, PMLR, 2025, pp. 3378–3400. arXiv: 2508.03645 [cs].
- [13] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA: MIT Press, 2016.
- [14] M. Assran et al., “Self-supervised learning from images with a joint-embedding predictive architecture,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv: 2301.08243 [cs].
- [15] S. Levine, C. Finn, T. Darrell, and P. Abbeel, “End-to-end training of deep visuomotor policies,” *Journal of Machine Learning Research*, vol. 17, no. 39, pp. 1–40, 2016.
- [16] Z. Ding, “Imitation learning,” in *Deep Reinforcement Learning: Fundamentals, Research and Applications*, H. Dong, Z. Ding, and S. Zhang, Eds. Singapore: Springer Singapore, 2020, pp. 273–306, ISBN: 978-981-15-4095-0. DOI: 10.1007/978-981-15-4095-0_8 [Online]. Available: https://doi.org/10.1007/978-981-15-4095-0_8
- [17] K. M. Lynch and F. C. Park, *Modern Robotics: Mechanics, Planning, and Control*. Cambridge University Press, 2017, ISBN: 978-1-107-15630-2.
- [18] Y. Zhou, C. Barnes, J. Lu, J. Yang, and H. Li, “On the continuity of rotation representations in neural networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 5745–5753. DOI: 10.1109/CVPR.2019.00589
- [19] N. Grossmann, E. Gröller, and M. Waldner, “Concept splatters: Exploration of latent spaces based on human interpretable concepts,” *Computers & Graphics*, vol. 105, pp. 73–84, 2022, ISSN: 0097-8493. DOI: 10.1016/j.cag.2022.04.013 [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0097849322000656>
- [20] A. Dawid and Y. LeCun, “Introduction to latent variable energy-based models: A path towards autonomous machine intelligence,” *Journal of Statistical Mechanics: Theory and Experiment*, p. 104 011, 2024. DOI: 10.1088/1742-5468/ad292b arXiv: 2306.02572 [cs].
- [21] D. P. Kingma and M. Welling, “Auto-encoding variational bayes,” in *2nd International Conference on Learning Representations (ICLR)*, 2014. arXiv: 1312.6114 [stat.ML].

- [22] K. Sohn, H. Lee, and X. Yan, “Learning structured output representation using deep conditional generative models,” in *Advances in Neural Information Processing Systems*, vol. 28, 2015.
- [23] A. Vaswani et al., “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017. DOI: 10.48550/arXiv.1706.03762 arXiv: 1706.03762 [cs.CL].
- [24] M. Assran et al., *V-JEPA 2: Self-supervised video models enable understanding, prediction and planning*, Jun. 2025. DOI: 10.48550/arXiv.2506.09985 arXiv: 2506.09985 [cs].
- [25] R. G. Goswami et al., *World models for learning dexterous hand-object interactions from human videos*, Dec. 2025. DOI: 10.48550/arXiv.2512.13644 arXiv: 2512.13644 [cs].
- [26] L. Maes, Q. Le Lidec, D. Scieur, Y. LeCun, and R. Balestriero, *Leworldmodel: Stable end-to-end joint-embedding predictive architecture from pixels*, Mar. 2026. arXiv: 2603.19312 [cs.LG].
- [27] G. Zhou, H. Pan, Y. LeCun, and L. Pinto, “DINO-WM: World models on pre-trained visual features enable zero-shot planning,” in *Proceedings of the 42nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 267, PMLR, 2025, pp. 79 115–79 135. arXiv: 2411.04983 [cs.LG].
- [28] X. Li et al., “Evaluating real-world robot manipulation policies in simulation,” in *Proceedings of the 8th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 270, PMLR, 2025, pp. 3705–3728. arXiv: 2405.05941 [cs].
- [29] Y. Li, Y. Zhu, J. Wen, C. Shen, and Y. Xu, *WorldEval: World model as real-world robot policies evaluator*, 2025. arXiv: 2505.19017 [cs].
- [30] S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Jozefowicz, and S. Bengio, “Generating sentences from a continuous space,” in *Proceedings of the 20th SIGNLL Conference on Computational Natural Language Learning (CoNLL)*, 2016, pp. 10–21. DOI: 10.18653/v1/K16-1002 arXiv: 1511.06349.
- [31] D. P. Kingma, T. Salimans, R. Jozefowicz, X. Chen, I. Sutskever, and M. Welling, “Improving variational inference with inverse autoregressive flow,” in *Advances in Neural Information Processing Systems 29 (NeurIPS)*, 2016. arXiv: 1606.04934 [cs.LG].

- [32] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, “Film: Visual reasoning with a general conditioning layer,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, 2018, pp. 3942–3951. DOI: 10.1609/aaai.v32i1.11671 arXiv: 1709.07871 [cs.CV].
- [33] M. Reuss, J. Pari, P. Agrawal, and R. Lioutikov, “Efficient diffusion transformer policies with mixture of expert denoisers for multitask learning,” in *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. arXiv: 2412.12953 [cs.LG].
- [34] A. Prasad, K. Lin, J. Wu, L. Zhou, and J. Bohg, “Consistency policy: Accelerated visuomotor policies via consistency distillation,” in *Proceedings of Robotics: Science and Systems XX*, 2024. DOI: 10.15607/RSS.2024.XX.071 arXiv: 2405.07503 [cs.R0].
- [35] K. Black, N. Brown, D. Driess, A. Esmail, M. Equi, C. Finn, et al., π_0 : A vision-language-action flow model for general robot control, 2024. arXiv: 2410.24164 [cs.LG].
- [36] T. Z. Zhao et al., “ALOHA Unleashed: A simple recipe for robot dexterity,” in *Proceedings of the 8th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 270, PMLR, 2025, pp. 1910–1924. arXiv: 2410.13126 [cs.R0].
- [37] S. Lee, Y. Wang, H. Etukuru, H. J. Kim, N. M. M. Shafullah, and L. Pinto, “Behavior generation with latent actions,” in *Proceedings of the 41st International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 235, PMLR, 2024, pp. 26 991–27 008. arXiv: 2403.03181 [cs.LG].
- [38] M. Drolet et al., “A comparison of imitation learning algorithms for bimanual manipulation,” *IEEE Robotics and Automation Letters*, vol. 9, no. 10, pp. 8579–8586, 2024. DOI: 10.1109/LRA.2024.3445630 arXiv: 2408.06536 [cs.R0].
- [39] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, “Mastering diverse control tasks through world models,” *Nature*, vol. 640, pp. 647–653, 2025. DOI: 10.1038/s41586-025-08744-2 arXiv: 2301.04104 [cs].
- [40] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, X. Chen, K. Choromanski, et al., “RT-2: Vision-language-action models transfer web knowledge to robotic control,” in *Proceedings of the 7th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 229, PMLR, 2023, pp. 2165–2183. arXiv: 2307.15818 [cs.R0].

- [41] M. J. Kim, K. Pertsch, S. Karamcheti, T. Xiao, A. Balakrishna, S. Nair, et al., “OpenVLA: An open-source vision-language-action model,” in *Proceedings of the 8th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 270, PMLR, 2025, pp. 2679–2713. arXiv: 2406.09246 [cs.R0].
- [42] O. Mees, L. Hermann, E. Rosete-Beas, and W. Burgard, “CALVIN: A benchmark for language-conditioned policy learning for long-horizon robot manipulation tasks,” *IEEE Robotics and Automation Letters*, vol. 7, no. 3, pp. 7327–7334, 2022. DOI: 10.1109/LRA.2022.3180108
- [43] J. Quevedo, A. K. Sharma, Y. Sun, V. Suryavanshi, P. Liang, and S. Yang, “WorldGym: World model as an environment for policy evaluation,” in *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026. arXiv: 2506.00613 [cs].
- [44] Q. Bu et al., “Closed-loop visuomotor control with generative expectation for robotic manipulation,” in *Advances in Neural Information Processing Systems 37 (NeurIPS)*, 2024. arXiv: 2409.09016 [cs].
- [45] S. Nair, A. Rajeswaran, V. Kumar, C. Finn, and A. Gupta, “R3M: A universal visual representation for robot manipulation,” in *Proceedings of the 6th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 205, PMLR, 2022, pp. 892–909. arXiv: 2203.12601 [cs].
- [46] A. Brohan et al., “RT-1: Robotics transformer for real-world control at scale,” in *Proceedings of Robotics: Science and Systems XIX*, 2023. DOI: 10.15607/RSS.2023.XIX.025 arXiv: 2212.06817 [cs.R0].
- [47] N. Lambert, B. Amos, O. Yadan, and R. Calandra, *Objective mismatch in model-based reinforcement learning*, 2020. arXiv: 2002.04523 [cs.LG].
- [48] A. Hu et al., *GAIA-1: A generative world model for autonomous driving*, 2023. arXiv: 2309.17080 [cs.CV].
- [49] J. Bruce, M. Dennis, A. Edwards, J. Parker-Holder, Y. Shi, E. Hughes, et al., “Genie: Generative interactive environments,” in *Proceedings of the 41st International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 235, PMLR, 2024, pp. 4603–4623. arXiv: 2402.15391 [cs.LG].
- [50] N. Tsagkas, A. Sochopoulos, D. Danier, C. X. Lu, and O. Mac Aodha, “The temporal trap: Entanglement in pre-trained visual representations for visuomotor policy learning,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2026. arXiv: 2502.03270 [cs].
- [51] Meta AI, *JEPA-WMs: Official JEPA world model checkpoints*, <https://huggingface.co/facebook/jepa-wms>, Model release; DROID world-model checkpoint, accessed July 2026.

- [52] B. Terver, T.-Y. Yang, J. Ponce, A. Bardes, and Y. LeCun, "What drives success in physical planning with joint-embedding predictive world models?" *Transactions on Machine Learning Research (TMLR)*, 2026. arXiv: 2512.24497 [cs.LG].
- [53] W. Huang et al., "Pointworld: Scaling 3d world models for in-the-wild robotic manipulation," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. arXiv: 2601.03782 [cs.RD].
- [54] R. Hoque, P. Huang, D. J. Yoon, M. Sivapurapu, and J. Zhang, "EgoDex: Learning dexterous manipulation from large-scale egocentric video," in *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026. arXiv: 2505.11709 [cs].
- [55] University of Malta, *M.AI.ESTRO: Mixed reality AI-enhanced skill transfer and robotic optimisation*,
<https://www.um.edu.mt/newspoint/news/2025/05/maiestro-mixed-reality-ai-enhanced-skill-transfer-robotic-optimisation>, University of Malta newspoint announcement, 2025.
- [56] J. Song, C. Meng, and S. Ermon, "Denoising diffusion implicit models," in *9th International Conference on Learning Representations (ICLR)*, 2021. arXiv: 2010.02502 [cs.LG].
- [57] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap* (Monographs on Statistics and Applied Probability). Chapman and Hall/CRC, 1993, vol. 57, ISBN: 978-0-412-04231-7.
- [58] A. C. Cameron, J. B. Gelbach, and D. L. Miller, "Bootstrap-based improvements for inference with clustered errors," *The Review of Economics and Statistics*, vol. 90, no. 3, pp. 414–427, 2008. DOI: 10.1162/rest.90.3.414
- [59] C. A. Field and A. H. Welsh, "Bootstrapping clustered data," *Journal of the Royal Statistical Society Series B: Statistical Methodology*, vol. 69, no. 3, pp. 369–390, 2007. DOI: 10.1111/j.1467-9868.2007.00593.x
- [60] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.
- [61] H. Fu, C. Li, X. Liu, J. Gao, A. Celikyilmaz, and L. Carin, "Cyclical annealing schedule: A simple approach to mitigating KL vanishing," in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 2019, pp. 240–250. DOI: 10.18653/v1/N19-1021 arXiv: 1903.10145.

- [62] R. Y. Rubinstein and D. P. Kroese, *The Cross-Entropy Method: A Unified Approach to Combinatorial Optimization, Monte-Carlo Simulation and Machine Learning*. Springer, 2004.
- [63] N. Hansen, H. Su, and X. Wang, “TD-MPC2: Scalable, robust world models for continuous control,” in *International Conference on Learning Representations (ICLR)*, 2024. arXiv: 2310.16828 [cs].
- [64] W. Peebles and S. Xie, “Scalable diffusion models with transformers,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023, pp. 4195–4205. arXiv: 2212.09748 [cs.CV].
- [65] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [66] L. Kuhn et al., *LeVJEPa: Efficient & scalable video pretraining without the heuristics*, 2026. arXiv: 2608.27395 [cs.CV]. [Online]. Available: <https://arxiv.org/abs/2608.27395>
- [67] S. Bai et al., *Qwen3-VL technical report*, Nov. 2025. DOI: 10.48550/arXiv.2511.21631 arXiv: 2511.21631 [cs].
- [68] Qwen Team, *Qwen3.5: Towards native multimodal agents*, <https://qwen.ai/blog?id=qwen3.5>, Qwen Team blog post, 2026.
- [69] H. Liu, C. Li, Y. Li, and Y. J. Lee, “Improved baselines with visual instruction tuning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv: 2310.03744.
- [70] H. Laurençon, A. Marafioti, V. Sanh, and L. Tronchon, *Building and better understanding vision-language models: Insights and future directions*, 2024. arXiv: 2408.12637. [Online]. Available: <https://arxiv.org/abs/2408.12637>
- [71] A. Marafioti et al., *SmoVLm: Redefining small and efficient multimodal models*, 2025. arXiv: 2504.05299. [Online]. Available: <https://arxiv.org/abs/2504.05299>
- [72] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, “Sigmoid loss for language image pre-training,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023, pp. 11 941–11 952. arXiv: 2303.15343 [cs.CV]. [Online]. Available: <https://arxiv.org/abs/2303.15343>
- [73] M. Cherti et al., “Reproducible scaling laws for contrastive language-image learning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023, pp. 2818–2829.

- [74] A. Radford et al., “Learning transferable visual models from natural language supervision,” in *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021. arXiv: 2103.00020.
- [75] S. Liu et al., “Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection,” in *Computer Vision – ECCV 2024*, ser. Lecture Notes in Computer Science, vol. 15105, Springer, 2024, pp. 38–55. DOI: 10.1007/978-3-031-72970-6_3 arXiv: 2303.05499 [cs.CV].
- [76] N. Ravi et al., “SAM 2: Segment anything in images and videos,” in *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. arXiv: 2408.00714 [cs.CV].

Appendix A Auxiliary Implemented Modules

This appendix describes two auxiliary modules that are fully implemented in the benchmark stack but excluded from the final fair benchmark: a semantic alignment pipeline and an object grounding stage. Both are labelled exploratory throughout this appendix. They produce useful artefacts and remain integrated with the broader system; their implementation is real and functional. Neither was part of the main validated experiment reported in Chapter 6, and neither carries a held-out evidentiary claim. Their absence from the main results chapters reflects a deliberate methodological choice, not a finding about their utility or correctness.

A.1 Why these modules are excluded, and what future work could do

The exclusion applies uniformly across every family this thesis evaluates, not only V-JEPA2-AC and LeWM. Semantic alignment artefacts (captions, tags, text embeddings) and object grounding artefacts (bounding boxes, segmentation masks) could in principle be added to any scoring pipeline through the semantic extraction path described below, but adding them to one family and not another would change the question from whether the world model improves selection under a shared protocol to whether extra semantic information helps a specific model, exactly the confound Section 7.5 avoids.

Two of the five evaluated families already carry an analogous, architecture-internal signal without needing either module here. PointWorld’s own point-cloud lifting and DexWM’s Hand Consistency loss (Section 7.5) each build in a form of scene or hand-geometry structure that V-JEPA2-AC and LeWM do not have. This is not evidence that the excluded modules are unnecessary for those two families specifically. It is a reminder that “no semantic or grounding input” does not mean “no structural prior at all” for every family in this comparison, only for the two held to the strict shared-input contract.

Both pipelines are implemented, validated at the scale reported below, and produce aligned artefacts stored alongside the canonical benchmark outputs. Testing them as a symmetric conditioning channel across every compared family, rather than as a privileged input to one, is proposed as future work in Chapter 9 (the *Selection design that targets the same alignment PPO gets by construction* and semantic-enrichment directions), which sets out the specific experiment this would require and what a positive or negative result would mean.

A.2 Semantic alignment pipeline

The semantic alignment stage converts indexed EgoDex windows into structured semantic artefacts for downstream conditioning and inspection. For each window it produces a natural-language caption, structured tags (contact object, secondary objects, observed action), a canonical semantic text string built from those tags, and a dense text embedding from a configurable backend. The primary downstream use is a semantically meaningful conditioning vector that could be concatenated with the visual latents used by the world-model scoring functions; the path exists but is not activated in the canonical benchmark.

The VLM component was evaluated across six candidate generators:

- Qwen3-VL-2B-Instruct and Qwen3-VL-8B-Instruct [67]
- Qwen3.5-4B [68]
- LLaVA-v1.6-Mistral-7B [69]
- Idefics3-8B-Llama3 [70]
- SmolVLM2-2.2B-Instruct [71]

Each was scored against two embedding backends, SigLIP [72] and OpenCLIP [73], on a shared 64-window sample (seed 42, stratified across task phases). Table A.1 reports the result.

Table A.1 Semantic-pipeline VLM comparison over the shared 64-window sample. SigLIP and OpenCLIP scores are cosine-similarity proxies computed independently within each embedding family and are not directly comparable across the two columns; lower seconds per window is faster.

Model	SigLIP score	OpenCLIP score	Sec./window
Qwen3-VL-2B-Instruct	0.365	0.409	1.432
Qwen3-VL-8B-Instruct	0.431	0.473	2.435
Qwen3.5-4B	0.434	0.478	5.763
LLaVA-v1.6-Mistral-7B	0.422	0.463	3.317
Idefics3-8B-Llama3	0.405	0.450	4.186
SmolVLM2-2.2B-Instruct	0.319	0.361	7.046

Both embedding backends rank the six generators in broadly the same order, so the comparison is not an artefact of one embedding space. Because SigLIP uses a pairwise sigmoid objective rather than a batch-contrastive one [72, 74], its similarity scores do not depend on batch composition, and it was kept as the operational backend as a convenient scoring choice for judging individual caption-frame pairs. Qwen/Qwen3-VL-2B-Instruct, paired with google/siglip-so400m-patch14-384 (SEMANTIC_EMBED_BACKEND=siglip_text), was the operational choice despite returning the lowest absolute SigLIP score (0.365) of the six. At 1.432 seconds per window it is

1.7 to 5 times faster than every stronger-scoring alternative, and a score at this level is still a reasonable semantic-alignment result for egocentric manipulation text-image pairs. For a pipeline intended to scale across large data extractions, that inference advantage outweighs the marginal quality gap. Neither score is checked against human-annotated ground truth; the critic and retry loop described next is what actually controls output quality at run time, and this comparison only chooses which model that loop runs on top of.

A naive approach of captioning once and trusting the output produces unreliable results on EgoDex frames, first-person egocentric views with partially occluded hands, cluttered surfaces, and rapidly changing visual content. A critic module scores the resulting caption against the input frames using a configurable similarity measure. If the score falls below a threshold, the pipeline retries the VLM call with adjusted prompting parameters, up to a configured maximum; if all retries fail or the score plateaus without improvement, the row is committed with `semantic_failure=true` rather than silently dropped or corrupted, so every window in the index gets a complete, inspectable output with an explicit quality flag.

The pipeline supports multi-GPU sharded execution (`semantic_shard_supervisor.py`, `semantic_merge_shards.py`), since captioning and embedding tens of thousands of windows at VLM inference speeds is infeasible on a single GPU, and writes a canonical merged artefact set (`captions.jsonl`, `tags.jsonl`, `embeddings.npy`, and a manifest recording the input fingerprint) that the downstream object grounding stage (Section A.3) reads directly to inherit the same window ordering, so semantic and grounding artefacts never drift out of alignment. The entrypoint is `src/cli/semantic_build_index.py`, orchestrated via `scripts/semantic.sh` or `scripts/run_stage.sh semantic`.

A.3 Object grounding and mask generation

The object grounding stage consumes semantic artefacts from the previous stage and produces two-dimensional bounding-box detections and box-conditioned segmentation masks for each window. It does not implement six-degree-of-freedom pose estimation, FoundationPose-style geometry reconstruction, or object identity tracking across frames, despite the name “object pose estimation” surviving in the codebase from an earlier, broader scope. For each window it reads the semantic row, samples frames, builds object prompts from the semantic tags (prioritising the `contact_object` tag, falling back to a task-derived hint when no tags are present, filtering generic labels such as “object” or “hand” that produce unreliable detections), runs the configured detector, then the configured masker on each detected box. This is

category-level prompting, not instance-level: the detector does not distinguish two instances of the same object class.

The production backends are Grounding DINO [75] for zero-shot open-vocabulary detection and SAM2 [76] for box-conditioned segmentation. Neither was finetuned on EgoDex for this study. Applied across all 26 task categories in `train_part1` (140,239 windows), the resulting pipeline produced at least one mask for 96.60% of windows, with the remaining 3.40% returning no detection for any prompt tried rather than a forced guess, consistent with Grounding DINO’s own open-vocabulary design; mean per-mask confidence across the 757,210 masks produced was 0.41.

Four limitations bound any future use of these outputs for manipulation conditioning: the 2D masks carry no depth or relative-pose information; small or visually similar objects can be merged or misassigned in cluttered egocentric views; there is no temporal tracking across frames within a window; and a missed or misidentified object in the upstream semantic stage propagates directly to an incorrect or absent detection here. Output artefacts (`masks.jsonl`, `detections.jsonl`, and a manifest) inherit their window ordering from the semantic build’s own manifest, so detection and mask rows always align to the same window boundaries used at every earlier stage.

A.4 Official-checkpoint world models: DINO-WM and JEPA-WM

DINO-WM and JEPA-WM entered the benchmark as exploratory official-checkpoint routes and are summarised only briefly in Chapter 3; this section records their literature detail.

DINO-WM performs latent world modelling over a frozen DINOv2 backbone that encodes each frame into a grid of spatial patch tokens, with a predictor advancing the patch-level latent state under action conditioning [27]. Like DexWM, it inherits the spatial resolution and generalisation benefits of DINOv2’s patch-level representations, at substantially higher computational cost than single-token designs such as LeWM, since the higher-dimensional latent state makes both training and planning slower, which is the trade-off LeWM targets with its reported $48\times$ planning speedup [26].

JEPA-WM refers to the officially released JEPA world-model checkpoints, including a model trained on the DROID robot manipulation dataset [51]. The release provides pretrained latent predictors, but their native action interface assumes a single-arm robot action space. Applying either model to the 18-dimensional bimanual wrist space of this benchmark therefore requires an action adapter, and results

obtained through that adapter are reported as validation-level exploratory evidence rather than headline comparisons.

Both models passed load, extraction, and scorer smoke tests, and were then carried through a bounded validation pilot (256 windows, `train_part1_2_3`) rather than the held-out evaluation used for the main comparison. Table A.2 reports the pilot's own numbers. Default DINO-WM and JEPA-WM improved their latent support scores in every cell but did not give a clean improvement on immediate trajectory metrics: total loss rose for both models, and endpoint translation worsened in every default and reweighted cell without exception. Validation-only reweighting recovered limited loss and rotation signals, especially for DINO-WM, whose reweighted rotation error improved by 0.57° against JEPA-WM's 0.37° , but translation and endpoint translation were neutral or worse under reweighting as well. No B2 reweighted run was carried out for either model, only B3; this is a gap in the pilot's own coverage rather than a withheld result. No held-out test claim is made from this route; the pilot is retained here as an implementation record rather than promoted to Chapter 6.

Table A.2 DINO-WM and JEPA-WM validation pilot, 256 windows, `train_part1_2_3`. Support penalty is the world model's own latent support score [unselected→selected], lower is better fit. The remaining four columns report the selected candidate's change from the same shared unselected (act) baseline, identical for every row since it is the same policy output regardless of which world model scores it: loss 0.0870, translation 0.1053, endpoint translation 0.1126, rotation 30.9° . No B2 reweighted run exists for either model.

Model	Cond.	Br.	Supp. pen.	Δ Loss	Δ Transl.	Δ Endpt.	Δ Rot. $^\circ$
DINO-WM	default	B2	853.6→819.8	+0.0011	-0.0003	+0.0002	+0.4
	default	B3	853.6→825.3	+0.0011	-0.0003	+0.0001	+0.4
	reweight.	B3	853.6→828.8	-0.0030	-0.0001	+0.0001	-0.5
JEPA-WM	default	B2	133.5→130.1	+0.0029	+0.0009	+0.0011	+0.7
	default	B3	133.5→130.6	+0.0034	+0.0029	+0.0030	+0.7
	reweight.	B3	133.5→132.5	-0.0016	+0.0010	+0.0007	-0.4

Appendix B Validity, Scaling, and Sensitivity Studies

This appendix reports the sensitivity and scaling studies that justify specific methodological choices fixed in Chapter 4 and used without further comment throughout Chapter 6: the choice of fixed-stride over state-change-targeted windowing, the restriction to three training splits, the candidate pool size, and the universal training-exposure budget. Each section states what was tested, what was found, and what that finding licenses or bounds in the main text. The last two sections depart slightly from this pattern: one reports CEMP’s compute cost, an implementation-level measurement rather than a methodological choice, kept here alongside this appendix’s other cost and scaling material rather than in Chapter 5; the other gives the full per-seed breakdown behind the pooled means reported throughout Chapter 6. Both are placed here as supporting evidence for the main text rather than as headline results in their own right.

B.1 Windowing strategy comparison: fixed-stride versus state-change-targeted

The canonical benchmark uses fixed-stride windowing (Strategy 1, Section 4.1.2) throughout. A state-change-targeted scheme (Strategy 2) was identified later, during pipeline iteration, as a way of correcting Strategy 1’s under-sampling of high-motion segments. It was not adopted for the canonical benchmark, but a standalone comparison was run to check whether it would have been worth adopting, using a parallel state-change-trained Diffusion Policy checkpoint, a matching support extract, and a test index of equal size to Strategy 1, scored under the same scorer, pool size, and evaluation contract.

The fixed-stride scheme samples the episode timeline uniformly, which causes it to under-sample high-motion events, since long near-static segments contribute windows at the same rate as active manipulation periods. This limitation motivated the revised strategy described below, adopted for this standalone comparison rather than the canonical benchmark. Rather than placing windows at fixed intervals, it computes smoothed hand-state change scores from the wrist transform and confidence sequences, selects candidate windows near high-change peaks, and applies episode-first task-stratified sampling with a fixed random seed (see Figure B.1). A relative-position fallback retains slower episodes that would otherwise be excluded, yielding one frozen window per episode stratified across all represented task

categories.




Figure B.1 State-change-targeted sampling pipeline yielding one frozen window per episode.

LeWM is included as the second scorer specifically because it is the weaker of the two, since its diagnostic scorer is already shown elsewhere in this thesis to fail its own ranking-validity gate at full test-set scale (Section 5.5.2), so it is a deliberately unfavourable second case, not a second endorsement. If a windowing change only ever helped, that would be equally consistent with “state-change windowing is generally better” and with “state-change windowing happens to suit V-JEPA2-AC’s specific representation.” A second, architecturally different, and already weaker scorer is the more informative test precisely because it can disagree; a scorer chosen for being reliably strong could not have told the two apart.

Table B.1 Candidate-capture rate under fixed-stride (Strategy 1) and state-change-targeted (Strategy 2) windowing, Diffusion Policy row, $K = 8$. Both scorer rows are pooled across 3 seeds. The bottom two rows are a separate, single-seed sub-analysis that splits V-JEPA2-AC’s own gain between the state-change test windows themselves and the accompanying state-change support extract, via a two-by-two crossing of window and support-extract strategy; this breakdown alone is not pooled.

Scorer	Strategy 1	Strategy 2	Change
V-JEPA2-AC	37.26% [32.46,40.33]	53.94% [52.85,55.88]	+16.68 pp
LeWM (action-prior)	60.08% [57.00,61.95]	57.92% [54.58,61.56]	−2.16 pp
<i>V-JEPA2-AC gain attribution (single-seed sub-analysis)</i>			
State-change windows + policy pairing	87% of gain		
State-change support extract	13% of gain		

Conclusion. State-change selection clearly helps the V-JEPA2-AC cell, a pooled gain whose confidence intervals do not overlap between strategies, and its direction is mechanistically located (mostly the windows and policy pairing, not the support extract). The LeWM cell moves the other way, but its two intervals overlap substantially once pooled, so this is read as consistent with a regression rather than an independently confirmed one. Either way, the effect is specific to which world model is scoring rather than a general property of the data-selection regime. That the sign reverses on the second scorer is itself the useful finding here, not a weakness of the comparison: it rules out "state-change windowing is a free improvement" and rules in "state-change windowing interacts with what a given scorer's representation is sensitive to." Given that LeWM is already established as the weaker scorer in this benchmark, its regression under Strategy 2 is not treated as a second data point against state-change windowing in general, only as evidence that the gain is not universal. As a result this scorer-dependent is not grounds to replace the study's canonical protocol, so Strategy 1 is retained throughout. State-change-aware sampling remains a promising direction for world models that benefit from it, worth testing across the full evaluation matrix in future work rather than adopted on the strength of one favourable cell.

B.2 Training-split scaling: performance across the three-part corpus

Section 4.1.3 restricts the canonical benchmark to the first three training parts of the EgoDex release (`train_part1`, `train_part1_2`, `train_part1_2_3`) rather than the full five-part release. Only the two policies, ACT and DP, were ever evaluated on all three cumulative splits; the three world-model scorers, V-JEPA2-AC, LeWM, and DexWM, are reported only at `train_part1_2_3`, the split used throughout the canonical benchmark. This is stated plainly here rather than implied by a complete table, since it means this appendix cannot show whether each family's result is sensitive to corpus breadth independent of exposure. It can only show that for the two policies.

The three scorers' missing cells were not all missing for the same reason. V-JEPA2-AC had a valid, exposure-budget-corrected checkpoint at both earlier splits; it had simply never been carried through the guard-evaluation stage, so those two cells were recoverable by running the existing evaluation pipeline rather than by retraining, and are now filled below. LeWM also has complete checkpoints at both earlier splits, but the support-bank, action-statistics, and test-export infrastructure that `train_part1_2_3` relies on for scoring was never built for these two splits, so filling those cells needs new pipeline work, not only a rerun. DexWM's checkpoints at both earlier splits predate the multi-GPU gradient-synchronisation fix that motivated the

dedicated `fair_training_fixed100k` retrain route used for the canonical split (Section 5.7), confirmed `world_size=4` and therefore affected, the same retired class as the pre-fix `train_part1_2_3` checkpoints; these two cells are left unfilled rather than reported from a confirmed-unsafe checkpoint, and would need a fresh retrain to fill safely.

Because the training-exposure budget (Section 4.5.1) is fixed at 100,000 effective items regardless of how many unique episodes the corpus contains, growing the corpus while holding exposure fixed does not mean more training. It means broader task coverage seen with less repetition per example. Table B.2 reports each policy’s own translation L_2 (candidate 0, no world-model selection) at every split, and the three world-model scorers’ selected-candidate translation L_2 against DP candidates at the one split where they were evaluated.

Table B.2 Translation L_2 by training split. Policy rows are each policy’s own unselected (candidate 0) prediction; world-model rows are Branch 2 selected-candidate translation L_2 against DP candidates, not a standalone reconstruction loss. LeWM and DexWM remain “not evaluated” at `train_part1` and `train_part1_2`, for the different reasons the paragraph above gives; V-JEPA2-AC’s cells at those splits are a fresh rerun, single canonical seed, matching this table’s existing single-seed convention.

Model	train_part1	train_part1_2	train_part1_2_3
ACT	0.1439	0.1427	0.1433
Diffusion Policy	0.1047	0.1092	0.1142
V-JEPA2-AC	0.1004	0.1033	0.1068
LeWM	not evaluated	not evaluated	0.1025
DexWM	not evaluated	not evaluated	0.1148

Broader corpus coverage under a fixed exposure budget does not produce a consistent improvement for either policy. ACT’s own translation error is roughly flat and non-monotonic across the three splits (0.1439, 0.1427, 0.1433). DP’s translation error rises from `train_part1` to `train_part1_2_3` (0.1047, 0.1092, 0.1142), the opposite of what an “additional data helps” story would predict. This is a real, sourced result, not smoothed over here. It does not, on its own, justify using `train_part1_2_3` as the canonical split on translation-accuracy grounds. The canonical split was fixed for the coverage reason stated in Section 4.1.3, breadth of task categories represented in training, rather than because this specific metric improves with more of the corpus. Whether broader coverage under fixed exposure trades average translation accuracy for something this offline metric does not capture, such as generalisation to task categories under-represented in `train_part1` alone, is not something this appendix’s data can answer and is not claimed here.

B.3 Candidate pool size (K) sweep

Section 4.2 fixes the candidate pool size at $K = 5$ for the canonical benchmark, on compute-budget grounds rather than a literature precedent, decided before any canonical result was scored. The sweep below was run afterward, once canonical results already existed at $K = 5$, specifically to check that fixed choice for robustness. It is therefore a post-hoc sensitivity analysis of an already-fixed experimental setting, not part of how $K = 5$ was originally selected, and it is presented and read that way throughout this appendix.

A common-random-number sweep was run over $K \in \{2, 3, 4, 5, 6, 8, 10, 12, 16, 20\}$, using nested candidate pools so that smaller pools are exact prefixes of larger ones, isolating the effect of pool size from any change in which candidates are available. Four of the five scorers, V-JEPA2-AC, V-JEPA2 no-AC, PointWorld, and LeWM action-prior, are pooled across all three training seeds, the same standard applied throughout this thesis; DexWM remains single-seed, on the same compute-cost grounds stated in Section 4.7. Two findings result.

First, the ranking of world-model scorers is identical at every pool size tested, so no comparative conclusion drawn in this work is an artefact of fixing K at five.

Second, the margin itself grows without a knee, following a logarithmic law across the swept range, so there is no sufficiency threshold at which additional candidates stop helping. This means the pool size cannot be defended as the point at which candidate diversity becomes adequate, and it is not defended that way in the main text. Reported margins are specific to $K = 5$ and would be larger at a larger pool; the comparisons between scorers, which are what the research question turns on, are not.

K is therefore treated throughout as a compute-budget variable, changing the size of the effect that is measured, but not the conclusion drawn from it.

Figure B.2 plots the reduction in translation error against K for all five DP-row scorers. For V-JEPA2-AC the pooled error reduction rises from 3.13% at $K = 2$ to 9.35% at $K = 20$, following $2.68 \ln K + 1.66$ with $R^2 = 0.988$. Each doubling of the pool buys roughly one and a half to two percentage points, tapering slightly at larger K , at twice the sampling cost. Random selection from the same pools was measured within noise of zero at every pool size, both signs appearing across seeds, confirming that the sampled candidates are exchangeable and that index 0 carries no advantage as a baseline. The entire margin is therefore attributable to ranking, not to pool composition.

The sweep covers the Diffusion Policy row only, and the reason is not the same at every seed. At seed 42, ACT’s checkpoint is CEMP-augmented (Section B.6.1), and CEMP candidates come from a planner that queries the world model during search




Figure B.2 Reduction in translation error against candidate pool size for the five world-model scorers on the Diffusion Policy row. All points are scored on nested pools, so differences along a curve reflect pool size alone. Pooled across 3 seeds except DexWM, which remains single-seed (Section 4.7). The vertical guide marks the operating point used throughout this dissertation.

rather than from independent sampling, so pool size there denotes the number of returned elites, not a proposal count comparable to K on the DP row; no common axis exists at this seed, and a sweep is not meaningful there regardless of budget. At seeds 43 and 44, ACT’s checkpoint is not collapsed, so its candidates are sampled natively, the same five-independent- z procedure DP uses, and a comparable K -sweep would in principle be well-defined there. This thesis does not run one. Doing so would only ever cover two of the three seeds, never the collapsed third, splitting the ACT row’s pool-size analysis into a native-sampling regime and a CEMP regime that cannot be pooled together, adding complexity this appendix judges not worth carrying for a row already treated as secondary corroboration throughout this chapter. This is a scope decision made on those grounds, not a fairness one: $K = 5$ itself is fixed identically for both policies regardless of which row’s sweep exists to check it, so the canonical comparison is unaffected either way. Pool-size invariance is therefore established for the Diffusion Policy row and assumed, not demonstrated, for the ACT row.

B.4 Training exposure convergence study

Section 4.5.1 fixes a universal training-exposure budget of $E = 100,000$ effective items, applied uniformly across all model families and all three training splits. This budget was

determined by a convergence study run prior to full training, rather than chosen arbitrarily.

Each of the seven implemented model families (ACT, Diffusion Policy, V-JEPA2-AC, V-JEPA2 without action conditioning, PointWorld, LeWM, and DexWM) was trained independently at six data scales, 1k, 5k, 10k, 25k, 50k, and 100k effective items, on a fixed stratified subset of `train_part1_2_3`. The two control families (V-JEPA2 without action conditioning and PointWorld) were included in the sweep for completeness but excluded from the budget-selection decision, since neither is designed to improve with additional data. Table B.3 reports the validation loss at each scale for all seven families, pooled across the 3 training seeds used throughout this thesis except DexWM (single-seed, Section 4.7); each row uses a different training objective, so raw values are not comparable across rows, and the relevant signal is each family’s own trend across scales.

Table B.3 Validation loss across six training scales for all seven model families, pooled across 3 seeds except DexWM (single-seed). Losses are objective-specific and not cross-family comparable. LeWM uses `fit/loss` for all three seeds, not the `validate/loss_epoch` metric used elsewhere for this family (the paragraph below explains why), and was trained on CPU rather than GPU throughout this specific study, unlike LeWM’s GPU training everywhere else in this thesis.

Model	1k	5k	10k	25k	50k	100k
ACT	0.1302	0.0758	0.0671	0.0660	0.0638	0.0620
Diffusion Policy	10.7765	5.5584	4.1532	3.0601	2.5960	2.2094
DexWM	4.314	2.414	1.995	1.899	1.693	1.356
LeWM	0.3059	0.1992	0.1600	0.0977	0.0847	0.0907
V-JEPA2-AC	1.3865	0.3234	0.2925	0.2455	0.2104	0.2062
V-JEPA2 (no AC)	0.4627	0.4683	0.3979	0.3699	0.3632	0.3781
PointWorld	0.6515	0.6713	0.6410	0.6347	0.6315	0.5879

ACT and LeWM’s own pooled trends are worth explaining in detail, since seed 42’s own trajectory does not tell the same story as the pooled mean for either family.

ACT’s own convergence-study checkpoint collapses at seed 42 specifically, confirmed directly by the same collapse diagnostic used in Section 5.4.1 (`collapsed_at_1e_4` fires by 100k, with pairwise candidate differences and z -injection sensitivity four to five orders of magnitude below seeds 43 and 44’s), the same seed-dependent CVAE posterior collapse established elsewhere in this thesis, appearing independently in a second part of the codebase rather than a new problem. Seed 42’s collapsed decoder regresses toward a near-fixed action from 5k onward regardless of how much more data it sees, while seeds 43 and 44 train normally and their own loss keeps dropping with scale; at 100k the pooled mean of 0.0620 sits

between seed 42’s own 0.1288 and seeds 43/44’s 0.0290 and 0.0283. The wide interval at every scale from 5k onward reports this genuine seed-level split honestly rather than smoothing it into a single number. Seed 42’s own trajectory looks like early saturation in isolation, but this is a collapsed checkpoint stuck at a floor, not evidence the family has converged.

LeWM’s own pooled trend declines steadily through 50k and then rises slightly at 100k, a small uptick present in all three seeds individually rather than sampling noise. No family in the pooled dataset meets a strict $< 0.5\%$ -improvement saturation criterion. V-JEPA2-AC comes closest, a 2.0 % improvement from 50k to 100k, but that step is not uniform across seeds either, seed 42 accounts for most of the pooled decrease while seeds 43 and 44 move by less or slightly the other way. Diffusion Policy and DexWM are still declining substantially at every scale tested, as Figure B.3 shows. The two controls show no consistent downward trend, consistent with their respective proxy designs having no learnable relationship with additional data.




Figure B.3 Normalised validation loss (each family’s value divided by its 1k loss) across six training scales, pooled across 3 seeds except DexWM (single-seed). Solid lines are primary models; dashed lines are controls.

Table B.4 summarises each primary family’s own 50k-to-100k trend, pooled the same way as Table B.3. None of the five shows a clearly earlier plateau. The universal budget is set to $E = 100,000$ as a compute-constrained lower bound rather than a true saturation point for any family, and every family receives the same budget by design, since the extra exposure cannot hurt a family still improving and it ensures every family is compared under the most favourable common condition reachable within the compute envelope.

This budget is a compute-constrained lower bound rather than a true saturation point for any family. The performance gaps between families reported in the results

Table B.4 50k-to-100k trend per primary model family, pooled across 3 seeds except DexWM (single-seed). The universal budget is $E = 100,000$ effective items regardless of this trend.

Family	50k→100k	Note
ACT	−2.8%	Still declining once pooled; seed 42 alone is collapsed and flat (Section 5.4.1)
LeWM	+7.1%	Small increase, present in all 3 seeds; not a plateau
V-JEPA2-AC	−2.0%	Smallest step of the five, but not uniform across seeds
Diffusion	−14.9%	Still declining steeply
DexWM	−19.9%	Still declining (single-seed)
Universal budget		$E = 100,000$

chapter may partly reflect differing distances from their respective convergence points rather than intrinsic architectural differences; this limitation is carried forward to Section 4.7.

B.5 CEMP compute cost

Table B.5 reports the compute cost of CEMP against the raw pool it replaces, Branch 2 only for clarity. Branch 3 timings follow the same pattern in every row. Every scorer costs more, since scoring 64 perturbations per iteration across three iterations plus the final pool means roughly 197 candidates per window instead of five. For the four lightweight scorers this is affordable, tens of seconds of extra runtime across the full test set. DexWM’s scorer is a full transformer over image patch tokens, so the same 40-fold increase in candidate count is a 40-fold increase in genuine forward-pass compute, not an overhead cost that could be optimised away. This cost was accepted rather than reducing CEMP’s search budget for DexWM alone, to keep the procedure identical across scorers. LeWM rollout’s raw-pool cell is empty because no raw, non-CEMP run of that configuration was ever launched on this row. Only its CEMP result was needed for the reported finding, so the direct comparison point was not collected. This is a compute-cost illustration, measured once, not a statistical result, so it is not pooled across seeds the way the guard-quality deltas elsewhere in this thesis are; a check against the equivalent seed-43/44 runs found the raw-pool timings closely reproduce these numbers, but the CEMP timings do not, some scorers cost noticeably more at those seeds for reasons not pinned down here, plausibly shared-cluster contention on the runs in question rather than anything about the scorers themselves. The relative comparison this table exists to make, raw pool against CEMP, holds

regardless.

Table B.5 Guard runtime, raw pool against CEMP, full 7,607-window test set, Branch 2 only. Branch 3 timings follow the same pattern.

Scorer	Raw pool	CEMP
V-JEPA2-AC	43.7s	118.3s
V-JEPA2 no-AC	41.7s	113.1s
PointWorld proxy	42.6s	106.7s
LeWM rollout	—	1,828.5s
DexWM	7,693s	169,873s (≈ 47.2 h)

B.6 Multi-seed validation

This section’s two subsections give the full per-seed picture behind the pooled means reported in Chapter 6: the primary EgoDex comparison first, then DiWA’s CALVIN reproduction, kept as a separate subsection since it is a secondary study on a different benchmark and is kept structurally distinct from the primary result throughout this thesis.

B.6.1 Primary EgoDex results

Section 6.1 and Section 6.2 report headline numbers pooled across three training seeds, not from any one seed alone; this section gives the fuller per-seed picture behind that pooled mean. Every family this thesis could afford to retrain (Section 4.5.3), DP, ACT, V-JEPA2-AC, both mechanistic controls, and LeWM’s rollout scorer, is trained and evaluated identically at seeds 42, 43, and 44, with no single seed treated as more authoritative than the others in the pooled statistic. The DP result is consistent across the three tested seeds; the ACT result is not, and remains affected by seed-dependent CVAE posterior collapse, discussed in full below. Seed 42 remains the reference point only for the small number of design decisions made before the multi-seed protocol existed, such as Table 6.3’s candidate-pool-size sweep, and this thesis is explicit wherever that narrower sense applies rather than leaving it ambiguous. Table B.8 and Table B.10 report Branch 2 translation change per seed, with the matching rotation-change table directly below each; Branch 3 follows the same pattern in every row and is not repeated here. DexWM is excluded from this picture entirely, on the compute-cost grounds already stated in Section 4.7, and is not plotted below.

Tables B.6 and B.7 show why ACT’s pooled Branch 1 baseline (Table 6.1) is not a single representative number. Seed 42 collapses, the same finding already established

Table B.6 Branch 1 policy-only baseline translation error by seed, both policies. Seed 42 is the canonical value reported in Table 6.1.

Policy	Seed 42	Seed 43	Seed 44
DP	0.11416	0.11777	0.11914
ACT	0.14329	0.02425	0.02659

Table B.7 Branch 1 policy-only baseline rotation error by seed, same policies and seeds as Table B.6.

Policy	Seed 42	Seed 43	Seed 44
DP	39.96°	39.91°	40.17°
ACT	41.48°	12.73°	12.83°

in Section 5.4.1, and its Branch 1 error is high accordingly; seeds 43 and 44 do not collapse, and their Branch 1 error is roughly a fifth of seed 42's on translation and a third on rotation. DP's own baseline, by contrast, is tight across all three seeds on both metrics, consistent with the stable headline result reported throughout this chapter.

Reading Table B.8 together with the tables above, the DP row's headline result holds at every seed tested: the same sign, a comparable magnitude, and no seed dominates the other two. This is the clearest confirmation available within this thesis's own protocol that the central finding is a property of the DP + V-JEPA2-AC pairing, not an artefact of one training run. V-JEPA2-AC and LeWM rollout show this same low, uneventful spread, so Section 6.1's interpretation of their pooled means carries over unchanged. PointWorld proxy is the one scorer on this row with a spread worth naming: its translation change ranges from +0.070% to +0.500% and its rotation change from -0.315° to -0.461° , a wider relative spread than the other three scorers show, consistent with this scorer's own pooled interval already crossing zero on translation (Section 6.1); the individual seeds do not contradict that finding, they are the reason for it.

Table B.8 DP-row Branch 2 translation change by training seed. Seed 42 is the canonical value reported in Table 6.4.

Scorer	Seed 42	Seed 43	Seed 44
V-JEPA2-AC	-6.406%	-6.094%	-5.458%
V-JEPA2 no-AC	0.000%	0.000%	0.000%
PointWorld proxy	+0.500%	+0.234%	+0.070%
LeWM rollout	-10.104%	-10.058%	-10.121%

ACT's row tells a different, more interesting story, and it is worth explaining in full rather than folding it into the same "stable" conclusion. Section 5.4.1 establishes that whether an ACT checkpoint's CVAE posterior collapses is itself seed-dependent.

Table B.9 DP-row Branch 2 rotation change by training seed, same scorers and seeds as Table B.8.

Scorer	Seed 42	Seed 43	Seed 44
V-JEPA2-AC	-1.008°	-0.833°	-1.038°
V-JEPA2 no-AC	0.000°	0.000°	0.000°
PointWorld proxy	-0.393°	-0.461°	-0.315°
LeWM rollout	-0.840°	-0.808°	-0.832°

Under this thesis’s canonical training configuration, one of three seeds tested collapses (seed 42) and two do not (seeds 43 and 44). Rather than force every seed’s candidates through the same pipeline regardless of whether that checkpoint actually needs it, this thesis’s evaluation harness checks each trained ACT checkpoint against the same collapse diagnostic used in Section 5.4.1, and applies CEMP only where the check fires. Seed 42’s checkpoint is collapsed, so its candidates come from CEMP’s search, exactly as reported in Table 6.5. Seeds 43 and 44 are not collapsed, so their candidates are sampled natively, the same five-independent- z procedure the original ACT design intends, with no CEMP augmentation at all. A diversity check on both native pools before scoring confirmed genuinely distinct candidates on every one of 2,000 sampled test windows, not merely candidates that clear the strict collapse threshold while remaining nearly degenerate in practice.

Seed 42 was not dropped in favour of a replacement seed once its checkpoint was found to collapse. The shorter route was available: report only seeds 43 and 44, or pick a fourth seed that happened not to collapse, and never surface the distinction between Section 6.1 and Section 6.2 at all. That route was not taken, because it would have hidden a real, reproducible property of ACT under this thesis’s training recipe rather than reported it, and would have erased the reason Sections 6.1 and 6.2 treat DP and ACT in separate sections at all. CEMP exists specifically because seed 42’s collapse was treated as a discovered problem to diagnose and correct for, not a result to quietly avoid: the adaptive diagnostic-then-augment rule described above was added to this thesis’s own evaluation methodology after the collapse was found, precisely so that every seed, including the one that collapses, could still be reported honestly rather than dropped from the study.

The seed-42/seed-43/44 split is this row’s one real outlier pattern, already explained above by the collapse diagnostic, not a second, unexplained source of variance. PointWorld proxy’s ACT-row numbers also cross sign, seed 42 positive, seeds 43 and 44 slightly negative, the same zero-crossing already established for its pooled result (Section 6.2); V-JEPA2-AC and V-JEPA2 no-AC otherwise show the low, consistent spread that lets Section 6.2’s pooled interpretation stand as given.

This per-seed breakdown supports Section 6.2’s conclusion: a small,

Table B.10 ACT-row Branch 2 translation change by training seed. Seed 42 uses CEMP (the checkpoint collapses); seeds 43 and 44 use native ACT sampling (neither checkpoint collapses), per the adaptive rule described above. Seed 42 values are the canonical ones reported in Table 6.5 and Table 6.6.

Scorer	Seed 42 (CEMP)	Seed 43 (native)	Seed 44 (native)
V-JEPA2-AC	−0.609%	−0.866%	−0.138%
V-JEPA2 no-AC	0.000%	0.000%	0.000%
PointWorld proxy	+0.503%	−0.236%	−0.000%
LeWM rollout	−2.159%	−1.478%	−0.476%

Table B.11 ACT-row Branch 2 rotation change by training seed, same scorers, seeds, and CEMP/native split as Table B.10.

Scorer	Seed 42 (CEMP)	Seed 43 (native)	Seed 44 (native)
V-JEPA2-AC	−0.049°	−0.011°	−0.003°
V-JEPA2 no-AC	0.000°	0.000°	0.000°
PointWorld proxy	−0.034°	−0.001°	−0.002°
LeWM rollout	−0.031°	−0.009°	−0.007°

consistently negative translation change at every seed once the adaptive rule is applied, with the same secondary-evidence caveat holding throughout, none of these numbers is promoted to the same standing as the DP row’s primary claim.

B.6.2 DiWA/CALVIN multi-seed validation

Every pooled mean reported above is a mean across three seeds, not a single run; this section gives the full per-seed picture behind those means, mirroring Appendix B.6.1 for the primary EgoDex results. C3a and C3b’s trainer-loop reading carry a seed-level bootstrap only, since that harness emits no per-episode trace to cluster on (Section 4.6.3); every other condition uses the same two-level clustered bootstrap as the primary results.

Table B.12 `close_drawer` success rate by seed, all conditions that apply to this task. C4b is excluded from this task by design.

Condition	Seed 42	Seed 43	Seed 44
C1	55.17%	65.52%	58.62%
C2	58.62%	58.62%	55.17%
C3a	72.41%	65.52%	55.17%
C3b (trainer-loop)	93.10%	93.10%	93.10%
C3b (shared harness)	93.10%	93.10%	93.10%
C4a	93.10%	93.10%	93.10%

Table B.13 turn_on_led success rate by seed, all conditions.

Condition	Seed 42	Seed 43	Seed 44
C1	58.62%	44.83%	55.17%
C2	48.28%	58.62%	44.83%
C3a	58.62%	62.07%	65.52%
C3b (trainer-loop)	89.66%	86.21%	96.55%
C3b (shared harness)	93.10%	96.55%	89.66%
C4a	79.31%	89.66%	96.55%
C4b	86.21%	93.10%	89.66%

Table B.14 move_slider_left success rate by seed, all conditions.

Condition	Seed 42	Seed 43	Seed 44
C1	59.38%	56.25%	56.25%
C2	53.12%	59.38%	50.00%
C3a	56.25%	56.25%	50.00%
C3b (trainer-loop)	84.38%	87.50%	75.00%
C3b (shared harness)	84.38%	84.38%	84.38%
C4a	84.38%	81.25%	75.00%
C4b	84.38%	81.25%	75.00%

Table B.15 turn_on_lightbulb success rate by seed, all conditions.

Condition	Seed 42	Seed 43	Seed 44
C1	69.70%	66.67%	66.67%
C2	75.76%	78.79%	66.67%
C3a	66.67%	60.61%	54.55%
C3b (trainer-loop)	96.97%	84.85%	81.82%
C3b (shared harness)	90.91%	87.88%	84.85%
C4a	90.91%	84.85%	81.82%
C4b	93.94%	84.85%	78.79%

Most of this picture is unremarkable. `close_drawer`'s C3b and C4a rows are flat at 93.10% across every seed regardless of harness, and `move_slider_left`'s shared-harness C3b, C4a, and C4b rows are similarly tight or exactly flat; for these, the pooled means already reported above are a fair summary and this table adds confirmation, not a new reading. Three genuine outliers are worth naming rather than folding into that same summary. First, C1's own `turn_on_led` baseline is volatile before any finetuning or selection is applied at all, 58.62% to 44.83%, a wider spread than C1 shows on any other task, which is why this thesis's own C2-versus-C1 comparison on that task carries a wide interval and should not be read as a precise measurement. Second, C3a's `close_drawer` reading ranges from 72.41% to 55.17%, the widest spread C3a shows on any task, ahead of any finetuning taking place. Third, `turn_on_lightbulb`'s trainer-loop C3b reading ranges from 96.97% at seed 42 down to 81.82% at seed 44, already discussed above as the reason this task's apparent closeness to DiWA's own reference checkpoint was partly a seed-specific artefact of reading the canonical seed alone. None of these three outliers changes which conclusion this section draws; they change how precisely that conclusion should be stated, which is why every headline number above is now reported with its own confidence interval rather than as a bare point estimate.

This per-seed breakdown supports Section 6.4's conclusion: the C1-versus-C2-versus-C3b comparison holds cleanly at every seed, and C4a/C4b's own inconclusive reading is not an artefact of reading a single seed either, the same wide, overlapping intervals appear once each seed is checked individually.