

A Comparative Study of Visuomotor Policies and Latent World Models in Dexterous Manipulation

Graham Pellegrini

Supervisor: Prof. Alexiei Dingli

[PLACEHOLDER: Submission Date]

*Submitted in partial fulfilment of the requirements
for the degree of Master of Science in Artificial Intelligence.*



L-Università ta' Malta
Faculty of Information &
Communication Technology

Plagiarism Declaration

I, the undersigned, hereby declare that the work contained in this dissertation is my own original work and that I have not previously in its entirety or in part submitted it at any university for a degree.

I declare that this submission is my own work and that, to the best of my knowledge and belief, it contains no material previously published or written by another person, except where due acknowledgement is made in the text.

I declare that the use of artificial intelligence tools in the preparation of this dissertation was limited to writing assistance, code suggestions, and literature search support. All research design, experimental decisions, implementation choices, result interpretation, and academic claims are my own.

Signature: _____

Name: Graham Pellegrini

Date: _____

Abstract

This dissertation investigates whether a learned world model can improve the quality of action predictions made by a robot imitation policy on a bimanual dexterous manipulation task. The central question is whether coupling a policy that learns to imitate expert demonstrations with a model that predicts the likely future consequences of candidate actions leads to trajectory predictions that more faithfully reproduce the demonstrated motion.

The implemented system, WACT (World Model Augmented Action Chunking Transformer), operates on a large collection of egocentric hand-manipulation demonstrations drawn from sixty-three task categories. It compares three candidate selection strategies that share the same underlying imitation policy: a baseline that applies the policy directly, a hybrid strategy in which the policy proposes candidate motions and a world model evaluates them, and a world-model-driven strategy in which the world model selects the best candidate independently of the policy's own preference. Two world model families are evaluated within this structure. The first draws on general-purpose visual representations learned from a large-scale video corpus before seeing any manipulation data. The second learns entirely from the task demonstrations themselves, without any external pretraining.

The main evaluation is conducted on a held-out test set of 7,607 evaluation windows spanning sixty-three task categories. The headline finding is that world model guidance substantially improves how closely predicted wrist trajectories match the demonstrated motion. Under world-model-driven selection with the pretrained family, mean translation error is reduced by 77 per cent compared to the policy-alone baseline, and this strategy outperforms the baseline in 91.5 per cent of head-to-head window comparisons. The hybrid strategy also improves substantially, achieving 86.6 per cent of comparisons, and shows greater consistency across individual test windows. Analysis of path shape reveals that the improvement is concentrated in spatial correctness and trajectory geometry rather than in the speed or timing of the predicted motion: path straightness improves to an 89.4 per cent win rate and trajectory smoothing to 93.1 per cent, while velocity and acceleration profiles remain near chance.

The task-data-trained world model achieves a 71.7 per cent win rate against the baseline despite learning entirely from scratch, confirming that world model guidance is not contingent on large-scale external pretraining, though the pretrained family is substantially stronger. Expanding the training dataset from thirty-nine to sixty-three task categories did not further improve the pretrained world model, suggesting it is near a capacity ceiling under the current policy architecture.

The dissertation also documents the benchmark development required to reach

defensible conclusions. This includes the correction of a data coverage defect that had limited one model family to 2.1 per cent of the intended test surface, and the identification and correction of a selection logic error that had caused two of the three branches to produce identical outputs. Both corrections were necessary before the final comparisons could be treated as scientifically meaningful. All conclusions are bounded to offline prediction quality on recorded demonstrations and do not constitute claims about deployed task success.

Acknowledgements

Placeholder for acknowledgements.

Contents

Plagiarism Declaration	i
Abstract	ii
Acknowledgements	iv
Contents	viii
List of Figures	x
List of Tables	xi
List of Abbreviations	xii
Glossary of Symbols	1
1 Introduction	1
1.1 Motivation	2
1.2 Research Question and Scope	2
1.3 Contributions	4
1.4 Dissertation Structure	4
2 Background	6
2.1 Visuomotor policies	6
2.1.1 Behaviour cloning and compounding error	6
2.1.2 Action chunking	7
2.2 Rigid body pose and the $SE(3)$ representation	7
2.3 Latent representations	9
2.4 Conditional variational autoencoders	10
2.5 Transformer networks and the attention mechanism	12
2.6 Self-supervised learning	13
3 Literature Review	15
3.1 Action Chunking with Transformers (ACT)	15

3.2	Diffusion Policy	16
3.3	World models	18
3.3.1	World models as policy evaluators	18
3.3.2	Latent predictive architectures	19
3.4	Offline evaluation	25
4	Methodology	27
4.1	Dataset and Experimental Setup	27
4.1.1	EgoDex: source and selection rationale	27
4.1.2	Windowing protocol	28
4.1.3	Training splits and test set	29
4.2	Three-Branch Experimental Framework	30
4.3	Policy Variants	32
4.4	World Model Families	32
4.4.1	V-JEPA2-AC — Pretrained Backbone Baseline	32
4.4.2	LeWM — Lightweight From-Scratch	32
4.4.3	DexWM — Domain-Specific Architecture	32
4.4.4	V-JEPA2 (No Action Conditioning) — Ablation Control	32
4.4.5	PointWorld — Out-of-Domain Negative Control	32
4.5	Fair Experimental Protocol	32
4.6	Evaluation Metrics and Score Harmonisation	32
4.7	Evaluation Scope and Limitations	32
5	Implementation	33
5.1	Data and windowing pipeline	33
5.2	ACT policy path	34
5.3	V-JEPA2-AC path and caching strategy	35
5.4	LeWM path	36
5.5	Policy guard implementation	37
5.6	Ensuring safe and fair evaluation	38
5.7	Implementation challenges and resolutions	39
5.8	Scaling and GPU management	41
5.9	Semantic alignment and object grounding	41
5.10	Benchmark reporting surface	42
6	Experimental Design and Metrics	43
6.1	Canonical evaluation surface	43
6.2	Frozen branch representatives	43
6.3	Prediction targets and their significance	44
6.4	Branch design and the five evaluated instances	45

6.5	Fair comparison rules	46
6.6	Metric harmonisation	46
6.7	Claims policy	46
6.8	Metric justification	47
6.9	Scaling experiment design	47
7	Results	48
7.1	Branch comparison	48
7.2	Branch authority: world-model-driven versus hybrid	50
7.3	World model family comparison	50
7.4	Path shape and trajectory geometry	51
7.5	Robustness across confidence strata and tasks	52
7.6	Scaling experiment	52
7.7	Why the canonical package supersedes the April 14 results	53
8	Discussion and Limitations	55
8.1	What the benchmark shows	55
8.2	The aggregate versus per-window tension	55
8.3	What the path shape results reveal	56
8.4	The family comparison and its architectural basis	57
8.5	Why benchmark evolution is part of the contribution	57
8.6	Why semantic alignment and object grounding were excluded	58
8.7	Main limitations	58
9	Conclusion	60
9.1	Summary of findings	60
9.2	Methodological contribution	60
9.3	Scope of conclusions	61
9.4	Future Work	61
A	Auxiliary Implemented Modules	67
A.1	Rationale for exclusion from the main benchmark	67
A.2	Semantic alignment pipeline	67
A.2.1	Purpose and outputs	67
A.2.2	Visual language model selection	68
A.2.3	Critic and retry loop	68
A.2.4	Multi-GPU sharding and output artefacts	69
A.2.5	Integration with the broader pipeline	70
A.3	Object grounding and mask generation	70
A.3.1	Scope and current implementation	70

A.3.2	Prompt construction	71
A.3.3	Detection and segmentation backends	71
A.3.4	Output artefacts and alignment	72
A.3.5	Known limitations	72
A.4	Future use under a symmetric design	73
B	Benchmark Evolution and Validity Record	74
B.1	The three benchmark eras	74
B.2	LeWM frame cap defect	75
B.2.1	Discovery	75
B.2.2	Effect on reported results	75
B.2.3	Correction	75
B.3	Branch 3 selection mode defect	76
B.3.1	Discovery	76
B.3.2	Effect on reported results	76
B.3.3	Correction	77
B.4	Oracle goal leakage	77
B.4.1	Discovery	77
B.4.2	Effect on reported results	78
B.4.3	Correction	78
B.5	Goal metric invalidity	78
B.5.1	Discovery	78
B.5.2	Correction	79
B.6	LeWM CPU inference defect	79
B.6.1	Discovery	79
B.6.2	Correction	79
B.7	HDF5 handle management defect	80
B.7.1	Discovery	80
B.7.2	Correction	80
B.8	Guard resumability	80
B.8.1	Motivation	80
B.8.2	Implementation	80
B.9	Rejected comparison inputs and superseded paths	81
B.9.1	Semantics and OPE as benchmark inputs	81
B.9.2	Veto and threshold selectors	81
B.9.3	Early Branch 3 shells	82
B.9.4	Conservative confidence gating	82
B.10	Summary of validity status	82

List of Figures

Figure 2.1	A visuomotor policy maps visual observations to motor actions. The input is a camera frame showing the hand and object, and the output is a proposed action for how the hand should move next.	6
Figure 2.2	Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [8].	7
Figure 2.3	Translation for each arm can be represented by a movement along each of the three spatial axis [12].	8
Figure 2.4	Rotation around the y-axis for each arm. The same circular movement can be imagined around the x and z axes; the total orientation of the wrist is the combined result of rotations around all three [12].	9
Figure 2.5	Example latent-space plot showing how structurally similar items cluster together in the learned representation. Structural similarity here refers to items that share shape, geometry, and likely physical behaviour [14].	10
Figure 2.6	Standard Variational Autoencoder (VAE) (a) and conditional Conditional Variational Autoencoder (CVAE) (b) encoder-decoder pipelines, adapted from [17].	11
Figure 2.7	VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation), adapted from [18].	12
Figure 2.8	Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [21].	13
Figure 2.9	Three machine learning paradigms: supervised, reinforcement, and self-supervised learning [23].	14
Figure 3.1	Architecture of the Action Chunking with Transformers (ACT) policy [7].	16
Figure 3.2	Diffusion Policy: general formulation (a), CNN variant (b), transformer variant (c) [9].	17
Figure 3.3	Classic closed-loop control (a) versus CLOVER's closed-loop visuomotor control (b) [29].	19
Figure 3.4	Three Self-Supervised Learning (SSL) prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [22].	20

Figure 3.5	Video Joint Embedding Predictive Architecture 2 (V-JEPA2) post-training pathways from internet-scale backbone to specialised applications [32]. . . .	21
Figure 3.6	Standard V-JEPA2 (left) versus action-conditioned Action-Conditioned Video Joint Embedding Predictive Architecture 2 (V-JEPA2-AC) (right) [32]. .	22
Figure 3.7	DexWM world model pipeline [34].	23
Figure 3.8	LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [35].	23
Figure 4.1	Fixed-stride windowing scheme (obs16_pred16_s128) used as the canonical training and evaluation index. Author-generated schematic.	28
Figure 4.2	State-change-targeted sampling pipeline yielding one frozen window per episode. Author-generated schematic.	29
Figure 4.3	Three-branch design space. B1 uses the policy’s own preference alone; B2 adds a world-model veto constrained by a margin threshold; B3 gives the world model full selection authority over the frozen candidate pool produced at B2 time. The policy is not re-run in B3.	31
Figure 7.1	Branch-level benchmark metrics. Left: mean translation L_2 . Right: mean rotation error in degrees. Both are chunk-level means over 16 prediction steps across 7,607 test windows. Lower is better.	49
Figure 7.2	Pairwise win rates from the final test package. Each bar shows the fraction of 7,607 windows where the listed branch outperforms the reference branch on translation L_2 . The dashed line marks 50 per cent.	49

List of Tables

Table 3.1	Average consecutive sub-task completions (<code>avg_len</code>) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [29].	19
Table 4.1	EgoDex splits under the <code>obs16_pred16_s128</code> windowing contract. Episode and window counts are derived from the canonical index.	29
Table 7.1	Final benchmark branch summary. Training split: <code>train_part1_2_3</code> (63 tasks, 195k episodes). Evaluation: 7,607 held-out test windows. Translation L_2 and rotation error are chunk-level means over all 16 prediction steps. . . .	48
Table 7.2	Selected pairwise win rates. Translation win rate: fraction of 7,607 windows where the candidate branch achieves lower translation L_2 than the reference branch.	49
Table 7.3	Path shape win rates for Branch 2 V-JEPA2-AC versus Branch 1 ACT-only (7,607 windows). Win rate: fraction of windows where the V-JEPA2-AC branch produces a trajectory with better shape by the given metric.	51
Table 7.4	Confidence stratification for Branch 2 V-JEPA2-AC versus Branch 1 ACT-only. High confidence: wrist confidence score ≥ 0.85 . Low confidence: wrist confidence score < 0.85	52
Table 7.5	Scaling experiment: Branch-level translation L_2 on the test split under two training scales. Lower is better.	53
Table B.1	Summary of benchmark validity concerns and their resolution status in the canonical final package.	83

List of Abbreviations

ACT Action Chunking Transformer.

BC Behaviour cloning.

CVAE Conditional variational autoencoder.

DDIM Denoising Diffusion Implicit Models, an accelerated inference scheme for diffusion models.

DP Diffusion Policy, a visuomotor policy that represents actions as a conditional denoising diffusion process.

I-JEPA Image Joint Embedding Predictive Architecture.

JEPA Joint Embedding Predictive Architecture.

LeWM LeWorldModel, a latent world model trained end-to-end from random initialisation.

SSL Self-supervised learning.

V-JEPA2 Video Joint Embedding Predictive Architecture 2.

V-JEPA2-AC Action-Conditioned Video Joint Embedding Predictive Architecture 2.

VAE Variational autoencoder.

WACT World Model Augmented Action Chunking Transformer.

1 Introduction

Imitation learning studies how computational systems can acquire behaviour from demonstrations rather than from manually specified rules. In manipulation settings, this is attractive because useful behaviour often depends on subtle motion patterns that are difficult to encode by hand. Behaviour cloning provides one direct route by learning a mapping from observations to demonstrated actions, while inverse reinforcement learning approaches the same problem by trying to infer the objectives that make a demonstration sensible [1, 2]. Both perspectives frame imitation as a way to learn from observed behaviour, but they do not by themselves resolve how a system should judge whether a predicted future motion is physically plausible.

This issue becomes sharper in dexterous manipulation. Human hand motion is temporally structured, object dependent, and sensitive to small changes in pose. This follows from the timing of natural prehension, the visuomotor transformation required for grasping objects, and the postural synergies used to shape the hand for tools and objects [3–5]. In imitation settings, this matters because small prediction errors can compound across a future sequence, so local agreement with a demonstration does not guarantee sequence-level alignment [6]. This thesis is therefore concerned with prediction and evaluation of future hand motion, rather than deployed control. The work remains relevant to robotic control because robotic manipulation systems ultimately require reliable motion selection, but the evidence in this dissertation is restricted to an offline computational benchmark.

Visuomotor policies are a common way to connect visual observations with future motion. In this thesis, the focused policy family is the Action Chunking Transformer (ACT), which predicts temporally extended action chunks rather than isolated next-step actions [7]. This chunked formulation is useful because dexterous motion unfolds over short sequences, not only through single-frame decisions. However, a policy that proposes a future chunk still requires a criterion for choosing between candidate futures. This creates the opening for world models.

World models learn predictive representations of how a state may evolve. The initial role considered in this thesis was conservative: a world model could evaluate candidate action chunks proposed by ACT and help select motions whose predicted consequences looked more consistent with the demonstrated future. As the implementation developed, the question broadened. The benchmark also examines whether a world model can take greater authority over candidate selection, moving from support for a policy toward a stronger world model driven selection role. This progression from assistance to selection motivates the comparative structure of the thesis.

This dissertation studies that structure through World Model Augmented Action Chunking Transformer (WACT), which is used here as a comparative benchmark framework for evaluating whether world model support changes the quality of future hand-motion prediction. The framework compares an ACT-only baseline with world model selected variants that use selected and ported world-model families, including V-JEPA2-AC and LeWM. The families themselves are not claimed as original contributions of this thesis. The contribution lies in the controlled comparison, the implemented evaluation surface, and the analysis of how these model families behave when attached to the same policy setting.

1.1 Motivation

The motivation for this work sits at the intersection of robotics and machine learning. Modern manipulation research increasingly depends on learned models that can process visual observations, predict future motion, and adapt to complex physical interaction. At the same time, the pace of open-source machine learning development can make it difficult to separate robust scientific progress from promising but weakly benchmarked demonstrations. This concern is especially important for industrially relevant settings, where unreliable claims about physical reasoning or manipulation capability can lead to poor engineering decisions.

The work was also shaped by the RSO I activity and the University of Malta M.AI.ESTRO project, where the convergence of robotics, machine learning, and Industry 4.0 raised a practical need for clearer evaluation of emerging model families. In that setting, world models are particularly relevant because they promise a form of predictive checking that could eventually support safer or more reliable action selection. This dissertation does not claim to deliver such a deployment system. It instead treats the perceived need from that work as motivation for a narrower research question: whether world models measurably improve dexterous manipulation prediction under a controlled offline benchmark.

1.2 Research Question and Scope

The central research question is:

Do world models improve dexterous manipulation prediction when compared with, and integrated alongside, an ACT visuomotor policy under a controlled offline benchmark?

This wording is a rescoping of the original proposal question. The proposal asked whether integrating a latent world model into an imitation learning policy would improve physical grounding and robustness, measured through latent agreement error and task success under perturbations. The final implementation does not evaluate deployed task success or perturbation robustness, and latent mean-squared error is used only as a training diagnostic for the world-model transition head rather than as the main thesis result. The finished thesis therefore frames the question around offline prediction quality: whether world model support selects future hand-motion chunks that are closer to demonstrated motion under a shared benchmark.

The benchmark focuses on future bimanual hand motion and transform state. In implementation terms, the main prediction surface is built around future left-hand and right-hand transform chunks, but the detailed representation is defined later in the methodology and implementation chapters. The Introduction therefore uses the more readable language of future hand motion and transform state, while later chapters specify the exact data contract, branch definitions, and metrics.

The scope is deliberately bounded. The thesis does not test closed-loop robotic deployment, real-world task completion, force interaction, contact stability, or general machine common sense. It also does not claim that any one world model family is universally superior. Instead, it asks whether selected world-model families improve an offline prediction benchmark when compared under a shared evaluation surface. The selected policy architecture is ACT, and the selected world-model families are V-JEPA2-AC and LeWM. These choices make the question concrete enough to evaluate while keeping the framework open to future model families and additional evaluation surfaces.

This scope also explains the structure of the comparison. The benchmark is organised around three branches:

- **ACT-only branch:** the policy baseline, where ACT provides the selected future motion without world model support.
- **Hybrid branch:** ACT proposes candidate chunks, and the world model evaluates or re-scores those candidates before selection.
- **World model driven branch:** no policy preference is used for selection. The selected future motion is chosen from the world model’s own assessment of predicted future-state consistency.

The purpose of this structure is to separate the value of the world model from the value of the underlying policy as far as possible within an offline benchmark.

1.3 Contributions

This thesis contributes a comparative benchmark framework for studying world model support in dexterous manipulation prediction. The framework evaluates a fixed ACT policy setting against world model selected alternatives, allowing the effect of world model support to be examined without changing the core policy family. This provides a controlled basis for asking whether predictive future-state models improve selected hand-motion chunks under the same offline evaluation conditions.

The thesis also contributes an implemented comparison between two selected world-model families, V-JEPA2-AC and LeWM. These families are used as benchmark representatives rather than presented as newly invented models. Their inclusion allows the dissertation to compare not only different levels of world model authority, but also different world-model strategies inside the same overall evaluation logic.

Finally, the thesis contributes a bounded empirical analysis of world model support for future bimanual hand-motion prediction. The analysis is limited to the selected dataset, branch definitions, prediction targets, and metrics. This boundedness is intentional. It keeps the claims scientifically defensible while leaving the framework open to later world models, alternative policy families, and evaluation designs that move closer to deployment.

1.4 Dissertation Structure

—To be updated at the end ———

The remainder of the dissertation is organised as follows:

- Chapter 2 introduces the technical background required to follow the dissertation, covering transformers, self-supervised learning, variational autoencoders, rigid body pose, visuomotor policies, and world models at a foundational level.
- Chapter 3 positions the work within the relevant literature on imitation learning, visuomotor policies, world models, and offline evaluation.
- Chapter 4 defines the benchmark design, scope, branch structure, and comparison logic.
- Chapter 5 describes the implemented system and the main engineering decisions needed to produce the benchmark.
- Chapter 6 explains the evaluation surface, selected metrics, and reporting rules.

- Chapter 7 presents the benchmark results.
- Chapter 8 interprets the findings and states the main limitations.
- Chapter 9 closes the dissertation and identifies future work.

2 Background

This chapter assumes familiarity with the basic mechanics of a neural network and introduces six concepts central to this work: visuomotor policies, rigid-body pose in $SE(3)$, latent representations, conditional VAEs, transformer attention, and self-supervised learning. These foundations support understanding of the specific architectures in Chapter 3.

2.1 Visuomotor policies

A visuomotor policy is a learned function that takes a visual observation as input and produces an action or sequence of future actions as output. In a manipulation setting, the visual observation is typically a single camera frame showing the hand and the object being manipulated, and the action describes how the hand should move next. The policy is called visuomotor because it connects vision directly to motor output, without requiring a hand-crafted intermediate representation of what the scene contains [1].




Figure 2.1 A visuomotor policy maps visual observations to motor actions. The input is a camera frame showing the hand and object, and the output is a proposed action for how the hand should move next.

Learning this mapping from data is appealing because useful hand behaviour depends on subtle visual cues that are difficult to specify by hand. The timing of a reaching movement, the grip shape adopted before contact, and the postural adjustments made in response to object geometry all depend on information visible in the camera image but difficult to encode as explicit rules [3–5]. A visuomotor policy learns these associations directly from recorded demonstrations, capturing the relationship between what the camera sees and what the hand does without requiring the programmer to describe it explicitly.

2.1.1 Behaviour cloning and compounding error

The most direct way to train a visuomotor policy from demonstrations is Behaviour Cloning (BC) [1]. Here an observed frame predicts the action taken by the expert at

that moment, and the policy is trained to minimise the difference between its predicted action and the recorded expert action across all frames in the dataset [6]. The problem arises when a small prediction error causes the hand to move to a state that is slightly different from any state in the training data. Which then results in a compounding error as the policy encounters states it has not seen during training, leading to larger and larger errors over time.




Figure 2.2 Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [8].

2.1.2 Action chunking

Action chunking addresses the compounding error problem by changing what the policy predicts. Instead of predicting a single next action at each step, the policy predicts a short sequence of future actions in one forward pass. This sequence is called an action chunk [7, 9]. Hence why Figure 2.1 shows the action sequence as a stack of multiple action blocks. By committing to a chunk of X future steps, the policy makes one prediction decision every X frames rather than one per frame. With fewer decision points, there are fewer opportunities for errors to accumulate, and the hand stays closer to the training distribution for longer [10].

2.2 Rigid body pose and the $SE(3)$ representation

To predict how a hand will move, a system needs a precise way to describe where the hand is and which way it is pointing. These two quantities together are called the pose of a rigid body. For the human wrist, pose is fully described by its position in space and its orientation relative to a fixed reference frame [11].

Translation – Position is represented as a vector of three numbers, one for each spatial dimension: how far the wrist is to the right or left, how far forward or backward, and how far up or down [3]. This vector $\mathbf{t} = [t_x, t_y, t_z]^\top$ is straightforward to predict with no special constraints on the output.




Figure 2.3 Translation for each arm can be represented by a movement along each of the three spatial axis [12].

Rotation – Orientation requires more care. A common representation uses three angles known as Euler angles, one per axis. These are intuitive but problematic for learning. At certain configurations the axes align and the system loses a degree of freedom, a singularity called gimbal lock. A subtler problem is discontinuity: a wrist rotated 359 degrees and one rotated 1 degree are almost physically identical, yet their angle values are numerically far apart. A network minimising prediction error will produce large errors near these boundaries even when the true motion is smooth [13].

A rotation matrix R avoids both issues: small physical rotations produce small numerical changes in the matrix entries, and there are no singularities.

The homogeneous transform – Translation and rotation are combined into a single 4×4 matrix, the standard representation of a pose in the special Euclidean group $SE(3)$:

$$T = \begin{bmatrix} R & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix} \quad (2.1)$$




Figure 2.4 Rotation around the y-axis for each arm. The same circular movement can be imagined around the x and z axes; the total orientation of the wrist is the combined result of rotations around all three [12].

2.3 Latent representations

A latent representation is a learned internal description of an observation. An image starts as a grid of pixel values, where each pixel stores three or four channel intensities (typically red, green, and blue, with an optional transparency value). A model does not usually reason with these raw numbers directly. Instead, it transforms them through several layers into a new set of values that no longer correspond to individual pixel coordinates or colours. These new values form a compact encoding that captures patterns in the image such as shape, position, motion, and spatial structure. The collection of these encodings is often called the latent space. Figure 2.5 shows an example in which similar items cluster together in the learned representation.




Figure 2.5 Example latent-space plot showing how structurally similar items cluster together in the learned representation. Structural similarity here refers to items that share shape, geometry, and likely physical behaviour [14].

The purpose of this transformation is to keep task-relevant structure while discarding unnecessary visual detail. For a manipulation scene, a useful latent state might encode hand position, object pose, contact progression, or motion trend without preserving every texture or lighting variation. A latent state is therefore not directly interpretable in the way an image is, but it can still preserve the information needed to reason about what happens next.

This is why latent prediction is often a better fit than pixel-level prediction for high-dimensional video observations. Predicting the next frame at the pixel level requires a model to account for every visual detail of the scene, including lighting changes, texture variation, and background movement that carry no information about task-relevant dynamics. A latent predictor instead operates on compact embeddings and concentrates predictive capacity on the structure that determines what actually happens next [15]. This design trade-off motivates the world model architectures reviewed in Chapter 3.

2.4 Conditional variational autoencoders

A VAE maps an input to a distribution of latent vectors described by a mean μ and standard deviation σ , rather than a single point. Sampling different latent vectors lets the decoder produce different but plausible outputs rather than one fixed answer. Figure 2.6a shows this standard encoder-latent-decoder flow [16].




Figure 2.6 Standard VAE (a) and conditional CVAE (b) encoder-decoder pipelines, adapted from [17].

For the policy to be trainable end-to-end, gradients must flow from the action prediction loss all the way back to the encoder. Sampling z directly from a distribution blocks this path. The reparameterisation trick resolves this by writing $z = \mu + \sigma\epsilon$, where ϵ is standard normal noise drawn independently of the learned parameters. Figure 2.7 shows the full encoder-latent-decoder pipeline alongside its training loss: a reconstruction term that rewards accurate output and a regularisation term that keeps the latent distribution close to a standard Gaussian, together forming the evidence lower bound [16].




Figure 2.7 VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation), adapted from [18].

A CVAE adds an extra conditioning input to both encoder and decoder, shown in Figure 2.6b. In a visuomotor policy, this condition is the current observation frame, so the latent variable only needs to capture how the future action varies within that scene [19]. The ACT architecture uses this to sample multiple latent vectors from one observation and decode each into a candidate action chunk [7].

2.5 Transformer networks and the attention mechanism

A transformer processes a sequence as a set of tokens. A token can represent a word, an image patch, a timestep in an action sequence, or any vector produced by an earlier model component. The architectural idea needed here is simple: each token is updated by looking across the other tokens and deciding which of them are most relevant. This relevance-weighted exchange of information is called attention [20].

Attention forms three learned views of every token:

- **Query** (Q) – what this token is looking for.
- **Key** (K) – what another token contains.
- **Value** (V) – what another token contributes if attended to.




Figure 2.8 Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [21].

Figure 2.8 shows how each token (highlighted) sends its query outward to every other token, receives weighted values in return, and produces an updated output. The bidirectional arrows indicate that every token attends to every other token simultaneously, and the crossing arrow beneath represents the weighted sum that combines those values into the final output. For each token, the transformer compares its query with the keys of all tokens. The resulting scores become weights, and those weights are used to combine the corresponding values. This operation is usually written as [20]:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V. \quad (2.2)$$

Transformers usually repeat this operation in several heads and add positional information so the model knows where each token sits in the sequence. Multiple heads allow different relationships to be represented at the same time, while positional information preserves order [20].

2.6 Self-supervised learning

Machine learning methods can be grouped by how they receive feedback during training, as illustrated in Figure 2.9. Supervised learning trains on data with explicit labels: the correct answer is provided for every input, and the model learns to reproduce it. BC from Section 2.1.1 is a direct example: the expert demonstration acts as the label. Reinforcement learning instead receives feedback through a reward signal after taking actions in an environment; the model learns which state-action sequences lead to good outcomes. SSL sits between these two: it trains on unlabelled data by

having the data itself provide the supervision signal. Part of the input is hidden, and the model learns to predict it from the remainder [22].




Figure 2.9 Three machine learning paradigms: supervised, reinforcement, and self-supervised learning [23].

Applied to video, SSL takes the following form: given a sequence of observed frames, hide one or more future frames and ask the model to predict them from the frames that came before. No human annotation is required because the video recording itself provides the hidden target.

Predicting raw pixels runs into a problem described in Section 2.3: when the future is ambiguous the model cannot commit to a single outcome, so it predicts the average of all possibilities and produces a blurred image that does not match any real frame. Predicting in latent space avoids this. Instead of reconstructing every pixel, the model predicts the compact embedding of the future frame. Background motion and fine texture are discarded at the encoding step, so the predictor focuses on the structure that determines what actually happens next [15]. When the predictor is also given the action the agent intends to take, the result is a world model: a learned function that predicts how the scene embedding will change in response to a proposed action. The architectures that build on this idea are reviewed in Chapter 3.

3 Literature Review

This chapter surveys the literature motivating the three central design choices of WACT: that a specialist imitation policy is a strong base for dexterous manipulation, that a world model can usefully evaluate candidate action chunks in an offline benchmark, and that latent predictive learning is the right approach for building such an evaluator. Sections 3.1 and 3.2 review the two candidate imitation policies, ACT and Diffusion Policy, examining both the challenges they address and the trade-offs between them. Section 3.3 introduces world models as learned environment predictors, surveys the evaluator architectures most relevant to this role, and covers the three latent predictive architectures implemented in the benchmark, together with a negative control. Section 3.4 covers offline evaluation frameworks and their validity as proxies for real-world performance. The chapter closes with a summary of the literature gap that the benchmark addresses. The chapter assumes familiarity with the technical foundations in Chapter 2; world models are introduced here.

3.1 Action Chunking with Transformers (ACT)

ACT was introduced by Zhao et al. to address the compounding error and multimodal distribution challenges that make dexterous bimanual manipulation demanding for imitation learning [7]. On the ALOHA bimanual platform it achieved success rates of up to 90% on precision tasks such as threading a cable tie, inserting a battery, and opening a translucent bag, outperforming behaviour cloning baselines and GATO by substantial margins [7]. In SE(3) wrist-pose prediction specifically, small deviations in predicted position or orientation alter contact geometry in ways that cascade into grasp failure [1], and a policy trained with a naïve regression loss produces averaged predictions that may not represent any plausible trajectory when the demonstrated distribution is multimodal [9]. ACT’s CVAE architecture addresses both of these failure modes directly.

ACT addresses both challenges through the use of a CVAE design (Section 2.4). The respective CVAE has an encoder-decoder structure depicted in Figure 3.1, where the encoder compresses an action sequence and joint observation into a style variable z that captures the multimodal distribution of plausible futures, and the decoder synthesizes images from multiple viewpoints, joint positions, and z with a transformer encoder, and predicts a sequence of actions with a transformer decoder. The encoder is discarded at test time, and z is simply set to the mean of the prior (i.e. zero) at test time [7].

At inference time, the CVAE decoder samples several latent vectors




Figure 3.1 Architecture of the Action Chunking with Transformers (ACT) policy [7].

independently and decodes each into a candidate action chunk. A temporal ensembling technique is then applied: consecutive chunks overlap in their predicted time horizons, and a recency-weighted average is taken across overlapping predictions to produce the final action. This averaging smooths the executed motion and reduces jerkiness caused by abrupt chunk transitions. The result is a structured pool of plausible futures, all conditioned on the same observation, from which the ensembled output is drawn [7].

ACT carries known limitations alongside these strengths. The Gaussian prior on the style variable z constrains how faithfully the CVAE can represent highly multimodal action distributions: at transitions where qualitatively different trajectories are equally plausible, the latent space can average across modes rather than sample cleanly from each one, producing trajectories that are plausible in aggregate but imprecise at the most ambiguous decision points [9]. Chunk length is a fixed hyperparameter that must be tuned per task; longer chunks improve smoothness but commit the policy to a trajectory over a horizon that may no longer be appropriate as the scene evolves. Temporal ensembling introduces a recency-weighted lag that can slow the policy’s response to sudden environmental changes [7]. These properties are well-documented in follow-on work comparing ACT with more expressive policy classes, and they motivate the inclusion of Diffusion Policy as the second candidate in the benchmark.

3.2 Diffusion Policy

Diffusion Policy offers a principled alternative approach to the same compounding error and multimodality challenges. Its core concept relies on starting from a noise distribution and iteratively refining action predictions through a diffusion process. By gradually denoising the predictions, Diffusion Policy captures complex multimodal distributions effectively, producing plausible trajectories that align with the demonstrated data [9]. Figure 3.2 illustrates the policy overview:

- a) General formulation: at timestep t , the policy takes the latest T steps of observation data O_t as input and outputs T steps of actions A_t .

- b) CNN-based Diffusion Policy: FiLM (Feature-wise Linear Modulation) [24] conditioning of the observation feature O_t is applied to every convolution layer, channel-wise. Starting from K_t drawn from Gaussian noise, the output of noise-prediction network ϵ_θ is subtracted, repeating K times to get A_t^0 , the denoised action sequence.
- c) Transformer-based Diffusion Policy: the embedding of observation O_t is passed into a multi-head cross-attention layer of each transformer decoder block. Each action embedding is constrained to only attend to itself and previous action embeddings (causal attention) using the attention mask illustrated [25, 26].




Figure 3.2 Diffusion Policy: general formulation (a), CNN variant (b), transformer variant (c) [9].

Chi et al. report that Diffusion Policy outperforms ACT on tasks with high action multimodality (such as multi-step object rearrangement and cup placement under visual ambiguity) while matching it on simpler reaching tasks [9]. The advantage arises because iterative denoising explores the full action distribution at each step rather than committing to a single latent mode as the CVAE prior must, making it more robust to the mode-averaging failure that affects ACT at ambiguous transitions.

The main limitation of Diffusion Policy is inference cost. Generating a single action chunk requires multiple sequential network evaluations: standard DDPM variants use on the order of 100 denoising steps compared with the single forward pass of ACT's decoder, which substantially reduces the achievable control frequency [9]. Efficient transformer variants address this through progressive noise scheduling and parameter sharing, recovering much of the inference speed at added architectural complexity [26]. Both policies are included in the benchmark to span this representational trade-off: ACT's structured latent sampling is computationally efficient but expressively constrained at high multimodality, while Diffusion Policy's iterative denoising is more expressive but more demanding at inference. Measuring the world model's effect across both allows the benchmark to determine whether evaluator gains are consistent across this axis.

3.3 World models

A world model is a learned predictor of how a system evolves over time. Given the current state and a proposed action, it estimates the state that is likely to follow. In a manipulation setting, the state may be a raw image, a latent representation of the scene (Section 2.3), or a structured description of hand and object configuration. Regardless of the representation chosen, the underlying principle is the same: the model internalises the dynamics of the environment and uses them to project forward from any given state [27].

World models become especially useful when a pool of candidate actions must be evaluated before any one is committed to. Each candidate is rolled forward through the model, producing a predicted future state; those predictions are ranked by predicted cost or consistency. The quality of this ranking depends heavily on the chosen representation: pixel-space world models must reconstruct every visual detail and suffer from mode-averaging at multimodal transitions, while models that operate in the compact learned latent space introduced in Section 2.3 can concentrate predictive capacity on the dynamics that matter for task outcomes [15]. The following subsections survey the literature supporting this evaluator design and describe the three world model architectures implemented in the benchmark.

3.3.1 World models as policy evaluators

WorldGym demonstrates that a learned world model used as a simulated environment for policy evaluation produces rankings that align with real-world task success with meaningful fidelity [28]. Candidate policies are evaluated not by deploying them physically but by rolling them out inside the model’s predicted environment, reducing the cost of policy comparison to a forward pass. CLOVER extends this to closed-loop control, as shown in Figure 3.3. A visual planner generates a sequence of sub-goal frames; the executor measures the embedding-space error between the predicted sub-goal and the current observation. When the error exceeds a threshold, the system triggers an adaptive transition and replans rather than committing to the original trajectory. The feedback loop, absent from open-loop imitation policies, is the mechanism that translates world model prediction accuracy into task completion gain [29]. Together, WorldGym and CLOVER establish that world model rankings are informative both before and during execution.

The broader question of whether such evaluator rankings transfer to physical outcomes, and the qualifications around out-of-distribution reliability, are examined in Section 3.4.

The practical value of this evaluator loop is measurable. Table 3.1 reproduces




Figure 3.3 Classic closed-loop control (a) versus CLOVER’s closed-loop visuomotor control (b) [29].

CLOVER’s real-world evaluation on a multi-step robot manipulation suite, reporting the average number of consecutive sub-tasks completed (`avg_len`) for representative baselines and CLOVER [29].

Table 3.1 Average consecutive sub-task completions (`avg_len`) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [29].

Method	<code>avg_len</code>
ACT [7]	0.6
R3M (pretrained visual repr.) [30]	0.7
RT-1 [31]	1.1
CLOVER (closed-loop)	2.1

Alongside the task completion gain, CLOVER reduces embedding-space prediction RMSE relative to open-loop baselines by correcting the policy’s course when the world model detects a divergence from the expected trajectory, demonstrating that prediction accuracy and task success move together under the evaluator framing.

3.3.2 Latent predictive architectures

The energy-based framework of LeCun and Dawid establishes why prediction should operate in latent rather than pixel space [15]: the same representational argument introduced in Section 2.3 applies directly to the prediction target, favouring compact embeddings over pixel reconstructions that force the model to account for irrelevant visual detail and average over all plausible continuations at ambiguous transitions.

Image Joint Embedding Predictive Architecture (I-JEPA) establishes the joint-embedding predictive architecture as a viable approach to learning rich visual




Figure 3.4 Three SSL prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [22].

representations without any pixel reconstruction objective [22]. Panel (c) of Figure 3.4 shows the data flow directly: a context block x from the image is passed through the context encoder to produce latent tokens z_x . A lightweight ViT predictor then takes z_x together with positional tokens indicating where the masked target regions y are located, and outputs predicted latent representations $\hat{s}_y$ for those regions. A separate target encoder, whose weights are a slow exponential moving average of the context encoder, independently encodes the actual target pixels into ground-truth latents s_y . The loss $\text{MSE}(\hat{s}_y, s_y)$ is computed entirely in latent space: no decoder, no pixel generation, no reconstruction of lighting or texture. The outcome is a context encoder whose representations capture semantically meaningful structure (object shape, spatial layout, and scene dynamics) because that is all the predictor ever needs to get the target embeddings right. Training is approximately 2.5 times faster than pixel-reconstruction methods as a result. V-JEPA2 extends this same joint-embedding predictive architecture from images to video sequences.

Four architectures are covered in the following subsections, each representing a distinct representational strategy. V-JEPA2 and V-JEPA2-AC share the same pretrained video backbone: V-JEPA2 is the internet-scale pretrained encoder, while V-JEPA2-AC extends it with an action-conditioned transition head and is the form used directly in the benchmark. DexWM uses a pretrained image backbone with specialised hand-geometry supervision. LeWorldModel (LeWM) learns all representations from scratch on the task data itself.

V-JEPA2

V-JEPA2 extends the I-JEPA framework from static images to video, pretraining on over one million hours of internet video [32]. The pretraining shifts from spatial to spatio-temporal masking: the context encoder receives past observation frames, and the predictor must recover future latent states. The training objective combines teacher-forcing with multi-step autoregressive rollout losses, encouraging accurate short-horizon predictions and consistent long-horizon trajectory representations. The result is a frozen ViT-Large encoder, approximately 300 million parameters, that encodes physical priors about hand and object interactions acquired from




Figure 3.5 V-JEPA2 post-training pathways from internet-scale backbone to specialised applications [32].

internet-scale pretraining. As shown in Figure 3.5, the pretrained backbone supports three post-training pathways: Language Alignment, which grounds visual embeddings in natural language; Attentive Probe, which adds a lightweight classification head; and Action Conditioning, which extends the predictor to take robot actions as input and forms the basis of V-JEPA2-AC.

V-JEPA2-AC

The frozen ViT-Large encoder from V-JEPA2 serves as the backbone of V-JEPA2-AC, producing a fixed latent embedding for each observation frame with no gradient updates to the backbone weights. A lightweight action-conditioned transition head is trained on top of these stored embeddings: it takes the current frame latent together with a proposed action chunk (robot joint states and end-effector poses) as the conditioning signal z , and predicts the latent representation of the future frame. Candidate action chunks are then scored by comparing the predicted future embedding against the actual observed future embedding under an L1 objective. This design, shown on the right of Figure 3.6, separates what the network must learn into two parts: the frozen encoder supplies rich visual features that generalise from internet-scale pretraining, and the small trained head learns only how actions propagate through those features.

An important qualification is raised by Tsagkas et al., who show that pretrained visual representations can encode spatial plausibility well while providing a weaker signal for fine-grained temporal dynamics [33]. A frozen encoder trained on internet video may represent hand shape and position accurately at each individual frame while providing a less reliable signal about whether the predicted trajectory is kinematically




Figure 3.6 Standard V-JEPA2 (left) versus action-conditioned V-JEPA2-AC (right) [32].

consistent across the full chunk horizon. This temporal limitation provides context for the path shape results in Chapter 8.

DexWM

DexWM [34] is a dexterous manipulation world model from Meta FAIR. As shown in Figure 3.7, its pipeline encodes each input frame with a frozen DINOv2 image encoder, producing a grid of spatial patch tokens rather than a single global embedding. These patch-level representations preserve fine-grained spatial information about hand positions and object locations, which gives the predictor strong generalisation across diverse environments and viewpoints. A Conditional Diffusion Transformer (CDiT) predictor then takes the spatial latent tokens together with action and camera-motion inputs to predict the next latent state, and a lightweight decoder maps the result back to image and keypoint space. Where V-JEPA2-AC uses a video-pretrained backbone, DexWM uses a backbone pretrained on static images; the comparison therefore isolates the contribution of video-specific temporal pretraining to world model quality.

A distinguishing feature of DexWM is the Hand Consistency (HC) loss: an auxiliary transformer head predicts fingertip and wrist heatmaps from the latent state, imposing a geometric prior that the predictor’s latent trajectory must remain consistent with plausible hand configurations. This explicit hand-geometry supervision is absent from both V-JEPA2-AC and LeWM, whose latent spaces are shaped only by prediction error or general regularisation. Goswami et al. report a 50% improvement over Diffusion Policy on grasping and placing tasks, with an 83% zero-shot real-world grasping success rate under latent planning [34].




Figure 3.7 DexWM world model pipeline [34].

LeWorldModel

LeWM is a latent world model trained end-to-end from random initialisation on raw pixel observations [35]. Its architecture consists of two components: an encoder that maps each frame o_t to a compact latent vector z_t , and a predictor that models environment dynamics by estimating the next latent embedding $\hat{z}_{t+1}$ from z_t and the action a_t . The encoder is a ViT-Tiny with approximately 15 million parameters. Each frame is represented as a single 192-dimensional token rather than a spatial patch grid, which keeps the latent space compact and makes planning significantly faster than models that produce higher-dimensional patch representations.




Figure 3.8 LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [35].

As shown in Figure 3.8, the training pipeline consists of an encoder that maps each observation frame to a compact latent vector, a dynamics predictor that autoregressively estimates the next latent given the current embedding and action, and

a SIGReg regulariser that enforces a Gaussian prior on the latent distribution by optimising normality along random univariate projections. The training objective reduces to two terms:

$$\mathcal{L}_{\text{LeWM}} = \mathcal{L}_{\text{pred}} + \lambda \text{SIGReg}(Z) \quad (3.1)$$

where $\mathcal{L}_{\text{pred}}$ penalises inaccurate next-embedding forecasts and $\text{SIGReg}(Z)$ stabilises end-to-end training without contrastive negative pairs or an exponential moving average target encoder [35]. This reduces the number of tunable loss hyperparameters from six, as required by prior Joint Embedding Predictive Architecture (JEPA)-style designs, to one.

DINO-WM, the closest comparable baseline, follows a similar latent prediction objective but uses a frozen DINOv2 backbone producing a grid of spatial patch tokens per frame [36]. Like DexWM, DINO-WM therefore carries the spatial resolution and generalisation benefits of DINOv2’s patch-level representations, but at substantially higher computational cost. LeWM replaces the frozen backbone with a fully learned ViT-Tiny encoder that collapses each frame to a single 192-dimensional token, trading spatial resolution for training speed. In practice, LeWM can be trained on a single GPU in a few hours and achieves a $48\times$ planning speedup over DINO-WM [35]. DexWM’s CDiT predictor operating on full patch grids is significantly slower to train and requires substantially more GPU memory; this difference in training cost is expected to widen further at larger dataset scales and is quantified in the implementation chapter. Despite its compact size, LeWM’s latent space encodes physically meaningful structure: the model detects physically implausible state transitions and captures spatial information such as agent and object locations across frames.

The three-way comparison across V-JEPA2-AC, DexWM, and LeWM is therefore not simply an architecture comparison. It is a comparison between fundamentally different representational strategies: video-pretrained transfer (V-JEPA2-AC), image-pretrained transfer with geometric supervision (DexWM), and self-supervised task-data learning from scratch (LeWM). This distinction makes the comparison principled and is the basis for the experimental design in Chapter 4.

PointWorld

PointWorld is a large pretrained 3D world model for in-the-wild rigid arm manipulation [37]. Unlike the latent-space models described above, it operates directly in 3D geometry: scene state is lifted from RGB-D images into a full-scene point cloud, and robot actions are represented as the 3D point flows traced by robot surface geometry through forward kinematics. This embodiment-agnostic representation allows the model to reason over physical interaction geometry rather than learned

abstract features.

Such a model is not designed to handle a 16-dimensional dexterous hand action space, where fine-grained finger articulation drives the manipulation outcome. The forward-kinematics representation assumes rigid arm geometry, making the model structurally mismatched when applied to dexterous bimanual hand motion.

3.4 Offline evaluation

A central concern for any offline benchmark is whether its metrics produce rankings that are meaningful proxies for real-world performance. Without this property, an offline evaluation framework reduces to a measure of in-distribution accuracy with no predictive validity for deployment.

At large scale, DreamerV3 demonstrates the effectiveness of latent-space candidate evaluation across diverse continuous control domains, including locomotion and manipulation [27]. DreamerV3 is an online reinforcement learning system, not an offline benchmark; it is included here because it provides large-scale empirical evidence that latent-space evaluation is an effective selection mechanism when properly calibrated, which motivates its use in the offline setting of WACT. Its world model maintains a compact recurrent latent state that transitions under action conditioning; candidate trajectories are rolled out entirely inside this imagined space and scored by a learned value function, with no return to the real environment between evaluations. TD-MPC2 confirms the principle with a scalable planner that optimises a temporal-difference value function defined over predicted latent trajectories rather than observed states [38]. Both systems show that ranking candidate actions by predicted latent-space outcomes is competitive with direct online evaluation across a range of motor control benchmarks.

A critical qualification is that latent-space scoring becomes unreliable outside the training distribution. MOPO, MOREL, MBOP, and MOPP each establish that learned models produce inaccurate predictions for out-of-distribution states, requiring conservative candidate selection to avoid trajectory degradation [39–42]. None of these works addresses dexterous manipulation specifically; they benchmark locomotion and simple control tasks, leaving open how well the principle transfers to the fine-grained wrist-pose prediction setting.

Two further bodies of work establish that offline rankings can transfer to physical outcomes when the evaluation domain is well-calibrated.

SimplerEnv establishes the correspondence between simulation-based evaluation and physical robot performance for visuomotor policies [43]. The study evaluates multiple manipulation tasks and policy types in simulation and then

physically, using a WidowX robot under matched visual and control conditions. Rankings produced from simulation-only evaluation closely track rankings from physical trials, with correlations holding across both policy families and task types. The key condition for transfer is domain calibration: visual appearance, camera viewpoint, and control interface must match between the simulation and the physical system. When these are well-matched, simulation rankings provide a reliable ordering of policies without requiring physical deployment.

WorldEval extends the validation to the world-model setting [44]. Rather than evaluating policies in a manually constructed simulator, WorldEval scores them using a learned world model that predicts future states from offline trajectory data. The world-model-based rankings correlate strongly with real-world task success: policies that score well in world-model evaluation also tend to succeed when deployed physically.

Both results carry a shared qualification: rankings hold under well-calibrated conditions, where the evaluation domain is consistent with the training distribution and the deployment context. When this alignment breaks down, offline rankings can degrade. The design implications of this constraint are discussed in Chapter 4.

Summary. The literature surveyed in this chapter establishes three supporting findings. First, world model rankings are informative evaluators both before and during execution, as demonstrated by WorldGym and CLOVER [28, 29]. Second, latent predictive objectives produce more efficient and better-calibrated representations than pixel-level reconstruction, as demonstrated by the JEPA family [22, 32]. Third, offline evaluation rankings can carry predictive validity for physical outcomes when the evaluation domain is well-calibrated, as demonstrated by SimplerEnv and WorldEval [43, 44].

The gap these findings leave open is twofold. First, there is no controlled offline benchmark that isolates the contribution of a world model evaluator to the trajectory quality of a fixed dexterous manipulation policy. The offline reinforcement learning literature establishes the need for principled candidate selection but benchmarks only locomotion and simple control tasks [39–42]. Second, the interaction between world model representational strategy and policy performance has not been examined: it is not known whether a pretrained backbone world model, a task-data-trained world model, and an out-of-domain negative control produce meaningfully different guidance when applied to the same policy on the same data. Both gaps require a benchmark that holds the policy and evaluation surface fixed while varying only the world model. The methodology for constructing and running such a benchmark is described in Chapter 4.

4 Methodology

4.1 Dataset and Experimental Setup

4.1.1 EgoDex: source and selection rationale

The benchmark is built on EgoDex [45], a large-scale egocentric dataset of bimanual dexterous manipulation collected using Meta Quest Pro head-mounted cameras. It spans 63 task categories including folding, slotting, pouring, and kneading, covering a wide range of fine-grained hand motions across more than 194,000 demonstration episodes and 1.3 TB of paired video and annotation data. The full release comprises five training parts and supplementary additional-data splits; this study uses the first three training parts. Each episode provides synchronised RGB video and full six-degree-of-freedom $SE(3)$ wrist annotations for both hands at approximately 30 frames per second.

The choice of dexterous manipulation as the evaluation domain is not incidental. This dissertation forms part of the University of Malta’s ongoing M.AI.ESTRO research project, which investigates the use of AI for scene and task prediction on dexterous manipulation tasks to support industrial transfer learning and human skill digitisation. The research domain was therefore set by this institutional context: the methodology operates on human hand demonstrations rather than robotic end-effector data because the target application is the transfer of tacit human dexterous skill. An evaluation framework built around robot arm trajectories would not address that objective. The introduction situates this motivation more broadly; the present chapter focuses on the methodological consequences of the domain choice.

Within the dexterous manipulation setting, EgoDex was selected on three properties. First, the benchmark requires a dexterous bimanual hand action space: fine-grained articulation across a sixteen-dimensional $SE(3)$ wrist pose representation drives the manipulation outcome, and large-scale corpora built around rigid gripper or arm contact do not provide this. Second, the egocentric viewpoint aligns with the intended downstream deployment setting of a wearable assistive device. Third, the dataset’s scale and task diversity are sufficient to support simultaneous training of five world model families and two policy architectures on declared non-overlapping splits while retaining a held-out test set of meaningful size.

EgoDex carries two acknowledged limitations. The demonstrations are produced by human subjects rather than a robotic system, which introduces a kinematic domain gap between the training distribution and any robotic deployment. No contact or force signals are present; the action representation is purely kinematic

SE(3). Both are accepted as known constraints rather than experimental variables and are discussed further in Section 4.7.

4.1.2 Windowing protocol

The dataset is stored as paired `.mp4` and `.hdf5` episode files. A deterministic index stage discovers all valid paired episodes, infers their lengths from HDF5 metadata, and writes canonical observation and prediction windows for each split. This shared index is the single source of truth for all downstream modules: policy training, world model extraction, and evaluation all consume the same window file for their respective split, ensuring that output differences across the experimental branches can be attributed to branch design rather than to differences in the data presented to each component. Two windowing strategies are used at different stages of the project.

Strategy 1: Fixed-stride sampling. Each window spans 16 observation frames followed by 16 prediction frames, with consecutive window starts separated by a stride of 128 frames. At 30 frames per second this places starts approximately 4.3 seconds apart, corresponding to an observation context and prediction horizon of approximately 0.53 seconds each. This contract is encoded as `obs16_pred16_s128` and forms the foundation for all policy training and world model training (see Figure 4.1).



Figure 4.1 Fixed-stride windowing scheme (`obs16_pred16_s128`) used as the canonical training and evaluation index. Author-generated schematic.

Strategy 2: State-change-targeted sampling. The fixed-stride scheme samples the episode timeline uniformly, which causes it to under-sample high-motion events: long near-static segments contribute windows at the same rate as active manipulation periods. As detailed in Chapter 5, this limitation motivated a revised strategy adopted for the placement perception study and qualitative development inspection. Rather than placing windows at fixed intervals, this approach computes smoothed hand-state change scores from the wrist transform and confidence sequences, selects candidate windows near high-change peaks, and applies episode-first task-stratified sampling with a fixed random seed (see Figure 4.2). A relative-position fallback retains slower episodes that would otherwise be excluded, yielding one frozen window per episode stratified across all represented task categories.

The canonical training and evaluation index uses the fixed-stride scheme across all experimental branches. The state-change-targeted subset was used exclusively for




Figure 4.2 State-change-targeted sampling pipeline yielding one frozen window per episode. Author-generated schematic.

the placement perception study and qualitative inspection during development; it does not form part of the primary benchmark evaluation.

4.1.3 Training splits and test set

Three cumulative training splits were assembled from the first three parts of the EgoDex release, each a strict superset of the previous; Table 4.1 summarises their composition. EgoDex ships five training parts plus supplementary additional-data splits; parts four and five and all additional-data splits are excluded from this study to bound the experimental scope. The held-out test split is the official EgoDex test set and is never used during training at any stage by any model or policy.

Table 4.1 EgoDex splits under the `obs16_pred16_s128` windowing contract. Episode and window counts are derived from the canonical index.

Split	Tasks	Episodes	Windows	Role
<code>train_part1</code>	26	46,091	140,239	Stage 1 training
<code>train_part1_2</code>	39	140,654	291,369	Stage 1+2 training
<code>train_part1_2_3</code>	63	194,333	430,962	Canonical training
<code>test</code>	111	3,229	7,607	Frozen evaluation

The episode counts reflect the uneven collection density across parts: the 13 task categories added in part two average approximately 7,300 episodes each, roughly four times the average of the 26 part-one tasks, because part two focuses on high-frequency foundational primitives such as pick-and-place and folding that were collected at greater volume.

The canonical results in Chapter 7 use `train_part1_2_3` across all trained

components. The test split is the official EgoDex held-out set spanning 111 task categories, which includes all 63 categories present in `train_part1_2_3` as well as 48 categories drawn from the excluded parts; evaluating on this broader set provides a measure of out-of-distribution generalisation. Its window index is frozen; no post-hoc filtering or re-sampling is applied. The incremental training curriculum across the three stages is described in Chapter ??.

4.2 Three-Branch Experimental Framework

The central question of this study is whether a latent world model improves candidate selection quality when paired with an imitation-learning policy, and whether that effect generalises across policy families. Answering both questions cleanly requires holding every other experimental factor constant. The benchmark therefore organises evaluation around three branches that share the same policy weights, the same evaluation windows, and the same candidate pool across B2 and B3; only the selection rule applied to the candidate pool differs (Figure 4.3). Any measured performance difference between branches is directly attributable to the selection mechanism rather than to training data, architecture, or window coverage.

Two policy families are evaluated within this framework: ACT, whose CVAE-Transformer decoder generates five candidate action chunks in a single forward pass, and Diffusion Policy (DP), which generates five candidates via five independent Denoising Diffusion Implicit Models (DDIM) denoising chains on the same observation. Both families are given equal standing; the candidate pool contract is identical for each, and all three branches apply the same selection rules to both.

Branch 1 (B1): Policy-only baseline. The policy generates its full candidate pool and selects the candidate it scores highest by its own internal preference. No world model is queried at runtime. B1 establishes the trajectory prediction quality the policy achieves acting on its own judgement alone and serves as the baseline against which both world-model branches are measured.

Branch 2 (B2): Hybrid with world-model veto. The policy generates the same five candidate action chunks. The world model scores all five by predicting forward in its latent space from the current observation and measuring the fidelity of each prediction. If the world model’s preferred candidate differs from the policy’s preferred candidate and its score exceeds the policy’s choice by a declared margin threshold, the world model overrides the selection; otherwise the policy’s preferred candidate stands. This conservative design confines world-model authority to windows where it has sufficient relative confidence. The mechanism is analogous in spirit to model-based offline reinforcement learning approaches such as MOREL [40] and MOPO [39], which

use a learned model’s uncertainty to constrain policy decisions to regions where the model is reliable.

Branch 3 (B3): World-model-driven selection. The margin threshold is removed entirely. The world model selects the candidate with the highest score unconditionally. Critically, B3 does not re-run the policy: it operates over the frozen candidate pool produced during the B2 pass for each window. The policy contributes its proposals at B2 time; B3 reuses them without generating any new candidates. This design ensures that comparing B2 against B3 isolates the effect of removing the threshold gate on identical inputs, with no other variable changing.

The isolation property holds at every level of the evaluation. All 7,607 test windows are fixed and shared across all branches and all five world model families. The five-candidate pool per window is generated once and reused across B2 and B3. World model weights are fixed at the end of training and not updated during evaluation. The only degree of freedom is the selection rule applied to the pool. Differences in measured trajectory prediction quality between branches therefore reflect the causal effect of world model involvement in selection.




Figure 4.3 Three-branch design space. B1 uses the policy’s own preference alone; B2 adds a world-model veto constrained by a margin threshold; B3 gives the world model full selection authority over the frozen candidate pool produced at B2 time. The policy is not re-run in B3.

4.3 Policy Variants

4.4 World Model Families

4.4.1 V-JEPA2-AC – Pretrained Backbone Baseline

4.4.2 LeWM – Lightweight From-Scratch

4.4.3 DexWM – Domain-Specific Architecture

4.4.4 V-JEPA2 (No Action Conditioning) – Ablation Control

4.4.5 PointWorld – Out-of-Domain Negative Control

4.5 Fair Experimental Protocol

4.6 Evaluation Metrics and Score Harmonisation

4.7 Evaluation Scope and Limitations

5 Implementation

The WACT implementation is organised as an executable research stack with a clear separation between code, data, and generated artefacts. The main source tree lives under `src/wact`, while subsystem entrypoints are provided under `scripts`. The stable shell surface consists of `policy.sh`, `jepa_wm.sh`, `lewm.sh`, `docs.sh`, and the higher-level orchestration wrapper `run_stage.sh`. Docker is the canonical runtime surface for all stages: datasets, caches, and run outputs are kept outside the repository checkout so that code, data, and generated artefacts remain separated. The most important implementation consequence for the thesis is that all reported results are outputs of this staged execution model, tied to frozen run artefacts rather than to manually reconstructed notebook states.

[FIGURE PLACEHOLDER: WACT repository structure overview — `src/wact` code tree, `scripts` shell surface, runs output directory tree, datasets mount point; arrows showing data flow between stages]

5.1 Data and windowing pipeline

The index stage is the foundation of the entire stack. It discovers valid paired EgoDex episodes (requiring both an `.mp4` video file and an `.hdf5` pose annotation file), infers episode lengths from the HDF5 metadata, and writes canonical observation and prediction windows. For an episode of length T , a window at position t_0 is valid when both the observation slice $[t_0, t_0 + \text{obs_len})$ and the prediction slice $[t_0 + \text{obs_len}, t_0 + \text{obs_len} + \text{pred_len})$ fall within the episode boundaries. The canonical contract used throughout this dissertation is `obs_len=16`, `pred_len=16`, `stride=128`, producing windows of 16 observed frames followed by 16 frames to predict, sampled every 128 frames along each episode.

[FIGURE PLACEHOLDER: window construction diagram — episode of length T , obs slice, pred slice, stride spacing, valid window boundaries, how overlapping windows are avoided by stride]

This shared window definition is architecturally important. ACT policy training, V-JEPA2-AC latent extraction, LeWM episode export, semantic alignment, and the final benchmark evaluation all consume the same window file. Comparisons between branches are therefore built on a common row space, not on branch-specific sample selections that could introduce confounds.

The index stage writes a resumable checkpoint under a hidden directory during construction, allowing an interrupted build to continue from the last completed episode rather than starting from scratch. For the largest split (`train_part1_2_3`,

approximately 1.3 TB and 195,000 episodes), this resumability was practically necessary because the indexing job ran over several hours.

5.2 ACT policy path

The ACT policy consumes indexed windows together with the RGB frame and bimanual wrist state at the observation boundary, and predicts five candidate action chunks for the left and right hands. Each chunk is a sequence of 16 SE(3) wrist transforms, one per prediction step, for each hand. The five candidates are generated by sampling the CVAE latent variable five times from the prior at evaluation time, giving a small pool of plausible trajectories from which the downstream guard stage can select.

The policy architecture uses a ResNet-18 visual backbone with approximately 51.6 million total parameters. The ResNet-18 processes the single observation RGB frame and produces a feature vector that is combined with the bimanual wrist state before being passed to the Transformer decoder. This is substantially heavier than either of the world model guard components at evaluation time. The V-JEPA2-AC transition head is a small multi-layer perceptron trained on stored embeddings and executes in milliseconds. The LeWM guard evaluates a 15-million-parameter ViT-Tiny end-to-end. ACT’s CNN and Transformer architecture therefore dominates the per-window compute budget at guard evaluation time.

[FIGURE PLACEHOLDER: ACT inference path — RGB frame through ResNet-18, bimanual state concatenated, CVAE prior sampled five times, five candidate chunks produced; contrast with V-JEPA2-AC guard’s lightweight MLP and LeWM guard’s ViT-Tiny rollout]

ACT training used the `train_part1_2_3` split for the canonical results. Training ran for 30 epochs over approximately 195,000 demonstration episodes and took approximately 18.5 GPU-hours on two A100 GPUs, using single-node distributed data-parallel (DDP) training launched via `torchrun`. The training loop writes a `best.pt` checkpoint (lowest validation loss), a `last.pt` checkpoint, per-epoch metrics in `metrics.jsonl`, a `normalization_stats.json` file that records the training-split statistics used to normalise wrist poses, and a frozen `manifest.json` that records the configuration and window file used. The normalisation statistics are critical for downstream stages: V-JEPA2-AC transition head training and the guard evaluation both normalise action representations using the same statistics to ensure consistency between training and evaluation.

ACT is heavier and slower to train than either world model component because it requires a full forward and backward pass through ResNet-18 and the Transformer decoder for every training window. The world model components operate either on

stored embeddings (V-JEPA2-AC head) or on compact HDF5 sequences (LeWM), which are cheaper to process per step.

5.3 V-JEPA2-AC path and caching strategy

The V-JEPA2-AC family has two distinct implemented phases. Their separation is a deliberate caching strategy that is central to the practical feasibility of the approach.

In Phase 1, indexed windows are processed by the frozen V-JEPA2 ViT-Large encoder in a single batched inference pass. The encoder takes 16-frame observation sequences and produces a latent embedding vector for each window. These embeddings are written to `embeddings.npy` on disk, one row per window, and a companion `window_ids.txt` records the exact identifier (split, task, episode, t_0) of each embedding. Phase 1 is not a training job: there is no backward pass, no optimiser, and no epoch loop. It is an offline feature extraction that runs once per split and stores the results permanently.

The motivation for this caching strategy is straightforward. V-JEPA2 is a 300-million-parameter ViT-Large that processes 16-frame sequences. Running it over every training window at each training epoch would be computationally prohibitive, and it is unnecessary because the backbone is frozen. Running it once and caching the result means the expensive backbone inference is amortised across all subsequent training and evaluation. The V-JEPA2-AC head then trains on the stored embeddings rather than on raw frames, which reduces AC head training from what would be many GPU-hours to approximately minutes.

For the `train_part1_2_3` split, Phase 1 extracted embeddings for approximately 430,962 training windows in a single pass requiring approximately 3.2 GPU-hours on two A100 GPUs. The extraction stage supports sharded multi-GPU execution: the window list is partitioned across GPUs and results are merged after all shards complete, enabling linear scaling of extraction time with GPU count.

[FIGURE PLACEHOLDER: V-JEPA2-AC caching pipeline – single extraction pass over all windows stores `embeddings.npy`; AC head trains on stored embeddings in minutes; guard stage loads frozen encoder once and AC head at runtime; contrast with a naive approach that would rerun the encoder per training step]

In Phase 2, the action-conditioned (AC) transition head is trained on the stored embeddings. The head is a small multi-layer perceptron that learns to predict the next embedding $z(t+1)$ from the current embedding $z(t)$ and an action conditioning vector derived from the bimanual wrist state delta. The `action_repr=policy_state_delta` conditioning aligns the action representation used in transition head training with the same bimanual state representation used by the ACT policy, ensuring that the world

model can score ACT candidate chunks consistently. The trained head is exported through an `act_adapter.json` contract that the guard stage loads at runtime.

At guard evaluation time, the frozen V-JEPA2 encoder processes each incoming observation frame and the AC head predicts the next latent, which is used to score each candidate chunk. The scoring function combines the predicted latent distance to a goal state with additional terms measuring how far the candidate falls from the training distribution (action out-of-distribution penalty, support distance). In the Branch 2 combined scoring mode (`selection_mode=best_score`), all scoring terms are combined and the candidate with the highest combined score is selected. In the Branch 3 world-model-only mode (`selection_mode=world_model_only`), only the world model's latent prediction quality is used and the policy's own candidate quality is ignored.

5.4 LeWM path

The LeWM family is implemented through three main stages: episode export, model fitting, and guard evaluation.

The export stage converts EgoDex episodes into contiguous HDF5 sequences that the LeWM training code can consume. Each exported sequence contains an array of pixel frames, the per-step bimanual wrist state, the per-step action (defined as the state delta between consecutive steps), per-episode length, and per-episode byte offset. This format is more compact than the raw EgoDex files because it pre-processes the sequences into a form that the LeWM training loop can read efficiently without re-opening individual episode files.

The fit stage runs the full LeWM training loop, jointly training the ViT-Tiny encoder, AR predictor, action embedder, and projector from random initialisation. This is genuine full-model backpropagation throughout, unlike the V-JEPA2-AC head which trains on stored embeddings. The `train_part1_2` LeWM fit ran for 2,048 steps to completion. This is a much shorter training loop than ACT training over the same data, because LeWM is a considerably smaller model (approximately 15 million parameters versus 51.6 million for ACT) and its training objective is next-embedding prediction rather than full-sequence imitation loss.

The guard evaluation stage scores candidates by encoding the current observation frame through the trained ViT-Tiny encoder, then running a rollout forward through the AR predictor using each candidate's action sequence, and ranking candidates by the predicted latent distance to a goal state.

[FIGURE PLACEHOLDER: LeWM pipeline — export stage converts episodes to HDF5 sequences; fit stage trains ViT-Tiny and predictor jointly from scratch; guard stage encodes current frame and rolls forward to score candidates]

Why LeWM performed differently from V-JEPA2-AC

The performance gap between V-JEPA2-AC and LeWM is structural rather than incidental. V-JEPA2-AC inherits a 300-million-parameter visual representation pretrained on over one million hours of internet video. This representation encodes rich spatiotemporal structure that transfers immediately to EgoDex frames. LeWM’s 15-million-parameter encoder is initialised randomly and must learn all of its visual understanding from EgoDex demonstrations alone. The EgoDex training data, while large for a manipulation dataset, is a fraction of the internet-scale video corpus from which V-JEPA2 learned. The representation capacity gap is therefore not merely a training-time limitation but a fundamental difference in the prior knowledge available to each model.

The self-supervised task-data strategy that LeWM represents has distinct practical merits. It requires no external pretraining corpus, no licensing constraints from a third-party model, and no compatibility assumptions between the pretraining domain and the task domain. It is also far smaller and faster to train. LeWM achieving a 71.7 per cent win rate against the ACT-only baseline from scratch is itself a meaningful finding; the comparison simply establishes that the pretrained backbone strategy has a substantial advantage at the current scale and training scope.

5.5 Policy guard implementation

The policy guard is the runtime component that implements the branch selection logic. For each test window, it receives the ACT policy’s five candidate chunks as input, scores them according to the branch-specific selection mode, and emits the selected chunk together with a decision record.

The Branch 2 guard (`policy_guard`) operates on the V-JEPA2-AC latent space. For each candidate chunk, the guard: (1) encodes the current observation frame through the frozen V-JEPA2 encoder; (2) converts the candidate’s bimanual pose sequence to a `policy_state_delta` action vector using the same normalisation statistics as AC head training; (3) rolls the current latent forward one step using the AC head; (4) scores the predicted latent by distance to the episode goal latent and by action out-of-distribution and support-distance penalties; (5) selects the highest-scoring candidate under the combined score. The Branch 3 guard follows the same path but omits all scoring terms except the world model’s own latent prediction quality.

The LeWM guard (`policy_guard_lewm`) follows the same structural pattern but uses the LeWM encoder and AR predictor rather than V-JEPA2 and the AC head. It consumes the pre-exported `episodes.h5` file to run the rollout prefix steps required by

the LeWM evaluation interface.

An important implementation discipline is that the guard emits a frozen manifest at the end of each run recording the selection mode, the candidate source, and aggregate decision statistics. The Branch 3 V-JEPA2-AC re-run on 18 April 2026 was verified by inspecting the manifest’s `override_reason_counts`:

`{'world_model_only_best_non_act': 7607}`, confirming that all 7,607 windows were decided by the world model rather than the policy. This manifest verification was an essential part of establishing that the corrected Branch 3 run was scientifically valid.

[FIGURE PLACEHOLDER: policy guard scoring diagram — five ACT candidate chunks enter the guard; each is scored by latent forward prediction; Branch 2 uses combined score; Branch 3 uses world-model score only; selected chunk emitted with decision record]

5.6 Ensuring safe and fair evaluation

Several specific safeguards were implemented and verified to maintain the fairness and scientific validity of the benchmark.

The oracle goal leak was the most serious early fairness failure. In early experiments, the goal latent used for scoring was derived from the ground-truth future frame of the episode, giving Branches 2 and 3 access to information that would not be available at deployment time and inflating their scores relative to Branch 1. This was corrected by setting `goal_source=none` in all canonical benchmark runs, removing the goal latent from the scoring function entirely. All results reported in the main results chapter use this corrected configuration.

The candidate pool invariant ensures that all branches score exactly the same set of ACT-generated candidates per window. Branch 3 uses a frozen manifest of the candidate sets generated during the Branch 2 run to guarantee that the same five candidates are evaluated under both selection modes. This prevents the Branch 3 and Branch 2 comparison from being confounded by differences in the candidate generation itself.

The shared `policy_chunk_compare` stage computes all cross-family metrics on the same prediction targets rather than relying on controller-native summaries, which measure different quantities for V-JEPA2-AC and LeWM.

The test split was held out throughout the project. All training, hyperparameter selection, and controller calibration used training and validation splits only. The test split was not consulted during any design decision.

5.7 Implementation challenges and resolutions

Several significant implementation challenges arose during the project and were resolved before the final benchmark could be produced.

LeWM coverage defect

The initial LeWM export used `max_frames_per_episode=256`. When the benchmark aligner attempted to reconstruct windows whose start positions lay beyond frame 256 of any episode, no matching data existed in the export and those windows were silently excluded. The result was that the March 2026 LeWM evaluation covered only 160 of 7,607 intended test windows, or 2.1 per cent of the evaluation surface.

The fix required computing the minimum safe frame cap from the actual benchmark window lists rather than using a fixed conservative value. The minimum required cap was 5,152 frames for the test split and 5,280 frames for `train_part1_2`. A safe contract of `max_frames_per_episode=5280` was adopted and the LeWM export and fit were rebuilt. This raised LeWM evaluation coverage from 2.1 per cent to 100 per cent of the test windows.

LeWM CPU inference

The initial LeWM runtime placed the model on CPU despite the container having GPU access. This occurred because the model loading code contained a hard-coded `.to("cpu")` call. Docker's `--gpus device=N` flag provides container GPU access but does not force model placement. Running 7,607 windows with five candidates each and multiple rollout prefix steps entirely on CPU was estimated to require approximately 54 hours per guard run.

The fix was to check CUDA availability at model load time and move the model and all tensors to the GPU accordingly. The corrected runtime achieved approximately 19.5 windows per second, representing approximately a 130-fold speedup over the CPU-bound configuration.

HDF5 file handling

The LeWM guard initially opened the exported `episodes.h5` file inside the per-window scoring function. With four rollout prefix steps required per window and 29,027 total windows to evaluate, this produced approximately 145,000 file open and close operations per guard run. Beyond being slow, this pattern risked file descriptor exhaustion on systems with low limits.

The fix was to open the HDF5 file once during runtime initialisation and store the file handle as a field of the runtime dataclass. This reduced file operations from 145,000 to one and eliminated the performance bottleneck.

Guard resumability

Guard runs had no crash recovery mechanism in the initial implementation. All per-window accumulations were held in memory and written only at the end of the run. An interrupted run of thousands of windows would require a complete restart.

The fix added per-window JSONL checkpointing to a `progress.jsonl` file, flushed after each window. A `--resume-dir` flag allows an interrupted run to reconstruct all accumulators from the checkpoint and continue from the last completed window.

Branch 3 selection mode error

The evaluation chain script `chain_eval_part1_2_3.sh` launched the Branch 3 V-JEPA2-AC guard without the `--selection-mode world_model_only` flag. The default selection mode is `best_score`, which combines ACT and world model scoring. Under `best_score` with the same candidate pool as Branch 2, Branch 3 produced outputs identical to Branch 2 in all 7,607 windows (0 per cent difference). The scientific purpose of Branch 3 requires the world model to make selections independently of the policy, which requires `world_model_only`. The error was detected through the pairwise comparison output showing a 0.000 win rate for Branch 3 against Branch 2, which indicated identity rather than competition between the two branches. The fix was to re-run the Branch 3 guard with the correct flag, verified by manifest inspection.

goal_l2 metric invalidity

The `goal_l2` metric in early benchmark reports was computed differently across branches. Branch 1 stored `translation_l2` in that field, while Branches 2 and 3 stored the JEPA latent distance (which is zero when `goal_mode=none`). This made the goal metric non-comparable across branches without being obviously wrong from the values alone.

The metric was replaced with `endpoint_translation_l2`, computed from the last quarter of each action chunk's translation predictions in a branch-agnostic way. This metric measures how close the predicted trajectory is to the target over the final portion of the prediction window, independently of how each branch scores or reports its selections internally.

[FIGURE PLACEHOLDER: summary of implementation challenges resolved — timeline showing which issue was discovered when, what evidence indicated the problem, and how it was corrected]

5.8 Scaling and GPU management

Three distinct training splits were used across the project’s development eras. `train_part1` (26 task categories, approximately 336 GB) was used for initial experiments and produced the first frozen three-branch benchmark. `train_part1_2` (39 task categories, approximately 683 GB) was used for the April 14 full-coverage package. `train_part1_2_3` (63 task categories, approximately 1.3 TB) was used for the final canonical package.

The three parts were assembled as a combined split using relative symbolic links rather than copying data, allowing the larger split to be constructed without duplicating the approximately 1 TB of earlier data on disk.

Training and extraction used two A100 GPUs throughout. ACT policy training used single-node DDP via `torchrun`, partitioning the training batch across both GPUs. V-JEPA2 extraction used sharded multi-GPU execution, partitioning the window list across both GPUs and merging results. LeWM fitting and all guard evaluation stages ran on a single GPU.

GPU indexing for the sharded V-JEPA2 extraction was managed by the extraction runner, which partitioned the canonical window list by index range and assigned each shard to a specific device ID. After all shards completed, a merge step concatenated the per-shard `embeddings.npy` files and wrote the final unified embedding array, verified by comparing the window count against the canonical window list.

5.9 Semantic alignment and object grounding

Two auxiliary modules are fully implemented in the repository and remain part of the WACT stack, although they were excluded from the final benchmark evaluation.

The semantic alignment stage converts indexed windows into captions, structured tags, and text embeddings. The pipeline uses a visual language model (Qwen3-VL) to generate free-form descriptions and task-relevant categorical tags for each window, followed by a separate embedding model that maps the text to a shared representation space. A critic and retry loop rejects low-quality outputs and requests regenerated captions, ensuring that semantic artefacts meet a minimum quality threshold.

The object grounding and mask generation stage consumes the semantic object prompts from the alignment stage and produces bounding box detections and segmentation masks for identified objects in the scene. The current implementation should be understood as two-dimensional grounding and mask generation rather than a full six-degree-of-freedom pose estimation pipeline. It provides pixel-space localisation of objects but does not estimate $SE(3)$ object pose.

Both modules were excluded from the final fair V-JEPA2-AC versus LeWM comparison because they could only be wired into one of the two families without the other having access to equivalent inputs, which would have violated the shared input contract. Their implementation is documented in Appendix A.

5.10 Benchmark reporting surface

The final benchmark preparation stage collects all frozen branch runs, the shared pairwise comparison run, and the harmonised branch summary into a single report bundle containing branch tables, pairwise tables, and figures. This is the primary reporting surface used in the results chapter.

The practical consequence of this design is that the dissertation no longer depends on a chain of historical notebook exports or exploratory summaries to establish its main branch comparison. Every quantitative claim in the results chapter is tied to a concrete frozen run or packaged report artefact whose inputs, configuration, and output files are all known. The same staged shell surface also builds the thesis PDF through Docker-backed documentation scripts, keeping the written report aligned with the executable repository workflow.

6 Experimental Design and Metrics

6.1 Canonical evaluation surface

The canonical evaluation surface for this dissertation is the final test benchmark package based on the `train_part1_2_3` training split, with Branch 3 V-JEPA2-AC correctly configured to use world-model-only candidate selection. This package evaluates all five branch instances on the held-out test split of 7,607 windows. All quantitative claims in the results chapter are drawn from this package.

Historical evidence from earlier packages remains useful for understanding how the final benchmark was reached and why certain design decisions were made. However, it is supporting material rather than the main results surface. The hierarchy described in Chapter 4 governs when earlier packages can be cited and prevents the thesis from treating bounded pilot comparisons as equivalent to the full evaluation surface.

6.2 Frozen branch representatives

The canonical package combines:

1. Branch 1: ACT policy trained on `train_part1_2_3`, evaluated on the test split without world model support.
2. Branch 2 V-JEPA2-AC: V-JEPA2-AC world model and AC transition head trained on `train_part1_2_3`, guard run with `selection_mode=best_score`.
3. Branch 3 V-JEPA2-AC: Same V-JEPA2-AC world model, guard re-run with `selection_mode=world_model_only`, verified by manifest inspection confirming all 7,607 windows decided by world model.
4. Branch 2 LeWM: LeWM model fitted on `train_part1_2` at `max_frames_per_episode=5280`, guard run in replay mode.
5. Branch 3 LeWM: Same LeWM model, guard run in candidate-world-model mode.

This construction is a frozen reporting surface, not a pure same-data ablation across all five branches. Branch comparisons should therefore be read as properties of the assembled package, with the understanding that the within-family matched comparisons (B2 versus B3 within V-JEPA2-AC; B2 versus B3 within LeWM) are the most internally controlled because they share the same underlying model weights.

6.3 Prediction targets and their significance

The benchmark evaluates predictions of bimanual wrist SE(3) poses over a 16-step horizon. Each prediction window targets the 16 frames following the observation window, representing approximately 0.53 seconds of future wrist motion at 30 frames per second.

Translation prediction measures how closely the predicted wrist position sequence matches the demonstrated wrist position sequence. Translation is the spatial position component of each SE(3) pose, expressed in world-frame coordinates. This is the primary metric because spatial position is the most directly relevant property for placing a hand at the correct location in a manipulation task. Translation L_2 is computed as the mean Euclidean distance between predicted and target positions across all 16 prediction steps and both hands.

Rotation prediction measures how closely the predicted wrist orientation sequence matches the demonstrated orientation. Rotation is expressed as the angular distance in degrees between predicted and target orientations (computed from the SE(3) rotation component). A policy that predicts the correct wrist position but with incorrect orientation may still fail to grasp or manipulate objects correctly. Rotation error therefore provides an important secondary view of trajectory quality.

Transforms (per-step 6-DoF poses) are available in the EgoDex annotation but are not used as a separate evaluation target in this benchmark. Translation and rotation are extracted and evaluated separately because they have different physical interpretations and different scales. Combining them into a single transform metric would require a weighting choice that is not theoretically grounded.

[FIGURE PLACEHOLDER: prediction targets illustration — 16-step observation window ending at $t_0 + \text{obs_len}$, 16-step prediction window, translation L_2 computed across both hands and all 16 steps, rotation error in degrees]

Shape metrics measure geometric properties of the predicted trajectory rather than pointwise distance to the target. These metrics assess whether the predicted wrist motion follows a plausible kinematic path, independently of whether it reaches exactly the right position. Five shape metrics are computed:

Path **straightness** measures whether the trajectory moves toward its endpoint without unnecessary detours. A highly straight path avoids oscillations and backtracking, which are common failure modes of chunked imitation policies that overfit to local features. A win on straightness means the predicted path is more direct than the baseline.

Oversmoothing (trajectory jitter) measures whether the trajectory avoids high-frequency oscillations around its mean path. Excessive jitter indicates that the policy is uncertain about the exact position at each step and oscillates between

plausible candidates. A win on oversmoothing means the predicted path is smoother.

Velocity consistency measures whether the speed of the predicted wrist motion across the 16 steps matches the demonstrated speed profile. A win on velocity means the predicted trajectory moves at a more appropriate speed.

Acceleration consistency measures whether changes in wrist speed are consistent with the demonstrated acceleration profile. This is distinct from velocity because a trajectory can have the right average speed but still accelerate and decelerate at the wrong moments.

Turning point count measures whether the number of direction reversals in the predicted trajectory matches the demonstration. Too many turning points indicates an oscillating trajectory; too few indicates a trajectory that is spatially smooth but misses important curvature events in the demonstrated motion.

Endpoint translation measures the translation L_2 error over only the final quarter of the prediction window (steps 12 through 16). This provides a measure of how accurately the trajectory arrives at the correct final position, which is often more relevant to task success than accuracy at the first prediction step.

[FIGURE PLACEHOLDER: shape metrics illustrated on an example trajectory — two wrist paths shown (predicted and target), annotations showing straightness, jitter, velocity profile, and turning points]

6.4 Branch design and the five evaluated instances

The five evaluated branch instances correspond to three branch structures and two world model families, as defined in Chapter 4. The key distinctions are:

Branch 1 uses `selection_mode=policy_only`: the ACT policy's internal scoring selects one of five candidates. This is the baseline.

Branch 2 uses `selection_mode=best_score`: a combined scoring function including the world model's predicted latent distance, an action out-of-distribution penalty, and a support distance term ranks the five candidates and selects the highest scorer. ACT's own preference enters through the combined scoring function.

Branch 3 uses `selection_mode=world_model_only`: only the world model's predicted latent distance ranks the candidates. ACT's internal candidate quality is not considered. This tests whether the world model alone, without any policy guidance in the selection step, can identify better trajectories.

[FIGURE PLACEHOLDER: branch selection logic comparison — same five ACT candidates enter all three branches; Branch 1 uses policy's own ranking; Branch 2 uses combined score; Branch 3 uses WM score only; arrows and decision labels]

6.5 Fair comparison rules

The main model family comparison between V-JEPA2-AC and LeWM follows a strict shared input contract. Both families receive the same raw RGB observation frame and the same bimanual wrist state as inputs. Oracle goal leakage is excluded from both families (`goal_source=none`). Both families are evaluated on exactly the same 7,607 test windows. The candidate pool used by Branch 3 is the same frozen set used by Branch 2, ensuring that differences in Branch 3 outcomes reflect differences in world model selection, not differences in the candidates available.

Cross-family metrics use the shared `policy_chunk_compare` surface, which recomputes chunk-level translation, rotation, and shape quantities on the same prediction targets. This is necessary because the native V-JEPA2-AC guard reports immediate-step translation error (first prediction step only) while the native LeWM guard reports mean chunk translation error, and a direct comparison of these two quantities would be invalid.

6.6 Metric harmonisation

The branch summary table in the results chapter is derived from the shared `policy_chunk_compare` surface rather than from controller-native metric logs. This harmonisation step is methodologically important: it means the branch table represents a common chunk-level reporting surface derived from paired branch and comparison artefacts, not a naive juxtaposition of incompatible native summaries.

The translation L_2 values in the branch table are therefore mean chunk translation errors (mean across all 16 prediction steps and both hands), computed on the same candidate chunks across all five branches. The rotation error values are mean angular distances across the same windows and steps.

6.7 Claims policy

The thesis adopts a strict claims policy governing what conclusions can and cannot be drawn from the benchmark results.

Claims that are supported by the evidence include: the implemented offline benchmark shows that world model support materially improves chunk-level trajectory quality; Branch 3 V-JEPA2-AC achieves the best aggregate mean translation error and the highest win rate against Branch 1; Branch 2 V-JEPA2-AC shows greater window-by-window consistency in the direct Branch 3 versus Branch 2 comparison; both LeWM branches improve on Branch 1; neither LeWM branch approaches the

V-JEPA2-AC family; the improvement is robust across confidence strata and task categories; path geometry improves substantially while temporal dynamics do not.

Claims that are not supported include: lower offline translation L_2 implies higher task completion rate; wrist trajectory quality is sufficient to characterise manipulation success; the WACT architecture generalises to non-EgoDex domains without retraining; semantic alignment and object grounding are unnecessary in general; the scaling finding is universal across architectures.

6.8 Metric justification

Translation L_2 is not established as the official evaluation metric in the EgoDex dataset paper. The justification for its use as a primary offline quality proxy rests on two lines of evidence from the manipulation evaluation literature. SimplerEnv [43] demonstrates that simulation-based offline evaluation reliably predicts real-world manipulation performance when control and visual disparities are managed, providing support for offline trajectory metrics as meaningful proxies for deployed quality. WorldEval [44] directly proposes world-model-based offline evaluation, showing strong correlation between offline world-model trajectory rankings and real-world task success. Together these establish a literature precedent for offline trajectory metrics as policy quality proxies in the manipulation domain. All results in this thesis are interpreted within the acknowledged limitation that no paper establishes this connection specifically for EgoDex wrist trajectories, and no claim of task completion rate improvement is made.

6.9 Scaling experiment design

The scaling experiment tests whether expanding from `train_part1_2` (39 tasks) to `train_part1_2_3` (63 tasks) improves the pretrained world model family. Both training splits were evaluated on the same held-out test split of 7,607 windows. The ACT policy, V-JEPA2-AC extractor, and LeWM model were all retrained from scratch on the larger split. No weights were transferred from the smaller split.

The results show that the `train_part1_2_3` training did not improve the V-JEPA2-AC family (Branch 2 V-JEPA2-AC: 0.0344 under `train_part1_2` versus 0.0354 under `train_part1_2_3`). The Branch 1 ACT-only baseline improved very slightly (0.1439 to 0.1429), while Branch 3 LeWM was essentially unchanged (0.0906 to 0.0909). These findings are interpreted as evidence of capacity saturation under the current architecture rather than a fundamental ceiling on world model guidance in general, as discussed in Chapter 8.

7 Results

All results in this chapter are drawn from the canonical final benchmark package based on the `train_part1_2_3` training split with Branch 3 V-JEPA2-AC correctly configured to use world-model-only selection. The evaluation surface is the held-out `test` split of 7,607 windows covering 63 task categories. Guard manifest verification confirmed that all 7,607 Branch 3 V-JEPA2-AC windows were decided by the world model rather than the policy (`override_reason_counts`: `{world_model_only_best_non_act: 7607}`).

The results are organised into four analytical perspectives: branch comparison, world model family comparison, path shape and trajectory quality, and scaling behaviour.

[FIGURE PLACEHOLDER: summary overview figure — four panels showing the four analytical perspectives, directing the reader to the subsections that follow]

7.1 Branch comparison

Table 7.1 presents the branch-level metrics for all five evaluated branch instances. Translation L_2 and rotation error are computed over all 16 prediction steps using the harmonised `policy_chunk_compare` surface, ensuring cross-family comparability.

Table 7.1 Final benchmark branch summary. Training split: `train_part1_2_3` (63 tasks, 195k episodes). Evaluation: 7,607 held-out `test` windows. Translation L_2 and rotation error are chunk-level means over all 16 prediction steps.

Branch	Translation $L_2 \downarrow$	Rotation ($^\circ$) $\downarrow$	Δ vs B1 (translation)
B1 – ACT only	0.1429	41.4 $^\circ$	—
B2 – V-JEPA2-AC hybrid	0.0354	15.8 $^\circ$	−75.2%
B3 – V-JEPA2-AC world-model-only	0.0328	15.2°	−77.0%
B2 – LeWM hybrid	0.1179	31.9 $^\circ$	−17.5%
B3 – LeWM	0.0909	27.5 $^\circ$	−36.4%

The table supports three immediate observations. First, both V-JEPA2-AC branches substantially improve on the ACT-only anchor across both translation and rotation, reducing mean translation L_2 by 75 and 77 per cent respectively. Second, both LeWM branches also improve on Branch 1, confirming that world model guidance is not contingent on a pretrained backbone. Third, the V-JEPA2-AC family is decisively stronger than the LeWM family across all matched branch comparisons.

Table 7.2 presents the pairwise win rates. These measure how often one branch produces a trajectory closer to the demonstration than the other on the same test




Figure 7.1 Branch-level benchmark metrics. Left: mean translation L_2 . Right: mean rotation error in degrees. Both are chunk-level means over 16 prediction steps across 7,607 test windows. Lower is better.

window.

Table 7.2 Selected pairwise win rates. Translation win rate: fraction of 7,607 windows where the candidate branch achieves lower translation L_2 than the reference branch.

Comparison	Windows	Translation win rate
B3 V-JEPA2-AC vs B1 ACT	7,607	91.5%
B2 V-JEPA2-AC vs B1 ACT	7,607	86.6%
B3 V-JEPA2-AC vs B2 V-JEPA2-AC	7,607	14.6%
B3 LeWM vs B1 ACT	7,607	71.7%
B3 LeWM vs B2 V-JEPA2-AC	7,607	17.4%
B2 LeWM vs B2 V-JEPA2-AC	7,607	12.9%




Figure 7.2 Pairwise win rates from the final test package. Each bar shows the fraction of 7,607 windows where the listed branch outperforms the reference branch on translation L_2 . The dashed line marks 50 per cent.

The pairwise results confirm the branch-level reading. Both V-JEPA2-AC branches dominate Branch 1 strongly and consistently. Branch 3 V-JEPA2-AC achieves the highest win rate against Branch 1 (91.5 per cent) while Branch 2 achieves 86.6 per cent. The LeWM family beats Branch 1 at 71.7 per cent (Branch 3) and substantially behind V-JEPA2-AC on matched comparisons (12.9 to 17.4 per cent).

7.2 Branch authority: world-model-driven versus hybrid

The direct comparison between Branch 3 V-JEPA2-AC and Branch 2 V-JEPA2-AC warrants careful interpretation. Branch 3 achieves a better aggregate mean translation L_2 (0.0328 versus 0.0354) and a higher win rate against Branch 1 (91.5 versus 86.6 per cent). However, in the direct head-to-head comparison, Branch 2 outperforms Branch 3 on 85.4 per cent of individual test windows (Branch 3 wins only 14.6 per cent).

This is not a contradiction. It arises because Branch 3 makes large improvements on a subset of windows where pure world model authority produces substantially better trajectories, while performing marginally worse than Branch 2 on the majority of windows where the combined hybrid score and the pure world model score point to different candidates. Branch 3 has a higher variance of per-window improvement: it wins less often but wins by larger margins when it does win, resulting in a better aggregate mean.

[FIGURE PLACEHOLDER: per-window difference distribution between B3 and B2 V-JEPA2-AC – histogram of (B3 translation L_2 minus B2 translation L_2) per window; showing the right-skewed distribution where B3 wins large and loses small]

The aggregate mean improvement favours Branch 3, while the window-by-window consistency favours Branch 2. Both are genuine properties of the two branches, and both are reported here. The broader implication is that full world model authority is beneficial when the world model is reliable enough to override the policy on the windows where they disagree, but the hybrid design offers more predictable per-window behaviour.

7.3 World model family comparison

The performance gap between the V-JEPA2-AC and LeWM families is substantial and consistent across branches. At the Branch 3 level, V-JEPA2-AC (0.0328) improves over LeWM (0.0909) by a factor of approximately 2.8 in mean translation L_2 . In direct head-to-head comparisons, the matched LeWM branch beats the matched V-JEPA2-AC branch in only 12.9 to 17.4 per cent of windows.

This gap reflects the structural difference between the two strategies. V-JEPA2-AC inherits a 300-million-parameter visual representation pretrained on over one million hours of internet video, providing rich spatiotemporal priors that transfer immediately to EgoDex frames. LeWM learns its 15-million-parameter representation entirely from EgoDex data. The approximately 20 percentage point gap in win rate against Branch 1 (91.5 versus 71.7 per cent) is consistent with the pretrained representation advantages reported by R3M [30] and MVP [46] in the manipulation

learning literature.

LeWM achieving 71.7 per cent against the ACT-only baseline from scratch is itself a meaningful result. It confirms that world model guidance can improve trajectory quality without any external pretraining corpus, and that LeWM is a viable world model family for this task. The gap with V-JEPA2-AC reflects the pretrained backbone advantage, not a failure of the from-scratch strategy.

[FIGURE PLACEHOLDER: V-JEPA2-AC versus LeWM family comparison — bar charts showing win rates and L2 values for all four world-model branches, grouped by family; dashed line at B1 baseline level]

7.4 Path shape and trajectory geometry

Beyond aggregate translation and rotation error, the benchmark records five path shape metrics for each evaluated window, comparing the predicted trajectory’s geometric properties against the ground-truth demonstration. Table 7.3 presents these results for the Branch 2 V-JEPA2-AC versus Branch 1 comparison, which has the largest sample size and the most consistent improvement.

Table 7.3 Path shape win rates for Branch 2 V-JEPA2-AC versus Branch 1 ACT-only (7,607 windows). Win rate: fraction of windows where the V-JEPA2-AC branch produces a trajectory with better shape by the given metric.

Shape metric	Win rate	Interpretation
Straightness	89.4%	V-JEPA2-AC strongly improves path linearity
Oversmoothing	93.1%	V-JEPA2-AC strongly reduces trajectory jitter
Velocity	56.5%	Near chance — temporal dynamics largely unchanged
Acceleration	50.1%	Near chance — temporal dynamics largely unchanged
Turning points	48.9%	Near chance — curvature events not improved

[FIGURE PLACEHOLDER: radar or bar chart of the five shape metric win rates for B2 V-JEPA2-AC vs B1 — showing high geometric metrics and near-chance temporal metrics]

The shape results reveal an important asymmetry in what world model guidance improves. V-JEPA2-AC selection strongly improves path geometry: trajectories are more spatially linear (89.4 per cent) and have substantially less jitter (93.1 per cent). However, world model guidance has almost no effect on temporal dynamics: velocity profiles (56.5 per cent) and acceleration profiles (50.1 per cent) are near chance, and turning point preservation (48.9 per cent) is slightly below chance.

The interpretation is that the V-JEPA2 latent space encodes spatial plausibility more reliably than temporal dynamics. The world model can identify which candidate

trajectory ends in the right place and follows a spatially coherent path, but cannot reliably identify which candidate moves at the right speed or decelerates at the right moments. This is consistent with the temporal entanglement limitation identified by Tsagkas et al. [33], who find that pretrained representations can conflate observations from different task stages.

7.5 Robustness across confidence strata and tasks

The benchmark records per-joint confidence scores for each EgoDex demonstration frame. Table 7.4 presents translation win rates for Branch 2 V-JEPA2-AC versus Branch 1 stratified by wrist annotation confidence.

Table 7.4 Confidence stratification for Branch 2 V-JEPA2-AC versus Branch 1 ACT-only.
High confidence: wrist confidence score ≥ 0.85 . Low confidence: wrist confidence score < 0.85 .

Window type	Count	% of test	Translation win rate
High confidence (≥ 0.85)	5,957	78%	88.0%
Low confidence (< 0.85)	1,650	22%	81.6%

The 6.4 percentage point gap between confidence strata is small relative to the overall improvement. V-JEPA2-AC guidance adds value across the full confidence distribution and is not inflated by being evaluated primarily on high-confidence, reliably-annotated frames. This is an important robustness finding: even when the ground-truth annotation is less reliable (low-confidence windows, where tracking may have been interrupted by fast motion or occlusion), the world model selection still produces substantially better trajectories than the ACT baseline.

Per-task analysis across all 111 unique task-category windows in the test split confirms that the improvement is broad rather than concentrated in a small number of tasks. The five worst-performing task categories still show win rates in the range of 62 to 66 per cent against Branch 1, well above chance. No single task cluster drives the 86.6 per cent aggregate headline.

[FIGURE PLACEHOLDER: per-task win rate distribution for B2 V-JEPA2-AC vs B1 — histogram of win rates across 111 task categories; vertical lines at 50%, 75%, and 86.6% (aggregate) to show spread and floor]

7.6 Scaling experiment

The `train_part1_2_3` training split was assembled after the April 14 package to test whether expanding from 39 to 63 task categories improved the pretrained world

model family. Table 7.5 presents the comparison.

Table 7.5 Scaling experiment: Branch-level translation L_2 on the test split under two training scales. Lower is better.

Branch	train_part1_2	train_part1_2_3	Δ
B1 ACT-only	0.1439	0.1429	-0.001
B2 V-JEPA2-AC hybrid	0.0344	0.0354	$+0.001$
B3 LeWM	0.0906	0.0909	≈ 0

The scaling experiment shows that expanding training data from 39 to 63 task categories did not improve the V-JEPA2-AC world model. Branch 2 V-JEPA2-AC regressed slightly (0.0344 to 0.0354), while Branch 1 ACT-only and Branch 3 LeWM were essentially unchanged. The slight V-JEPA2-AC regression is consistent with the additional 24 task categories in `train_part3` introducing mild distributional shift relative to the test split’s 26 represented tasks, without providing enough additional overlap to compensate. Crucially, the win rate against Branch 1 remains above 86 per cent across all three training splits (approximately 91 per cent on `train_part1` and `train_part1_2`, and 86.6 per cent on `train_part1_2_3`), confirming that the world model benefit is not an artefact of any single training data volume.

[FIGURE PLACEHOLDER: scaling experiment bar charts — B1, B2 JEPA, B3 LeWM translation L_2 under `train_part1_2` and `train_part1_2_3`; showing flat or slightly negative scaling for V-JEPA2-AC]

This finding suggests that the V-JEPA2-AC family is near a capacity ceiling at the current policy architecture scale. The pretrained V-JEPA2 backbone already provides rich representations that do not benefit substantially from seeing more EgoDex task variety. The lightweight AC head that learns from these frozen embeddings may similarly be near its capacity given the current policy size. A larger policy backbone (for example, replacing ResNet-18 with a ViT-Small or larger) might yield different scaling behaviour and is identified as future work in Chapter 9.

7.7 Why the canonical package supersedes the April 14 results

The April 14 full-coverage package was the first evaluation to achieve 7,607-window coverage for the LeWM family and constituted a significant advance over the March 2026 package. However, it cannot be treated as the final canonical results surface for two reasons.

First, Branch 3 V-JEPA2-AC in the April 14 package ran with `selection_mode=best_score` rather than `world_model_only`, making it scientifically

identical to Branch 2. The three-branch comparison was effectively a two-branch comparison. The thesis cannot claim a genuine Branch 3 result without the world-model-only correction.

Second, the April 14 package used `train_part1_2` trained models. The scaling experiment demonstrates that the transition to `train_part1_2_3` produced a small but real change in V-JEPA2-AC performance. Reporting only the April 14 numbers would require explaining why a later, larger training run was excluded.

The canonical results therefore use the `train_part1_2_3` package with the Branch 3 correction as the definitive reporting surface. The April 14 package remains useful as a reference point for the scaling experiment and for understanding the development trajectory of the benchmark.

8 Discussion and Limitations

8.1 What the benchmark shows

The central finding of the dissertation is that world model guidance materially and consistently improves offline trajectory quality in a controlled bimanual manipulation benchmark. This improvement holds across both evaluated world model families, both levels of world model authority, and across the full breadth of the test surface including lower-confidence windows and the worst-performing task categories.

The improvement is not marginal. The V-JEPA2-AC world model reduces mean translation L_2 by 77 per cent relative to the policy-only baseline and achieves a 91.5 per cent win rate in head-to-head window comparisons under world-model-only selection. Even the task-data-trained LeWM family, which has no access to any external pretraining, achieves a 71.7 per cent win rate from scratch. The effect of world model guidance is therefore robust to world model architecture.

At the same time, the results establish that the structure of the guidance matters. The two world model families differ by approximately 20 percentage points in win rate against the baseline, and the degree of world model authority (Branch 2 versus Branch 3) affects both the aggregate mean improvement and the window-by-window consistency of that improvement.

8.2 The aggregate versus per-window tension

The most interesting interpretive tension in the canonical results is between Branch 3 V-JEPA2-AC and Branch 2 V-JEPA2-AC. Branch 3 achieves a better aggregate mean translation L_2 (0.0328 versus 0.0354) and a higher win rate against Branch 1 (91.5 versus 86.6 per cent). However, in the direct head-to-head, Branch 2 outperforms Branch 3 on 85.4 per cent of individual windows.

This pattern arises because Branch 3 makes larger improvements on a smaller number of windows, while Branch 2 makes smaller but more consistent improvements across a larger fraction of windows. Branch 3's aggregate advantage comes from winning by large margins on the windows where full world model authority produces substantially better trajectories; on the majority of windows, the combined hybrid score in Branch 2 is a more reliable guide.

Neither metric is more correct than the other. The aggregate mean measures the typical improvement in expected prediction error. The pairwise win rate measures how often one branch produces a closer trajectory than another on a given window. Both are reported because they provide complementary views of the same phenomenon.

The broader implication is that full world model authority is beneficial when the world model is reliable enough to systematically prefer better candidates. In the V-JEPA2-AC case, the pretrained backbone provides rich enough representations that world-model-only selection outperforms hybrid selection in aggregate, even though the hybrid design is more consistent window-by-window. A world model with weaker representations (such as LeWM) shows the opposite pattern: Branch 3 LeWM achieves a better aggregate mean than Branch 2 LeWM (0.0909 versus 0.1179), again because concentrated improvements on a subset of windows more than compensate for losses on others.

8.3 What the path shape results reveal

The shape metric results in Table 7.3 provide a mechanistic interpretation of what the world model is and is not correcting. V-JEPA2-AC guidance strongly improves spatial trajectory geometry: path straightness (89.4 per cent win rate) and oversmoothing reduction (93.1 per cent). This indicates that the world model’s latent space reliably encodes which candidate follows a spatially plausible path toward the goal.

By contrast, temporal dynamics are essentially unchanged. Velocity (56.5 per cent), acceleration (50.1 per cent), and turning point preservation (48.9 per cent) are all near chance. The world model selects trajectories that end in better places and follow straighter paths, but it does not select trajectories with better timing or curvature event profiles.

This is consistent with the temporal entanglement limitation identified by Tsagkas et al. [33]. Pretrained visual representations trained on diverse video data can encode where a scene is likely to lead without reliably encoding the precise dynamics of how it gets there. The V-JEPA2 backbone, pretrained on internet-scale video, provides strong spatial priors but weaker temporal dynamics priors. The AC head trained on top of these frozen embeddings inherits both the spatial strength and the temporal limitation.

This finding is practically informative. A downstream system that needs spatially correct wrist trajectories (for example, to ensure that the hand reaches the right region of space) would benefit substantially from V-JEPA2-AC guidance. A system that also needs temporally correct velocity profiles (for example, to synchronise with a moving target) would need additional mechanisms beyond the current world model architecture.

8.4 The family comparison and its architectural basis

The V-JEPA2-AC versus LeWM comparison is a comparison of two coherent research strategies rather than a fair ablation of model quality at fixed scale. V-JEPA2-AC uses a 300-million-parameter backbone pretrained on over one million hours of video, while LeWM trains a 15-million-parameter model from scratch on EgoDex. The pretrained backbone strategy carries a structural advantage that is not explained away by noting the architectural difference; the difference is the point.

The approximately 20 percentage point gap in win rate against the baseline (91.5 versus 71.7 per cent) is consistent with the advantage reported by R3M and MVP for pretrained egocentric video representations over task-specific training in robot manipulation [30, 46]. The gap is therefore not a WACT-specific artefact.

LeWM achieving 71.7 per cent from scratch is a meaningful result in its own right. It establishes that a compact, task-data-trained world model can provide useful guidance even without external pretraining, and that the LeWM architecture is a viable option in deployment contexts where large pretrained backbones are unavailable or unsuitable. The question of whether a pretrained variant of LeWM could close the gap is a valid future direction but would require architectural changes that would make the result non-comparable to the current benchmark.

8.5 Why benchmark evolution is part of the contribution

An important part of the dissertation’s contribution is the development of a defensible evaluation surface, not only the final numerical rankings. The project progressed through three distinct benchmark eras: early bounded pilot slices used for controller selection, the March 2026 frozen three-branch bundle with limited LeWM coverage, and the final canonical package with full coverage and correct Branch 3 configuration.

Two corrections were necessary before the final results could be trusted as scientifically meaningful. The first was the LeWM coverage fix, which required recomputing the frame cap from actual benchmark window lists and rebuilding the export contract. The second was the Branch 3 selection mode correction, which required identifying through pairwise manifest inspection that two nominally distinct branches had produced identical outputs, diagnosing the root cause as a missing script flag, re-running with the correct flag, and verifying the result through manifest inspection.

Without both corrections, the reported conclusions would have been weaker and potentially misleading. An uncorrected LeWM result based on 160 windows would have left open the question of whether the V-JEPA2-AC advantage over LeWM was a

coverage artefact. An uncorrected Branch 3 result would have made the three-branch comparison a two-branch comparison and missed the finding that full world model authority produces a better aggregate mean than the hybrid design for the V-JEPA2-AC family.

8.6 Why semantic alignment and object grounding were excluded

The decision to exclude semantic alignment and object grounding from the final benchmark deserves explicit discussion because it might appear to weaken the result by removing potentially useful information from the evaluation.

Both modules are fully implemented and produce useful artefacts. They were excluded from the main benchmark because incorporating either into only one world model family’s scoring would have given that family privileged inputs not available to the other. The V-JEPA2-AC family could in principle be conditioned on object grounding masks through the semantic extraction path, while the LeWM family has no equivalent semantic conditioning in its current architecture. Including such conditioning for V-JEPA2-AC alone would change the nature of the comparison from “which world model family provides better guidance given the same inputs” to “does additional semantic information help V-JEPA2-AC.” These are distinct scientific questions. The thesis addresses only the former.

It remains entirely plausible that semantic conditioning would improve V-JEPA2-AC performance beyond the results reported here. That question is left as future work under a symmetric design where both families receive equivalent semantic inputs.

8.7 Main limitations

The first limitation is scope. The benchmark is offline and trajectory-level. It measures how closely predicted wrist motions match expert demonstrations. It does not measure task completion, object manipulation success, or any deployed behavioural outcome. The connection between offline trajectory L_2 and deployed task success is not established by this dissertation.

The second limitation is the wrist-only evaluation scope. EgoDex provides 138 joint transform sequences per frame, but the benchmark uses only the left and right wrist poses as both policy targets and evaluation targets. Wrist trajectory is a reasonable proxy for manipulation intent but does not capture finger articulation,

elbow configuration, or body posture. A policy that achieves lower wrist L_2 might still exhibit poor finger configuration or elbow alignment.

The third limitation is the offline-to-online gap. Even if offline trajectory quality is a reliable proxy for deployment success (as suggested by `SimplerEnv` and `WorldEval`), there is no deployment evaluation in the present system. The GRASP mixed-reality deployment project is designed to address this gap but its results are outside the scope of this dissertation.

The fourth limitation concerns the scaling finding. The scaling experiment tests only one expansion, from 39 to 63 task categories. Whether a larger backbone architecture would also saturate at this training scale, or whether the saturation is specific to the ResNet-18 and AC head configuration, cannot be determined from the current results alone.

The fifth limitation is metric coverage. The benchmark uses translation L_2 , rotation error, and five path shape metrics. It does not yet include a single scalar notion of overall prediction quality that is theoretically grounded for non-linear manipulation trajectories. The multi-metric bundle provides useful triangulation but does not collapse to one interpretable number.

9 Conclusion

9.1 Summary of findings

This dissertation asked to what extent world model support improves an ACT policy across three implemented branches: policy-only control, hybrid ACT plus world model support, and world-model-driven candidate selection. The answer is clear and supported by a fully corrected, full-coverage evaluation surface.

World model guidance substantially improves offline trajectory quality. Under world-model-driven selection with the pretrained V-JEPA2-AC family, mean translation error is reduced by 77 per cent compared to the policy-only baseline, and the world model selects a closer trajectory than the policy alone in 91.5 per cent of head-to-head window comparisons across the 7,607-window test set. The hybrid selection strategy achieves 86.6 per cent, with greater window-by-window consistency. Both results hold across 63 task categories, across the full range of annotation confidence, and with no single task cluster driving the aggregate headline.

World model guidance does not depend on a pretrained backbone. The task-data-trained LeWM family achieves a 71.7 per cent win rate against the policy-only baseline despite learning all of its visual representations from EgoDex demonstrations alone. The pretrained backbone strategy nevertheless carries a decisive advantage of approximately 20 percentage points, consistent with the literature on pretrained visual representations for robot learning.

The shape analysis reveals that world model guidance improves where trajectories end and how spatially coherent their paths are, but does not substantially alter their temporal dynamics. Path straightness and oversmoothing improve to win rates of 89.4 and 93.1 per cent respectively, while velocity and acceleration profiles remain near chance. This points to a clear direction for future improvement.

Scaling the training data from 39 to 63 task categories did not further improve the pretrained world model, suggesting that the current ResNet-18 policy backbone and lightweight transition head are near a capacity ceiling at this training scope.

9.2 Methodological contribution

The second contribution of the thesis is the documented development of a defensible evaluation surface. Two significant benchmark defects were identified, diagnosed, and corrected during the project: a data coverage defect that limited one model family to 2.1 per cent of the intended test surface, and a selection logic error that caused two nominally distinct branches to produce identical outputs across all 7,607 windows.

Both corrections required non-trivial implementation work and altered the scientific interpretation of the results.

The thesis contribution therefore includes benchmark curation and the establishment of a trustworthy evidence hierarchy, not only the final numerical rankings. The process of identifying and correcting evaluation defects is part of the scientific work. A result that depends on a broken evaluation surface is not a valid result, and the effort required to ensure validity is as important to report as the result itself.

9.3 Scope of conclusions

The conclusions of this dissertation are explicitly bounded. They concern offline prediction quality on a held-out set of recorded demonstrations, not deployed task success. The system does not execute actions in any physical or simulated environment, and no task completion rates are reported. The improvement in trajectory quality relative to the policy-only baseline is a finding about offline imitation fidelity, which is a meaningful proxy for deployment quality under appropriate conditions but is not equivalent to it.

The thesis does not claim solved dexterous manipulation, solved action chunk fidelity, or closed-loop deployment success. It establishes a reproducible offline benchmark with a real policy and world model interface and a set of bounded empirical results that can support more targeted future work.

9.4 Future Work

The results and limitations identified in this dissertation suggest several concrete directions for future work.

The most immediate next step is online evaluation through the GRASP mixed-reality deployment project. GRASP deploys an ACT-based control loop in a live extended reality environment, providing the first opportunity to measure whether the improvements in offline trajectory quality found in this dissertation translate to improved task completion rates in a physical system. The connection between offline L_2 improvement and deployed task success is the primary open question left by the present work.

A larger policy backbone is the most directly motivated architectural extension. The scaling experiment shows that the current ResNet-18 backbone saturates with 63 training tasks. Replacing ResNet-18 with a ViT-Small or ViT-Base backbone would

substantially increase the visual capacity of the policy and might allow the AC transition head to learn more detailed action-conditioned dynamics from the larger training split.

Temporal dynamics improvement is suggested by the shape analysis. The world model currently improves spatial geometry but not velocity or acceleration profiles. A transition head that is explicitly trained to predict temporal dynamics (rather than only spatial latent distance) might address this. One approach would be to add a temporal consistency loss to the AC head training, penalising transitions that change velocity profile unexpectedly.

Symmetric semantic conditioning would allow a direct test of whether semantic alignment and object grounding improve world model guidance. The current benchmark excluded both modules to maintain fairness, but a future experiment with symmetric semantic inputs across both model families would measure their effect independently of the fairness constraint.

A pretrained LeWM variant using a standard ViT-Small or DINOv2 backbone would allow a more direct comparison between the pretrained and from-scratch strategies at comparable model capacity. The current LeWM result reflects both the from-scratch training and the smaller model scale; a pretrained variant at the same parameter count would isolate the effect of initialisation from the effect of capacity.

Wrist-only evaluation could be extended to include additional joint targets from EgoDex’s 138-joint annotation. Including finger articulation or elbow configuration as additional evaluation targets would test whether the world model improvements in wrist trajectory extend to the full kinematic configuration, or whether wrist improvements mask degradation elsewhere.

Finally, the DAgger-style distillation pathway [6] offers a route to closing the loop between world model guidance and policy quality. World-model-accepted trajectories could be used as a teacher signal to retrain the ACT policy, progressively aligning the policy’s own output distribution with the trajectories that the world model rates most highly. If the world model reliably identifies better trajectories, distillation should produce a policy that generates better candidates without requiring world model evaluation at every step.

References

- [1] B. D. Argall, S. Chernova, M. Veloso, and B. Browning, "A survey of robot learning from demonstration," *Robotics and Autonomous Systems*, vol. 57, no. 5, pp. 469–483, 2009. DOI: 10.1016/j.robot.2008.10.024
- [2] A. Y. Ng and S. J. Russell, "Algorithms for inverse reinforcement learning," in *Proceedings of the Seventeenth International Conference on Machine Learning*, Morgan Kaufmann, 2000, pp. 663–670.
- [3] M. Jeannerod, "The timing of natural prehension movements," *Journal of Motor Behavior*, vol. 16, no. 3, pp. 235–254, 1984. DOI: 10.1080/00222895.1984.10735319
- [4] M. Jeannerod, M. A. Arbib, G. Rizzolatti, and H. Sakata, "Grasping objects: The cortical mechanisms of visuomotor transformation," *Trends in Neurosciences*, vol. 18, no. 7, pp. 314–320, 1995. DOI: 10.1016/0166-2236(95)93921-J
- [5] M. Santello, M. Flanders, and J. F. Soechting, "Postural hand synergies for tool use," *Journal of Neuroscience*, vol. 18, no. 23, pp. 10 105–10 115, 1998. DOI: 10.1523/JNEUROSCI.18-23-10105.1998
- [6] S. Ross, G. Gordon, and D. Bagnell, "A reduction of imitation learning and structured prediction to no-regret online learning," in *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 15, 2011.
- [7] T. Z. Zhao, V. Kumar, S. Levine, and C. Finn, *Learning fine-grained bimanual manipulation with low-cost hardware*, Apr. 2023. DOI: 10.48550/arXiv.2304.13705 arXiv: 2304.13705 [cs].
- [8] Z. Ding, "Imitation learning," in *Deep Reinforcement Learning: Fundamentals, Research and Applications*, H. Dong, Z. Ding, and S. Zhang, Eds. Singapore: Springer Singapore, 2020, pp. 273–306, ISBN: 978-981-15-4095-0. DOI: 10.1007/978-981-15-4095-0_8 [Online]. Available: https://doi.org/10.1007/978-981-15-4095-0_8
- [9] C. Chi et al., *Diffusion policy: Visuomotor policy learning via action diffusion*, Mar. 2024. DOI: 10.48550/arXiv.2303.04137 arXiv: 2303.04137 [cs].
- [10] S. Jain et al., *DiWA: Diffusion policy adaptation with world models*, ArXiv preprint, Aug. 2025. arXiv: 2508.03645 [cs].
- [11] K. M. Lynch and F. C. Park, *Modern Robotics: Mechanics, Planning, and Control*. Cambridge University Press, 2017, ISBN: 978-1-107-15630-2.

- [12] Omitram, *Left-hand vs. right-hand coordinate systems: The 3d developer's guide* - omitram – omitram.com, <https://omitram.com/3d-coordinate-systems-left-vs-right-hand/>, 2024.
- [13] Y. Zhou, C. Barnes, J. Lu, J. Yang, and H. Li, "On the continuity of rotation representations in neural networks," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 5745–5753. DOI: 10.1109/CVPR.2019.00589
- [14] N. Grossmann, E. Gröller, and M. Waldner, "Concept splatters: Exploration of latent spaces based on human interpretable concepts," *Computers & Graphics*, vol. 105, pp. 73–84, 2022, ISSN: 0097-8493. DOI: <https://doi.org/10.1016/j.cag.2022.04.013> [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0097849322000656>
- [15] A. Dawid and Y. LeCun, "Introduction to latent variable energy-based models: A path towards autonomous machine intelligence," *Journal of Statistical Mechanics: Theory and Experiment*, 2023. arXiv: 2306.02572 [cs].
- [16] D. P. Kingma and M. Welling, *Auto-encoding variational bayes*, 2013. DOI: 10.48550/arXiv.1312.6114 arXiv: 1312.6114 [stat.ML].
- [17] L. Li and R. Betti, "A machine learning-based data augmentation strategy for structural damage classification in civil infrastructure system," *Journal of Civil Structural Health Monitoring*, vol. 13, pp. 1–21, May 2023. DOI: 10.1007/s13349-023-00705-5
- [18] AnotherDatum. "Variational autoencoders explained," Accessed: Apr. 24, 2026. [Online]. Available: <https://anotherdatum.com/vae.html>
- [19] K. Sohn, H. Lee, and X. Yan, "Learning structured output representation using deep conditional generative models," in *Advances in Neural Information Processing Systems*, vol. 28, 2015.
- [20] A. Vaswani et al., "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017. DOI: 10.48550/arXiv.1706.03762 arXiv: 1706.03762 [cs.CL].
- [21] S. Panda, *Decoding attention types and strategies for effective nlp models*, <https://www.linkedin.com/pulse/decoding-attention-types-strategies-effective-nlp-models-swagat-panda-erhef/>, 2024.
- [22] M. Assran et al., "Self-supervised learning from images with a joint-embedding predictive architecture," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv: 2301.08243 [cs].

- [23] IBM Developer, *Models for machine learning*, <https://developer.ibm.com/articles/cc-models-machine-learning/>, 2017. Accessed: May 14, 2026.
- [24] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, *Film: Visual reasoning with a general conditioning layer*, 2017. arXiv: 1709.07871 [cs.CV]. [Online]. Available: <https://arxiv.org/abs/1709.07871>
- [25] *MTDP: Modulated Transformer Diffusion Policy Model*, <https://arxiv.org/html/2502.09029v1>. Accessed: Apr. 26, 2026.
- [26] M. Reuss, J. Pari, P. Agrawal, and R. Lioutikov, *Efficient diffusion transformer policies with mixture of expert denoisers for multitask learning*, 2024. arXiv: 2412.12953 [cs.LG]. [Online]. Available: <https://arxiv.org/abs/2412.12953>
- [27] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, *Mastering diverse domains through world models*, *Nature*, 2025, 2023. arXiv: 2301.04104 [cs].
- [28] J. Quevedo, A. K. Sharma, Y. Sun, V. Suryavanshi, P. Liang, and S. Yang, *WorldGym: World model as an environment for policy evaluation*, Sep. 2025. DOI: 10.48550/arXiv.2506.00613 arXiv: 2506.00613 [cs].
- [29] Q. Bu et al., *Closed-loop visuomotor control with generative expectation for robotic manipulation*, Oct. 2024. DOI: 10.48550/arXiv.2409.09016 arXiv: 2409.09016 [cs].
- [30] S. Nair, A. Rajeswaran, V. Kumar, C. Finn, and A. Gupta, “R3M: A universal visual representation for robot manipulation,” in *Conference on Robot Learning (CoRL)*, 2022. arXiv: 2203.12601 [cs].
- [31] A. Brohan et al., *RT-1: Robotics transformer for real-world control at scale*, 2022. arXiv: 2212.06817 [cs.R0]. [Online]. Available: <https://arxiv.org/abs/2212.06817>
- [32] M. Assran et al., *V-JEPA 2: Self-supervised video models enable understanding, prediction and planning*, Jun. 2025. DOI: 10.48550/arXiv.2506.09985 arXiv: 2506.09985 [cs].
- [33] N. Tsagkas, A. Sochopoulos, D. Danier, C. X. Lu, and O. Mac Aodha, *When pre-trained visual representations fall short: Limitations in visuo-motor robot learning*, 2025. arXiv: 2502.03270 [cs].
- [34] R. G. Goswami et al., *World models can leverage human videos for dexterous manipulation*, Dec. 2025. DOI: 10.48550/arXiv.2512.13644 arXiv: 2512.13644 [cs].

- [35] L. Maes, Q. Le Lidec, D. Scieur, Y. LeCun, and R. Balestriero, *Leworldmodel: Stable end-to-end joint-embedding predictive architecture from pixels*, Verified from local PDF metadata and first page, Mar. 2026. arXiv: 2603.19312 [cs.LG].
- [36] G. Zhou, H. Pan, Y. LeCun, and L. Pinto, *DINO-WM: World models on pre-trained visual features enable zero-shot planning*, 2024. arXiv: 2411.04983 [cs.LG]. [Online]. Available: <https://arxiv.org/abs/2411.04983>
- [37] W. Huang et al., *Pointworld: Scaling 3d world models for in-the-wild robotic manipulation*, 2026. arXiv: 2601.03782 [cs.R0]. [Online]. Available: <https://arxiv.org/abs/2601.03782>
- [38] N. Hansen, H. Su, and X. Wang, "TD-MPC2: Scalable, robust world models for continuous control," in *International Conference on Learning Representations (ICLR)*, 2024. arXiv: 2310.16828 [cs].
- [39] T. Yu et al., "MOPO: Model-based offline policy optimization," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv: 2005.13239 [cs].
- [40] R. Kidambi, A. Rajeswaran, P. Netrapalli, and T. Joachims, "MOREL: Model-based offline reinforcement learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv: 2005.05951 [cs].
- [41] A. Argenson and G. Dulac-Arnold, "Model-based offline planning," in *International Conference on Learning Representations (ICLR)*, 2021. arXiv: 2008.05556 [cs].
- [42] Z. Liu, T. Swann, S. Garg, S. Levine, and A. Kumar, *Model-based offline planning with trajectory pruning*, 2021. arXiv: 2105.07351 [cs].
- [43] X. Li et al., *Evaluating real-world robot manipulation policies in simulation*, CoRL 2024, 2024. arXiv: 2405.05941 [cs].
- [44] Y. Li, Y. Zhu, J. Wen, C. Shen, and Y. Xu, *WorldEval: World model as real-world robot policies evaluator*, 2025. arXiv: 2505.19017 [cs].
- [45] R. Hoque, P. Huang, D. J. Yoon, M. Sivapurapu, and J. Zhang, *EgoDex: Learning dexterous manipulation from large-scale egocentric video*, ICLR 2026, 2025. DOI: 10.48550/arXiv.2505.11709 arXiv: 2505.11709 [cs].
- [46] T. Xiao, I. Radosavovic, T. Darrell, and J. Malik, *Masked visual pre-training for motor control*, 2022. arXiv: 2203.06173 [cs].

Appendix A Auxiliary Implemented Modules

This appendix describes the two auxiliary modules that are fully implemented in the WACT stack but were excluded from the final fair benchmark. The semantic alignment pipeline and the object grounding stage each produce useful artefacts and remain integrated with the broader system. Their absence from the main results chapters reflects a deliberate methodological choice rather than a finding about their utility or correctness.

A.1 Rationale for exclusion from the main benchmark

The canonical benchmark evaluates two world model families under a shared input contract: the same raw RGB observation frame and the same bimanual wrist state sequence are provided to both V-JEPA2-AC and LeWM. No additional conditioning is admitted on either side.

Semantic alignment artefacts (captions, tags, text embeddings) and object grounding artefacts (bounding boxes, segmentation masks) could in principle be incorporated into the V-JEPA2-AC scoring pipeline through the semantic extraction path. The LeWM architecture, however, has no equivalent semantic conditioning module. Including semantic or grounding conditioning for V-JEPA2-AC alone would change the nature of the comparison from a world model family comparison under equal information to a mixed comparison that conflates world model architecture with access to additional input channels. The two scientific questions are distinct, and the thesis addresses only the former.

The exclusion applies symmetrically across all five evaluated branch instances. Both the V-JEPA2-AC family and the LeWM family run without semantic or grounding inputs in the final benchmark. Future work under a symmetric design, in which both families receive equivalent semantic conditioning, could isolate the effect of this information channel independently of the model family comparison.

A.2 Semantic alignment pipeline

A.2.1 Purpose and outputs

The semantic alignment stage converts indexed EgoDex windows into structured semantic artefacts for downstream conditioning and inspection. For each window, the pipeline produces:

- A natural language caption describing the observable manipulation activity in the selected frames.
- Structured tags organised into categories, including the contact object, secondary objects, and observed actions.
- A canonical semantic text string constructed from the structured tags, suitable for embedding.
- A dense text embedding vector produced by a configurable embedding backend.

The primary downstream use is to provide a semantically meaningful conditioning vector that can be concatenated with or appended to the visual latent representations used by the world model scoring functions. In the current implementation, this conditioning path is available but is not activated in the canonical benchmark.

A.2.2 Visual language model selection

The VLM component was evaluated across several model and embedding combinations before the operational configuration was stabilised. The semantic model comparison table records these evaluations and their relative quality on caption coherence and tag precision. The operationally selected configuration uses Qwen/Qwen3-VL-2B-Instruct as the visual language model for caption and tag generation, with the google/siglip-so400m-patch14-384 model as the default text embedding backend (`WACT_EMBED_BACKEND=siglip_text`).

The Qwen3-VL-2B-Instruct model was selected for practical reasons: it is compact enough to run on a single GPU alongside other pipeline stages, it produces grammatically coherent captions from single-frame inputs, and its structured tag output is consistent enough to support deterministic canonical text construction. The SigLIP embedding backend was selected because it provides a joint image-text embedding space that is directly compatible with the V-JEPA2-AC semantic scoring path, allowing future similarity computations between caption embeddings and visual latents without an additional projection layer.

OpenCLIP and VLM-text hidden-state backends are also implemented and configurable. The SigLIP-backed Qwen3-VL run represents the strongest row in the comparison table for the current stack configuration.

A.2.3 Critic and retry loop

A key engineering feature of the semantic pipeline is an integrated critic and retry loop. The naive approach of captioning once and trusting the output produces unreliable

results on EgoDex frames, which are first-person egocentric views with partially occluded hands, cluttered surfaces, and rapidly changing visual content. The critic module scores the resulting caption against the input frames using a configurable similarity measure. If the score falls below a threshold or plateaus across retries, the pipeline either retries the VLM call with adjusted prompting parameters or marks the row as `semantic_failure=true` and writes it with a failure flag.

This design means that the semantic stage produces a complete, aligned output for every window in the index, with explicit quality flags rather than silently dropping or corrupting rows on difficult windows. The failure-flagged rows remain usable for inspection and can be filtered before any downstream conditioning step.

The implemented retry logic is:

1. Score the caption against the sampled frame using the configured score mode.
2. If the score exceeds the acceptance threshold, commit the row.
3. If the score is below threshold, retry with a revised prompt up to the configured maximum number of retries.
4. If all retries fail or the score plateaus without improvement, commit the row with `semantic_failure=true`.

A.2.4 Multi-GPU sharding and output artefacts

The semantic pipeline supports multi-GPU sharded execution through a shard supervisor (`semantic_shard_supervisor.py`) and a deterministic shard merge step (`semantic_merge_shards.py`). This is necessary because captioning and embedding 7,607 test windows or tens of thousands of training windows at VLM inference speeds is infeasible on a single GPU within a reasonable run budget.

The canonical merged output layout is:

```
/runs/w-act/egodex/semantic/<split>/<model_id>/<build_id>/
  embeddings.npy
  window_ids.txt
  meta.jsonl
  tags.jsonl
  captions.jsonl
  manifest.json
  index.npz
```

During a multi-GPU run, intermediate shard directories (`shard_0/`, `shard_1/`, and so on) accumulate partial outputs. The merge step validates row counts before

producing the canonical root artefact set, ensuring that partial or interrupted shard outputs are not silently used as complete builds. The manifest records the input fingerprint, allowing downstream stages to detect cache invalidation if the source windows file is replaced.

A.2.5 Integration with the broader pipeline

The semantic stage is positioned between the data indexing stage and the object grounding stage. The OPE stage (Section A.3) reads semantic manifests to obtain the canonical window ordering and to construct object prompts from the semantic tags produced by the VLM. This tight coupling is intentional: it ensures that semantic and grounding artefacts are always aligned to the same window index, preventing silent join drift if the windows file is regenerated.

The entrypoint for the semantic stage is `src/wact/cli/semantic_build_index.py`, with orchestration handled by `scripts/semantic.sh` or the stage-runner wrapper `scripts/run_stage.sh semantic`. The notebook companion `notebooks/semantic_aligner.ipynb` provides interactive inspection of caption quality, tag distributions, and embedding nearest-neighbour retrievals.

A.3 Object grounding and mask generation

A.3.1 Scope and current implementation

The object grounding stage takes semantic artefacts from the previous stage and produces spatially grounded detection and segmentation outputs for each window. The current implementation provides two-dimensional bounding-box detection and box-conditioned mask generation. It does not implement canonical six degree-of-freedom pose estimation, FoundationPose-style geometry reconstruction, or object identity tracking across frames. The name “object pose estimation” in the codebase reflects the intended scope of a future extension rather than the current capability.

For each window, the implemented pipeline:

1. Reads the semantic row for the window from the upstream semantic build.
2. Samples one or more frames from the window according to the configured frame sampling strategy.
3. Constructs object prompts from the structured semantic tags.

4. Runs the configured detector to produce bounding boxes.
5. Runs the configured masker on each detected box to produce segmentation masks.
6. Writes aligned detection and mask artefacts alongside a progress checkpoint.

A.3.2 Prompt construction

Object prompts are constructed from semantic tags using a priority scheme implemented in `src/wact/ope/prompting.py`. The logic prioritises the `contact_object` tag, which the VLM assigns to the primary object being actively manipulated, and then appends deduplicated entries from the `objects` tag list. Generic labels such as “object”, “thing”, “tool”, and “hand” are filtered before the prompt is passed to the detector, as these terms produce unreliable or overspecific detections in the grounding model. When the VLM returns no object tags, the prompt construction falls back to a task-derived hint derived from the EgoDex task label.

This is category-level prompting rather than instance-level prompting. The detector does not distinguish between two instances of the same object class, and the resulting masks identify regions containing the prompted category rather than a specific tracked object instance.

A.3.3 Detection and segmentation backends

The production backends, configured through `configs/stage_defaults.env` and accessible via the stage-runner, are:

- Detector: `IDEA-Research/grounding-dino-base` (Grounding DINO), a zero-shot open-vocabulary detector that accepts free-text prompts.
- Masker: `facebook/sam2-hiera-large` (Segment Anything Model 2), a promptable segmentation model that produces high-quality binary masks from bounding box inputs.

The combination of Grounding DINO for detection and SAM2 for box-conditioned segmentation is a well-established zero-shot grounding pipeline. It does not require object-specific training data from EgoDex and generalises across the 194 task categories in the dataset without fine-tuning.

The bare `scripts/ope.sh` script defaults to dummy backends for development and debugging. Production runs must use the stage-runner or explicitly override the backend flags to select the real Grounding DINO and SAM2 models.

A.3.4 Output artefacts and alignment

The canonical output layout is:

```
/runs/w-act/egodex/pose/<split>/<pipeline_id>/<build_id>/
  masks.jsonl
  detections.jsonl
  window_ids.txt
  errors.jsonl
  progress.json
  manifest.json
```

The `masks.jsonl` file is the primary downstream artefact: it contains a row per window with the detected object categories, bounding box coordinates, and run-length-encoded or polygon-format segmentation masks. The `detections.jsonl` file records the raw detection outputs before masking and is retained for audit purposes. The `progress.json` file supports resumable execution through the shard supervisor.

Alignment between the OPE output and upstream artefacts is enforced by inheriting the `window_ids.txt` ordering from the semantic build. The OPE stage reads the semantic manifest to resolve the windows pickle that was used to produce the semantic index, ensuring that detection and mask rows correspond to the same window boundaries used at every earlier pipeline stage.

A.3.5 Known limitations

The current implementation has several limitations that are relevant to any future use of OPE outputs for manipulation conditioning:

1. The 2D masks identify spatial regions in the image frame but do not provide three-dimensional geometry information. Object depth, relative pose, or spatial relationship between manipulated objects are not recoverable from the mask output alone.
2. Small or visually similar objects can confuse the Grounding DINO detector, particularly in the cluttered egocentric field of view typical of EgoDex demonstrations. When two similar objects are present, the detector may merge them or select the wrong instance.
3. There is no temporal tracking across frames within a window. Each sampled frame is processed independently, and there is no mechanism to maintain object identity across the prediction horizon.

4. The zero-shot grounding approach depends on the quality of the upstream semantic tag output. A missed or misidentified contact object in the semantic stage will propagate to an incorrect or absent detection in the OPE stage.

These limitations are consistent with the current scope of two-dimensional grounding and mask generation. They would need to be addressed before OPE outputs could serve as reliable conditioning for world model scoring in a deployed system.

A.4 Future use under a symmetric design

The exclusion of semantics and object grounding from the final benchmark is a fairness decision, not a statement that these modules lack value. Both pipelines are implemented, tested on the training split, and produce aligned artefacts that are stored alongside the canonical benchmark outputs.

The most natural future experiment would introduce semantic conditioning symmetrically across both world model families. This would require either adapting the LeWM architecture to accept a text embedding as an additional conditioning channel, or training a new LeWM variant with semantic inputs from the outset. Under a symmetric design, the contribution of semantic conditioning to world model scoring quality could be measured independently of the model family comparison, and the effect of caption quality, embedding backend selection, and object grounding granularity could each be evaluated in isolation.

Similarly, the object grounding masks could be used to produce spatially localised visual features by masking the RGB input before the visual backbone processes it. Whether this localisation improves world model guidance in the manipulation setting, or whether the pretrained backbone’s attention mechanism already implicitly grounds relevant objects without explicit masking, is an open empirical question.

[FIGURE PLACEHOLDER: semantic pipeline overview – input window to VLM to caption plus tags to embedding backend to output artefacts; critic/retry loop shown as a feedback arc between caption and score threshold]

[FIGURE PLACEHOLDER: object grounding example – sampled EgoDex frame with Grounding DINO bounding box and SAM2 mask overlaid on detected contact object; prompt text shown]

Appendix B Benchmark Evolution and Validity Record

This appendix documents the development trajectory of the WACT benchmark and the validity concerns that were identified and resolved before the final canonical results package was established. Its purpose is to provide an auditable record of the decisions and corrections that shaped the final evaluation surface, and to explain why certain earlier result sets were superseded or rejected.

The record is presented in roughly chronological order of discovery. Each section identifies the concern, describes its effect on the results as reported at the time, and explains the correction applied.

B.1 The three benchmark eras

The dissertation results passed through three distinct evaluation eras before the canonical package was established. Understanding the transition between eras is necessary for interpreting references to earlier result numbers in project documentation.

Era 1: Bounded pilot slices. The earliest quantitative comparisons were conducted on small, handpicked windows selected to represent the diversity of EgoDex task categories. These were useful for rapid controller iteration and for identifying which architectural choices were clearly harmful, but they could not support generalisable claims because the selection was not held-out from the development process. All Era 1 results were explicitly demoted from the thesis claims surface once full held-out evaluation became practical.

Era 2: March 2026 frozen three-branch bundle with limited LeWM coverage. The first evaluation to use the held-out test split was the March 2026 frozen package. This package provided valid Branch 1 and Branch 2 V-JEPA2-AC results, but its LeWM coverage was severely limited by a frame cap configuration error (described in Section B.2). The Branch 3 selection mode error (described in Section B.3) also meant that Branch 3 V-JEPA2-AC was not independently evaluated. Both defects required correction before the three-branch comparison could be treated as scientifically valid.

Era 3: Canonical final package (April 2026). The April 14, 2026 package achieved full 7,607-window coverage for the LeWM family after the frame cap fix. A further correction on April 18, 2026 re-ran Branch 3 V-JEPA2-AC with the correct `selection_mode=world_model_only` flag, producing the first genuine Branch 3 evaluation. The `train_part1_2_3` training split was used for all five branch instances. This package is the canonical results surface for all quantitative claims in the thesis.

B.2 LeWM frame cap defect

B.2.1 Discovery

The LeWM export pipeline requires a `max_frames_per_episode` parameter that controls how many frames from each EgoDex episode are included when building the training corpus. This parameter has a direct impact on the coverage of the resulting evaluation surface, because the guard can only evaluate windows whose constituent frames were included in the export.

During the March 2026 benchmark packaging, the frame cap was set to `max_frames_per_episode=256`. This value was not derived from the actual window length distribution of the benchmark window set; it was a development default carried over from early pipeline testing. When the benchmark results were inspected, the LeWM guard had produced predictions for only 160 of the 7,607 test windows, representing 2.1 per cent of the intended evaluation surface.

The coverage failure was identified by comparing the count of LeWM guard output rows against the expected window count. The root cause was confirmed by recomputing the maximum frame index required to cover all windows in the benchmark window list and finding that the actual maximum was 5,280 frames, far above the 256-frame cap in use.

B.2.2 Effect on reported results

Any LeWM result reported from the March 2026 package is based on 160 windows, not 7,607. The 160 covered windows are not a random sample of the test split; they are the windows whose episodes happened to fall within the 256-frame cap. These windows are concentrated in episodes with early-episode actions and do not represent the full range of task categories or annotation confidence levels in the test split.

The practical effect is that the March 2026 LeWM numbers cannot be compared directly to Branch 1 or Branch 2 V-JEPA2-AC numbers, because the LeWM evaluation is on a different, non-representative slice of the test data. Any conclusion drawn from these numbers about the relative merit of the LeWM family would be confounded by the coverage difference.

B.2.3 Correction

The frame cap was recomputed by iterating over the benchmark window list, reading the episode frame indices for each window, and taking the maximum observed frame index across all windows. The resulting cap of 5,280 frames was applied to the export

rebuild. The LeWM model was not retrained; only the export was rebuilt with the corrected parameter. The guard was then re-run, achieving full 7,607-window coverage on the April 14, 2026 package.

The corrected LeWM numbers in the canonical package are therefore not comparable to the March 2026 LeWM numbers on a like-for-like basis. The March 2026 numbers are retained in project documentation for audit purposes but are not cited in any thesis claims.

B.3 Branch 3 selection mode defect

B.3.1 Discovery

The three-branch design requires Branch 3 to use `selection_mode=world_model_only`, meaning that only the world model's predicted latent distance ranks the candidates, without any contribution from the ACT policy's own scoring function. This is the experimental condition that tests whether the world model alone, without policy guidance in the selection step, can identify better trajectories.

After the April 14, 2026 package was assembled, the per-window pairwise comparison between Branch 3 V-JEPA2-AC and Branch 2 V-JEPA2-AC was computed as a validation check. The result was a 0.0 per cent win rate for Branch 3 against Branch 2 across all 7,607 windows. That is, Branch 3 did not produce a single window outcome that differed from Branch 2.

This was immediately identified as inconsistent with any genuine difference in selection logic between the two branches. Inspection of the Branch 3 guard manifest revealed that the selection mode was recorded as `best_score`, not `world_model_only`. The chain evaluation script (`chain_eval_part1_2_3.sh`) had been written to launch both Branch 2 and Branch 3 guards, but the Branch 3 launch did not pass the `--selection-mode world_model_only` flag. Both guards therefore ran with `selection_mode=best_score` on the same frozen candidate pool, producing identical outputs.

B.3.2 Effect on reported results

Until the defect was corrected, the "Branch 3 V-JEPA2-AC" entry in the April 14 package was scientifically identical to Branch 2 V-JEPA2-AC. The three-branch comparison was effectively a two-branch comparison. Any conclusion about whether world-model-only selection outperforms hybrid selection under V-JEPA2-AC could not be drawn from this data.

The defect was non-obvious from the output metrics alone. The reported Branch 3 translation error (0.0354, identical to Branch 2) was not implausible; one might expect Branch 3 and Branch 2 to produce similar numbers if the world model’s ranking correlates with the combined score’s ranking. The defect was only detectable through pairwise comparison, which revealed the zero win rate that is impossible if the two branches genuinely differ.

B.3.3 Correction

The Branch 3 V-JEPA2-AC guard was re-run on April 18, 2026 with the correct flag:

```
--selection-mode world_model_only
--candidate-source frozen_manifest
--candidate-manifest-run-dir <B2_JEPA_run_dir>
```

Using the frozen candidate pool from Branch 2 ensured that the Branch 3 result reflects only the difference in selection logic, not any difference in the candidates available. Guard manifest verification after the run confirmed that all 7,607 windows were decided by the world model (`override_reason_counts: {world_model_only_best_non_act: 7607}`), confirming that the corrected selection mode was in effect for every window.

The corrected Branch 3 V-JEPA2-AC result (translation $L_2 = 0.0328$, 91.5 per cent win rate against Branch 1) is the first genuine evaluation of world-model-only selection in the WACT stack.

B.4 Oracle goal leakage

B.4.1 Discovery

In early V-JEPA2-AC guard configurations, the world model scoring function was provided with the terminal latent state of the evaluation episode as a goal signal. The scoring function measured the predicted final latent distance to this goal and used it as part of the candidate ranking. This goal signal was derived from the ground-truth future frames of the demonstration, information that would not be available in a deployed system.

The oracle goal leak was identified during the fair comparison audit when it was noted that the V-JEPA2-AC guard configuration differed from the ACT-only baseline in the information it received at evaluation time. The ACT policy generates candidates from the current observation alone and has no access to future frames. Providing the world model with the ground-truth terminal latent created an asymmetry: the world

model scoring function was effectively guided toward candidates that reached the demonstrated endpoint, while the ACT policy’s own ranking had no equivalent oracle signal.

B.4.2 Effect on reported results

Any result from a guard run using oracle goal conditioning overstates the benefit of world model guidance, because part of the improvement is attributable to the oracle signal rather than to the world model’s intrinsic representation quality. The magnitude of the overstatement cannot be precisely determined from the available data, but it is substantial enough to make oracle-conditioned results non-comparable to policy-only results.

Early exploration of V-JEPA2-AC guidance showed results that were more impressive than the final canonical numbers. The oracle goal was a contributing factor. Removing the oracle signal and rerunning under the fair contract produced lower but more defensible numbers. The final canonical results (77.0 per cent reduction in mean translation error) are therefore conservative relative to the earliest reported V-JEPA2-AC results.

B.4.3 Correction

The fair evaluation contract was established by setting `goal_source=none` for all branch evaluations. Under this setting, the world model scoring function does not receive any goal latent derived from future frames. The candidate ranking is based solely on the world model’s predicted latent dynamics given the observed context and the proposed action chunk.

This correction was applied to all five branch instances in the canonical benchmark. The change from oracle goal to no goal required validating that the world model scoring function still produced meaningful candidate rankings without the terminal goal signal, which was confirmed by the final translation improvement results.

B.5 Goal metric invalidity

B.5.1 Discovery

The V-JEPA2-AC guard natively reports a metric called `goal_12` alongside the translation error. In the original implementation, this metric stored different quantities in Branch 1 and Branches 2 and 3. In Branch 1, it stored the translation L_2 error for the

selected candidate. In Branches 2 and 3, it stored the JEPA latent-space distance between the predicted final state and the goal latent.

When `goal_source=none` was adopted as the fair evaluation contract, there was no longer a meaningful goal latent for Branches 2 and 3. The `goal_12` metric for these branches consequently reported a latent distance of zero for every window, making the metric uninformative. Comparing `goal_12` across branches was therefore comparing a real translation error in Branch 1 against a vacuous zero in Branches 2 and 3.

B.5.2 Correction

The `goal_12` metric was replaced with `endpoint_translation_12`, defined as the mean translation L_2 error computed over only the final quarter of the prediction window (prediction steps 12 through 16). This metric measures how accurately the predicted trajectory arrives at the correct final position, which is often more relevant to task success than early-window prediction accuracy. The endpoint metric is computed identically across all branch instances and is not affected by the presence or absence of a goal latent. It is included in the canonical benchmark shape metric tables reported in the results chapter.

B.6 LeWM CPU inference defect

B.6.1 Discovery

During the first full test-split guard run for LeWM after the frame cap fix, the estimated completion time for 7,607 windows was approximately 54 hours. This was implausibly slow relative to the expected GPU inference throughput for a 15-million-parameter model. Inspection of the guard inference loop revealed a hard-coded `.to("cpu")` call in the LeWM model forward pass that overrode any device specification passed at runtime. The model was running on CPU regardless of whether a CUDA device was available.

B.6.2 Correction

The hard-coded device assignment was replaced with a dynamic device detection that checks for CUDA availability and falls back to CPU only if no GPU is detected. After this correction, the LeWM guard ran at approximately 19.5 windows per second on a single GPU, completing the full 7,607-window test evaluation in roughly 6.5 minutes. The effective speedup was approximately 130-fold relative to the CPU-bound run.

This defect is worth recording because it illustrates a failure mode common in research code developed iteratively: a debugging convenience (forcing CPU to avoid CUDA memory errors during development) was never removed before production use. The correction required no changes to the model architecture or training, only to the device placement logic in the inference wrapper.

B.7 HDF5 handle management defect

B.7.1 Discovery

The data loading loop in an early version of the guard infrastructure opened and closed the HDF5 data file once per evaluation window. On a test set of 7,607 windows, this produced approximately 145,000 open-close cycles (one open and one close per window, plus frame-level re-opens within each window). The overhead from repeated file handle acquisition was a significant fraction of total evaluation wall time on large test splits.

B.7.2 Correction

The data loading implementation was refactored to open the HDF5 file once at the start of the evaluation run, hold the handle in memory, and close it only when the evaluation completes. This change reduced per-window I/O overhead to a small seek operation rather than a full open-close cycle and improved guard evaluation throughput noticeably on the full test split.

B.8 Guard resumability

B.8.1 Motivation

Full test-split guard runs covering 7,607 windows take several minutes on a GPU but can be interrupted by node preemption, out-of-memory errors, or external compute scheduler events. Without checkpointing, an interrupted run required restarting from window zero, discarding all partial results.

B.8.2 Implementation

A per-window JSONL checkpointing system was added to the guard infrastructure. After each window is evaluated, a row is appended to a checkpoint file in the output

directory. When a run is invoked with `--resume-dir` pointing to an existing output directory, the guard reads the checkpoint file, identifies which window IDs have already been completed, and skips them. The final manifest and aggregate metrics are computed from the complete set of per-window rows, including both the resumed and newly computed windows.

This design ensures that an interrupted run can be resumed without any manual intervention and without requiring the window evaluation order to be deterministic. The checkpoint is written atomically per row, so a crash mid-window produces at worst one incomplete row that the resume logic can detect and skip.

B.9 Rejected comparison inputs and superseded paths

B.9.1 Semantics and OPE as benchmark inputs

An early hypothesis held that incorporating semantic alignment and object grounding into the V-JEPA2-AC scoring function would produce the most informative world model guidance, and that this enriched V-JEPA2-AC configuration should be the primary benchmark representative for the V-JEPA2-AC family. This hypothesis was rejected for the reasons described in Appendix A and in Section 8.6 of the discussion chapter: the enriched configuration gives V-JEPA2-AC privileged inputs not available to LeWM, confounding the family comparison.

The decision to reject this path was not based on the performance of the enriched configuration. The lean RGB-plus-trajectory contract was adopted because it is the only contract that supports a scientifically unconfounded comparison between the two world model families.

B.9.2 Veto and threshold selectors

Before the current scoring function was stabilised, several candidate selection mechanisms based on explicit threshold vetos were implemented and evaluated. These included the `act_first_veto` design, which rejected world model candidates whose support distance exceeded a fixed threshold, and the `veto_m10` selector, which applied a margin-based veto to prevent the world model from selecting candidates that were far from the ACT policy’s preferred candidate even if they scored well on the world model metric alone.

These mechanisms were useful for early exploration of the policy-guard interface and helped establish that naive world model selection could be dominated by retrieval artefacts when the candidate pool contained low-quality outliers. They were superseded by the structured scoring function in Branch 2 (`best_score`: combined

latent distance, action out-of-distribution penalty, and support distance term) and the clean world-model-only design in Branch 3. Neither the veto selectors nor the margin-based threshold approaches are used in the canonical benchmark.

B.9.3 Early Branch 3 shells

Several experimental Branch 3 designs preceded the final frozen representative. The first explicit world-model-only shell for LeWM was evaluated on the fair held-out slice and was found to be weaker than the ACT-only baseline on aggregate slice metrics, showing that world-model-only selection requires a world model of sufficient quality before it improves over the policy alone. Neutral-seed variants of Branch 3, which removed the ACT seed from the candidate pool, also collapsed on the fair slice, demonstrating that the bottleneck was the controller structure rather than any single parameter choice. These negative results were important for understanding the design space and for eventually arriving at a Branch 3 configuration that genuinely improves over Branch 2 in aggregate.

B.9.4 Conservative confidence gating

A further rejected path explored whether applying more conservative confidence gating to the LeWM guard, requiring higher override margins before the world model was permitted to override the policy candidate, could make LeWM competitive with V-JEPA2-AC on the shared benchmark. The bounded sweep showed that increasing the override margin made LeWM more conservative but did not produce a competitive replacement for V-JEPA2-AC on chunk-level comparison metrics. This was recorded as a useful negative result: the performance gap between the two world model families is not primarily explained by calibration of the override threshold, and therefore cannot be closed by threshold tuning alone.

B.10 Summary of validity status

Table B.1 summarises the seven validity concerns identified during the benchmark development cycle and the resolution status of each at the time the canonical package was assembled.

Concerns 1, 3, and 6 were resolved by changes to the benchmark artefacts or analysis code. Concern 2 is addressed by consistent framing throughout the thesis: translation L_2 is positioned as an offline diagnostic proxy rather than a task success predictor, and the connection to deployed outcomes is explicitly left as future work. Concern 4 is partially addressed by the per-task win rate distribution reported in the

Table B.1 Summary of benchmark validity concerns and their resolution status in the canonical final package.

#	Concern	Severity	Status in canonical package
1	Branch 3 V-JEPA2-AC identity defect (<code>selection_mode</code> bug)	Critical	Resolved: re-run April 18, 2026
2	<code>translation_12</code> not paper-backed as primary metric	Moderate	Documented: framed as offline diagnostic proxy
3	Shape metrics suppressed from reported summary	Moderate	Resolved: shape metrics added to results chapter
4	No per-task win rate breakdown	Moderate	Partially resolved: aggregate confirmed broad; worst-5 reported
5	Wrist-only evaluation scope	Low–Moderate	Acknowledged design constraint; documented in limitations
6	Confidence scores unused in evaluation	Low	Resolved: stratification analysis added to results
7	Offline-to-online gap	Moderate	Acknowledged: explicitly bounded in scope section

results chapter, which confirms that improvement is broad and not driven by a small number of tasks. Concerns 5 and 7 are acknowledged design constraints that are transparent in the scope of conclusions and limitations sections.

The final canonical results are drawn from a benchmark in which all critical and moderate concerns have been either resolved or explicitly bounded. No result cited in the thesis claims chapter relies on a defective or uncorrected evaluation artefact.