


Learning to Select Actions for Complex Hand Tasks: A Comparative Study of Visuomotor Policies and Latent World Models

Graham Pellegrini

Supervisor: Prof. Alexiei Dingli

[PLACEHOLDER: Submission Date]

*Submitted in partial fulfilment of the requirements
for the degree of Master of Science in Artificial Intelligence.*



L-Università ta' Malta

Faculty of Information &
Communication Technology

Plagiarism Declaration

I, the undersigned, hereby declare that the work contained in this dissertation is my own original work and that I have not previously in its entirety or in part submitted it at any university for a degree.

I declare that this submission is my own work and that, to the best of my knowledge and belief, it contains no material previously published or written by another person, except where due acknowledgement is made in the text.

I declare that the use of artificial intelligence tools in the preparation of this dissertation was limited to writing assistance, code suggestions, and literature search support. All research design, experimental decisions, implementation choices, result interpretation, and academic claims are my own.

Signature: _____

Name: Graham Pellegrini

Date: _____

Abstract

This dissertation investigates latent world models: learned models that predict how a scene will evolve in a compact representation space rather than in raw pixels. World models are a young and rapidly developing area of machine learning, and their promise for robotics lies in judging a proposed action by its predicted consequences before it is executed. The guiding idea behind this work casts the imitation policy as an actor that proposes future motions and the world model as a brain that evaluates them, so that the pair selects better actions than the policy alone.

The research question is whether this pairing improves the selection of future bimanual hand motions under a controlled offline benchmark. The study deliberately favours quality of evidence over quantity of tests. One pairing is established as a clean, fully controlled positive result, and every other implemented route is followed far enough to explain in detail why it falls short of the same standard.

The benchmark evaluates 7,607 held-out windows from the EgoDex egocentric manipulation dataset. Each window pairs 16 observed frames with a 16-frame prediction target, and actions are represented as 18-dimensional bimanual wrist-motion chunks. Policy and world model are interfaced through three selection branches: the policy alone, world-model scoring with the policy's preferred candidate retained, and world-model scoring with that candidate excluded. In the strongest branch, Diffusion Policy candidates scored by V-JEPA2-AC reduce first-step translation error by 6.62 per cent and rotation error by 1.08 degrees relative to the policy alone. The remaining routes (ACT, DexWM, LeWM, DINO-WM, and JEPA-WM) return negative or bounded outcomes, each traced to a concrete caveat such as candidate collapse or an unavailable pretrained checkpoint.

Beyond the headline numbers, the dissertation contributes a reproducible protocol for separating genuine world-model guidance from misleading shortcuts, and a documented account of what open model releases must provide before their claims can be independently tested. All conclusions concern offline wrist-trajectory prediction on recorded demonstrations.

Acknowledgements

Placeholder for acknowledgements.

Contents

Plagiarism Declaration	i
Abstract	ii
Acknowledgements	iii
Contents	viii
List of Figures	xi
List of Tables	xiii
List of Abbreviations	xiv
List of Notations	xvi
1 Introduction	1
1.1 Motivation	2
1.2 Research Question and Scope	2
1.3 Contributions	4
1.4 Dissertation Structure	4
2 Background	6
2.1 Visuomotor policies	6
2.1.1 Behaviour cloning and compounding error	7
2.1.2 Action chunking	7
2.2 Rigid body pose and the $SE(3)$ representation	8
2.3 Latent representations	9
2.4 Conditional variational autoencoders	10
2.5 Transformer networks and the attention mechanism	12
2.6 Self-supervised learning	13
3 Literature Review	15
3.1 Action Chunking with Transformers (ACT)	17

3.2	Diffusion Policy	18
3.3	World models	19
3.3.1	World models as policy evaluators	20
3.3.2	Latent predictive architectures	21
3.4	Offline evaluation	27
4	Methodology	30
4.1	Dataset and Experimental Setup	30
4.1.1	EgoDex: source and selection rationale	30
4.1.2	Windowing protocol	31
4.1.3	Training splits and test set	32
4.2	Three-Branch Experimental Framework	33
4.3	Policy Variants	36
4.4	World Model Families	36
4.5	Fair Experimental Protocol	38
4.5.1	Training exposure budget	38
4.5.2	Evaluation contract	39
4.6	Evaluation Metrics	40
4.6.1	Offline trajectory prediction setting	40
4.6.2	Primary metrics	40
4.6.3	Candidate selection comparison	41
4.7	Evaluation Scope and Limitations	41
5	Implementation	43
5.1	Repository and execution model	43
5.2	Data pipeline and windowing	45
5.2.1	Index construction	45
5.2.2	Incremental training curriculum	45
5.3	Training configuration	45
5.4	Policy training	47
5.5	World model training	51
5.5.1	V-JEPA2-AC: extraction and transition head	51
5.5.2	LeWM: end-to-end training	52
5.5.3	DexWM: from-scratch training	53
5.5.4	DINO-WM and JEPA-WM exploratory selectors	54
5.5.5	DiWA/CALVIN reproduction infrastructure	55
5.5.6	V-JEPA2 (no action conditioning): ablation control	55
5.5.7	PointWorld: action-degradation control	56
5.6	Convergence study and universal training budget	56

5.6.1	Loss curves and plateau detection	57
5.6.2	Universal budget decision	57
5.7	B2 margin threshold calibration	59
5.8	Guard evaluation	59
5.8.1	Branch 1: policy-only baseline	60
5.8.2	Branch 2: world-model-inclusive selection	60
5.8.3	Branch 3: world-model-exclusive selection	60
5.8.4	Manifest and checkpoint design	60
5.9	Evaluation safeguards	60
5.10	Implementation challenges and resolutions	61
6	Results	63
6.1	Overview of the frozen evidence	63
6.2	Diffusion Policy guard results	64
6.3	ACT guard results	65
6.4	DexWM and LeWM closeout	66
6.5	Overall interpretation	67
7	Discussion and Limitations	69
7.1	What the benchmark shows	69
7.2	Why Branch 2 is stronger than Branch 3	69
7.3	Interpreting the model-family differences	70
7.4	Reproducibility and open model releases	70
7.5	Why benchmark evolution is part of the contribution	71
7.6	Why semantic alignment and object grounding were excluded	71
7.7	Main limitations	72
8	Conclusion	73
8.1	Summary of findings	73
8.2	Methodological contribution	73
8.3	Reproducibility lesson	74
8.4	Scope of conclusions	74
8.5	Future Work	74
A	Auxiliary Implemented Modules	80
A.1	Rationale for exclusion from the main benchmark	80
A.2	Semantic alignment pipeline	80
A.2.1	Purpose and outputs	80
A.2.2	Visual language model selection	81
A.2.3	Critic and retry loop	81

A.2.4	Multi-GPU sharding and output artefacts	82
A.2.5	Integration with the broader pipeline	83
A.3	Object grounding and mask generation	83
A.3.1	Scope and current implementation	83
A.3.2	Prompt construction	84
A.3.3	Detection and segmentation backends	84
A.3.4	Output artefacts and alignment	85
A.3.5	Known limitations	85
A.4	Future use under a symmetric design	86
B	Benchmark Evolution and Validity Record	87
B.1	The three benchmark eras	87
B.2	LeWM frame cap defect	88
B.2.1	Discovery	88
B.2.2	Effect on reported results	88
B.2.3	Correction	88
B.3	Branch 3 selection mode defect	89
B.3.1	Discovery	89
B.3.2	Effect on reported results	89
B.3.3	Correction	90
B.4	Oracle goal leakage	90
B.4.1	Discovery	90
B.4.2	Effect on reported results	91
B.4.3	Correction	91
B.5	Goal metric invalidity	91
B.5.1	Discovery	91
B.5.2	Correction	92
B.6	LeWM CPU inference defect	92
B.6.1	Discovery	92
B.6.2	Correction	92
B.7	HDF5 handle management defect	93
B.7.1	Discovery	93
B.7.2	Correction	93
B.8	Guard resumability	93
B.8.1	Motivation	93
B.8.2	Implementation	93
B.9	Rejected comparison inputs and superseded paths	94
B.9.1	Semantics and OPE as benchmark inputs	94
B.9.2	Veto and threshold selectors	94

B.9.3	Early Branch 3 shells	95
B.9.4	Conservative confidence gating	95
B.10	Summary of validity status	95
C	Auxiliary Implementation Modules	97
D	Notes for Future Revision	98
D.1	Strategy 2 comparative evaluation	98

List of Figures

Figure 2.1	A visuomotor policy maps an EgoDex-style camera frame to future bimanual wrist actions in the thesis setting.	6
Figure 2.2	Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [8].	7
Figure 2.3	Translation for each arm can be represented by a movement along each of the three spatial axis [12].	8
Figure 2.4	Rotation around the y-axis for each arm. The same circular movement can be imagined around the x and z axes; the total orientation of the wrist is the combined result of rotations around all three [12].	9
Figure 2.5	Example latent-space plot showing how structurally similar items cluster together in the learned representation. Structural similarity here refers to items that share shape, geometry, and likely physical behaviour [14].	10
Figure 2.6	Standard Variational Autoencoder (VAE) (a) and conditional Conditional Variational Autoencoder (CVAE) (b) encoder-decoder pipelines [17]. . .	11
Figure 2.7	VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation) [16].	11
Figure 2.8	Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [18].	12
Figure 2.9	Three machine learning paradigms: supervised learning (explicit labels), reinforcement learning (reward signal), and self-supervised learning (unlabelled data provides its own supervision signal). The third panel is labelled unsupervised learning in the source figure; this thesis uses the narrower term self-supervised learning for prediction tasks constructed from the data itself.	13
Figure 3.1	Architecture of the Action Chunking with Transformers (ACT) policy [7].	17
Figure 3.2	Diffusion Policy: general formulation (a), Convolutional Neural Network (CNN) variant (b), transformer variant (c) [9].	19
Figure 3.3	Classic closed-loop control (a) versus CLOVER's closed-loop visuomotor control (b) [31].	20

Figure 3.4	Three Self-Supervised Learning (SSL) prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [19].	21
Figure 3.5	Video Joint Embedding Predictive Architecture 2 (V-JEPA2) post-training pathways from internet-scale backbone to specialised applications [20]. . . .	22
Figure 3.6	Standard V-JEPA2 (left) versus action-conditioned Action-Conditioned Video Joint Embedding Predictive Architecture 2 (V-JEPA2-AC) (right) [20]. .	24
Figure 3.7	High-level layout of the DexWM world model pipeline, illustrating how predicted keypoints are incorporated alongside the encoder, Conditional Diffusion Transformer (CDiT) predictor, and decoder stages [21].	25
Figure 3.8	LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [22].	25
Figure 4.1	Fixed-stride windowing scheme (obs16_pred16_s128) used as the canonical training and evaluation index. Author-generated schematic.	31
Figure 4.2	State-change-targeted sampling pipeline yielding one frozen window per episode. Author-generated schematic.	32
Figure 4.3	Three-branch evaluation framework (Section 4.2). ACT row uses CEMP candidates due to CVAE collapse; DP row follows the ideal design.	35
Figure 5.1	Shared ResNet-18 visual backbone preceding both policy decoders. Adapted from [44].	48
Figure 5.2	CVAE posterior collapse in the trained Action Chunking Transformer (ACT) checkpoint. Three independently sampled latent vectors z_1, z_2, z_3 are passed to the CVAE decoder (dashed arrows indicate the latent path the decoder effectively ignores). All inputs produce an identical output $\hat{a}$; the decoder has learned to reconstruct the action from the visual observation alone, making the latent redundant.	49
Figure 5.3	Cross-Entropy Method with Policy (CEMP) candidate augmentation. Three iterations of the Cross-Entropy Method initialised from the ACT B1 output refine SE(3) perturbations; the standard deviation decays by $0.7\times$ per iteration. The four highest-scoring candidates form the B2/B3 evaluation pool.	50
Figure 5.4	DDIM subsampling correction. The ten-step inference schedule spans $t \in \{90, 80, \dots, 0\}$. At each reverse step the correct $\bar{\alpha}_{t_{\text{prev}}}$ is the <i>next timestep in the subsampled schedule</i> ; the bug substituted the consecutive training-schedule timestep $\bar{\alpha}_{t-1}$, miscalibrating the $\hat{x}_0$ weighting at every step.	51
Figure 5.5	DexWM pipeline as implemented, comprising a frozen DINOv2 encoder, CDiT predictor, lightweight decoder, and HC loss supervision path. Adapted from [21].	54

Figure 5.6 Normalised validation loss (each family’s value divided by its 1k loss) across six training scales. Solid lines are primary models; dashed lines are controls. Large filled circles mark saturation points matching the bold entries in Table 5.5. Families with no marker are still declining at 100k. 58

List of Tables

Table 3.1	Literature roles in the WACT benchmark design.	16
Table 3.2	Average consecutive sub-task completions (<code>avg_len</code>) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [31].	21
Table 4.1	EgoDex splits under the <code>obs16_pred16_s128</code> windowing contract. Episode and window counts are derived from the canonical index.	33
Table 4.2	Evidence role assigned to each implemented or investigated world-model path.	38
Table 5.1	Implementation ownership and evidence status.	44
Table 5.2	Training configuration for all model families at $E = 100,000$. Approximate step time is shown for trainable phases only; Phase 1 of V-JEPA2-AC is a one-time offline extraction with no optimiser step. V-JEPA2 (no AC) and PointWorld reuse the V-JEPA2-AC Phase 1 embeddings and require no additional training (--- denotes not applicable).	46
Table 5.3	Model input/output contracts used by the implemented routes.	47
Table 5.4	ACT CVAE posterior collapse: three successive retraining attempts. <code>kl_weight</code> , cyclical annealing schedule, and observation dropout were varied in sequence; all runs produced a mean pairwise candidate difference of zero on every training split.	49
Table 5.5	Validation loss across six training scales for all seven model families. Bold marks each family’s first scale of saturation; losses are objective-specific and not cross-family comparable.	57
Table 5.6	Plateau budget per primary model family. The universal budget is $E = 100,000$ effective items.	58
Table 5.7	Material implementation challenges and resolution status.	62
Table 6.1	Frozen evaluation status used in the Results chapter.	64
Table 6.2	Primary Diffusion Policy (DP) guard results on the frozen 7,607-window EgoDex test set.	64

Table 6.3	ACT guard results. Values are numerically zero because the candidate pool collapses.	65
Table 6.4	DexWM paired causal results against the frozen DP candidate-0 baseline.	66
Table 6.5	LeWM action-prior causal results against the frozen DP candidate-0 baseline.	67
Table B.1	Summary of benchmark validity concerns and their resolution status in the canonical final package.	96

List of Abbreviations

ACT Action Chunking Transformer.

ALOHA A Low-cost Open-source Hardware System for Bimanual Teleoperation.

BC Behaviour Cloning.

CDiT Conditional Diffusion Transformer.

CEMP Cross-Entropy Method with Policy.

CNN Convolutional Neural Network.

CVAE Conditional Variational Autoencoder.

DDIM Denoising Diffusion Implicit Models.

DDPM Denoising Diffusion Probabilistic Models.

DexWM Dexterous Manipulation World Model.

DiT Diffusion Transformer.

DP Diffusion Policy.

GATO Generalist Agent.

HDF5 Hierarchical Data Format 5.

I-JEPA Image Joint Embedding Predictive Architecture.

JEPA Joint Embedding Predictive Architecture.

LeWM LeWorldModel.

RMSE Root Mean Square Error.

SSL Self-Supervised Learning.

V-JEPA2 Video Joint Embedding Predictive Architecture 2.

V-JEPA2-AC Action-Conditioned Video Joint Embedding Predictive Architecture 2.

VAE Variational Autoencoder.

ViT Vision Transformer.

VLM Vision Language Model.

List of Notations

I Identity matrix; gives the isotropic unit Gaussian covariance.

$\mathcal{N}(0, I)$ Standard isotropic Gaussian prior.

ε Reparameterisation noise, $\varepsilon \sim \mathcal{N}(0, I)$.

μ, σ Encoder-predicted mean and standard deviation.

z CVAE latent variable sampled from the prior $\mathcal{N}(0, I)$.

$\bar{\alpha}_t$ Cumulative noise-schedule product at diffusion timestep t .

ε_θ Noise-prediction network parameterised by θ .

t_{prev} Next step index in a subsampled inference schedule.

$\hat{x}_0$ Predicted clean sample in a DDIM reverse step.

x_T Noise sample at the diffusion schedule endpoint T .

B_{global} Effective global batch size; equals $B \times G \times A$ where B is the per-device batch size, G the gradient-accumulation steps, and A the number of accelerator devices.

E Effective training exposure, measured as total training items consumed, equal to optimisation steps multiplied by effective global batch size.

K Number of sampled candidate action chunks scored for one observation window.

N Number of evaluated samples or windows, depending on the metric context; pairwise comparisons report the number of held-out windows or candidate pairs explicitly.

S Number of optimisation steps.

$\mathbf{a}_{t:t+15}$ Sixteen-step future action chunk predicted from sixteen observed frames; each step has 18 action dimensions.

a_t Action vector at timestep t ; in this thesis each wrist-action step is an 18-dimensional bimanual vector, with left and right hands each represented by 3 translation and 6 rotation coordinates.

z_t Latent embedding of observation frame or video state at timestep t ; dimensionality depends on the encoder family.

1 Introduction

Imitation learning studies how computational systems can acquire behaviour from demonstrations rather than from manually specified rules. In manipulation settings, this is attractive because useful behaviour often depends on subtle motion patterns that are difficult to encode by hand. Behaviour cloning provides one direct route by learning a mapping from observations to demonstrated actions, while inverse reinforcement learning approaches the same problem by trying to infer the objectives that make a demonstration sensible [1, 2]. Both perspectives frame imitation as a way to learn from observed behaviour, but they do not by themselves resolve how a system should judge whether a predicted future motion is physically plausible.

This issue becomes sharper in dexterous manipulation. Human hand motion is temporally structured, object dependent, and sensitive to small changes in pose. This follows from the timing of natural prehension, the visuomotor transformation required for grasping objects, and the postural synergies used to shape the hand for tools and objects [3–5]. In imitation settings, this matters because small prediction errors can compound across a future sequence, so local agreement with a demonstration does not guarantee sequence-level alignment [6]. This thesis is therefore concerned with prediction and evaluation of future hand motion, rather than deployed control. The work remains relevant to robotic control because robotic manipulation systems ultimately require reliable motion selection, but the evidence in this dissertation is restricted to an offline computational benchmark.

Visuomotor policies are a common way to connect visual observations with future motion. In this thesis, the focused policy family is the ACT, which predicts temporally extended action chunks rather than isolated next-step actions [7]. This chunked formulation is useful because dexterous motion unfolds over short sequences, not only through single-frame decisions. However, a policy that proposes a future chunk still requires a criterion for choosing between candidate futures. This creates the opening for world models.

World models learn predictive representations of how a state may evolve. The initial role considered in this thesis was conservative: a world model could evaluate candidate action chunks proposed by ACT and help select motions whose predicted consequences looked more consistent with the demonstrated future. As the implementation developed, the question broadened. The benchmark also examines whether a world model can take greater authority over candidate selection, moving from support for a policy toward a stronger world model driven selection role. This progression from assistance to selection motivates the comparative structure of the thesis.

This dissertation studies that structure through WACT, which is used here as a comparative benchmark framework for evaluating whether world-model support changes the quality of future hand-motion prediction. The work began as a broad comparison across several world-model architectures. During implementation, that plan exposed a second research problem: published world-model systems are not equally reproducible, action-compatible, or fair to compare under one offline benchmark. The final contribution is therefore not a simple leaderboard. It is a controlled benchmark with one clean positive pairing, several boundary cases, and a documented account of why some promising model routes could not support headline claims.

1.1 Motivation

The motivation for this work sits at the intersection of robotics and machine learning. Modern manipulation research increasingly depends on learned models that can process visual observations, predict future motion, and adapt to complex physical interaction. At the same time, the pace of open-source machine learning development can make it difficult to separate robust scientific progress from promising but weakly benchmarked demonstrations. This concern is especially important for industrially relevant settings, where unreliable claims about physical reasoning or manipulation capability can lead to poor engineering decisions.

The work was also shaped by the RSO I activity and the University of Malta M.AI.ESTRO project, where the convergence of robotics, machine learning, and Industry 4.0 raised a practical need for clearer evaluation of emerging model families. Recent systems such as DiWA motivate this direction by treating world models as mechanisms for improving diffusion policies through imagined experience. WACT asks a related but different question: can a world model improve the action selected at inference time without retraining the policy? This dissertation does not claim to deliver a deployment system. It uses that practical motivation to ask a narrower benchmark question about offline candidate selection on recorded hand-motion demonstrations.

1.2 Research Question and Scope

The central research question is:

Do world models improve dexterous manipulation prediction when compared with, and integrated alongside, an ACT visuomotor policy under a controlled offline benchmark?

This wording is a rescoping of the original proposal question. The proposal asked whether integrating a latent world model into an imitation learning policy would improve physical grounding and robustness, measured through latent agreement error and task success under perturbations. The final implementation does not evaluate deployed task success or perturbation robustness, and latent mean-squared error is used only as a training diagnostic for the world-model transition head rather than as the main thesis result. The finished thesis therefore frames the question around offline prediction quality: whether world model support selects future hand-motion chunks that are closer to demonstrated motion under a shared benchmark.

The benchmark focuses on future bimanual hand motion and transform state. In implementation terms, the main prediction surface is built around future left-hand and right-hand transform chunks, but the detailed representation is defined later in Chapter 4 and Chapter 5. The Introduction therefore uses the more readable language of future hand motion and transform state, while later chapters specify the exact data contract, branch definitions, and metrics.

The scope is deliberately bounded. The thesis does not test closed-loop robotic deployment, real-world task completion, force interaction, contact stability, or general machine common sense. It also does not claim that any one world model family is universally superior. Instead, it asks whether selected world-model families improve an offline prediction benchmark when compared under a shared evaluation surface. The final evidence focuses on Diffusion Policy candidates scored by V-JEPA2-AC as the clean positive pairing. ACT, DexWM, LeWM, DINO-WM, JEPA-WM, PointWorld proxy controls, and DiWA/CALVIN infrastructure are used to define the boundaries around that result.

This scope also explains the structure of the comparison. The benchmark is organised around three branches:

- **ACT-only branch:** the policy baseline, where ACT provides the selected future motion without world model support.
- **Hybrid branch:** ACT proposes candidate chunks, and the world model evaluates or re-scores those candidates before selection.
- **World model driven branch:** no policy preference is used for selection. The selected future motion is chosen from the world model’s own assessment of predicted future-state consistency.

The purpose of this structure is to separate the value of the world model from the value of the underlying policy as far as possible within an offline benchmark.

1.3 Contributions

This thesis contributes a comparative benchmark framework for studying world-model support in dexterous manipulation prediction. The framework fixes the policy candidates, separates training support from held-out test windows, and records which information a selector is allowed to consume. This provides a controlled basis for asking whether predictive future-state models improve selected hand-motion chunks under the same offline evaluation conditions.

The thesis also contributes an implemented evidence hierarchy. V-JEPA2-AC is the main positive semantic world-model scorer. LeWM action-prior scoring is a secondary action-regularity result. DexWM, DINO-WM, JEPA-WM, ACT, and the PointWorld proxy expose limitations involving checkpoint availability, pretraining scale, action-space compatibility, candidate diversity, and evaluation fairness.

Finally, the thesis contributes a reproducibility account. Several model routes were implemented far enough to identify why they could not support headline claims. These negative and mixed outcomes are not treated as failures to hide. They are part of the benchmark contribution because they show what must be true before world-model guidance becomes valid scientific evidence.

1.4 Dissertation Structure

—To be updated at the end ———

The remainder of the dissertation is organised as follows:

- Chapter 2 introduces the technical background required to follow the dissertation, covering transformers, self-supervised learning, variational autoencoders, rigid body pose, visuomotor policies, and world models at a foundational level.
- Chapter 3 positions the work within the relevant literature on imitation learning, visuomotor policies, world models, and offline evaluation.
- Chapter 4 defines the benchmark design, scope, branch structure, leakage controls, metrics, and reporting rules.
- Chapter 5 describes the implemented system, the model-specific routes, and the reproduction constraints encountered during the study.
- Chapter 6 presents the benchmark results.
- Chapter 7 interprets the findings and states the main limitations.

- Chapter 8 closes the dissertation and identifies future work.

2 Background

This chapter assumes familiarity with the basic mechanics of a neural network and introduces six concepts central to this work: visuomotor policies, rigid-body pose in $SE(3)$, latent representations, conditional VAEs, transformer attention, and self-supervised learning. These foundations support understanding of the specific architectures in Chapter 3 and underpin the candidate selection mechanism and world model evaluator design examined in Chapter 4. Together, they explain the thesis question: WACT asks whether a policy can propose bimanual action chunks and whether a latent world model can score those chunks before any action is executed.

2.1 Visuomotor policies

A visuomotor policy is a learned function that takes a visual observation as input and produces an action or sequence of future actions as output. In a manipulation setting, the visual observation is typically a single camera frame showing the hand and the object being manipulated, and the action describes how the hand should move next. The policy is called visuomotor because it connects vision directly to motor output, without requiring a hand-crafted intermediate representation of what the scene contains [1].




Figure 2.1 A visuomotor policy maps an EgoDex-style camera frame to future bimanual wrist actions in the thesis setting.

Learning this mapping from data is appealing because useful hand behaviour depends on subtle visual cues that are difficult to specify by hand. The timing of a reaching movement, the grip shape adopted before contact, and the postural adjustments made in response to object geometry all depend on information visible in the camera image but difficult to encode as explicit rules [3–5]. A visuomotor policy learns these associations directly from recorded demonstrations, capturing the relationship between what the camera sees and what the hand does without requiring the programmer to describe it explicitly.

2.1.1 Behaviour cloning and compounding error

The most direct way to train a visuomotor policy from demonstrations is Behaviour Cloning (BC) [1]. Here an observed frame predicts the action taken by the expert at that moment, and the policy is trained to minimise the difference between its predicted action and the recorded expert action across all frames in the dataset [6]. The problem arises when a small prediction error causes the hand to move to a state that is slightly different from any state in the training data. This results in a compounding error: as the policy encounters increasingly unfamiliar states, its prediction errors grow larger over time.




Figure 2.2 Compounding error in behaviour cloning. The policy trajectory drifts progressively further away at each step. The growing gap between the two paths illustrates how small errors can accumulate over time [8].

2.1.2 Action chunking

Action chunking addresses the compounding error problem by changing what the policy predicts. Instead of predicting a single next action at each step, the policy predicts a short sequence of future actions in one forward pass. This sequence is called an action chunk [7, 9]. In WACT, the policy observes 16 frames and predicts a 16-frame future action chunk. Each predicted step is an 18-dimensional vector: two hands, each with a six-dimensional rotation representation and a three-dimensional translation. This is why Figure 2.1 shows the action sequence as a stack of multiple action blocks. By committing to a chunk of k future steps, where k is a hyperparameter set per task, the policy makes one prediction decision every k frames rather than one per frame. With fewer decision points, there are fewer opportunities for errors to accumulate, and the hand stays closer to the training distribution for longer [10].

2.2 Rigid body pose and the $SE(3)$ representation

To predict how a hand will move, a system needs a precise way to describe where the hand is and which way it is pointing. These two quantities together are called the pose of a rigid body. For the human wrist, pose is fully described by its position in space and its orientation relative to a fixed reference frame [11].

Translation. Position is represented as a vector of three numbers, one for each spatial dimension: how far the wrist is to the right or left, how far forward or backward, and how far up or down [3]. This vector $\mathbf{t} = [t_x, t_y, t_z]^\top$ is straightforward to predict with no special constraints on the output.




Figure 2.3 Translation for each arm can be represented by a movement along each of the three spatial axis [12].

Rotation. Orientation requires more care. A common representation uses three angles known as Euler angles, one per axis. These are intuitive but problematic for learning. At certain configurations the axes align and the system loses a degree of freedom, a singularity called gimbal lock. A subtler problem is discontinuity: a wrist rotated 359 degrees and one rotated 1 degree are almost physically identical, yet their angle values are numerically far apart. A network minimising prediction error will produce large errors near these boundaries even when the true motion is smooth [13].

A rotation matrix R avoids both issues: small physical rotations produce small numerical changes in the matrix entries, and there are no singularities.

Homogeneous transform. Translation and rotation are combined into a single 4×4 matrix, the standard representation of a pose in the special Euclidean group $SE(3)$:

$$T = \begin{bmatrix} R & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix} \quad (2.1)$$




Figure 2.4 Rotation around the y-axis for each arm. The same circular movement can be imagined around the x and z axes; the total orientation of the wrist is the combined result of rotations around all three [12].

The benchmark does not predict the full human body. It predicts left- and right-wrist motion. For each future step, WACT stores one rotation and one translation per wrist, giving the 18-dimensional bimanual action vector used by the policy and world-model scorers.

2.3 Latent representations

A latent representation is a learned internal description of an observation. An image starts as a grid of pixel values, where each pixel stores three or four channel intensities (typically red, green, and blue, with an optional transparency value). A model does not usually reason with these raw numbers directly. Instead, it transforms them through several layers into a new set of values that no longer correspond to individual pixel coordinates or colours. These new values form a compact encoding that captures patterns in the image such as shape, position, motion, and spatial structure. The collection of these encodings is often called the latent space. Figure 2.5 shows an example in which similar items cluster together in the learned representation.




Figure 2.5 Example latent-space plot showing how structurally similar items cluster together in the learned representation. Structural similarity here refers to items that share shape, geometry, and likely physical behaviour [14].

The purpose of this transformation is to keep task-relevant structure while discarding unnecessary visual detail. For a manipulation scene, a useful latent state might encode hand position, object pose, contact progression, or motion trend without preserving every texture or lighting variation. A latent state is therefore not directly interpretable in the way an image is, but it can still preserve the information needed to reason about what happens next.

This is why latent prediction is often a better fit than pixel-level prediction for high-dimensional video observations. Predicting the next frame at the pixel level requires a model to account for every visual detail of the scene, including lighting changes, texture variation, and background movement that carry no information about task-relevant dynamics. A latent predictor instead operates on compact embeddings and concentrates predictive capacity on the structure that determines what actually happens next [15]. This design trade-off motivates the world model architectures reviewed in Chapter 3.

2.4 Conditional variational autoencoders

A VAE maps an input to a distribution of latent vectors described by a mean μ and standard deviation σ , rather than a single point. Sampling different latent vectors lets the decoder produce different but plausible outputs rather than one fixed answer. Figure 2.6a shows this standard encoder-latent-decoder flow [16].




Figure 2.6 Standard VAE (a) and conditional CVAE (b) encoder-decoder pipelines [17].

For the policy to be trainable end-to-end, gradients must flow from the action prediction loss all the way back to the encoder. Sampling z directly from a distribution blocks this path. The reparameterisation trick resolves this by writing $z = \mu + \sigma\epsilon$, where ϵ is standard normal noise drawn independently of the learned parameters. Figure 2.7 shows the full encoder-latent-decoder pipeline alongside its training loss: a reconstruction term that rewards accurate output and a regularisation term that keeps the latent distribution close to a standard Gaussian, together forming the evidence lower bound [16]:

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}[\log p_{\theta}(\mathbf{x} | \mathbf{z})] - D_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z})). \quad (2.2)$$




Figure 2.7 VAE encoder-decoder architecture and two-part training loss (reconstruction and regularisation) [16].

A CVAE adds an extra conditioning input to both encoder and decoder, shown in Figure 2.6b. In a visuomotor policy, this condition is the current observation frame, so the latent variable only needs to capture how the future action varies within that scene [17]. The ACT architecture uses this to sample multiple latent vectors from one observation and decode each into a candidate action chunk [7].

2.5 Transformer networks and the attention mechanism

A transformer processes a sequence as a set of tokens. A token can represent a word, an image patch, a timestep in an action sequence, or any vector produced by an earlier model component. The architectural idea needed here is simple: each token is updated by looking across the other tokens and deciding which of them are most relevant. This relevance-weighted exchange of information is called attention [18].




Figure 2.8 Transformer attention block: each token queries all others with Q , gathers weighted values V via keys K , and the weighted sum yields the updated output token [18].

Attention forms three learned views of every token:

- **Query** (Q) – what this token is looking for.
- **Key** (K) – what another token contains.
- **Value** (V) – what another token contributes if attended to.

Figure 2.8 shows how each token (highlighted) sends its query outward to every other token, receives weighted values in return, and produces an updated output. The bidirectional arrows indicate that every token attends to every other token simultaneously, and the crossing arrow beneath represents the weighted sum that combines those values into the final output. For each token, the transformer compares

its query with the keys of all tokens. The resulting scores become weights, and those weights are used to combine the corresponding values. This operation is usually written as [18]:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right) V. \quad (2.3)$$

Transformers usually repeat this operation in several heads and add positional information so the model knows where each token sits in the sequence. Multiple heads allow different relationships to be represented at the same time, while positional information preserves order [18].

2.6 Self-supervised learning

Machine learning methods can be grouped by how they receive feedback during training, as illustrated in Figure 2.9. Supervised learning trains on data with explicit labels: the correct answer is provided for every input, and the model learns to reproduce it. BC from Section 2.1.1 is a direct example: the expert demonstration acts as the label. Reinforcement learning instead receives feedback through a reward signal after taking actions in an environment; the model learns which state-action sequences lead to good outcomes. SSL sits between these two: it trains on unlabelled data by having the data itself provide the supervision signal. Part of the input is hidden, and the model learns to predict it from the remainder [19].




Figure 2.9 Three machine learning paradigms: supervised learning (explicit labels), reinforcement learning (reward signal), and self-supervised learning (unlabelled data provides its own supervision signal). The third panel is labelled unsupervised learning in the source figure; this thesis uses the narrower term self-supervised learning for prediction tasks constructed from the data itself.

Applied to video, SSL takes the following form: given a sequence of observed

frames, hide one or more future frames and ask the model to predict them from the frames that came before. No human annotation is required because the video recording itself provides the hidden target.

Predicting raw pixels runs into a problem described in Section 2.3: when the future is ambiguous the model cannot commit to a single outcome, so it predicts the average of all possibilities and produces a blurred image that does not match any real frame. Predicting in latent space avoids this. Instead of reconstructing every pixel, the model predicts the compact embedding of the future frame. Background motion and fine texture are discarded at the encoding step, so the predictor focuses on the structure that determines what actually happens next [15]. This makes the prediction target, the latent embedding of the future frame, a compact and meaningful signal for evaluating whether a proposed action is likely to lead to a desirable outcome. When the predictor is also given the action the agent intends to take, the result is a world model: a learned function that predicts how the scene embedding will change in response to a proposed action. For example, if a policy proposes five possible wrist-motion chunks, a latent world model can predict the future embedding for each chunk and assign the lowest cost to the candidate whose predicted future is most plausible under its learned dynamics. This connects the concepts in this chapter: visuomotor policies produce action chunks, $SE(3)$ describes the wrist motion being predicted, latent representations provide the scoring space, and self-supervised prediction supplies the world-model training signal. The architectures that build on this idea are reviewed in Chapter 3.

3 Literature Review

This chapter reviews the literature that leads directly to the two research questions. RQ1 asks whether a world-model evaluator can improve action-chunk selection for bimanual EgoDex wrist-action prediction, and RQ2 asks which evaluator conditions make that improvement reliable. The reviewed work therefore has three roles: first, to justify ACT and Diffusion Policy as the candidate policy families; second, to motivate latent, action-conditioned world models as candidate selectors; and third, to bound what an offline benchmark can and cannot claim without physical deployment.

The review is also selective about evidential role. Some papers motivate the final positive test, especially Diffusion Policy and V-JEPA2-AC; others explain why negative or exploratory routes are still scientifically useful, especially DexWM, LeWorldModel (LeWM), DINO-WM, and JEPA-WM. DiWA is treated as the closest adjacent policy-improvement study because it uses a world model to improve Diffusion Policy performance, but its training-time adaptation setting differs from this thesis’s inference-time candidate selection setting. Table 3.1 summarises how each literature group maps to the benchmark design.

The chapter assumes familiarity with the technical foundations in Chapter 2; world models are introduced here.

Table 3.1 Literature roles in the WACT benchmark design.

Work	Relevant strength	Relevance and limitation for this thesis
ACT [7]	Efficient bimanual action chunking with a structured latent variable.	Provides the transformer policy baseline, but its sampled candidates collapsed in the final WACT implementation, limiting its use as a selector stress test.
Diffusion Policy [9]	Strong multimodal visuomotor policy with diverse sampled action chunks.	Main policy family for the headline WACT result, with higher inference cost than ACT.
V-JEPA2 and V-JEPA2-AC [20]	Internet-scale video representation and action-conditioned latent prediction.	Main positive world-model evaluator; the action-conditioned head gives the most direct match to candidate scoring.
DexWM [21]	Dexterous manipulation world model with hand-consistency supervision.	Important comparator, but no public pretrained checkpoint was available and full upstream-scale pretraining fell outside the fair local budget.
LeWM [22]	Lightweight end-to-end latent dynamics model.	Useful for action-prior evidence, but the visual rollout route did not support the main semantic-selection claim.
DINO-WM and JEPA-WM [23]	Official checkpoint route for image/JEP-style latent prediction.	Exploratory official-checkpoint comparators; action-interface mismatch required an adapter and results remain appendix-level.
DiWA [10]	Uses imagined experience to improve Diffusion Policy.	Closest related study, but it modifies training rather than selecting candidates at test time and could not be fully reproduced without missing artifacts.
SimplerEnv and WorldEval [24, 25]	Evidence that offline/simulated rankings can correlate with real-world outcomes.	Supports the offline-evaluation premise, while requiring careful claims about calibration and deployment validity.

3.1 Action Chunking with Transformers (ACT)

ACT was introduced by Zhao et al. to address the compounding error and multimodal distribution challenges that make dexterous bimanual manipulation demanding for imitation learning [7]. On the A Low-cost Open-source Hardware System for Bimanual Teleoperation (ALOHA) bimanual platform it achieved success rates of up to 90% on precision tasks such as threading a cable tie, inserting a battery, and opening a translucent bag, outperforming behaviour cloning baselines and Generalist Agent (GATO) by substantial margins [7]. In $SE(3)$ wrist-pose prediction specifically, small deviations in predicted position or orientation alter contact geometry in ways that cascade into grasp failure [1], and a policy trained with a naïve regression loss produces averaged predictions that may not represent any plausible trajectory when the demonstrated distribution is multimodal [9]. ACT’s CVAE architecture addresses both of these failure modes directly.

ACT addresses both challenges through the use of a CVAE design (Section 2.4). The respective CVAE has an encoder-decoder structure depicted in Figure 3.1, where the encoder compresses an action sequence and joint observation into a style variable z (List of Notations) that captures the multimodal distribution of plausible futures, and the decoder takes visual observations and joint positions, together with z , through a transformer encoder, and predicts a sequence of future action steps as an action chunk with a transformer decoder. The encoder is discarded at test time. In the original formulation, z is set to the mean of the prior (zero) to produce a single deterministic action chunk per observation [7]. When a pool of candidate chunks is required, multiple values of z are sampled independently from the prior and each is decoded into a separate candidate, all conditioned on the same observation.




Figure 3.1 Architecture of the Action Chunking with Transformers (ACT) policy [7].

At inference time, the CVAE decoder produces each candidate action chunk by decoding an independently sampled z vector. A temporal ensembling technique is then applied: consecutive chunks overlap in their predicted time horizons, and a recency-weighted average is taken across overlapping predictions to produce the final action. This averaging smooths the executed motion and reduces jerkiness caused by

abrupt chunk transitions. The result is a structured pool of plausible futures, all conditioned on the same observation, from which the ensembled output is drawn [7].

ACT carries known limitations alongside these strengths. The Gaussian prior on the style variable z constrains how faithfully the CVAE can represent highly multimodal action distributions: at transitions where qualitatively different trajectories are equally plausible, the latent space can average across modes rather than sample cleanly from each one, producing trajectories that are plausible in aggregate but imprecise at the most ambiguous decision points [9]. Chunk length is a fixed hyperparameter that must be tuned per task; longer chunks improve smoothness but commit the policy to a trajectory over a horizon that may no longer be appropriate as the scene evolves. Temporal ensembling introduces a recency-weighted lag that can slow the policy’s response to sudden environmental changes [7]. These properties are well-documented in follow-on work comparing ACT with more expressive policy classes, and they motivate the inclusion of Diffusion Policy as the second candidate in the benchmark.

3.2 Diffusion Policy

Diffusion Policy offers a principled alternative approach to the same compounding error and multimodality challenges. Its core concept relies on starting from a noise distribution and iteratively refining action predictions through a diffusion process. By gradually denoising the predictions, Diffusion Policy captures complex multimodal distributions effectively, producing plausible trajectories that align with the demonstrated data [9]. Figure 3.2 illustrates the policy overview:

- a) General formulation: at timestep t , the policy takes the latest T steps of observation data O_t as input and outputs T steps of actions A_t .
- b) CNN-based Diffusion Policy: FiLM (Feature-wise Linear Modulation) [26] conditioning of the observation feature O_t is applied to every convolution layer, channel-wise. Starting from K_t drawn from Gaussian noise (List of Notations), the output of noise-prediction network ϵ_θ is subtracted, repeating K times to get A_t^0 , the denoised action sequence.
- c) Transformer-based Diffusion Policy: the embedding of observation O_t is passed into a multi-head cross-attention layer of each transformer decoder block. Each action embedding is constrained to only attend to itself and previous action embeddings (causal attention) using the attention mask illustrated [27, 28].

Chi et al. report that Diffusion Policy outperforms ACT on tasks with high action multimodality (such as multi-step object rearrangement and cup placement under



Figure 3.2 Diffusion Policy: general formulation (a), CNN variant (b), transformer variant (c) [9].

visual ambiguity) while matching it on simpler reaching tasks [9]. The advantage arises because iterative denoising explores the full action distribution at each step rather than committing to a single latent mode as the CVAE prior must, making it more robust to the mode-averaging failure that affects ACT at ambiguous transitions.

The main limitation of Diffusion Policy is inference cost. Generating a single action chunk requires multiple sequential network evaluations: standard Denoising Diffusion Probabilistic Models (DDPM) variants use on the order of 100 denoising steps compared with the single forward pass of ACT’s decoder, which substantially reduces the achievable control frequency [9]. Efficient transformer variants address this through progressive noise scheduling and parameter sharing, recovering much of the inference speed at added architectural complexity [28]. Both policies are included in the benchmark to span this representational trade-off: ACT’s structured latent sampling is computationally efficient but expressively constrained at high multimodality, while Diffusion Policy’s iterative denoising is more expressive but more demanding at inference. Measuring the world model’s effect across both allows the benchmark to determine whether evaluator gains are consistent across this axis.

3.3 World models

A world model is a learned predictor of how a system evolves over time. Given the current state and a proposed action, it estimates the state that is likely to follow. In a manipulation setting, the state may be a raw image, a latent representation of the scene (Section 2.3), or a structured description of hand and object configuration. Regardless of the representation chosen, the underlying principle is the same: the model internalises the dynamics of the environment and uses them to project forward from any given state [29].

World models become especially useful when a pool of candidate actions must be evaluated before any one is committed to. Each candidate is rolled forward through the model, producing a predicted future state; those predictions are ranked by

predicted cost or consistency. The quality of this ranking depends heavily on the chosen representation: pixel-space world models must reconstruct every visual detail and suffer from mode-averaging at multimodal transitions, while models that operate in the compact learned latent space introduced in Section 2.3 can concentrate predictive capacity on the dynamics that matter for task outcomes [15]. The following subsections survey the literature supporting this evaluator design and describe the three world model architectures implemented in the benchmark.

3.3.1 World models as policy evaluators

WorldGym demonstrates that a learned world model used as a simulated environment for policy evaluation produces rankings that align with real-world task success with meaningful fidelity [30]. Candidate policies are evaluated not by deploying them physically but by rolling them out inside the model’s predicted environment, reducing the cost of policy comparison to a forward pass. CLOVER extends this to closed-loop control, as shown in Figure 3.3. A visual planner generates a sequence of sub-goal frames; the executor measures the embedding-space error between the predicted sub-goal and the current observation. When the error exceeds a threshold, the system triggers an adaptive transition and replans rather than committing to the original trajectory. The feedback loop, absent from open-loop imitation policies, is the mechanism that translates world model prediction accuracy into task completion gain [31]. Together, WorldGym and CLOVER establish that world model rankings are informative both before and during execution.

The broader question of whether such evaluator rankings transfer to physical outcomes, and the qualifications around out-of-distribution reliability, are examined in Section 3.4.




Figure 3.3 Classic closed-loop control (a) versus CLOVER’s closed-loop visuomotor control (b) [31].

The practical value of this evaluator loop is measurable. Table 3.2 reproduces CLOVER’s real-world evaluation on a multi-step robot manipulation suite, reporting the average number of consecutive sub-tasks completed (`avg_len`) for representative baselines and CLOVER [31].

Table 3.2 Average consecutive sub-task completions (`avg_len`) on a real-world multi-step manipulation suite. Higher is better. CLOVER’s closed-loop feedback via world model prediction nearly doubles the completion depth of the strongest open-loop baseline. Data from [31].

Method	<code>avg_len</code>
ACT [7]	0.6
R3M (pretrained visual repr.) [32]	0.7
RT-1 [33]	1.1
CLOVER (closed-loop)	2.1

Alongside the task completion gain, CLOVER reduces embedding-space prediction Root Mean Square Error (RMSE) relative to open-loop baselines by correcting the policy’s course when the world model detects a divergence from the expected trajectory, demonstrating that prediction accuracy and task success move together under the evaluator framing.

3.3.2 Latent predictive architectures

The energy-based framework of LeCun and Dawid establishes why prediction should operate in latent rather than pixel space [15]: the same representational argument introduced in Section 2.3 applies directly to the prediction target, favouring compact embeddings over pixel reconstructions that force the model to account for irrelevant visual detail and average over all plausible continuations at ambiguous transitions.




Figure 3.4 Three SSL prediction paradigms: joint-embedding (a), generative (b), and joint-embedding predictive (c) [19].

Image Joint Embedding Predictive Architecture (I-JEPA) establishes the joint-embedding predictive architecture as a viable approach to learning rich visual representations without any pixel reconstruction objective [19]. Panel (c) of Figure 3.4 shows the data flow directly: a context block x from the image is passed through the context encoder to produce latent tokens z_x . A lightweight ViT predictor then takes z_x

together with positional tokens indicating where the masked target regions y are located, and outputs predicted latent representations $\hat{s}_y$ for those regions. A separate target encoder, whose weights are a slow exponential moving average of the context encoder, independently encodes the actual target pixels into ground-truth latents s_y . The loss $\text{MSE}(\hat{s}_y, s_y)$ is computed entirely in latent space: no decoder, no pixel generation, no reconstruction of lighting or texture. The outcome is a context encoder whose representations capture semantically meaningful structure (object shape, spatial layout, and scene dynamics) because that is all the predictor ever needs to get the target embeddings right. Training is approximately 2.5 times faster than pixel-reconstruction methods as a result. V-JEPA2 extends this same joint-embedding predictive architecture from images to video sequences.

The following subsections cover the architectures that shaped the benchmark, but they should not be read as equal final-result routes. V-JEPA2-AC is the main evaluator because it directly conditions latent prediction on candidate actions. DexWM is the closest dexterous-manipulation comparator, but its missing public pretrained checkpoint turned the thesis implementation into a bounded fixed-budget test rather than a full upstream reproduction. LeWM is relevant because it offers a compact learned latent space, although the final thesis evidence uses it primarily as an action-prior route. DINO-WM and JEPA-WM provide official-checkpoint latent predictors, but their action interface mismatch makes them exploratory rather than headline comparators.

V-JEPA2




Figure 3.5 V-JEPA2 post-training pathways from internet-scale backbone to specialised applications [20].

V-JEPA2 extends the I-JEPA framework from static images to video, pretraining

on over one million hours of internet video [20]. The pretraining shifts from spatial to spatio-temporal masking: the context encoder receives past observation frames, and the predictor must recover future latent states. The training objective combines teacher-forcing with multi-step autoregressive rollout losses, encouraging accurate short-horizon predictions and consistent long-horizon trajectory representations. The result is a frozen ViT-Large encoder, approximately 300 million parameters, that encodes physical priors about hand and object interactions acquired from internet-scale pretraining. As shown in Figure 3.5, the pretrained backbone supports three post-training pathways: Language Alignment, which grounds visual embeddings in natural language; Attentive Probe, which adds a lightweight classification head; and Action Conditioning, which extends the predictor to take robot actions as input and forms the basis of V-JEPA2-AC.

V-JEPA2-AC

The frozen ViT-Large encoder from V-JEPA2 serves as the backbone of V-JEPA2-AC, producing a fixed latent embedding for each observation frame with no gradient updates to the backbone weights. A lightweight action-conditioned transition head is trained on top of these stored embeddings: it takes the current frame latent together with a proposed action chunk (robot joint states and end-effector poses) as the conditioning signal z , and predicts the latent representation of the future frame. Candidate action chunks are then scored by comparing the predicted future embedding against the actual observed future embedding under an L1 objective. This design, shown on the right of Figure 3.6, separates what the network must learn into two parts: the frozen encoder supplies rich visual features that generalise from internet-scale pretraining, and the small trained head learns only how actions propagate through those features.

An important qualification is raised by Tsagkas et al., who show that pretrained visual representations can encode spatial plausibility well while providing a weaker signal for fine-grained temporal dynamics [34]. A frozen encoder trained on internet video may represent hand shape and position accurately at each individual frame while providing a less reliable signal about whether the predicted trajectory is kinematically consistent across the full chunk horizon. This temporal limitation provides context for the path shape results in Chapter 7.

DexWM

DexWM [21] is a dexterous manipulation world model from Meta FAIR. As shown in Figure 3.7, its pipeline encodes each input frame with a frozen DINOv2 image encoder, producing a grid of spatial patch tokens rather than a single global embedding. These




Figure 3.6 Standard V-JEPA2 (left) versus action-conditioned V-JEPA2-AC (right) [20].

patch-level representations preserve fine-grained spatial information about hand positions and object locations, which gives the predictor strong generalisation across diverse environments and viewpoints. A CDiT predictor then takes the spatial latent tokens together with action and camera-motion inputs to predict the next latent state, and a lightweight decoder maps the result back to image and keypoint space. Where V-JEPA2-AC uses a video-pretrained backbone, DexWM uses a backbone pretrained on static images; the comparison therefore isolates the contribution of video-specific temporal pretraining to world model quality.

A distinguishing feature of DexWM is the Hand Consistency (HC) loss: an auxiliary transformer head predicts fingertip and wrist heatmaps from the latent state, imposing a geometric prior that the predictor’s latent trajectory must remain consistent with plausible hand configurations. This explicit hand-geometry supervision is absent from both V-JEPA2-AC and LeWM, whose latent spaces are shaped only by prediction error or general regularisation. Goswami et al. report a 50% improvement over Diffusion Policy on grasping and placing tasks, with an 83% zero-shot real-world grasping success rate under latent planning [21].

LeWorldModel

LeWM is a latent world model trained end-to-end from random initialisation on raw pixel observations [22]. Its architecture consists of two components: an encoder that maps each frame o_t to a compact latent vector z_t , and a predictor that models environment dynamics by estimating the next latent embedding $\hat{z}_{t+1}$ from z_t and the action a_t . The encoder is a ViT-Tiny with approximately 15 million parameters. Each frame is represented as a single 192-dimensional token rather than a spatial patch grid,




Figure 3.7 High-level layout of the DexWM world model pipeline, illustrating how predicted keypoints are incorporated alongside the encoder, CDiT predictor, and decoder stages [21].

which keeps the latent space compact and makes planning significantly faster than models that produce higher-dimensional patch representations.




Figure 3.8 LeWM training pipeline: encoder, dynamics predictor, and SIGReg regulariser [22].

As shown in Figure 3.8, the training pipeline consists of an encoder that maps each observation frame to a compact latent vector, a dynamics predictor that autoregressively estimates the next latent given the current embedding and action, and a SIGReg regulariser that enforces a Gaussian prior on the latent distribution by optimising normality along random univariate projections. The training objective reduces to two terms:

Z	batch of latent vectors produced by the encoder
λ	regularisation weight balancing prediction and SIGReg losses
$\text{SIGReg}(Z)$	Gaussian regulariser on random univariate projections of Z

$$\mathcal{L}_{\text{LeWM}} = \mathcal{L}_{\text{pred}} + \lambda \text{SIGReg}(Z) \quad (3.1)$$

where $\mathcal{L}_{\text{pred}}$ penalises inaccurate next-embedding forecasts and $\text{SIGReg}(Z)$ stabilises end-to-end training without contrastive negative pairs or an exponential moving average target encoder [22]. This reduces the number of tunable loss hyperparameters from six, as required by prior Joint Embedding Predictive Architecture (JEPA)-style designs, to one.

DINO-WM, the closest comparable baseline, follows a similar latent prediction objective but uses a frozen DINOv2 backbone producing a grid of spatial patch tokens per frame [23]. Like DexWM, DINO-WM therefore carries the spatial resolution and generalisation benefits of DINOv2’s patch-level representations, but at substantially higher computational cost. LeWM replaces the frozen backbone with a fully learned ViT-Tiny encoder that collapses each frame to a single 192-dimensional token, trading spatial resolution for training speed. In practice, LeWM can be trained on a single GPU in a few hours and achieves a $48\times$ planning speedup over DINO-WM [22]. DexWM’s CDiT predictor operating on full patch grids is significantly slower to train and requires substantially more GPU memory; this difference in training cost is expected to widen further at larger dataset scales and is quantified in Chapter 5. Despite its compact size, LeWM’s latent space encodes physically meaningful structure: the model detects physically implausible state transitions and captures spatial information such as agent and object locations across frames.

The three-way comparison across V-JEPA2-AC, DexWM, and LeWM is therefore not simply an architecture comparison. It is a comparison between fundamentally different representational strategies: video-pretrained transfer (V-JEPA2-AC), image-pretrained transfer with geometric supervision (DexWM), and self-supervised task-data learning from scratch (LeWM). This distinction makes the comparison principled and is the basis for the experimental design in Chapter 4.

PointWorld

PointWorld is a large pretrained 3D world model for in-the-wild rigid arm manipulation [35]. Unlike the latent-space models described above, it operates directly in 3D geometry: scene state is lifted from RGB-D images into a full-scene point cloud, and robot actions are represented as the 3D point flows traced by robot surface geometry through forward kinematics. This embodiment-agnostic representation allows the model to reason over physical interaction geometry rather than learned abstract features.

Such a model is not designed to handle an 18-dimensional bimanual wrist-pose action space. Its forward-kinematics representation assumes rigid arm geometry with a

discrete gripper aperture, making it structurally mismatched when applied to continuous SE(3) wrist transforms from human hand demonstrations.

3.4 Offline evaluation

A central concern for any offline benchmark is whether its metrics produce rankings that are meaningful proxies for real-world performance. Without this property, an offline evaluation framework reduces to a measure of in-distribution accuracy with no predictive validity for deployment.

At large scale, DreamerV3 demonstrates the effectiveness of latent-space candidate evaluation across diverse continuous control domains, including locomotion and manipulation [29]. DreamerV3 is an online reinforcement learning system, not an offline benchmark; it is included here because it provides large-scale empirical evidence that latent-space evaluation is an effective selection mechanism when properly calibrated, which motivates its use in the offline setting examined in this study. Its world model maintains a compact recurrent latent state that transitions under action conditioning; candidate trajectories are rolled out entirely inside this imagined space and scored by a learned value function, with no return to the real environment between evaluations. TD-MPC2 confirms the principle with a scalable planner that optimises a temporal-difference value function defined over predicted latent trajectories rather than observed states [36]. Both systems show that ranking candidate actions by predicted latent-space outcomes is competitive with direct online evaluation across a range of motor control benchmarks.

A critical qualification is that latent-space scoring becomes unreliable outside the training distribution. MOPO, MOREL, MBOP, and MOPP each establish that learned models produce inaccurate predictions for out-of-distribution states, requiring conservative candidate selection to avoid trajectory degradation [37–40]. None of these works addresses dexterous manipulation specifically; they benchmark locomotion and simple control tasks, leaving open how well the principle transfers to the fine-grained wrist-pose prediction setting.

Two further bodies of work establish that offline rankings can transfer to physical outcomes when the evaluation domain is well-calibrated.

SimplerEnv establishes the correspondence between simulation-based evaluation and physical robot performance for visuomotor policies [24]. The study evaluates multiple manipulation tasks and policy types in simulation and then physically, using a WidowX robot under matched visual and control conditions. Rankings produced from simulation-only evaluation closely track rankings from physical trials, with correlations holding across both policy families and task types. The

key condition for transfer is domain calibration: visual appearance, camera viewpoint, and control interface must match between the simulation and the physical system. When these are well-matched, simulation rankings provide a reliable ordering of policies without requiring physical deployment.

WorldEval extends the validation to the world-model setting [25]. Rather than evaluating policies in a manually constructed simulator, WorldEval scores them using a learned world model that predicts future states from offline trajectory data. The world-model-based rankings correlate strongly with real-world task success: policies that score well in world-model evaluation also tend to succeed when deployed physically.

Both results carry a shared qualification: rankings hold under well-calibrated conditions, where the evaluation domain is consistent with the training distribution and the deployment context. When this alignment breaks down, offline rankings can degrade. The design implications of this constraint are discussed in Chapter 4.

Summary. The literature surveyed in this chapter establishes three supporting findings. First, world model rankings are informative evaluators both before and during execution, as demonstrated by WorldGym and CLOVER [30, 31]. Second, latent predictive objectives produce more efficient and better-calibrated representations than pixel-level reconstruction, as demonstrated by the JEPA family [19, 20]. Third, offline evaluation rankings can carry predictive validity for physical outcomes when the evaluation domain is well-calibrated, as demonstrated by SimplerEnv and WorldEval [24, 25].

These findings leave two questions unanswered, which this dissertation addresses directly.

RQ1. Does augmenting a fixed imitation-learning policy with a latent world model evaluator improve candidate selection quality in an offline dexterous manipulation benchmark, and does any improvement hold across policy families?

The offline reinforcement learning literature establishes the need for principled candidate selection but benchmarks only locomotion and simple control tasks [37–40]. No controlled study isolates the world model evaluator’s contribution on a fixed dexterous manipulation policy.

RQ2. Does the world model’s representational strategy (pretrained video backbone, task-data-trained from scratch, or out-of-domain modality) produce meaningfully different selection guidance when applied to the same policy on the same evaluation data?

The interaction between world model architecture and policy performance has not been examined in this setting. Answering RQ2 requires a benchmark that holds the policy and evaluation surface fixed while varying only the world model family. The

methodology for constructing and running such a benchmark is described in Chapter 4.

4 Methodology

This chapter defines the protocol used to decide which world-model results are claimable in the thesis. The final study is not a broad leaderboard over every attempted model. It is a controlled candidate-selection benchmark with one headline positive pairing and several boundary results. The design therefore specifies the dataset and action space (Section 4.1), the three branch definitions (Section 4.2), the policy and candidate-generation routes (Section 4.3), the evidence roles assigned to each world-model family (Section 4.4), and the fairness and leakage rules used to separate held-out evidence from exploratory implementation work (Section 4.5). Evaluation metrics and scope limitations are addressed in Sections 4.6 and 4.7.

4.1 Dataset and Experimental Setup

4.1.1 EgoDex: source and selection rationale

The benchmark is built on EgoDex [41], a large-scale egocentric dataset of bimanual dexterous manipulation collected using Meta Quest Pro head-mounted cameras. It spans more than 100 task categories across its full release, including folding, slotting, pouring, and kneading, covering a wide range of fine-grained hand motions. The full release comprises five training parts and supplementary additional-data splits; this study uses the first three training parts, which together comprise more than 194,000 demonstration episodes and approximately 1.3 TB of paired video and annotation data. Each episode provides synchronised RGB video and full six-degree-of-freedom SE(3) wrist annotations for both hands at approximately 30 frames per second.

The choice of dexterous manipulation as the evaluation domain is not incidental. This dissertation forms part of the University of Malta’s ongoing M.AI.ESTRO research project [42], which investigates the use of AI for scene and task prediction on dexterous manipulation tasks to support industrial transfer learning and human skill digitisation. The research domain was therefore set by this institutional context: the methodology operates on human hand demonstrations rather than robotic end-effector data because the target application is the transfer of tacit human dexterous skill. An evaluation framework built around robot arm trajectories would not address that objective. The introduction situates this motivation more broadly; the present chapter focuses on the methodological consequences of the domain choice.

Within the dexterous manipulation setting, EgoDex was selected on three properties:

- **Bimanual hand action space.** The benchmark requires fine-grained articulation

across an eighteen-dimensional bimanual wrist pose representation; large-scale corpora built around rigid gripper or arm contact do not provide this.

- **Egocentric viewpoint.** The first-person perspective aligns with the intended downstream deployment setting of a wearable assistive device.
- **Dataset scale and diversity.** The corpus is sufficient to support simultaneous training of five world model families and two policy architectures on declared non-overlapping splits while retaining a held-out test set of meaningful size.

EgoDex carries two acknowledged limitations. The demonstrations are produced by human subjects rather than a robotic system, which introduces a kinematic domain gap between the training distribution and any robotic deployment. No contact or force signals are present; the action representation is purely kinematic $SE(3)$. Both are accepted as known constraints rather than experimental variables and are discussed further in Section 4.7.

4.1.2 Windowing protocol

The dataset is stored as paired `.mp4` and `.hdf5` episode files. A deterministic index stage discovers all valid paired episodes, infers their lengths from HDF5 metadata, and writes canonical observation and prediction windows for each split. This shared index is the single source of truth for all downstream modules: policy training, world model extraction, and evaluation all consume the same window file for their respective split, ensuring that output differences across the experimental branches can be attributed to branch design rather than to differences in the data presented to each component. The benchmark uses a fixed-stride windowing scheme as the canonical contract; a supplementary state-change-targeted scheme was developed during early exploration and used exclusively for qualitative inspection.

Strategy 1: Fixed-stride sampling. Each window spans 16 observation frames followed by 16 prediction frames, with consecutive window starts separated by a stride of 128 frames. At 30 frames per second this places starts approximately 4.3 seconds apart, corresponding to an observation context and prediction horizon of approximately 0.53 seconds each. This contract is encoded as `obs16_pred16_s128` and forms the foundation for all policy training and world model training (see Figure 4.1).




Figure 4.1 Fixed-stride windowing scheme (`obs16_pred16_s128`) used as the canonical training and evaluation index. Author-generated schematic.

Strategy 2: State-change-targeted sampling. The fixed-stride scheme samples the episode timeline uniformly, which causes it to under-sample high-motion events: long near-static segments contribute windows at the same rate as active manipulation periods. As detailed in Chapter 5, this limitation motivated a revised strategy adopted for a subsidiary perceptual inspection task and qualitative development review. Rather than placing windows at fixed intervals, this approach computes smoothed hand-state change scores from the wrist transform and confidence sequences, selects candidate windows near high-change peaks, and applies episode-first task-stratified sampling with a fixed random seed (see Figure 4.2). A relative-position fallback retains slower episodes that would otherwise be excluded, yielding one frozen window per episode stratified across all represented task categories.




Figure 4.2 State-change-targeted sampling pipeline yielding one frozen window per episode. Author-generated schematic.

The canonical training and evaluation index uses the fixed-stride scheme across all experimental branches. The state-change-targeted subset was used exclusively for a subsidiary perceptual inspection task and qualitative development review; it does not form part of the primary benchmark evaluation.

4.1.3 Training splits and test set

Three cumulative training splits were assembled from the first three parts of the EgoDex release, each a strict superset of the previous; Table 4.1 summarises their composition. EgoDex ships five training parts plus supplementary additional-data splits; parts four and five and all additional-data splits are excluded from this study to bound the experimental scope. The held-out test split is the official EgoDex test set and is never used during training at any stage by any model or policy. For the Strategy 1 benchmark, the same fixed-stride test index is used for quantitative

evaluation and for the qualitative review examples, keeping the reported comparisons aligned with the training contract in Table 4.1.

Table 4.1 EgoDex splits under the `obs16_pred16_s128` windowing contract. Episode and window counts are derived from the canonical index.

Split	Tasks	Episodes	Windows	Role
<code>train_part1</code>	26	46,091	140,239	Stage 1 training
<code>train_part1_2</code>	39	140,654	291,369	Stage 1+2 training
<code>train_part1_2_3</code>	63	194,333	430,962	Canonical training
<code>test</code>	111	3,229	7,607	Frozen evaluation

The episode counts reflect the uneven collection density across parts: the 13 task categories added in part two average approximately 7,300 episodes each, roughly four times the average of the 26 part-one tasks, because part two focuses on high-frequency foundational primitives such as pick-and-place and folding that were collected at greater volume.

The canonical results in Chapter 6 use `train_part1_2_3` across all trained components. The test split is the official EgoDex held-out set spanning 111 task categories, which includes all 63 categories present in `train_part1_2_3` as well as 48 categories drawn from the excluded parts; evaluating on this broader set provides a measure of out-of-distribution generalisation. Its window index is frozen; no post-hoc filtering or re-sampling is applied. The incremental training curriculum across the three stages is described in Chapter 5.

4.2 Three-Branch Experimental Framework

The central question of this study is whether a latent world model improves candidate selection quality when paired with an imitation-learning policy, and whether that effect generalises across policy families. Answering both questions cleanly requires holding every other experimental factor constant. The benchmark therefore organises evaluation around three branches that share the same policy weights and the same evaluation windows; the candidate pool composition and selection rule differ across branches (Figure 4.3). Any measured performance difference between branches is intended to reflect the selection mechanism, since training data, architecture, and window coverage are held constant across branches.

Branch 1 (B1): Policy-only baseline. The policy generates its full candidate pool and selects by its own internal preference; no world model is queried. The ideal design produces several structurally independent candidates so the downstream branches can compare meaningfully different trajectories. B1 establishes the trajectory prediction

quality each policy achieves acting alone and is the baseline against which B2 and B3 are measured.

Branch 2 (B2): World-model-inclusive selection. The world model scores all candidates in the pool and selects the highest-scoring one unconditionally; no margin threshold is applied. The candidate pool is the same as B1. Scoring proceeds by predicting forward in the world model’s latent space from the current observation and measuring prediction fidelity against each candidate action. The policy’s own output is not privileged; it is selected only if the world model assigns it the highest score. For V-JEPA2 without action conditioning, the transition head required to produce candidate-specific latent predictions is absent; if the model assigns indistinguishable scores to all candidates, B2 degenerates to B1, directly testing whether action conditioning is the operative mechanism.

Branch 3 (B3): World-model-exclusive selection. The policy’s preferred output is excluded from the evaluation pool; the world model selects the highest-scoring remaining candidate unconditionally. The ideal design ensures the remaining candidates are structurally independent so that exclusion meaningfully changes the selection outcome. B3 tests whether the world model maintains selection quality when the policy’s top choice is unavailable. All five world model families apply this protocol.

Two policy families are evaluated: ACT and DP (Figure 4.3).

- **DP.** The three-branch progression follows the ideal description above. Five independent reverse-diffusion chains seeded from distinct noise samples serve as candidates. B1 uses DP’s policy-internal argmax over those five; B2 passes the same pool to the world model scorer; B3 excludes DP’s preferred sample and scores the remaining four independent chains.
- **ACT.** The intended ACT design samples multiple CVAE latent vectors to produce several candidates. In the trained checkpoint used for the final causal matrix, these samples collapse numerically to the same action. B1 is therefore the deterministic policy candidate, B2 evaluates candidate 0 together with four sampled variants that are effectively identical, and B3 excludes candidate 0 but still has no meaningful action diversity. The ACT rows are retained as a candidate diversity diagnostic rather than as a positive or negative test of world-model ranking over a useful pool.

Candidate generation details are in Section 4.3.




Figure 4.3 Three-branch evaluation framework (Section 4.2). ACT row uses CEMP candidates due to CVAE collapse; DP row follows the ideal design.

4.3 Policy Variants

The benchmark evaluates two imitation-learning policy families. Their architectures and training objectives are reviewed in Chapter 3; the properties relevant to this study are how each generates its candidate pool and what that implies for world-model scoring.

- **Action Chunking Transformer** (Section 3.1). ACT is architecturally designed to generate diverse candidates by sampling latent vectors $z \sim \mathcal{N}(0, I)$ from its CVAE prior and decoding each into a 16-step bimanual SE(3) action chunk. In the trained checkpoint used here, CVAE posterior collapse is observed: the decoder is insensitive to z at inference time, producing identical outputs regardless of the sampled vector (maximum output deviation under $z=10$ is 3×10^{-7}). ACT therefore contributes one deterministic candidate per window ($N=1$). The single forward pass requires one network evaluation. Three retraining attempts to resolve the collapse and the subsequent adoption of CEMP for B2/B3 candidate diversity are documented in Section 5.4.
- **Diffusion Policy** (Section 3.2). DP generates each candidate through an independent reverse-diffusion chain using a Denoising Diffusion Implicit Models (DDIM) schedule of ten steps (Section 3.2) [43]. Producing five candidates requires five separate chains and approximately fifty network evaluations per window, roughly an order of magnitude more than ACT. The Diffusion Transformer (DiT) backbone is used. Candidates are structurally independent because each starts from a fresh noise sample; diversity arises from stochastic initialisation rather than from a shared latent manifold.

Evaluating both families under the same three-branch protocol allows the study to ask whether world-model augmentation improves candidate selection in general or only under a particular generation mechanism. A result that holds across both constitutes stronger evidence for the generality of the approach.

4.4 World Model Families

The model families have different evidence roles in the final thesis. They are not treated as equal candidates for the headline claim. A result is promoted to main evidence only when it uses frozen candidates, training-only support, held-out test windows, and a scorer whose action representation is compatible with the candidate pool.

Headline scorer: V-JEPA2-AC. The architecture is reviewed in Section 3.3.2. A frozen V-JEPA2 [20] Vision Transformer (ViT)-Large backbone provides a 1024-dimensional

visual embedding, and a lightweight V-JEPA2-AC transition head predicts a future embedding from the current embedding and an 18-dimensional bimanual action chunk. The distance between the predicted latent and the candidate’s reference latent is used as a selection score. This is the primary semantic world-model scorer because it combines a large pretrained video representation with explicit action conditioning.

Mechanistic controls. V-JEPA2 without action conditioning reuses the same visual representation but removes candidate-specific action information. PointWorld proxy scoring keeps the V-JEPA2-AC visual pathway but degrades the bimanual action chunk into an incompatible robot-arm-style representation. These controls test whether the positive result depends on meaningful action conditioning rather than on a strong visual backbone alone.

Secondary action-regularity scorer: LeWM action-prior. LeWM [22] was initially implemented as a visual rollout and latent-support world-model scorer. That route is not promoted to canonical semantic evidence because the causal diagnostic failed its action-support and ranking-validity gates. A narrower LeWM-derived action-prior scorer is retained as a secondary result. It uses training-only action statistics to prefer candidates closer to the training action distribution. This is a valid action-regularity signal, but it is not interpreted as successful visual rollout imagination.

Fixed-budget negative comparator: DexWM. Dexterous Manipulation World Model (DexWM) [21] is included because it is a close architecture-level comparator for dexterous world modelling. Public pretrained weights for the upstream EgoDex+DROID recipe could not be sourced. The thesis therefore evaluates a reproducible fixed-budget WACT implementation and treats its result as evidence only for that implementation. Full official pretraining is excluded from the main protocol because it requires large DROID raw-data onboarding, substantial preprocessing and inference cost, and a different training scale from the frozen fairness contract.

Exploratory official-checkpoint comparators: DINO-WM and JEPA-WM. DINO-WM and JEPA-WM were implemented later as official pretrained checkpoint routes through the JEPA-WMS codebase. They require an adapter from WACT’s 18-dimensional bimanual wrist chunks to the 7-dimensional DROID-style action interface expected by those checkpoints. Because the validation pilots were mixed and the best settings required validation-only reweighting, they are not promoted to held-out test evidence. They belong to the implementation and appendix record as fair exploratory attempts.

Related reproduction infrastructure: DiWA/CALVIN. DiWA is treated as close related work rather than EgoDex evidence. It studies world-model-guided diffusion-policy improvement through imagined training, whereas WACT studies test-time candidate selection without retraining the policy. A CALVIN reproduction path was prepared, but

missing official policy checkpoints, normalisation statistics, and evaluation metadata prevent a thesis claim. The methodology therefore keeps DiWA/CALVIN separate from the EgoDex branch definitions.

Table 4.2 Evidence role assigned to each implemented or investigated world-model path.

Path	Evidence role	Main claim status
V-JEPA2-AC	Headline semantic scorer	Held-out test claim
V-JEPA2 no-AC	Action-conditioning ablation	Held-out control
PointWorld proxy	Action-degradation control	Held-out mixed control
LeWM visual rollout	Diagnostic scorer	No-go for semantic rollout claim
LeWM action-prior	Training-action regularity scorer	Secondary held-out claim
DexWM fixed-budget	Domain-specific comparator	Held-out negative/null claim
DINO-WM / JEPA-WM	Official-checkpoint exploratory routes	Validation-only, appendix
DiWA/CALVIN	Related reproduction infrastructure	No thesis result

4.5 Fair Experimental Protocol

Comparing the five world model families and two policy architectures against each other requires holding every non-design variable constant. Two aspects of fairness are addressed: training exposure and evaluation contract.

4.5.1 Training exposure budget

Different model families have different computational costs per training step and different sensitivities to data volume. Without a principled exposure budget, a comparison could conflate the effect of architecture with the effect of receiving more or fewer training samples. To determine an appropriate shared budget, a convergence study was conducted prior to full training: each model family was trained independently at six data scales (1k, 5k, 10k, 25k, 50k, and 100k effective items) on a fixed subset of the canonical training data, and the resulting loss curves were examined for plateau behaviour. The study and its results are described in detail in Chapter 5; the finding relevant to the experimental design is summarised here.

The effective training exposure E is defined using the following symbols (see the List of Notations):

E	effective training exposure (items consumed)
S	number of optimisation steps
B	per-device batch size
G	gradient-accumulation steps
A	number of accelerator devices
B_{global}	effective global batch size

$$E = S \times B_{\text{global}} \quad (4.1)$$

where:

$$B_{\text{global}} = B \times G \times A \quad (4.2)$$

This formulation counts the total training items consumed regardless of how steps are distributed across hardware, allowing models with different batch sizes and GPU counts to be normalised to a common scale.

The convergence study, described in Chapter 5, established a universal budget of $E = 100,000$ effective items. This figure is applied uniformly across all model families and all three training splits, ensuring that any measured performance difference reflects architecture and data rather than exposure volume.

4.5.2 Evaluation contract

All models receive the same RGB input at evaluation time: the observation frames from the canonical test window, without additional depth, point cloud, or semantic metadata. This ensures that any performance difference between world model families reflects their internal representation and prediction quality rather than access to a richer input signal. All five families, including PointWorld, operate on the same V-JEPA2 visual latent representations; PointWorld’s distinctive characteristic is its degraded action conditioning rather than a different observation modality.

The evaluation windows are fixed at indexing time and shared across all branches and all world model families. No window is selected, filtered, or re-sampled after training. Policy weights and world model weights are frozen before evaluation begins and are not updated. The only variable at evaluation time is the selection rule applied to the candidate pool, isolating the contribution of each branch design and each world model family from all other sources of variance.

The causal validator checks four leakage boundaries before a run can be treated as evidence. First, support episodes must not overlap with evaluation episodes. Second, support windows must not temporally overlap the held-out test window. Third, task labels may not be used to retrieve test support unless the same restriction is declared for every compared family. The final canonical runs use global retrieval rather than test-task-scoped retrieval. Fourth, repeated-subject metadata are preserved when

available but not used as a selection feature. If a model route cannot provide these guarantees, it remains exploratory or appendix material rather than a headline result.

4.6 Evaluation Metrics

4.6.1 Offline trajectory prediction setting

EgoDex does not define an official offline evaluation metric. The original benchmark measures policy effectiveness via online task success rate on a physical robot deployment [41], which is outside the scope of this study. The evaluation setting here differs fundamentally: rather than measuring task execution, the study evaluates offline visuomotor trajectory prediction, measuring how closely each configuration’s predicted wrist pose sequences align with ground-truth demonstrated trajectories in the held-out test set. Results should therefore be interpreted as trajectory prediction quality under each branch and world model configuration, not as claims about downstream task success. The relationship between offline trajectory prediction quality and online task performance is addressed in Section 3.4 of the literature review and in Section 4.7.

4.6.2 Primary metrics

The action representation throughout this study is a 16-step sequence of bimanual wrist actions. At step t , hand $h \in \{L, R\}$ has predicted translation $\hat{\mathbf{p}}_{t,h}$, target translation $\mathbf{p}_{t,h}$, predicted rotation $\hat{R}_{t,h}$, and target rotation $R_{t,h}$.

Translation L2 error is the mean Euclidean wrist-position error over both hands and all prediction steps:

$$e_{\text{trans}} = \frac{1}{2T} \sum_{t=1}^T \sum_{h \in \{L, R\}} \|\hat{\mathbf{p}}_{t,h} - \mathbf{p}_{t,h}\|_2, \quad T = 16. \quad (4.3)$$

It is reported in metres and is the primary spatial fidelity metric.

Rotation error is the mean geodesic distance on $\text{SO}(3)$, reported in degrees:

$$e_{\text{rot}} = \frac{1}{2T} \sum_{t=1}^T \sum_{h \in \{L, R\}} \cos^{-1} \left(\frac{\text{tr}(\hat{R}_{t,h}^\top R_{t,h}) - 1}{2} \right) \frac{180}{\pi}. \quad (4.4)$$

This avoids the discontinuities of Euler-angle differences and reflects wrist orientation fidelity.

Root-mean-square error is used for diagnostics where a squared aggregate is

more appropriate than a mean absolute distance:

$$\text{RMSE} = \sqrt{\frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2}. \quad (4.5)$$

Here M denotes the number of scalar components being compared.

Endpoint translation uses Equation 4.3 over the final four prediction steps only. Velocity and acceleration diagnostics apply the same paired-error logic after first and second finite differences of the wrist position sequence.

4.6.3 Candidate selection comparison

For each window w , let $m(c, w)$ be the trajectory error of candidate c and let $c_0(w)$ be the policy-only candidate. A branch selecting $c_s(w)$ records a pairwise win when:

$$\mathbb{K}_{\text{win}}(w) = \begin{cases} 1, & m(c_s, w) < m(c_0, w), \\ 0, & \text{otherwise.} \end{cases} \quad (4.6)$$

The win rate is:

$$\hat{p}_{\text{win}} = \frac{1}{W} \sum_{w=1}^W \mathbb{K}_{\text{win}}(w), \quad W = 7607. \quad (4.7)$$

Mean deltas are reported as selected candidate minus candidate 0, so negative values indicate improvement. Confidence intervals are computed by clustered bootstrap over evaluation episodes, and task-clustered intervals are used as a robustness check where task-level dependence is the relevant uncertainty source. Small win-rate differences are not interpreted without their clustered interval or a corresponding paired delta analysis.

4.7 Evaluation Scope and Limitations

The following scope constraints and known limitations apply to the experimental design.

Offline evaluation only. No physical robot is used and no task-completion signal is available; all metrics are computed against held-out demonstrated trajectories (Section 4.6.1). Results should be read as evidence about candidate selection quality in the offline trajectory prediction setting, not as claims about downstream robot task success.

Kinematic action space and embodiment. The action representation is the bimanual $\text{SE}(3)$ wrist space described in Section 4.4, with two consequences for scope:

- *No finger, contact, or force signals.* Finger-level articulation, contact forces, and haptic feedback are absent; tasks whose outcome depends on fine-grained in-hand manipulation or contact dynamics lie outside the evaluation scope.
- *Human-to-robot kinematic gap.* Demonstrations are produced by human subjects, whose wrist kinematics differ from robot manipulators in range of motion and joint compliance. Downstream deployment would require bridging this embodiment gap.

Training budget as a compute-constrained lower bound. The universal 100k-item budget (Section 4.5.1) is not a saturation point for all families; some were still improving at this scale. Performance differences may partly reflect differing distances from respective saturation budgets rather than intrinsic architectural differences.

DexWM pretraining abstention. Published DexWM weights are unavailable [21]. A fixed-budget WACT implementation was evaluated, but the full upstream EgoDex+DROID pretraining recipe was abstained from. The reasons are practical and methodological: DROID raw onboarding is storage- and preprocessing-heavy, DexWM inference and training are substantially more expensive than the other routes, the fixed-budget result did not suggest a likely headline gain, and full upstream-scale pretraining would move outside the common fairness contract used for the main benchmark.

5 Implementation

This chapter describes what was actually implemented and what each route can support as evidence. The implementation does not end with a uniform comparison where every model family becomes a headline result. Instead, it produces an evidence hierarchy: V-JEPA2-AC is the main semantic world-model scorer, LeWM action-prior is a secondary action-regularity scorer, DexWM is a fixed-budget negative comparator, DINO-WM and JEPA-WM are exploratory official-checkpoint routes, and DiWA/CALVIN remains reproduction infrastructure.

The chapter is organised around six concerns: repository ownership and reproducibility (Section 5.1); data and windowing (Section 5.2); training and hardware configuration (Section 5.3); policy candidate generation (Section 5.4); model-specific implementation paths (Section 5.5); and implementation challenges that changed the scientific interpretation of the project (Section 5.10).

5.1 Repository and execution model

The research stack is organised as a single Python package under `src/`, with subsystem entry points exposed under `scripts/`. The main repository is `/home/graham/projects/w-act`. The thesis repository is separate at `/home/graham/dossier`. Reported run artefacts live under `/home/graham/runs/wact`, so code, generated data, and thesis text remain auditable as distinct layers.

Docker is the canonical runtime surface for WACT training and evaluation. The local image is `w-act:latest`, built from `docker/Dockerfile` and `requirements.txt`. The repository README records the intended container as PyTorch with CUDA 12.4 support. Individual routes also use isolated caches when upstream projects impose separate dependency constraints, such as the JEPA-WMS checkout under `/home/graham/cache/jepa-wms` and the DiWA checkout under `/home/graham/cache/diwa/diwa`.

Table 5.1 gives the ownership and reproducibility status of the main implementation paths. Green means original WACT code or a fully controlled adapter. Amber means official upstream code or weights were used but required non-trivial local adaptation. Red means the route was prepared or investigated but does not produce thesis evidence.

The run output directory follows a split-first, then model layout. Each training job writes a `result.json`, model checkpoints, metrics, and a manifest. Canonical evidence additionally records candidate hashes, support-role metadata, window counts, and fair-evaluation receipts. These manifests are treated as part of the result,

Table 5.1 Implementation ownership and evidence status.

Path	Status	Meaning for thesis evidence
DP candidate generation	Green	WACT-controlled candidate pool; five independent DDIM chains used for the headline result.
V-JEPA2-AC scorer	Green/Amber	WACT transition head on official V-JEPA2 features; claimable held-out result.
LeWM action-prior	Green	WACT-derived action regularity scorer; secondary claim, not visual rollout.
DexWM fixed-budget	Amber	Local implementation of published architecture; valid negative/null result for this budget only.
DINO-WM / JEPA-WM	Amber	Official checkpoints loaded through adapted action interface; validation-only mixed results.
DiWA/CALVIN	Red	Infrastructure and artifact checks only; no reproduction or numerical thesis result.

not as optional logs.

5.2 Data pipeline and windowing

5.2.1 Index construction

The index stage is the foundation of the entire stack. It discovers valid paired EgoDex episodes, requiring both a `.mp4` video file and a `.hdf5` pose annotation file, infers episode lengths from the Hierarchical Data Format 5 (HDF5) metadata, and writes canonical observation and prediction windows. The fixed-stride index is configured as `obs16_pred16_s128` following the window protocol described in Section 4.1.2: 16-frame observation windows followed by 16-frame prediction windows, sampled every 128 frames. All training, world model extraction, and guard evaluation in this study consume this windowing contract. The index stage writes resumable checkpoints to allow interrupted builds to continue from the last completed episode.

As established in Section 4.2, the shared window definition ensures any measured performance difference is attributable to branch design or architecture rather than to differences in the data each component received. The canonical index contains 194,333 training episodes and 430,962 training windows for `train_part1_2_3`, plus 3,229 held-out test episodes and 7,607 held-out test windows. The data pipeline records discovered, paired, and rejected episodes before window generation; any unpaired video or annotation file is excluded before training rather than handled downstream.

5.2.2 Incremental training curriculum

Three cumulative training splits were assembled from the first three parts of the EgoDex release, as detailed in Table 4.1 of the methodology. Each split is a strict superset of the previous: the larger splits were assembled using relative symbolic links rather than copying data, allowing incremental construction without duplicating approximately 1 TB of earlier data on disk.

This curriculum allows the data-scaling axis of the benchmark to be evaluated: all five world model families and both policy architectures are trained independently on each of the three splits at the universal budget of 100,000 effective items (Section 5.6), enabling a direct measure of how additional training data diversity affects prediction quality independent of training exposure volume.

5.3 Training configuration

Table 5.2 consolidates the hardware and batch configuration for all model families at the universal budget $E = 100,000$. Settings marked (*shared*) are kept identical across all

trainable models. The effective-item budget is verified for each trainable family via $E = S \times B_{\text{global}}$, where S is the number of optimisation steps and B_{global} is the effective batch per step. V-JEPA2 (no AC) and PointWorld require no model training and reuse the Phase 1 extraction shared with V-JEPA2-AC.

Table 5.2 Training configuration for all model families at $E = 100,000$. Approximate step time is shown for trainable phases only; Phase 1 of V-JEPA2-AC is a one-time offline extraction with no optimiser step. V-JEPA2 (no AC) and PointWorld reuse the V-JEPA2-AC Phase 1 embeddings and require no additional training (--- denotes not applicable).

Family	GPUs	Batch/device	Steps S	B_{global}	s/step
ACT	2×A100	32	1,562	64	≈0.5
Diffusion Policy	1×A100	32	3,125	32	≈2.5
V-JEPA2-AC Ph. 1 (extraction) [†]	4×A100	—	—	—	—
V-JEPA2-AC Ph. 2 (AC head)	4×A100	32	781	128	<0.01
LeWM	4×A100	4	6,250	16	≈0.047
DexWM	3×L40S	4	8,333	12	≈28
V-JEPA2 (no AC)	—	—	—	—	—
PointWorld	—	—	—	—	—

[†] Phase 1 is a one-time offline extraction (4×A100, sharded, 27.4 windows/s, no backward pass). Measured wall times: train_part1 ≈0.75 h, train_part1_2 ≈1.4 h, train_part1_2_3 ≈2.2 h. Phase 2 head training completes in under ten seconds per split (≈0.01 s/step on cached embeddings). DexWM total wall time across all three training splits was approximately 130 h (≈28 s/step × 8,333 steps × 3 splits); LeWM total wall time was approximately 15 minutes (≈0.047 s/step), substantially faster than DexWM despite the identical effective budget.

Across trainable WACT components, random seed 42 is the default unless a run explicitly records a different seed. Training manifests record optimiser steps, effective exposure, device count, batch size, and checkpoint paths. V-JEPA2-AC uses AdamW for the transition head; the JEPA-WMS adapter calibration also uses AdamW with learning rate 10^{-3} . Checkpoint selection is by validation loss where a validation surface exists; otherwise the frozen run manifest and hash identify the exact artefact used for evaluation.

Table 5.3 Model input/output contracts used by the implemented routes.

Route	Input	Output / target	Loss or score
ACT	RGB frame plus bimanual wrist state	16-step, 18D action chunk	Reconstruction plus KL training loss
Diffusion Policy	RGB frame plus wrist state	16-step, 18D action chunk	Diffusion denoising objective
V-JEPA2-AC	V-JEPA2 embedding plus 18D action chunk	Future latent embedding	Latent prediction distance
LeWM rollout	RGB/action sequence	Predicted compact latent	Diagnostic latent-support score
LeWM action-prior	Candidate 18D action chunk	Scalar regularity cost	Training-action z-score prior
DexWM fixed-budget	DINOv2 patch tokens plus pose actions	Future patch/key-point latent	CDiT prediction and HC training losses
DINO-WM / JEPA-WM	Official visual latent plus adapted 7D action	Future official latent	Support/action-policy weighted score

5.4 Policy training

Both policy families share the same observation encoder: a ResNet-18 visual backbone [44] processes the single observation RGB frame (96×96 pixels) through four residual stages (64, 128, 256, and 512 channels), producing a 512-dimensional feature vector that is concatenated with the bimanual wrist state before being passed to the policy-specific decoder (Figure 5.1). Image resizing to 96×96 is applied at loading time and is consistent across all branches. Wrist pose sequences are normalised to zero mean and unit variance using statistics computed from the training split; these statistics are frozen after training and consumed by the world model scoring stages to ensure consistency.

Action Chunking Transformer

The ACT architecture is reviewed in Section 3.1 and illustrated in Figure 3.1. The implementation follows the architecture and training procedure of Zhao et al. [7] with no dependency on an external codebase. Custom components include the EgoDex HDF5 window loader, the 18-dimensional bimanual wrist action space adapter, and a checkpoint serialisation layer that freezes the CVAE configuration and normalisation statistics alongside the model weights. The CVAE encoder-decoder and transformer decoder are realised directly from the paper specification. At training time, the CVAE encoder compresses a ground-truth action chunk together with the current




Figure 5.1 Shared ResNet-18 visual backbone preceding both policy decoders. Adapted from [44].

observation into a style variable $z \sim \mathcal{N}(\mu, \sigma^2)$; the transformer decoder reconstructs the chunk from z and the observation. The encoder is discarded at evaluation time, and z is drawn from the standard Gaussian prior $\mathcal{N}(0, I)$. The design intent is to generate up to N diverse candidates by sampling independent z vectors and decoding each under the same observation. Training uses two A100 GPUs with a per-device batch of 32 (global batch 64): $E = 1,562 \times 64 \approx 100,000$ (Table 5.2). Each run writes a best .pt checkpoint selected by lowest validation loss, a last.pt checkpoint, per-epoch metrics in metrics.jsonl, and a frozen manifest.json recording the configuration and window file consumed.

CVAE posterior collapse: investigation and retraining. The trained ACT checkpoint exhibits CVAE posterior collapse: at inference time the decoder is latent-insensitive, and sampling five independent z vectors from $\mathcal{N}(0, I)$ produces five identical action chunks. A diagnostic confirms the extent of the collapse: injecting $z = 10\sigma$ (ten standard deviations from the prior) via a forward pre-hook on the latent projection layer shifts the output by approximately 3×10^{-7} in absolute value, which is effectively zero. The mean pairwise difference across five independently drawn candidates is 0.000000 on all three training splits.

This is a well-documented failure mode of conditional VAE models rather than a mis-implementation [45]. When the conditioning signal (the visual observation) is sufficiently informative to reconstruct the target without consulting z , the decoder learns to ignore z entirely. The KL divergence term in the training objective regularises the *encoder* distribution toward the prior but cannot compel the *decoder* to use the latent; the collapse is therefore not resolved by increasing the KL weight. Three

retraining runs were conducted to test this:

- v1: $\text{kl_weight}=10$ (default): collapse unchanged.
- v2: $\text{kl_weight}=100$: weighted KL contribution ($100 \times 0.25 \approx 25$ nats) exceeded reconstruction loss (~ 15), yet collapse persisted.
- v3: cyclical KL annealing [46] ($\text{kl_cycle_batches}=200, \text{kl_cycle_ramp_frac}=0.5$) and 30 % observation dropout to force decoder reliance on z : collapse persisted.

All three runs produced mean pairwise differences indistinguishable from zero on every split (Table 5.4). The EgoDex visual observations are sufficiently informative that the decoder finds a loss minimum that ignores z regardless of training pressure, and correcting the collapse without an architectural change (discrete latents, vector-quantisation, or an explicit diversity head) is outside the scope of this thesis.

Table 5.4 ACT CVAE posterior collapse: three successive retraining attempts. kl_weight , cyclical annealing schedule, and observation dropout were varied in sequence; all runs produced a mean pairwise candidate difference of zero on every training split.

Run	kl_weight	Cyclical KL	Obs. dropout	Mean pairwise diff
v1	10 (default)	no	0 %	0.000000
v2	100	no	0 %	0.000000
v3	100	yes (200-batch cycle, 50 % ramp)	30 %	0.000000




Figure 5.2 CVAE posterior collapse in the trained ACT checkpoint. Three independently sampled latent vectors z_1, z_2, z_3 are passed to the CVAE decoder (dashed arrows indicate the latent path the decoder effectively ignores). All inputs produce an identical output $\hat{a}$; the decoder has learned to reconstruct the action from the visual observation alone, making the latent redundant.

The collapse is treated as an empirical finding. It motivates the use of an external candidate generator, described below, and is reported in Section 6.3 as a characteristic of the ACT policy on this dataset.

CEMP candidate augmentation. To restore the five-candidate evaluation pool, ACT-based guards use a CEMP planner: an adaptation of the Cross-Entropy Method [47] initialised from the policy’s output rather than a uniform prior, following

the model-based planning approach of TD-MPC2 [36]. A perturbation is a small SE(3) displacement applied to the base action chunk, shifting each wrist pose by a sampled translation and rotation offset while preserving temporal coherence across the 16-step chunk. Starting from the collapsed ACT B1 output, three iterations each draw 64 such perturbations around the current mean, score them with the world model’s latent prediction function, and retain the top 25 % (16 elites); the standard deviation decays by $0.7\times$ per iteration. The four highest-scoring candidates across all iterations are returned (Figure 5.3).

For Branch 2 the pool contains these four alternatives plus the ACT B1 output ($N=5$); for Branch 3 the B1 output is excluded, leaving the four alternatives only. Because the same world model scores candidates inside CEMP and performs the final selection, the B2 and B3 results measure the quality of the combined world-model-plus-planner system rather than the world model’s scoring ability in isolation.




Figure 5.3 CEMP candidate augmentation. Three iterations of the Cross-Entropy Method initialised from the ACT B1 output refine SE(3) perturbations; the standard deviation decays by $0.7\times$ per iteration. The four highest-scoring candidates form the B2/B3 evaluation pool.

Diffusion Policy

The Diffusion Policy implementation follows the transformer-backbone variant (DP-T) described in Section 3.2 and illustrated in Figure 3.2. The implementation follows Chi et al. [9] with no external codebase dependency. Custom components include the 18-dimensional bimanual wrist action space adapter, the EgoDex HDF5 window loader, and the five-candidate DDIM evaluation protocol (five independent noise seeds decoded under the shared conditioning signal). The same ResNet-18 observation encoder produces a conditioning vector `obs_cond` passed to a DiT denoising backbone via cross-attention. At evaluation time, five independent DDIM denoising chains are run on the same observation to produce five candidate action chunks, each starting from independently sampled Gaussian noise and converging to a distinct trajectory under the shared conditioning signal. Training uses one A100 GPU with a per-device batch of 32: $E = 3,125 \times 32 = 100,000$ (Table 5.2). The same checkpoint, metrics, and manifest artefacts as ACT are written at the end of each run.

DDIM subsampling correction.

$\bar{\alpha}_t$	cumulative noise schedule product at step t (List of Notations)
t_{prev}	next step in the subsampled inference schedule (List of Notations)
$\hat{x}_0$	predicted clean sample (List of Notations)
x_T	Gaussian noise at the schedule endpoint (List of Notations)

An implementation audit identified an error in the reverse-diffusion sampler [43]: the original implementation passed $\bar{\alpha}_{t-1}$ (the consecutive index in the training schedule) instead of $\bar{\alpha}_{t_{\text{prev}}}$:

$$\text{Bug: } \bar{\alpha}_{t-1} \longrightarrow \text{Fix: } \bar{\alpha}_{t_{\text{prev}}} \quad (5.1)$$

For a ten-step schedule $t \in \{90, 80, \dots, 0\}$ this substitutes $\bar{\alpha}_{89}$ for the correct $\bar{\alpha}_{80}$ at the first denoising step, miscalibrating the $\hat{x}_0$ weighting throughout. The diversity guarantee is unaffected: each candidate begins from an independent $x_T \sim \mathcal{N}(0, I)$, so the collapsed-latent failure mode present in ACT cannot arise. The fix supplies the explicit next-in-schedule index to `ddim_step`; all three evaluation splits are re-run under the corrected sampler and the corrected results supersede the initial runs. The correction did not change the qualitative conclusion that Diffusion Policy supplied a diverse candidate pool, but it improved the integrity of the candidate source used by all downstream world-model selectors. The before/after audit is retained in the implementation record rather than used as a separate thesis result.




Figure 5.4 DDIM subsampling correction. The ten-step inference schedule spans $t \in \{90, 80, \dots, 0\}$. At each reverse step the correct $\bar{\alpha}_{t_{\text{prev}}}$ is the *next timestep in the subsampled schedule*; the bug substituted the consecutive training-schedule timestep $\bar{\alpha}_{t-1}$, miscalibrating the $\hat{x}_0$ weighting at every step.

5.5 World model training

5.5.1 V-JEPA2-AC: extraction and transition head

The V-JEPA2-AC architecture is reviewed in Section 3.3.2. Figure 3.6 illustrates the standard V-JEPA2 backbone alongside the action-conditioned variant that forms Phase 2. The frozen ViT-Large backbone is loaded via PyTorch Hub

(facebookresearch/vjepa2) and cached locally; the action-conditioned transition head is implemented and trained from scratch. Training separates into two phases.

In Phase 1, the frozen ViT-Large encoder processes all indexed windows in parallel across four A100 GPUs via a sharded extraction pipeline. The encoder takes 16-frame observation sequences and produces a latent embedding vector per window; these are written to `embeddings.npy` together with a companion `window_ids.txt` recording the exact identifier (split, task, episode, t_0) of each row. Phase 1 is a one-time offline extraction: there is no backward pass, no optimiser, and no epoch loop. A critical bug in an early run applied a 2,000-window default cap producing only approximately 1,000 training pairs from a split of 140,000 windows; this is documented in Section 5.10.

In Phase 2, a small MLP transition head is trained on the stored embeddings to predict the next embedding z_{t+1} from the current embedding z_t and an action conditioning vector derived from the bimanual wrist state delta (`action_repr=policy_state_delta`). Training runs on four A100 GPUs: $E = 781 \times 128 = 99,968 \approx 100,000$ (Table 5.2). Wall-clock time is under ten seconds because each step is a small MLP forward-backward pass with no backbone inference. The trained head is exported via an `act_adapter.json` contract loaded by the guard stage at runtime.

5.5.2 LeWM: end-to-end training

LeWM is reviewed in Section 3.3.2 and illustrated in Figure 3.8. The architecture is implemented following the published specification [22] and trained entirely from scratch with no external weights. Training jointly optimises the ViT-Tiny encoder, the autoregressive predictor, the action embedder, and the projector from random initialisation using the next-embedding prediction loss and the SIGReg regulariser (Equation 3.1) that enforces a Gaussian prior on the latent distribution [22].

EgoDex episodes are exported into contiguous HDF5 sequences containing pixel frames, per-step bimanual wrist state, per-step action (state delta), and per-episode length metadata. Two distinct export paths exist with different frame caps. The training export (`_extractions/lewm_export/`) uses `max_frames_per_episode=128` (the default in `lewm_export_episodes.py`). The guard/evaluation export (`evaluation/lewm_test_export/`) uses `max_frames_per_episode=5280`, derived from the maximum window start position across splits (up to frame 5,280 in `train_part1_2`). An earlier evaluation-export cap of 256 frames caused 97.9 per cent of test windows to be silently excluded; the full resolution of this defect is documented in Section 5.10.

LeWM training runs on four A100 GPUs with a per-device batch of 4 (global batch 16, no gradient accumulation): $E = 6,250 \times 16 = 100,000$ (Table 5.2). Wall-clock

training time on `train_part1_2_3` was approximately 4.9 minutes per run ($\approx 0.047 \text{ s/step} \times 6,250 \text{ steps}$), for a total of approximately 15 minutes across all three training splits, reflecting the lightweight ViT-Tiny + MLP architecture relative to DexWM.

A normalisation statistics defect discovered during the evaluation audit caused the guard stage to load `action_mean` and `action_std` from the test-set export rather than the training-set export, introducing a $1.74\times$ scale mismatch in the LeWM cost function across all initial guard runs. The full description and resolution are documented in Section 5.10.

The final positive LeWM route is not this visual rollout scorer. After the causal diagnostic showed that rollout/support scoring did not rank candidates reliably, a narrower action-prior scorer was implemented. It skips visual rollout and scores each candidate by its normalised distance from training-only action statistics. This implementation is computationally light and causally clean, but its interpretation is intentionally narrower: it measures action regularity, not future visual prediction.

5.5.3 DexWM: from-scratch training

DexWM is reviewed in Section 3.3.2. Its pipeline pairs a frozen DINOv2 image encoder with a CDiT predictor and incorporates a hand-consistency loss that supervises fingertip and wrist heatmap predictions (Figure 5.5). The DINOv2 ViT-Large encoder is loaded from the `facebookresearch/dinov2` HuggingFace repository; the CDiT predictor and HC head are implemented from the published architecture [21] and trained from scratch.

Checkpoint sourcing and abstained full pretraining

The published DexWM architecture [21] exposes architecture code and dataset instructions, but no public pretrained checkpoint could be sourced. The associated HuggingFace repository contains trajectory data rather than weights, and the GitHub repository describes pretraining on EgoDex and DROID without linking a downloadable model checkpoint. This matters because the published result depends on a large upstream training recipe that is not reproduced merely by having the code.

The WACT implementation therefore uses a fixed-budget local route. The frozen DINOv2 ViT-Large encoder provides 1,024-dimensional patch tokens; the CDiT predictor and hand-consistency head are trained from random initialisation on the declared WACT training splits. Training runs on three L40S GPUs with a per-device batch of 4. With 100,000 extracted windows per split, one epoch spans $\lfloor 100,000 / (3 \times 4) \rfloor = 8,333$ optimiser steps, yielding $E \approx 100,000$. Wall-clock training




Figure 5.5 DexWM pipeline as implemented, comprising a frozen DINOv2 encoder, CDiT predictor, lightweight decoder, and HC loss supervision path. Adapted from [21].

time was approximately 43 hours per split, substantially more expensive than LeWM and V-JEPA2-AC head training.

A later preflight tested whether the full official EgoDex+DROID pretraining suite should be reopened. The preflight passed the EgoDex-side checks but reported that the DROID raw root was missing. Onboarding the full DROID raw data would require substantial storage, conversion, indexing, preprocessing, and training time. More importantly, such a run would no longer match the fairness contract used by the main WACT comparison. Since the fixed-budget DexWM result was already valid but negative on the primary metric, full upstream pretraining was abstained from rather than launched as a new open-ended experiment.

This is a reproducibility limitation of the available release as well as a project-scope decision. Releasing code is valuable, but code without the pretrained checkpoint leaves independent researchers to repeat a large-scale pretraining programme before they can evaluate the published model under a new benchmark.

5.5.4 DINO-WM and JEPA-WM exploratory selectors

DINO-WM and JEPA-WM were added after the initial closeout as bounded official checkpoint comparators. The implementation uses the official JEPA-WMS code path and public Meta checkpoints. DINO-WM DROID loaded directly. JEPA-WM DROID also required the official gated DINOv3 ViT-L/16 backbone, which was downloaded

after access was accepted and converted into the native format expected by the upstream loader.

The central implementation issue is action compatibility. WACT candidates are 16-step chunks in an 18-dimensional bimanual wrist action space. The official DROID/RoboCasa JEPA-WMS checkpoints expect a 7-dimensional single-arm robot action. A direct truncation would invent semantics and discard one hand, so the implementation uses a train-only linear adapter from EgoDex action chunks into the expected 7D interface. This keeps the route auditable, but it also weakens the claim: the scorer is an adapted transfer experiment, not a native EgoDex world model.

Both models passed load, extraction, and scorer smoke tests. Validation pilots were mixed. Default DINO-WM and JEPA-WM improved their latent support scores but did not give a clean improvement on immediate trajectory metrics. Validation-only reweighting recovered limited loss or rotation signals, especially for DINO-WM, but translation and endpoint translation were neutral or worse. No held-out test claim is therefore reported from this route; the details belong in the appendix and implementation record.

5.5.5 DiWA/CALVIN reproduction infrastructure

DiWA is the closest training-time comparator to WACT. It studies improving a diffusion policy through imagined world-model rollouts, whereas WACT studies inference-time candidate selection without retraining the policy. A separate DiWA/CALVIN checkout was prepared to test whether that comparison could be reproduced or extended.

The implementation work remained infrastructure only. Public artifacts were downloaded privately, hashed, and checked for archive safety. The close-drawer world model and reward classifier loaded on CPU and in one bounded GPU timing check. A C0–C3 experimental contract was written to separate the untouched upstream reproduction lane from a WACT-style inference-time guidance lane. However, the official base-policy checkpoint, fine-tuned policy checkpoint, `statistics_minmax.yaml`, and exact sampler/evaluation metadata were not available. No policy, optimiser, selector episode, or CALVIN reproduction run is therefore reported as thesis evidence.

5.5.6 V-JEPA2 (no action conditioning): ablation control

The methodological role of this ablation is described in Section 4.4. In implementation, the variant is realised by reusing the Phase 1 embeddings extracted for V-JEPA2-AC; no additional training is required (Table 5.2). At guard time, the scorer passes a zero action vector to the transition model for all candidates, producing action-blind scores

that reflect visual plausibility alone. A dedicated result artefact records the persistence loss (mean MSE plus $0.1 \times$ cosine distance between consecutive frame embeddings, using the same weighted objective as the AC head training loss) as a diagnostic metric. This loss is computed directly as the weighted distance between consecutive stored frame embeddings from `embeddings.npy`. No transition model is invoked for this diagnostic, confirming the world model makes no action-differentiated prediction when conditioned on a zero action vector.

5.5.7 PointWorld: action-degradation control

The methodological role of the PointWorld control is described in Section 4.4. The implementation uses the same frozen V-JEPA2-AC backbone and transition head as the V-JEPA2-AC family, requiring no additional model training (Table 5.2). The sole difference is that candidate action chunks are degraded before being passed to the transition head.

A `HandToArmActionAdapter` extracts the six left-wrist rotation dimensions (`rot6d`, dims 0–5) and computes a scalar summary as the mean of the remaining twelve dimensions (left translation, right rotation, right translation) via the `wrist_plus_gripper` strategy, then zero-pads the result back to 18 dimensions. The right-hand pose is collapsed into the scalar summary dimension and not recoverable from the output; it does not survive as a distinguishable signal in the degraded representation. The transition head therefore receives the correct visual features but systematically corrupted action conditioning.

The convergence study diagnostic additionally computes a proxy loss:

$$\mathcal{L}_{\text{PointWorld-proxy}} = \mathcal{L}_{\text{persist}} + 0.5 \cdot \text{mismatch} \quad (5.2)$$

where $\mathcal{L}_{\text{persist}}$ is the frame-to-frame embedding distance and $\text{mismatch} = 1 - \|\mathbf{arm}\|/\|\mathbf{hand}\|$ is the fraction of action-vector norm discarded by the projection (clipped to $[0, 1]$). This diagnostic confirms that the adapter introduces genuine information loss: the bimanual hand action space does not map faithfully to the 7-DOF arm representation, as expected for a degradation control.

5.6 Convergence study and universal training budget

As described in Section 4.5.1 of the methodology, a convergence study was conducted prior to full fair training to determine an appropriate shared exposure budget. Each of the five primary model families was trained independently at six data scales (1k, 5k, 10k, 25k, 50k, and 100k effective items) on a fixed stratified subset of

train_part1_2_3. The two control families (V-JEPA2 no-AC and PointWorld) were included in the sweep but excluded from the budget selection decision.

5.6.1 Loss curves and plateau detection

Table 5.5 reports the validation loss at each scale for all seven families. Each row uses a different training objective so raw values are not comparable across families; the relevant signal is each family’s own trend across scales. Bold entries mark the earliest scale at which a family’s loss saturates (improvement below 0.5 per cent relative to the next scale).

Table 5.5 Validation loss across six training scales for all seven model families. Bold marks each family’s first scale of saturation; losses are objective-specific and not cross-family comparable.

Family	1k	5k	10k	25k	50k	100k
ACT	0.134	0.134	0.129	0.133	0.130	0.129
Diffusion	10.86	5.22	4.25	3.11	2.57	2.14
DexWM	4.31	2.41	2.00	1.90	1.69	1.36
LeWM	0.335	0.209	0.165	0.145	0.130	0.130
V-JEPA2-AC	1.796	0.375	0.308	0.286	0.257	0.227
V-JEPA2 (no AC)	0.460	0.477	0.373	0.352	0.353	0.377
PointWorld	0.655	0.679	0.628	0.629	0.632	0.592

The bold entries identify the two families that saturate within the sweep: ACT is essentially flat from 1k onwards (a 4 per cent improvement over a 100-fold data increase), and LeWM plateaus between 50k and 100k (0.5 per cent improvement). V-JEPA2-AC, Diffusion Policy, and DexWM carry no bold entry because all three are still declining at 100k, as Figure 5.6 shows. The two controls show no consistent downward trend, consistent with their respective proxy designs having no learnable relationship with additional data.

5.6.2 Universal budget decision

The plateau budget per primary family is summarised in Table 5.6. Given that Diffusion Policy, DexWM, and V-JEPA2-AC all continue to improve at 100k, the universal budget is set to $E = 100,000$ as a compute-constrained lower bound. ACT and LeWM, which plateau earlier, receive the same 100k budget by design: the extra exposure cannot hurt them and ensures all families are compared under the most favourable common condition reachable within the compute envelope.




Figure 5.6 Normalised validation loss (each family’s value divided by its 1k loss) across six training scales. Solid lines are primary models; dashed lines are controls. Large filled circles mark saturation points matching the bold entries in Table 5.5. Families with no marker are still declining at 100k.

Table 5.6 Plateau budget per primary model family. The universal budget is $E = 100,000$ effective items.

Family	Plateau budget	Evidence
ACT	1k	Flat from 1k; 4% drop over $100\times$ more data
LeWM	50k	50k $\rightarrow$ 100k improvement of 0.5%; effectively converged
V-JEPA2-AC	$>100k$	11% drop at 50k $\rightarrow$ 100k; still declining
Diffusion	$>100k$	17% drop at 50k $\rightarrow$ 100k; still declining steeply
DexWM	$>100k$	20% drop at 50k $\rightarrow$ 100k; still declining
Universal budget		$E = 100,000$

This budget is a compute-constrained lower bound rather than a true saturation point for all families. The performance gaps between families reported in the results chapter may partly reflect differing distances from their respective saturation budgets rather than intrinsic architectural differences. This limitation is acknowledged in Section 4.7 of the methodology.

5.7 B2 margin threshold calibration

The Branch 2 guard includes an optional margin threshold δ : the world model's override is accepted only when its preferred candidate outscores the policy output by at least δ . Setting $\delta = 0$ yields unconditional override on every disagreeing window; higher values suppress overrides and collapse B2 toward B1.

A development sweep over $\delta \in \{0.25, 0.50, 1.00\}$ was run on the `train_part1_2` validation split using retrieval-based candidate pools for the V-JEPA2-AC guard. The results followed a clear monotone pattern: at $\delta = 0.25$, approximately 21% of windows were overridden with the largest loss improvement relative to B1; at $\delta = 1.00$, the override rate fell below 0.2% and only 6% of potential gain was recovered.

The canonical evaluation adopts $\delta = 0$ (`selection_mode: best_score`, `override_margin=0.0`), driven by two considerations:

1. **CEMP renders a post-hoc gate partially redundant** (Section 5.9). The internal CEM scoring loop already pre-filters candidates toward world-model-preferred actions before the guard selects; adding a margin gate would reintroduce a hyperparameter that varies across model families and pool types.
2. **Unconditional argmax isolates discriminative quality directly.** Measuring the world model's preference without a threshold suppression layer produces a cleaner experimental signal and avoids per-family calibration.

The sweep results are retained for transparency but do not contribute to any reported result. DexWM was excluded from the sweep: its initial guard runs were found to be vacuous due to a candidate-loading defect (Section 5.10); the canonical DexWM evaluation uses the dedicated `guards_native_dexwm/` pipeline (Section 5.8).

5.8 Guard evaluation

The three-branch framework is defined in Section 4.2; the implementation realises it without bespoke deviation.

5.8.1 Branch 1: policy-only baseline

No world model is called. For ACT, CVAE posterior collapse (Section 5.4) yields a single deterministic chunk, which is recorded directly. For DP, B1 selects the chain with the lowest denoising residual across the five independent samples.

5.8.2 Branch 2: world-model-inclusive selection

The guard scores the five-candidate pool with `selection_mode: best_score` and no margin override. Each world model uses its canonical latent scorer: DexWM via its CDiT patch-token predictor; LeWM by encoding a rollout prefix to establish latent context before scoring each candidate; PointWorld via the proxy function described in Section 5.5.7. V-JEPA2 without action conditioning cannot produce candidate-specific predictions and reduces to B1 by default.

An evaluation audit (Section 5.10) found that V-JEPA2-AC, V-JEPA2 (no AC), and PointWorld initially used retrieval-based pools rather than CEMP; all three were re-evaluated with `candidate_source=planner` to enforce a consistent candidate source across all families.

5.8.3 Branch 3: world-model-exclusive selection

All five world model families use the same generic pipeline: four CEMP alternatives per window with `selection_mode: world_model_only`. Manifest correctness is enforced by verifying that every B3 run records all 7,607 windows under `world_model_only_best_non_act`, confirming the policy chunk is absent. This invariant was used to detect and correct the selection-mode error documented in Section 5.10.

5.8.4 Manifest and checkpoint design

Every guard run emits a frozen manifest at completion recording the selection mode, candidate source, aggregate score statistics, and per-reason override counts. Guard runs also write per-window JSONL checkpoints flushed after each window, allowing interrupted runs to be resumed from the last completed window via a `--resume-dir` flag rather than restarted from scratch.

5.9 Evaluation safeguards

Oracle goal elimination. In early experiments, the latent goal used for scoring was derived from the ground-truth future frame of each episode, giving Branches 2 and 3

access to information unavailable at deployment time. This was corrected by setting `goal_source=none` in all canonical benchmark runs. All results in the results chapter use this corrected configuration.

Candidate pool invariant. Both Branch 2 and Branch 3 score five-candidate pools generated independently per window by the CEMP planner. Branch 2 pools contain the ACT output plus four planned alternatives; Branch 3 pools contain four planned alternatives only. The key controlled variable between the two branches is pool membership (policy output present vs. absent), not the number of candidates, which is five for both. The pool-size correctness of every B3 run is verified through the guard manifest: all 7,607 override-reason records must read `world_model_only_best_non_act` with zero policy fallbacks. This invariant was used to detect the Branch 3 selection-mode error documented in Section 5.10.

Shared comparison surface. The `policy_chunk_compare` stage computes all cross-branch and cross-family metrics on the same prediction targets rather than relying on controller-native summaries, ensuring that translation L2 and rotation error values are computed identically across all configurations. The evaluation window-count shortfall documented in Section 5.10 was detected by checking `num_windows` in each manifest against the expected 7,607.

Test-split isolation. The test split was held out throughout all stages of the project. All training, hyperparameter selection, threshold calibration, and qualitative development work used training and validation splits only.

Candidate source standardisation. An audit conducted after the initial evaluation runs found that V-JEPA2-AC, V-JEPA2 (no AC), and PointWorld used retrieval-based candidate pools (`candidate_source=retrieval`) while LeWM used CEMP-planned alternatives (`candidate_source=planner`). Because the candidate source is a controlled variable in the B2/B3 comparison, a source asymmetry confounds cross-model comparisons: differences between families could reflect candidate quality rather than world model quality. All five families were therefore re-run with `candidate_source=planner` before any results were reported. The original retrieval-based runs are archived under `*_retrieval_archived/` subdirectories and are not used in this thesis. The full audit is documented in Section 5.10.

5.10 Implementation challenges and resolutions

The implementation challenges are part of the contribution because several early runs looked plausible but were not valid evidence. Table 5.7 records the material issues, their impact, the corrective action, rerun status, and remaining risk. Longer forensic

notes are retained in the project archive and appendices.

Table 5.7 Material implementation challenges and resolution status.

Issue	Impact	Fix	Rerun status	Risk
DDIM subsampling error	Miscalibrated reverse-diffusion updates for DP candidates.	Passed the explicit previous timestep t_{prev} into <code>ddim_step</code> .	DP evaluations rerun; corrected candidates used downstream.	Low.
Branch 3 selection-mode error	B3 initially duplicated B2 because candidate 0 remained eligible.	Relaunched with <code>world_model_only</code> ; manifest checks require zero policy fallbacks.	Corrected B3 runs replace earlier outputs.	Low.
Score-cache memory collision	Tensor memory-address reuse could silently corrupt scale factors after garbage collection.	Re-keyed cache on explicit run-scoped identifiers.	Affected runs excluded or rerun.	Low.
Evaluation window shortfall	Default guard limits evaluated 787 rather than 7,607 windows.	Added full-evaluation flags and manifest count checks.	Full 7,607-window runs used.	Low.
LeWM wrong normalization source	Test-set action statistics distorted candidate costs.	Loaded training-only statistics and reran diagnostics.	Rollout scorer still failed; action-prior route separated.	Medium for rollout only.
Candidate-source asymmetry	Earlier families used different candidate sources, confounding comparison.	Final causal matrix uses shared frozen candidates and explicit candidate hashes.	Superseded runs archived.	Low.
DexWM missing checkpoint	Published pretrained model unavailable; zero-score stub risk found.	Stub replaced by hard failure; fixed-budget local route evaluated.	Valid negative/null for local route only.	Medium for general DexWM claims.
DINO-WM/JEPA-WM action mismatch	Official 7D DROID action interface does not match 18D bimanual WACT chunks.	Train-only adapter and validation-only pilots.	Kept exploratory; no held-out claim.	High for headline use.
DiWA missing artifacts	Full CALVIN reproduction blocked by absent policy checkpoints and statistics.	Compatibility infrastructure and artifact checks only.	No thesis result.	High for reproduction.

The remaining risk is concentrated in routes that were not promoted to main claims. The headline V-JEPA2-AC result, LeWM action-prior result, DexWM fixed-budget negative result, and ACT collapse diagnostic all use frozen candidate artifacts, held-out windows, and manifest-level validation.

6 Results

6.1 Overview of the frozen evidence

This chapter reports the final frozen EgoDex evidence used for the dissertation. All canonical generic guard runs use the held-out test split of the `train_part1_2_3` configuration: 7,607 windows from 3,229 episodes and 111 task categories. The split, candidate pools, support sets, and branch authority rules were fixed before the reported evaluations were interpreted. No model was updated during evaluation, and no test episode was used as world model support.

The central result is narrow but positive. With DP candidates, V-JEPA2-AC improves the policy-only candidate on both first-step translation and rotation error. Branch 2 improves translation by 6.62 per cent and rotation by 1.08 degrees. Branch 3 improves translation by 5.70 per cent and rotation by 0.95 degrees. Episode-clustered confidence intervals exclude zero for both metrics. Branch 2 is the stronger version because it can keep the policy candidate when the world-model score is unhelpful.

The other families mainly define the claim boundary. The zero-action ablation preserves the baseline in Branch 2 and therefore confirms that action conditioning is needed for selection. The PointWorld proxy produces mixed behaviour: translation becomes worse while rotation improves. DexWM is a valid negative result because it slightly worsens translation under the same frozen DP candidate protocol. LeWM separates into two findings. Its visual rollout and latent-support diagnostic did not provide a defensible canonical test scorer, but a LeWM-derived action-prior scorer did improve frozen DP candidate selection. This makes V-JEPA2-AC the only strong semantic world-model result in the matrix, while LeWM contributes a secondary positive result about action-distribution regularity. ACT is also a valid null result, but for a different reason: its sampled candidates collapse numerically to the same action.

What the results do not claim. The reported numbers are offline wrist-trajectory errors. They do not measure robot task success, human preference, or closed-loop deployment. A negative result does not prove that a model family is generally useless. It means that under this frozen WACT protocol, this checkpoint and candidate pool did not produce a defensible improvement on the held-out EgoDex metric.

Table 6.1 Frozen evaluation status used in the Results chapter.

Family	Policy	B2	B3	Status and interpretation
V-JEPA2-AC	DP	valid	valid	Headline positive causal result
V-JEPA2 no-AC	DP	valid	valid	Zero-action ablation; confirms action conditioning
PointWorld proxy	DP	valid	valid	Mixed action-degradation control; not published PointWorld
DexWM	DP	valid	valid	Negative/null causal result under frozen protocol
V-JEPA2-AC	ACT	valid	valid	Null because sampled ACT candidates collapse
V-JEPA2 no-AC	ACT	valid	valid	Expected null under collapsed candidates
PointWorld proxy	ACT	valid	valid	Expected null under collapsed candidates
LeWM action-prior	DP	valid	valid	Secondary positive action-regularity result; not visual rollout evidence

6.2 Diffusion Policy guard results

Table 6.2 gives the main DP results. Each delta is the selected candidate minus candidate 0, so negative values mean the guard selected a better trajectory than the policy-only baseline. The candidate-0 selection column is also important. In Branch 2, the guard can keep candidate 0. In Branch 3, candidate 0 is excluded by design, so the selection rate is structurally zero.

Table 6.2 Primary DP guard results on the frozen 7,607-window EgoDex test set.

Scorer	Branch	Candidate 0 kept	Translation change	Translation delta, episode 95% CI	Rotation delta, episode 95% CI
V-JEPA2-AC	B2	23.1%	-6.620%	-0.007557 [-0.008043, -0.007062]	-1.081° [-1.205, -0.960]
V-JEPA2-AC	B3	0.0%	-5.702%	-0.006509 [-0.007040, -0.005965]	-0.946° [-1.084, -0.810]
V-JEPA2 no-AC	B2	100.0%	0.000%	0.000000 [0.000000, 0.000000]	0.000° [0.000, 0.000]
V-JEPA2 no-AC	B3	0.0%	-0.032%	-0.000037 [-0.000580, 0.000495]	-0.148° [-0.288, -0.015]
PointWorld proxy	B2	19.9%	+0.500%	+0.000571 [+0.000063, +0.001069]	-0.393° [-0.521, -0.260]
PointWorld proxy	B3	0.0%	+0.500%	+0.000571 [+0.000038, +0.001112]	-0.382° [-0.529, -0.240]
LeWM action-prior	B2	20.0%	-10.236%	-0.011685 [-0.012162, -0.011227]	-0.849° [-0.972, -0.726]
LeWM action-prior	B3	0.0%	-9.009%	-0.010284 [-0.010792, -0.009790]	-0.741° [-0.885, -0.603]

The V-JEPA2-AC result supports the core WACT hypothesis in its most bounded form: a semantic world model can improve inference-time selection when it scores a shared policy-generated candidate pool. The improvement is not only a mean shift. Branch 2 wins on translation in 50.7 per cent of windows, ties candidate 0 in 23.1 per cent, and loses in 26.2 per cent. The tie rate is exactly the set of windows where

Branch 2 keeps the policy candidate. This matters because it shows why Branch 2 is preferable to Branch 3. Branch 2 gains from world-model ranking while retaining a fallback to the policy’s own best sample.

The controls sharpen the interpretation. V-JEPA2 no-AC assigns tied scores in Branch 2 and therefore returns the baseline unchanged. The PointWorld proxy selects candidates that improve rotation but worsen translation. This is useful as a degradation control, but it is not evidence that the published PointWorld model was evaluated. In this dissertation, PointWorld proxy means the WACT adapter that reuses the same visual latent surface with deliberately degraded action conditioning. The LeWM action-prior rows are also not semantic world-model evidence in the same sense as V-JEPA2-AC. They show that a training-action regularity signal can select better DP candidates, even when the LeWM visual rollout objective itself was not reliable enough for a canonical test scorer.

6.3 ACT guard results

The ACT guard matrix is valid but not informative as positive selection evidence. The shared ACT candidate pool contains candidate 0 from the deterministic policy pass and four seeded CVAE samples. In practice these samples collapse to almost the same action. The mean root-mean-square difference from the deterministic candidate is approximately 2.04×10^{-8} , and the mean maximum absolute difference is approximately 6.45×10^{-8} . Additional sampling-parameter tests did not materially increase candidate separation, so the null result is attributed to limited candidate diversity under this ACT checkpoint rather than to a failure of the guard evaluator.

This collapse explains Table 6.3. The scorer can change the selected index, but the selected action is numerically identical to candidate 0 for practical purposes. The correct interpretation is therefore “no usable candidate diversity”, not “world-model scoring failed for a diverse ACT pool”.

Table 6.3 ACT guard results. Values are numerically zero because the candidate pool collapses.

Scorer	Branch	Candidate 0 kept	Translation change	Rotation delta
V-JEPA2-AC	B2	58.1%	approximately 0%	approximately 0°
V-JEPA2-AC	B3	0.0%	approximately 0%	approximately 0°
V-JEPA2 no-AC	B2	100.0%	0.000%	0.000°
V-JEPA2 no-AC	B3	0.0%	approximately 0%	approximately 0°
PointWorld proxy	B2	69.8%	approximately 0%	approximately 0°
PointWorld proxy	B3	0.0%	approximately 0%	approximately 0°

The practical consequence is that the ACT rows should not be used to rank world models. They diagnose a policy-candidate problem. For this checkpoint, ACT’s

stochastic candidate generator does not create action alternatives with enough numerical separation for an inference-time selector to matter.

6.4 DexWM and LeWM closeout

DexWM was evaluated through its native causal path using the same frozen DP candidate pool. The result is a valid negative or null result. Branch 2 increased translation error by 0.570 per cent and endpoint translation error by 0.922 per cent. Branch 3 increased translation error by 0.563 per cent and endpoint translation error by 0.913 per cent. The rotation deltas were slightly negative, but their episode-clustered intervals included zero. The primary claim is therefore that DexWM did not improve the frozen offline trajectory metric in this setting.

Table 6.4 DexWM paired causal results against the frozen DP candidate-0 baseline.

Com- pari- son	Metric	Relative change	Mean delta	Episode 95% CI
B2 vs B1	Translation L2	+0.570%	+0.000651	[+0.000135, +0.001164]
B2 vs B1	Endpoint translation L2	+0.922%	+0.001099	[+0.000456, +0.001713]
B2 vs B1	Rotation error	-0.129%	-0.051°	[-0.187, +0.077]
B3 vs B1	Translation L2	+0.563%	+0.000642	[+0.000098, +0.001184]
B3 vs B1	Endpoint translation L2	+0.913%	+0.001088	[+0.000434, +0.001745]
B3 vs B1	Rotation error	-0.210%	-0.084°	[-0.230, +0.060]

LeWM reached a more nuanced endpoint. Historical LeWM evaluations used future-goal or test-derived support assumptions and are not used as causal evidence. A bounded causal diagnostic was implemented instead. It removed future-goal access, used training-only support, reused frozen DP candidates, and converted policy-state chunks into LeWM’s trained action-delta representation. The original rollout and latent-support scorer failed its progression gates: 1,283 converted action-delta values exceeded the training action bounds, and the score-margin versus translation-improvement rank association was negative at -0.0207.

The useful LeWM signal came from a narrower source. When the scorer was reduced to a normalized action-prior cost, using only training-action statistics, it passed the validation diagnostic and was then evaluated on the frozen test candidate pool.

Branch 2 improved translation by 10.236 per cent and endpoint translation by 7.806 per cent. Branch 3 improved translation by 9.009 per cent and endpoint translation by 6.649 per cent. Both branches also reduced mean rotation error. Table 6.5 reports the paired test deltas.

Table 6.5 LeWM action-prior causal results against the frozen DP candidate-0 baseline.

Com- parison	Metric	Relative change	Mean delta	Episode 95% CI
B2 vs B1	Translation L2	-10.236%	-0.011685	[-0.012162, -0.011227]
B2 vs B1	Endpoint translation L2	-7.806%	-0.009303	[-0.009880, -0.008709]
B2 vs B1	Rotation error	-2.124%	-0.849°	[-0.972, -0.726]
B3 vs B1	Translation L2	-9.009%	-0.010284	[-0.010792, -0.009790]
B3 vs B1	Endpoint translation L2	-6.649%	-0.007924	[-0.008582, -0.007287]
B3 vs B1	Rotation error	-1.855%	-0.741°	[-0.885, -0.603]

This result should not be read as evidence that LeWM visual imagination solved the selection problem. It shows that the action distribution learned during LeWM training is useful as a regularizer over DP samples. The broader interpretation is therefore conservative: V-JEPA2-AC is the only strong semantic world-model scorer in this dissertation, while DexWM and the LeWM rollout path expose training and alignment limitations of the available implementations.

This distinction matters for DexWM as well. Public pretrained DexWM weights were not available for this dissertation. The implemented checkpoint therefore tests a reproducible fixed-budget DexWM training route, not the full original training recipe at its intended scale. The negative DexWM result is valid for the evaluated implementation, but it should not be interpreted as a general disproof of DexWM-style representations for WACT.

6.5 Overall interpretation

The frozen results support one main empirical claim. Under a causal offline protocol, V-JEPA2-AC improves DP candidate selection on held-out EgoDex wrist-trajectory errors. This is the strongest world-model result in the dissertation because it combines a pretrained visual representation with explicit action conditioning. The effect is clearer for Branch 2 than for Branch 3 because Branch 2 combines world-model ranking with

policy fallback. This supports the policy-first version of WACT more strongly than the world-model-authority version.

The remaining results prevent overclaiming. Zero-action scoring does not help. PointWorld proxy is mixed and cannot be presented as a published PointWorld result. DexWM is valid but negative on the primary metric. ACT does not provide a meaningful branch comparison because its candidate pool collapses. LeWM visual rollout did not pass the causal diagnostic needed for a semantic world-model test claim, although its training-action prior did improve DP candidate selection. Taken together, the evidence shows that world-model guidance can help, but only when the scorer has action-conditioned information, sufficient representation quality, and a diverse candidate pool that matches its training support.

7 Discussion and Limitations

7.1 What the benchmark shows

The final benchmark gives a narrower and more defensible conclusion than the earlier project plan. World-model guidance can improve offline trajectory selection, but the positive semantic result is concentrated in Diffusion Policy candidates scored by V-JEPA2-AC. Branch 2 reduces first-step translation error by 6.62 per cent and rotation error by 1.08 degrees relative to the policy-only candidate. Branch 3 also improves both metrics, but less strongly. This supports the policy-first version of WACT: keep the policy candidate available, then use the world model to choose among diverse policy-generated alternatives.

The corrected evidence does not support a general claim that adding any world model improves action selection. DexWM is a valid negative/null result for the fixed-budget WACT implementation. The PointWorld proxy is mixed, with worse translation but better rotation. The zero-action ablation confirms that action conditioning is necessary. ACT mainly diagnoses candidate collapse. LeWM visual rollout did not become a valid semantic scorer, although LeWM action-prior scoring gives a secondary positive result through training-action regularity.

This creates a useful scientific result. WACT works when three conditions align: the candidate pool is diverse, the scorer receives meaningful action-conditioned information, and the candidate actions remain compatible with the scorer’s training support. When one of these conditions fails, the world-model label alone is not enough.

7.2 Why Branch 2 is stronger than Branch 3

The Branch 2 and Branch 3 comparison is central to the interpretation. Branch 3 is the stronger world-model-authority condition because candidate 0 is excluded. If the world model is reliable enough, this should let it escape the policy’s preferred action more often. In the final Diffusion Policy plus V-JEPA2-AC result, that extra authority does not help. Branch 2 improves translation by 6.62 per cent, while Branch 3 improves it by 5.70 per cent.

The reason is practical. Branch 2 keeps candidate 0 on 23.1 per cent of windows. Those cases are not failures of WACT; they are cases where the policy candidate is already competitive or the world-model score does not justify changing it. Branch 3 has no such fallback. It must choose from candidates 1 to 4 even when candidate 0 is the safer option. The result therefore favours world-model-guided

selection over unrestricted world-model authority.

7.3 Interpreting the model-family differences

The model-family comparison should be read as an evidence hierarchy, not as a simple leaderboard. V-JEPA2-AC is the only family that gives a clear positive semantic world-model result on the headline Diffusion Policy evaluation. This is consistent with the role of its pretrained video backbone and action-conditioned head. The zero-action version removes that action signal and collapses to the baseline in Branch 2, confirming that the action-conditioned component is doing real work.

LeWM reached a different endpoint. Its rollout and latent-support scorer failed the causal diagnostic once oracle and transductive shortcuts were removed. The valid LeWM result is narrower: an action-prior scorer based on training-only action statistics improves frozen Diffusion Policy candidate selection. This is important, but it should not be called successful visual imagination. It shows that action regularity can be useful even when the visual rollout world-model claim is not supported.

DexWM is important because it prevents a biased narrative. It was implemented and evaluated under a fixed-budget WACT route, but it did not improve the primary translation metric. The thesis should not interpret this as a general failure of DexWM-style models, because the public pretrained checkpoint for the published EgoDex+DROID recipe was not available and full pretraining was outside the fairness contract. The correct interpretation is narrower: under the local fixed-budget implementation, DexWM did not provide a positive WACT selector.

DINO-WM and JEPA-WM add a further boundary. They are official pretrained checkpoint routes, but their DROID action interface is not native to WACT's 18-dimensional bimanual wrist chunks. The validation pilots show that a model can improve its own latent support score while failing to improve the trajectory metric that matters for this thesis. This reinforces the central point: world models must be action-compatible and metric-aligned before they can become useful evaluators.

7.4 Reproducibility and open model releases

The DexWM investigation raises a broader reproducibility issue. Open-source code for recent world-model systems is valuable. It allows independent researchers to inspect architectures, reuse data loaders, and understand how a result was constructed. That openness matters in a field where increasingly capable AI systems risk being concentrated among organisations with unusually large compute and data resources.

At the same time, code alone is not equivalent to a reproducible model release when the published claim depends on large-scale pretraining. In the DexWM case, the architecture and data recipe are available, but the pretrained checkpoint for the reported model was not. Reproducing the full recipe would require large DROID raw-data onboarding, preprocessing, storage, and compute. For a Master's thesis with a fixed fairness contract, that is not a reasonable prerequisite for checking whether the model works as a WACT selector.

The point is not that incomplete releases have no value. They do. The point is that benchmark papers and downstream theses must distinguish between open code, open data, open checkpoints, and fully reproducible evidence. These are different levels of openness, and they lead to different claim boundaries.

7.5 Why benchmark evolution is part of the contribution

A major contribution of this dissertation is the construction of a defensible evaluation surface. Several earlier results looked stronger but were not clean enough for final claims. Some used test-derived support. Some allowed future-goal access. Some used inconsistent candidate generation. The final result set is weaker in headline size but stronger academically because it removes these shortcuts.

This matters because the thesis is not only asking whether WACT can produce a large number. It is asking whether world-model involvement can be evaluated fairly as an inference-time selection mechanism. The final protocol fixes the candidate pool, separates training support from test windows, blocks future-frame access, and reports negative outcomes when the evidence does not support a claim. That process is part of the scientific contribution.

7.6 Why semantic alignment and object grounding were excluded

Semantic alignment and object grounding remain useful implemented modules, but they are excluded from the final benchmark claims. Adding them only to V-JEPA2-AC would give one model family privileged inputs not available to LeWM, DexWM, or the controls. That would change the question from whether the world model improves selection under a shared protocol to whether extra semantic information helps a specific model.

A future experiment can test semantic conditioning directly, but it should do so symmetrically. Either every compared family receives equivalent semantic inputs, or

the experiment should be labelled as an ablation of the V-JEPA2-AC path alone.

7.7 Main limitations

The first limitation is scope. The benchmark is offline and trajectory-level. It measures how closely selected wrist-motion chunks match expert demonstrations. It does not execute the action in a simulator or on a robot, and it does not measure task success.

The second limitation is candidate diversity. The ACT results show that a valid evaluation can still be uninformative if the candidate generator does not produce meaningfully different actions. WACT needs both a scorer and a useful candidate set.

The third limitation is model coverage. LeWM visual rollout, DexWM full upstream-scale pretraining, DINO-WM, JEPA-WM, and DiWA/CALVIN all reached bounded implementation states but not headline evidence states. These routes are valuable because they define limits, but they do not become additional positive claims.

The fourth limitation is metric coverage. First-step translation and rotation errors are meaningful for offline imitation fidelity, but they are incomplete. They do not measure grasp stability, finger articulation, object interaction, or long-horizon recovery after an error.

8 Conclusion

8.1 Summary of findings

This dissertation asked whether world models can improve action selection for complex bimanual hand-motion prediction at inference time. The final answer is conditional. World-model guidance helps in the strongest clean setting tested: Diffusion Policy candidates scored by V-JEPA2-AC. Branch 2 reduces first-step translation error by 6.62 per cent and rotation error by 1.08 degrees on the held-out EgoDex test set. Branch 3 also improves both metrics, but less strongly.

The result supports the policy-first WACT design. The best-performing branch is not full world-model authority. It is the branch that lets the world model rank policy-generated candidates while still retaining the policy’s own preferred candidate when that candidate is best.

The broader result is not that every world model helps. DexWM is a valid negative/null result for the evaluated fixed-budget route. LeWM visual rollout did not pass the causal diagnostic required for a semantic rollout claim, while LeWM action-prior scoring provides a secondary action-regularity result. DINO-WM and JEPA-WM were implemented as official-checkpoint exploratory routes but did not produce clean validation evidence for a held-out test claim. The ACT evaluation is null because the sampled candidate pool collapses numerically. These outcomes define when WACT is reliable and when it is not.

8.2 Methodological contribution

The thesis contributes a reproducible evaluation protocol for separating causal world-model guidance from misleading shortcuts. The final protocol fixes the test windows, uses frozen candidate pools, blocks future-goal access, uses training-only support, and keeps branch authority explicit. It also reports negative or no-go outcomes instead of retuning after test exposure.

This protocol changed the scientific interpretation of the project. Earlier historical results appeared stronger, but some relied on transductive support, oracle access, inconsistent candidate generation, or incomplete branch alignment. The final evidence is smaller in headline magnitude but stronger as academic evidence.

8.3 Reproducibility lesson

A secondary conclusion concerns reproducibility. The study benefited from open research releases, but it also exposed the difference between open code and complete reproducible evidence. DexWM is the clearest case. The architecture and training recipe are public, but the pretrained checkpoint for the reported model was not available. Repeating the full EgoDex+DROID pretraining recipe would have required substantial data onboarding, storage, preprocessing, and compute, and would have moved outside the fairness contract of this thesis.

The conclusion is not that code releases are unimportant. They are essential for independent research and for avoiding a field where only large organisations can inspect or build on frontier systems. The more careful point is that model papers should make checkpoint availability and reproduction requirements clear. When a published result depends on pretraining that most researchers cannot repeat, downstream work must treat the missing checkpoint as a claim boundary.

8.4 Scope of conclusions

The conclusions concern offline wrist-trajectory prediction on EgoDex. They do not claim closed-loop robot success, physical task completion, or solved dexterous manipulation. The benchmark measures whether a selected candidate is closer to the demonstrated wrist motion than the policy-only candidate.

Within that scope, the conclusion is clear. World-model scoring can improve inference-time action selection when the candidate pool is diverse and the scorer has meaningful action-conditioned representations. It does not help merely by being added to the system.

8.5 Future Work

The most immediate next step is online evaluation. A simulator or robot study is needed to test whether the offline wrist-error improvements translate into task success.

A second direction is candidate generation. The ACT collapse shows that WACT needs diverse candidate actions. Future work should improve policy sampling or use a candidate generator whose diversity can be measured before world-model scoring.

A third direction is stronger causal LeWM development. Future LeWM work should separate visual rollout evidence from action-prior evidence from the start, using validation-only development and a fresh untouched test resource.

A fourth direction is full reproduction of close related systems such as DiWA. The infrastructure prepared in this thesis can support such work if official policy checkpoints, normalisation statistics, and evaluation metadata become available.

Finally, WACT guidance could be distilled back into a policy. If the world model reliably identifies better candidates, accepted candidates could become a teacher signal for policy retraining. That would move WACT from test-time selection toward training-time policy improvement.

References

- [1] B. D. Argall, S. Chernova, M. Veloso, and B. Browning, "A survey of robot learning from demonstration," *Robotics and Autonomous Systems*, vol. 57, no. 5, pp. 469–483, 2009. DOI: 10.1016/j.robot.2008.10.024
- [2] A. Y. Ng and S. J. Russell, "Algorithms for inverse reinforcement learning," in *Proceedings of the Seventeenth International Conference on Machine Learning*, Morgan Kaufmann, 2000, pp. 663–670.
- [3] M. Jeannerod, "The timing of natural prehension movements," *Journal of Motor Behavior*, vol. 16, no. 3, pp. 235–254, 1984. DOI: 10.1080/00222895.1984.10735319
- [4] M. Jeannerod, M. A. Arbib, G. Rizzolatti, and H. Sakata, "Grasping objects: The cortical mechanisms of visuomotor transformation," *Trends in Neurosciences*, vol. 18, no. 7, pp. 314–320, 1995. DOI: 10.1016/0166-2236(95)93921-J
- [5] M. Santello, M. Flanders, and J. F. Soechting, "Postural hand synergies for tool use," *Journal of Neuroscience*, vol. 18, no. 23, pp. 10 105–10 115, 1998. DOI: 10.1523/JNEUROSCI.18-23-10105.1998
- [6] S. Ross, G. Gordon, and D. Bagnell, "A reduction of imitation learning and structured prediction to no-regret online learning," in *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 15, 2011.
- [7] T. Z. Zhao, V. Kumar, S. Levine, and C. Finn, *Learning fine-grained bimanual manipulation with low-cost hardware*, Apr. 2023. DOI: 10.48550/arXiv.2304.13705 arXiv: 2304.13705 [cs].
- [8] Z. Ding, "Imitation learning," in *Deep Reinforcement Learning: Fundamentals, Research and Applications*, H. Dong, Z. Ding, and S. Zhang, Eds. Singapore: Springer Singapore, 2020, pp. 273–306, ISBN: 978-981-15-4095-0. DOI: 10.1007/978-981-15-4095-0_8 [Online]. Available: https://doi.org/10.1007/978-981-15-4095-0_8
- [9] C. Chi et al., *Diffusion policy: Visuomotor policy learning via action diffusion*, Mar. 2024. DOI: 10.48550/arXiv.2303.04137 arXiv: 2303.04137 [cs].
- [10] S. Jain et al., *DiWA: Diffusion policy adaptation with world models*, ArXiv preprint, Aug. 2025. arXiv: 2508.03645 [cs].
- [11] K. M. Lynch and F. C. Park, *Modern Robotics: Mechanics, Planning, and Control*. Cambridge University Press, 2017, ISBN: 978-1-107-15630-2.

- [12] Omitram, *Left-hand vs. right-hand coordinate systems: The 3d developer's guide* - omitram – omitram.com, <https://omitram.com/3d-coordinate-systems-left-vs-right-hand/>, 2024.
- [13] Y. Zhou, C. Barnes, J. Lu, J. Yang, and H. Li, "On the continuity of rotation representations in neural networks," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 5745–5753. DOI: 10.1109/CVPR.2019.00589
- [14] N. Grossmann, E. Gröller, and M. Waldner, "Concept splatters: Exploration of latent spaces based on human interpretable concepts," *Computers & Graphics*, vol. 105, pp. 73–84, 2022, ISSN: 0097-8493. DOI: <https://doi.org/10.1016/j.cag.2022.04.013> [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0097849322000656>
- [15] A. Dawid and Y. LeCun, "Introduction to latent variable energy-based models: A path towards autonomous machine intelligence," *Journal of Statistical Mechanics: Theory and Experiment*, 2023. arXiv: 2306.02572 [cs].
- [16] D. P. Kingma and M. Welling, *Auto-encoding variational bayes*, 2013. DOI: 10.48550/arXiv.1312.6114 arXiv: 1312.6114 [stat.ML].
- [17] K. Sohn, H. Lee, and X. Yan, "Learning structured output representation using deep conditional generative models," in *Advances in Neural Information Processing Systems*, vol. 28, 2015.
- [18] A. Vaswani et al., "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017. DOI: 10.48550/arXiv.1706.03762 arXiv: 1706.03762 [cs.CL].
- [19] M. Assran et al., "Self-supervised learning from images with a joint-embedding predictive architecture," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. arXiv: 2301.08243 [cs].
- [20] M. Assran et al., *V-JEPA 2: Self-supervised video models enable understanding, prediction and planning*, Jun. 2025. DOI: 10.48550/arXiv.2506.09985 arXiv: 2506.09985 [cs].
- [21] R. G. Goswami et al., *World models can leverage human videos for dexterous manipulation*, Dec. 2025. DOI: 10.48550/arXiv.2512.13644 arXiv: 2512.13644 [cs].
- [22] L. Maes, Q. Le Lidec, D. Scieur, Y. LeCun, and R. Balestriero, *Leworldmodel: Stable end-to-end joint-embedding predictive architecture from pixels*, Verified from local PDF metadata and first page, Mar. 2026. arXiv: 2603.19312 [cs.LG].

- [23] G. Zhou, H. Pan, Y. LeCun, and L. Pinto, *DINO-WM: World models on pre-trained visual features enable zero-shot planning*, 2024. arXiv: 2411.04983 [cs.LG]. [Online]. Available: <https://arxiv.org/abs/2411.04983>
- [24] X. Li et al., *Evaluating real-world robot manipulation policies in simulation*, CoRL 2024, 2024. arXiv: 2405.05941 [cs].
- [25] Y. Li, Y. Zhu, J. Wen, C. Shen, and Y. Xu, *WorldEval: World model as real-world robot policies evaluator*, 2025. arXiv: 2505.19017 [cs].
- [26] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, *Film: Visual reasoning with a general conditioning layer*, 2017. arXiv: 1709.07871 [cs.CV]. [Online]. Available: <https://arxiv.org/abs/1709.07871>
- [27] *MTDP: Modulated Transformer Diffusion Policy Model*, <https://arxiv.org/html/2502.09029v1>. Accessed: Apr. 26, 2026.
- [28] M. Reuss, J. Pari, P. Agrawal, and R. Lioutikov, *Efficient diffusion transformer policies with mixture of expert denoisers for multitask learning*, 2024. arXiv: 2412.12953 [cs.LG]. [Online]. Available: <https://arxiv.org/abs/2412.12953>
- [29] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, *Mastering diverse domains through world models*, Nature, 2025, 2023. arXiv: 2301.04104 [cs].
- [30] J. Quevedo, A. K. Sharma, Y. Sun, V. Suryavanshi, P. Liang, and S. Yang, *WorldGym: World model as an environment for policy evaluation*, Sep. 2025. DOI: 10.48550/arXiv.2506.00613 arXiv: 2506.00613 [cs].
- [31] Q. Bu et al., *Closed-loop visuomotor control with generative expectation for robotic manipulation*, Oct. 2024. DOI: 10.48550/arXiv.2409.09016 arXiv: 2409.09016 [cs].
- [32] S. Nair, A. Rajeswaran, V. Kumar, C. Finn, and A. Gupta, “R3M: A universal visual representation for robot manipulation,” in *Conference on Robot Learning (CoRL)*, 2022. arXiv: 2203.12601 [cs].
- [33] A. Brohan et al., *RT-1: Robotics transformer for real-world control at scale*, 2022. arXiv: 2212.06817 [cs.R0]. [Online]. Available: <https://arxiv.org/abs/2212.06817>
- [34] N. Tsagkas, A. Sochopoulos, D. Danier, C. X. Lu, and O. Mac Aodha, *When pre-trained visual representations fall short: Limitations in visuo-motor robot learning*, 2025. arXiv: 2502.03270 [cs].
- [35] W. Huang et al., *Pointworld: Scaling 3d world models for in-the-wild robotic manipulation*, 2026. arXiv: 2601.03782 [cs.R0]. [Online]. Available: <https://arxiv.org/abs/2601.03782>

- [36] N. Hansen, H. Su, and X. Wang, "TD-MPC2: Scalable, robust world models for continuous control," in *International Conference on Learning Representations (ICLR)*, 2024. arXiv: 2310.16828 [cs].
- [37] T. Yu et al., "MOPO: Model-based offline policy optimization," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv: 2005.13239 [cs].
- [38] R. Kidambi, A. Rajeswaran, P. Netrapalli, and T. Joachims, "MOREL: Model-based offline reinforcement learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. arXiv: 2005.05951 [cs].
- [39] A. Argenson and G. Dulac-Arnold, "Model-based offline planning," in *International Conference on Learning Representations (ICLR)*, 2021. arXiv: 2008.05556 [cs].
- [40] Z. Liu, T. Swann, S. Garg, S. Levine, and A. Kumar, *Model-based offline planning with trajectory pruning*, 2021. arXiv: 2105.07351 [cs].
- [41] R. Hoque, P. Huang, D. J. Yoon, M. Sivapurapu, and J. Zhang, *EgoDex: Learning dexterous manipulation from large-scale egocentric video*, ICLR 2026, 2025. DOI: 10.48550/arXiv.2505.11709 arXiv: 2505.11709 [cs].
- [42] University of Malta, *M.AI.ESTRO: Mixed reality AI-enhanced skill transfer and robotic optimisation*,
<https://www.um.edu.mt/newspoint/news/2025/05/maiestro-mixed-reality-ai-enhanced-skill-transfer-robotic-optimisation>, University of Malta
 newspoint announcement, 2025.
- [43] J. Song, C. Meng, and S. Ermon, *Denoising diffusion implicit models*, Oct. 2020. DOI: 10.48550/arXiv.2010.02502 arXiv: 2010.02502 [cs.LG].
- [44] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.
- [45] S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Jozefowicz, and S. Bengio, "Generating sentences from a continuous space," *arXiv preprint arXiv:1511.06349*, 2016. arXiv: 1511.06349.
- [46] H. Fu, C. Li, X. Liu, J. Gao, A. Celikyilmaz, and L. Carin, "Cyclical annealing schedule: A simple approach to mitigating KL vanishing," *arXiv preprint arXiv:1903.10145*, 2019. arXiv: 1903.10145.
- [47] R. Y. Rubinstein and D. P. Kroese, *The Cross-Entropy Method: A Unified Approach to Combinatorial Optimization, Monte-Carlo Simulation and Machine Learning*. Springer, 2004.

Appendix A Auxiliary Implemented Modules

This appendix describes the two auxiliary modules that are fully implemented in the WACT stack but were excluded from the final fair benchmark. The semantic alignment pipeline and the object grounding stage each produce useful artefacts and remain integrated with the broader system. Their absence from the main results chapters reflects a deliberate methodological choice rather than a finding about their utility or correctness.

A.1 Rationale for exclusion from the main benchmark

The canonical benchmark evaluates two world model families under a shared input contract: the same raw RGB observation frame and the same bimanual wrist state sequence are provided to both V-JEPA2-AC and LeWM. No additional conditioning is admitted on either side.

Semantic alignment artefacts (captions, tags, text embeddings) and object grounding artefacts (bounding boxes, segmentation masks) could in principle be incorporated into the V-JEPA2-AC scoring pipeline through the semantic extraction path. The LeWM architecture, however, has no equivalent semantic conditioning module. Including semantic or grounding conditioning for V-JEPA2-AC alone would change the nature of the comparison from a world model family comparison under equal information to a mixed comparison that conflates world model architecture with access to additional input channels. The two scientific questions are distinct, and the thesis addresses only the former.

The exclusion applies symmetrically across all five evaluated branch instances. Both the V-JEPA2-AC family and the LeWM family run without semantic or grounding inputs in the final benchmark. Future work under a symmetric design, in which both families receive equivalent semantic conditioning, could isolate the effect of this information channel independently of the model family comparison.

A.2 Semantic alignment pipeline

A.2.1 Purpose and outputs

The semantic alignment stage converts indexed EgoDex windows into structured semantic artefacts for downstream conditioning and inspection. For each window, the pipeline produces:

- A natural language caption describing the observable manipulation activity in the selected frames.
- Structured tags organised into categories, including the contact object, secondary objects, and observed actions.
- A canonical semantic text string constructed from the structured tags, suitable for embedding.
- A dense text embedding vector produced by a configurable embedding backend.

The primary downstream use is to provide a semantically meaningful conditioning vector that can be concatenated with or appended to the visual latent representations used by the world model scoring functions. In the current implementation, this conditioning path is available but is not activated in the canonical benchmark.

A.2.2 Visual language model selection

The VLM component was evaluated across several model and embedding combinations before the operational configuration was stabilised. The semantic model comparison table records these evaluations and their relative quality on caption coherence and tag precision. The operationally selected configuration uses Qwen/Qwen3-VL-2B-Instruct as the visual language model for caption and tag generation, with the google/siglip-so400m-patch14-384 model as the default text embedding backend (`WACT_EMBED_BACKEND=siglip_text`).

The Qwen3-VL-2B-Instruct model was selected for practical reasons: it is compact enough to run on a single GPU alongside other pipeline stages, it produces grammatically coherent captions from single-frame inputs, and its structured tag output is consistent enough to support deterministic canonical text construction. The SigLIP embedding backend was selected because it provides a joint image-text embedding space that is directly compatible with the V-JEPA2-AC semantic scoring path, allowing future similarity computations between caption embeddings and visual latents without an additional projection layer.

OpenCLIP and VLM-text hidden-state backends are also implemented and configurable. The SigLIP-backed Qwen3-VL run represents the strongest row in the comparison table for the current stack configuration.

A.2.3 Critic and retry loop

A key engineering feature of the semantic pipeline is an integrated critic and retry loop. The naive approach of captioning once and trusting the output produces unreliable

results on EgoDex frames, which are first-person egocentric views with partially occluded hands, cluttered surfaces, and rapidly changing visual content. The critic module scores the resulting caption against the input frames using a configurable similarity measure. If the score falls below a threshold or plateaus across retries, the pipeline either retries the VLM call with adjusted prompting parameters or marks the row as `semantic_failure=true` and writes it with a failure flag.

This design means that the semantic stage produces a complete, aligned output for every window in the index, with explicit quality flags rather than silently dropping or corrupting rows on difficult windows. The failure-flagged rows remain usable for inspection and can be filtered before any downstream conditioning step.

The implemented retry logic is:

1. Score the caption against the sampled frame using the configured score mode.
2. If the score exceeds the acceptance threshold, commit the row.
3. If the score is below threshold, retry with a revised prompt up to the configured maximum number of retries.
4. If all retries fail or the score plateaus without improvement, commit the row with `semantic_failure=true`.

A.2.4 Multi-GPU sharding and output artefacts

The semantic pipeline supports multi-GPU sharded execution through a shard supervisor (`semantic_shard_supervisor.py`) and a deterministic shard merge step (`semantic_merge_shards.py`). This is necessary because captioning and embedding 7,607 test windows or tens of thousands of training windows at VLM inference speeds is infeasible on a single GPU within a reasonable run budget.

The canonical merged output layout is:

```
/runs/w-act/egodex/semantic/<split>/<model_id>/<build_id>/
  embeddings.npy
  window_ids.txt
  meta.jsonl
  tags.jsonl
  captions.jsonl
  manifest.json
  index.npz
```

During a multi-GPU run, intermediate shard directories (`shard_0/`, `shard_1/`, and so on) accumulate partial outputs. The merge step validates row counts before

producing the canonical root artefact set, ensuring that partial or interrupted shard outputs are not silently used as complete builds. The manifest records the input fingerprint, allowing downstream stages to detect cache invalidation if the source windows file is replaced.

A.2.5 Integration with the broader pipeline

The semantic stage is positioned between the data indexing stage and the object grounding stage. The OPE stage (Section A.3) reads semantic manifests to obtain the canonical window ordering and to construct object prompts from the semantic tags produced by the VLM. This tight coupling is intentional: it ensures that semantic and grounding artefacts are always aligned to the same window index, preventing silent join drift if the windows file is regenerated.

The entrypoint for the semantic stage is `src/wact/cli/semantic_build_index.py`, with orchestration handled by `scripts/semantic.sh` or the stage-runner wrapper `scripts/run_stage.sh semantic`. The notebook companion `notebooks/semantic_aligner.ipynb` provides interactive inspection of caption quality, tag distributions, and embedding nearest-neighbour retrievals.

A.3 Object grounding and mask generation

A.3.1 Scope and current implementation

The object grounding stage takes semantic artefacts from the previous stage and produces spatially grounded detection and segmentation outputs for each window. The current implementation provides two-dimensional bounding-box detection and box-conditioned mask generation. It does not implement canonical six degree-of-freedom pose estimation, FoundationPose-style geometry reconstruction, or object identity tracking across frames. The name “object pose estimation” in the codebase reflects the intended scope of a future extension rather than the current capability.

For each window, the implemented pipeline:

1. Reads the semantic row for the window from the upstream semantic build.
2. Samples one or more frames from the window according to the configured frame sampling strategy.
3. Constructs object prompts from the structured semantic tags.

4. Runs the configured detector to produce bounding boxes.
5. Runs the configured masker on each detected box to produce segmentation masks.
6. Writes aligned detection and mask artefacts alongside a progress checkpoint.

A.3.2 Prompt construction

Object prompts are constructed from semantic tags using a priority scheme implemented in `src/wact/ope/prompting.py`. The logic prioritises the `contact_object` tag, which the VLM assigns to the primary object being actively manipulated, and then appends deduplicated entries from the `objects` tag list. Generic labels such as “object”, “thing”, “tool”, and “hand” are filtered before the prompt is passed to the detector, as these terms produce unreliable or overspecific detections in the grounding model. When the VLM returns no object tags, the prompt construction falls back to a task-derived hint derived from the EgoDex task label.

This is category-level prompting rather than instance-level prompting. The detector does not distinguish between two instances of the same object class, and the resulting masks identify regions containing the prompted category rather than a specific tracked object instance.

A.3.3 Detection and segmentation backends

The production backends, configured through `configs/stage_defaults.env` and accessible via the stage-runner, are:

- Detector: `IDEA-Research/grounding-dino-base` (Grounding DINO), a zero-shot open-vocabulary detector that accepts free-text prompts.
- Masker: `facebook/sam2-hiera-large` (Segment Anything Model 2), a promptable segmentation model that produces high-quality binary masks from bounding box inputs.

The combination of Grounding DINO for detection and SAM2 for box-conditioned segmentation is a well-established zero-shot grounding pipeline. It does not require object-specific training data from EgoDex and generalises across the 194 task categories in the dataset without fine-tuning.

The bare `scripts/ope.sh` script defaults to dummy backends for development and debugging. Production runs must use the stage-runner or explicitly override the backend flags to select the real Grounding DINO and SAM2 models.

A.3.4 Output artefacts and alignment

The canonical output layout is:

```
/runs/w-act/egodex/pose/<split>/<pipeline_id>/<build_id>/
  masks.jsonl
  detections.jsonl
  window_ids.txt
  errors.jsonl
  progress.json
  manifest.json
```

The `masks.jsonl` file is the primary downstream artefact: it contains a row per window with the detected object categories, bounding box coordinates, and run-length-encoded or polygon-format segmentation masks. The `detections.jsonl` file records the raw detection outputs before masking and is retained for audit purposes. The `progress.json` file supports resumable execution through the shard supervisor.

Alignment between the OPE output and upstream artefacts is enforced by inheriting the `window_ids.txt` ordering from the semantic build. The OPE stage reads the semantic manifest to resolve the windows pickle that was used to produce the semantic index, ensuring that detection and mask rows correspond to the same window boundaries used at every earlier pipeline stage.

A.3.5 Known limitations

The current implementation has several limitations that are relevant to any future use of OPE outputs for manipulation conditioning:

1. The 2D masks identify spatial regions in the image frame but do not provide three-dimensional geometry information. Object depth, relative pose, or spatial relationship between manipulated objects are not recoverable from the mask output alone.
2. Small or visually similar objects can confuse the Grounding DINO detector, particularly in the cluttered egocentric field of view typical of EgoDex demonstrations. When two similar objects are present, the detector may merge them or select the wrong instance.
3. There is no temporal tracking across frames within a window. Each sampled frame is processed independently, and there is no mechanism to maintain object identity across the prediction horizon.

4. The zero-shot grounding approach depends on the quality of the upstream semantic tag output. A missed or misidentified contact object in the semantic stage will propagate to an incorrect or absent detection in the OPE stage.

These limitations are consistent with the current scope of two-dimensional grounding and mask generation. They would need to be addressed before OPE outputs could serve as reliable conditioning for world model scoring in a deployed system.

A.4 Future use under a symmetric design

The exclusion of semantics and object grounding from the final benchmark is a fairness decision, not a statement that these modules lack value. Both pipelines are implemented, tested on the training split, and produce aligned artefacts that are stored alongside the canonical benchmark outputs.

The most natural future experiment would introduce semantic conditioning symmetrically across both world model families. This would require either adapting the LeWM architecture to accept a text embedding as an additional conditioning channel, or training a new LeWM variant with semantic inputs from the outset. Under a symmetric design, the contribution of semantic conditioning to world model scoring quality could be measured independently of the model family comparison, and the effect of caption quality, embedding backend selection, and object grounding granularity could each be evaluated in isolation.

Similarly, the object grounding masks could be used to produce spatially localised visual features by masking the RGB input before the visual backbone processes it. Whether this localisation improves world model guidance in the manipulation setting, or whether the pretrained backbone’s attention mechanism already implicitly grounds relevant objects without explicit masking, is an open empirical question.

[FIGURE PLACEHOLDER: semantic pipeline overview – input window to VLM to caption plus tags to embedding backend to output artefacts; critic/retry loop shown as a feedback arc between caption and score threshold]

[FIGURE PLACEHOLDER: object grounding example – sampled EgoDex frame with Grounding DINO bounding box and SAM2 mask overlaid on detected contact object; prompt text shown]

Appendix B Benchmark Evolution and Validity Record

This appendix documents the development trajectory of the WACT benchmark and the validity concerns that were identified and resolved before the final canonical results package was established. Its purpose is to provide an auditable record of the decisions and corrections that shaped the final evaluation surface, and to explain why certain earlier result sets were superseded or rejected.

The record is presented in roughly chronological order of discovery. Each section identifies the concern, describes its effect on the results as reported at the time, and explains the correction applied.

B.1 The three benchmark eras

The dissertation results passed through three distinct evaluation eras before the canonical package was established. Understanding the transition between eras is necessary for interpreting references to earlier result numbers in project documentation.

Era 1: Bounded pilot slices. The earliest quantitative comparisons were conducted on small, handpicked windows selected to represent the diversity of EgoDex task categories. These were useful for rapid controller iteration and for identifying which architectural choices were clearly harmful, but they could not support generalisable claims because the selection was not held-out from the development process. All Era 1 results were explicitly demoted from the thesis claims surface once full held-out evaluation became practical.

Era 2: March 2026 frozen three-branch bundle with limited LeWM coverage. The first evaluation to use the held-out test split was the March 2026 frozen package. This package provided valid Branch 1 and Branch 2 V-JEPA2-AC results, but its LeWM coverage was severely limited by a frame cap configuration error (described in Section B.2). The Branch 3 selection mode error (described in Section B.3) also meant that Branch 3 V-JEPA2-AC was not independently evaluated. Both defects required correction before the three-branch comparison could be treated as scientifically valid.

Era 3: Canonical final package (April 2026). The April 14, 2026 package achieved full 7,607-window coverage for the LeWM family after the frame cap fix. A further correction on April 18, 2026 re-ran Branch 3 V-JEPA2-AC with the correct `selection_mode=world_model_only` flag, producing the first genuine Branch 3 evaluation. The `train_part1_2_3` training split was used for all five branch instances. This package is the canonical results surface for all quantitative claims in the thesis.

B.2 LeWM frame cap defect

B.2.1 Discovery

The LeWM export pipeline requires a `max_frames_per_episode` parameter that controls how many frames from each EgoDex episode are included when building the training corpus. This parameter has a direct impact on the coverage of the resulting evaluation surface, because the guard can only evaluate windows whose constituent frames were included in the export.

During the March 2026 benchmark packaging, the frame cap was set to `max_frames_per_episode=256`. This value was not derived from the actual window length distribution of the benchmark window set; it was a development default carried over from early pipeline testing. When the benchmark results were inspected, the LeWM guard had produced predictions for only 160 of the 7,607 test windows, representing 2.1 per cent of the intended evaluation surface.

The coverage failure was identified by comparing the count of LeWM guard output rows against the expected window count. The root cause was confirmed by recomputing the maximum frame index required to cover all windows in the benchmark window list and finding that the actual maximum was 5,280 frames, far above the 256-frame cap in use.

B.2.2 Effect on reported results

Any LeWM result reported from the March 2026 package is based on 160 windows, not 7,607. The 160 covered windows are not a random sample of the test split; they are the windows whose episodes happened to fall within the 256-frame cap. These windows are concentrated in episodes with early-episode actions and do not represent the full range of task categories or annotation confidence levels in the test split.

The practical effect is that the March 2026 LeWM numbers cannot be compared directly to Branch 1 or Branch 2 V-JEPA2-AC numbers, because the LeWM evaluation is on a different, non-representative slice of the test data. Any conclusion drawn from these numbers about the relative merit of the LeWM family would be confounded by the coverage difference.

B.2.3 Correction

The frame cap was recomputed by iterating over the benchmark window list, reading the episode frame indices for each window, and taking the maximum observed frame index across all windows. The resulting cap of 5,280 frames was applied to the export

rebuild. The LeWM model was not retrained; only the export was rebuilt with the corrected parameter. The guard was then re-run, achieving full 7,607-window coverage on the April 14, 2026 package.

The corrected LeWM numbers in the canonical package are therefore not comparable to the March 2026 LeWM numbers on a like-for-like basis. The March 2026 numbers are retained in project documentation for audit purposes but are not cited in any thesis claims.

B.3 Branch 3 selection mode defect

B.3.1 Discovery

The three-branch design requires Branch 3 to use `selection_mode=world_model_only`, meaning that only the world model's predicted latent distance ranks the candidates, without any contribution from the ACT policy's own scoring function. This is the experimental condition that tests whether the world model alone, without policy guidance in the selection step, can identify better trajectories.

After the April 14, 2026 package was assembled, the per-window pairwise comparison between Branch 3 V-JEPA2-AC and Branch 2 V-JEPA2-AC was computed as a validation check. The result was a 0.0 per cent win rate for Branch 3 against Branch 2 across all 7,607 windows. That is, Branch 3 did not produce a single window outcome that differed from Branch 2.

This was immediately identified as inconsistent with any genuine difference in selection logic between the two branches. Inspection of the Branch 3 guard manifest revealed that the selection mode was recorded as `best_score`, not `world_model_only`. The chain evaluation script (`chain_eval_part1_2_3.sh`) had been written to launch both Branch 2 and Branch 3 guards, but the Branch 3 launch did not pass the `--selection-mode world_model_only` flag. Both guards therefore ran with `selection_mode=best_score` on the same frozen candidate pool, producing identical outputs.

B.3.2 Effect on reported results

Until the defect was corrected, the "Branch 3 V-JEPA2-AC" entry in the April 14 package was scientifically identical to Branch 2 V-JEPA2-AC. The three-branch comparison was effectively a two-branch comparison. Any conclusion about whether world-model-only selection outperforms hybrid selection under V-JEPA2-AC could not be drawn from this data.

The defect was non-obvious from the output metrics alone. The reported Branch 3 translation error (0.0354, identical to Branch 2) was not implausible; one might expect Branch 3 and Branch 2 to produce similar numbers if the world model’s ranking correlates with the combined score’s ranking. The defect was only detectable through pairwise comparison, which revealed the zero win rate that is impossible if the two branches genuinely differ.

B.3.3 Correction

The Branch 3 V-JEPA2-AC guard was re-run on April 18, 2026 with the correct flag:

```
--selection-mode world_model_only
--candidate-source frozen_manifest
--candidate-manifest-run-dir <B2_JEPA_run_dir>
```

Using the frozen candidate pool from Branch 2 ensured that the Branch 3 result reflects only the difference in selection logic, not any difference in the candidates available. Guard manifest verification after the run confirmed that all 7,607 windows were decided by the world model (`override_reason_counts: {world_model_only_best_non_act: 7607}`), confirming that the corrected selection mode was in effect for every window.

The corrected Branch 3 V-JEPA2-AC result (translation $L_2 = 0.0328$, 91.5 per cent win rate against Branch 1) is the first genuine evaluation of world-model-only selection in the WACT stack.

B.4 Oracle goal leakage

B.4.1 Discovery

In early V-JEPA2-AC guard configurations, the world model scoring function was provided with the terminal latent state of the evaluation episode as a goal signal. The scoring function measured the predicted final latent distance to this goal and used it as part of the candidate ranking. This goal signal was derived from the ground-truth future frames of the demonstration, information that would not be available in a deployed system.

The oracle goal leak was identified during the fair comparison audit when it was noted that the V-JEPA2-AC guard configuration differed from the ACT-only baseline in the information it received at evaluation time. The ACT policy generates candidates from the current observation alone and has no access to future frames. Providing the world model with the ground-truth terminal latent created an asymmetry: the world

model scoring function was effectively guided toward candidates that reached the demonstrated endpoint, while the ACT policy’s own ranking had no equivalent oracle signal.

B.4.2 Effect on reported results

Any result from a guard run using oracle goal conditioning overstates the benefit of world model guidance, because part of the improvement is attributable to the oracle signal rather than to the world model’s intrinsic representation quality. The magnitude of the overstatement cannot be precisely determined from the available data, but it is substantial enough to make oracle-conditioned results non-comparable to policy-only results.

Early exploration of V-JEPA2-AC guidance showed results that were more impressive than the final canonical numbers. The oracle goal was a contributing factor. Removing the oracle signal and rerunning under the fair contract produced lower but more defensible numbers. The final canonical results (77.0 per cent reduction in mean translation error) are therefore conservative relative to the earliest reported V-JEPA2-AC results.

B.4.3 Correction

The fair evaluation contract was established by setting `goal_source=none` for all branch evaluations. Under this setting, the world model scoring function does not receive any goal latent derived from future frames. The candidate ranking is based solely on the world model’s predicted latent dynamics given the observed context and the proposed action chunk.

This correction was applied to all five branch instances in the canonical benchmark. The change from oracle goal to no goal required validating that the world model scoring function still produced meaningful candidate rankings without the terminal goal signal, which was confirmed by the final translation improvement results.

B.5 Goal metric invalidity

B.5.1 Discovery

The V-JEPA2-AC guard natively reports a metric called `goal_12` alongside the translation error. In the original implementation, this metric stored different quantities in Branch 1 and Branches 2 and 3. In Branch 1, it stored the translation L_2 error for the

selected candidate. In Branches 2 and 3, it stored the JEPA latent-space distance between the predicted final state and the goal latent.

When `goal_source=none` was adopted as the fair evaluation contract, there was no longer a meaningful goal latent for Branches 2 and 3. The `goal_12` metric for these branches consequently reported a latent distance of zero for every window, making the metric uninformative. Comparing `goal_12` across branches was therefore comparing a real translation error in Branch 1 against a vacuous zero in Branches 2 and 3.

B.5.2 Correction

The `goal_12` metric was replaced with `endpoint_translation_12`, defined as the mean translation L_2 error computed over only the final quarter of the prediction window (prediction steps 12 through 16). This metric measures how accurately the predicted trajectory arrives at the correct final position, which is often more relevant to task success than early-window prediction accuracy. The endpoint metric is computed identically across all branch instances and is not affected by the presence or absence of a goal latent. It is included in the canonical benchmark shape metric tables reported in the results chapter.

B.6 LeWM CPU inference defect

B.6.1 Discovery

During the first full test-split guard run for LeWM after the frame cap fix, the estimated completion time for 7,607 windows was approximately 54 hours. This was implausibly slow relative to the expected GPU inference throughput for a 15-million-parameter model. Inspection of the guard inference loop revealed a hard-coded `.to("cpu")` call in the LeWM model forward pass that overrode any device specification passed at runtime. The model was running on CPU regardless of whether a CUDA device was available.

B.6.2 Correction

The hard-coded device assignment was replaced with a dynamic device detection that checks for CUDA availability and falls back to CPU only if no GPU is detected. After this correction, the LeWM guard ran at approximately 19.5 windows per second on a single GPU, completing the full 7,607-window test evaluation in roughly 6.5 minutes. The effective speedup was approximately 130-fold relative to the CPU-bound run.

This defect is worth recording because it illustrates a failure mode common in research code developed iteratively: a debugging convenience (forcing CPU to avoid CUDA memory errors during development) was never removed before production use. The correction required no changes to the model architecture or training, only to the device placement logic in the inference wrapper.

B.7 HDF5 handle management defect

B.7.1 Discovery

The data loading loop in an early version of the guard infrastructure opened and closed the HDF5 data file once per evaluation window. On a test set of 7,607 windows, this produced approximately 145,000 open-close cycles (one open and one close per window, plus frame-level re-opens within each window). The overhead from repeated file handle acquisition was a significant fraction of total evaluation wall time on large test splits.

B.7.2 Correction

The data loading implementation was refactored to open the HDF5 file once at the start of the evaluation run, hold the handle in memory, and close it only when the evaluation completes. This change reduced per-window I/O overhead to a small seek operation rather than a full open-close cycle and improved guard evaluation throughput noticeably on the full test split.

B.8 Guard resumability

B.8.1 Motivation

Full test-split guard runs covering 7,607 windows take several minutes on a GPU but can be interrupted by node preemption, out-of-memory errors, or external compute scheduler events. Without checkpointing, an interrupted run required restarting from window zero, discarding all partial results.

B.8.2 Implementation

A per-window JSONL checkpointing system was added to the guard infrastructure. After each window is evaluated, a row is appended to a checkpoint file in the output

directory. When a run is invoked with `--resume-dir` pointing to an existing output directory, the guard reads the checkpoint file, identifies which window IDs have already been completed, and skips them. The final manifest and aggregate metrics are computed from the complete set of per-window rows, including both the resumed and newly computed windows.

This design ensures that an interrupted run can be resumed without any manual intervention and without requiring the window evaluation order to be deterministic. The checkpoint is written atomically per row, so a crash mid-window produces at worst one incomplete row that the resume logic can detect and skip.

B.9 Rejected comparison inputs and superseded paths

B.9.1 Semantics and OPE as benchmark inputs

An early hypothesis held that incorporating semantic alignment and object grounding into the V-JEPA2-AC scoring function would produce the most informative world model guidance, and that this enriched V-JEPA2-AC configuration should be the primary benchmark representative for the V-JEPA2-AC family. This hypothesis was rejected for the reasons described in Appendix A and in Section 7.6 of the discussion chapter: the enriched configuration gives V-JEPA2-AC privileged inputs not available to LeWM, confounding the family comparison.

The decision to reject this path was not based on the performance of the enriched configuration. The lean RGB-plus-trajectory contract was adopted because it is the only contract that supports a scientifically unconfounded comparison between the two world model families.

B.9.2 Veto and threshold selectors

Before the current scoring function was stabilised, several candidate selection mechanisms based on explicit threshold vetos were implemented and evaluated. These included the `act_first_veto` design, which rejected world model candidates whose support distance exceeded a fixed threshold, and the `veto_m10` selector, which applied a margin-based veto to prevent the world model from selecting candidates that were far from the ACT policy’s preferred candidate even if they scored well on the world model metric alone.

These mechanisms were useful for early exploration of the policy-guard interface and helped establish that naive world model selection could be dominated by retrieval artefacts when the candidate pool contained low-quality outliers. They were superseded by the structured scoring function in Branch 2 (`best_score`: combined

latent distance, action out-of-distribution penalty, and support distance term) and the clean world-model-only design in Branch 3. Neither the veto selectors nor the margin-based threshold approaches are used in the canonical benchmark.

B.9.3 Early Branch 3 shells

Several experimental Branch 3 designs preceded the final frozen representative. The first explicit world-model-only shell for LeWM was evaluated on the fair held-out slice and was found to be weaker than the ACT-only baseline on aggregate slice metrics, showing that world-model-only selection requires a world model of sufficient quality before it improves over the policy alone. Neutral-seed variants of Branch 3, which removed the ACT seed from the candidate pool, also collapsed on the fair slice, demonstrating that the bottleneck was the controller structure rather than any single parameter choice. These negative results were important for understanding the design space and for eventually arriving at a Branch 3 configuration that genuinely improves over Branch 2 in aggregate.

B.9.4 Conservative confidence gating

A further rejected path explored whether applying more conservative confidence gating to the LeWM guard, requiring higher override margins before the world model was permitted to override the policy candidate, could make LeWM competitive with V-JEPA2-AC on the shared benchmark. The bounded sweep showed that increasing the override margin made LeWM more conservative but did not produce a competitive replacement for V-JEPA2-AC on chunk-level comparison metrics. This was recorded as a useful negative result: the performance gap between the two world model families is not primarily explained by calibration of the override threshold, and therefore cannot be closed by threshold tuning alone.

B.10 Summary of validity status

Table B.1 summarises the seven validity concerns identified during the benchmark development cycle and the resolution status of each at the time the canonical package was assembled.

Concerns 1, 3, and 6 were resolved by changes to the benchmark artefacts or analysis code. Concern 2 is addressed by consistent framing throughout the thesis: translation L_2 is positioned as an offline diagnostic proxy rather than a task success predictor, and the connection to deployed outcomes is explicitly left as future work. Concern 4 is partially addressed by the per-task win rate distribution reported in the

Table B.1 Summary of benchmark validity concerns and their resolution status in the canonical final package.

#	Concern	Severity	Status in canonical package
1	Branch 3 V-JEPA2-AC identity defect (<code>selection_mode</code> bug)	Critical	Resolved: re-run April 18, 2026
2	<code>translation_12</code> not paper-backed as primary metric	Moderate	Documented: framed as offline diagnostic proxy
3	Shape metrics suppressed from reported summary	Moderate	Resolved: shape metrics added to results chapter
4	No per-task win rate breakdown	Moderate	Partially resolved: aggregate confirmed broad; worst-5 reported
5	Wrist-only evaluation scope	Low–Moderate	Acknowledged design constraint; documented in limitations
6	Confidence scores unused in evaluation	Low	Resolved: stratification analysis added to results
7	Offline-to-online gap	Moderate	Acknowledged: explicitly bounded in scope section

results chapter, which confirms that improvement is broad and not driven by a small number of tasks. Concerns 5 and 7 are acknowledged design constraints that are transparent in the scope of conclusions and limitations sections.

The final canonical results are drawn from a benchmark in which all critical and moderate concerns have been either resolved or explicitly bounded. No result cited in the thesis claims chapter relies on a defective or uncorrected evaluation artefact.

Appendix C Auxiliary Implementation Modules

Two modules are fully implemented in the repository but were excluded from the primary benchmark evaluation. They are documented here as reference material for any continuation of this line of work toward a general framework for semantic feature enrichment, policy candidate generation, and world model selection.

Semantic alignment. This stage converts indexed windows into captions, structured tags, and text embeddings. A Vision Language Model (VLM) generates free-form descriptions and task-relevant categorical tags for each window; a separate embedding model maps the text to a shared representation space. A critic-and-retry loop rejects low-quality outputs and requests regenerated captions, providing a quality filter before downstream use.

Object grounding and mask generation. This stage consumes semantic object prompts from the alignment stage and produces bounding box detections and segmentation masks for identified objects. The current implementation provides pixel-space localisation; SE(3) object pose estimation is not yet included.

Both modules were excluded from the primary benchmark because a shared input contract across all five world model families could not be maintained: the semantic and grounding outputs could not be wired into the frozen DINOv2 and ViT-Large encoder paths used by DexWM, V-JEPA2-AC, and related families without violating the evaluation fairness requirement (Section 4.5.2). The initial intuition motivating these modules — that frozen encoder pre-extraction followed by lightweight head training is productive — has since been validated by the DexWM and DINOv2-based world model architectures, each of which adopts an analogous two-phase structure. Extending this framework to incorporate semantic and spatial grounding signals remains a natural direction for future work.

Appendix D Notes for Future Revision

This appendix collects scope notes identifying parts of the thesis that would require revision if certain extensions were carried out after submission. They are gathered here to keep the main chapters clean while preserving a complete record for any follow-on work.

D.1 Strategy 2 comparative evaluation

All benchmark results in this thesis use the fixed-stride windowing contract (Strategy 1, `obs16_pred16_s128`) exclusively. The state-change-targeted alternative (Strategy 2, Section 4.1.2) is fully implemented and has been used for qualitative development inspection, but a systematic Strategy 1 vs. Strategy 2 comparison has not been conducted and is noted as a natural future direction.

If a full Strategy 2 evaluation were completed and the two contracts compared as co-primary evaluation designs, the following sections would require revision:

- **Section 4.1.2** (index construction): restore the Strategy 1 / Strategy 2 bold-paragraph pattern. Add a paragraph describing the state-change-targeted index in full (peak detection, task-stratified sampling, one window per episode with a fixed seed, no temporal overlap). Clarify that both strategies enforce the shared-window invariant independently.
- **Section 4.1.3** (training splits): extend Table 4.1 to include Strategy 2 window counts per split and task; update the test-index description to distinguish the two contracts.
- **Section 4.6** (evaluation metrics): add explicit strategy labels to all result tables and figures; add a dedicated analysis subsection for cross-strategy comparison if results differ materially.
- **Results chapter**: all quantitative results would need split-by-strategy labelling to make clear which sampling contract each row corresponds to.