

M.AI.ESTRO-Twin Vision Pipeline Placement Report

Graham Pellegrini

0352804L

University of Malta

graham.pellegrini.22@um.edu.mt

ARI5907 Placement in AI

Placement Title:

M.AI.ESTRO-Twin: Automated Digital Twin Synthesis for Industrial Skill Transfer

Supervisor:

Prof. Alexiei Dingli

Abstract

This report presents the final outcome of the ARI5907 placement, which contributed a vision perception pipeline to the M.AI.ESTRO-Twin project. The placement addressed the problem of generating structured perception signals automatically from egocentric manipulation demonstrations without manual prompt entry. Six vision language models were benchmarked on a shared 64-window sample drawn from the EgoDex dataset, with SigLIP and OpenCLIP compared as semantic quality backends. Qwen3-VL-2B-Instruct paired with SigLIP was selected as the production configuration, achieving a semantic quality score of 0.365 at 1.432 seconds per window, the best latency adjusted result across all candidates. The semantic stage was extended into a downstream grounding path in which object lists from the semantic manifest became automated prompts for Grounding DINO and SAM 2, producing aligned bounding boxes and binary masks. The resulting pipeline yields structured manifests containing captions, tags, embeddings, spatial detections, and segmentation masks that serve as a conditioning layer for subsequent world model and policy training stages. The 6 DoF stage was not carried forward because the semantic and 2D grounding stages had already reached a useful prototype state, whereas a geometry aware extension would have introduced substantially more complexity with limited immediate value. The primary contribution of this placement is a reusable, inspectable AI vision layer that advances the M.AI.ESTRO-Twin pipeline from raw egocentric video to machine readable task state representations.

1 Introduction

The placement was carried out alongside Research Support Officer duties linked to Project M.AI.ESTRO, an active Malta based research programme developing automated digital twin technology for industrial skill transfer. The declared placement title remained M.AI.ESTRO-Twin: Automated Digital Twin Synthesis for Industrial Skill Transfer, and the motivating problem was how to generate structured perception signals automatically from expert demonstrations without the need for manual prompt entry or annotation.

The one page progress report submitted at the midpoint of the placement established that a working semantic alignment pipeline had been implemented and that an initial vision language model comparison had been completed. This final report continues from that checkpoint by documenting how the semantic stage was refined, how downstream grounding and mask stages were added, and why the placement delivered a 2D perception and conditioning layer rather than the full four stage pipeline originally proposed.

The aim of this placement was to develop a reusable AI vision pipeline that transforms raw

egocentric manipulation data into machine readable representations of task relevant state, suitable for conditioning downstream learning modules within M.AI.ESTRO. To achieve this aim, five objectives were pursued:

1. Establish a semantic alignment pipeline capable of generating structured object and action descriptors from EgoDex manipulation windows.
2. Evaluate and select a vision language model and embedding backend appropriate for automated semantic quality assessment.
3. Extend the semantic stage into a downstream grounding pipeline using open-set object detection and segmentation.
4. Produce aligned structured manifests containing captions, tags, bounding boxes, and masks suitable for world model conditioning.
5. Document the design decisions and model selection rationale in a form that supports future M.AI.ESTRO pipeline development.

The central research question guiding the implementation was: Can vision language model out-

puts serve as reliable automated prompt sources for downstream open-set object grounding in ego-centric manipulation data? The original proposed stages were semantic understanding, bounding box grounding, binary mask generation, and 6 DoF pose estimation. As the implementation matured, the scope was deliberately narrowed to the first three stages. The report is structured as follows: Section 2 describes the methodology, Section 3 presents results, Section 4 reports placement effort, and Section 5 draws conclusions and identifies future work.

2 Methodology

All data used in this placement was drawn from EgoDex, a publicly released dataset; no human participants were recruited, no personal data were collected, and no ethical approval beyond standard academic data use was required. EgoDex is a large scale egocentric manipulation dataset comprising approximately 1.1 million demonstration clips across 194 object categories designed for dexterous interaction learning [1]. The implemented pipeline operated on indexed windows extracted from the dataset rather than processing every frame independently. Since the source videos run at 30 frames per second, frame-by-frame semantic inference would have introduced an impractical amount of redundant computation. A fixed windowing strategy shown in figure 1 was therefore adopted as the first operational approach, in which each sample contained 16 observation frames and 16 prediction frames at a stride of 128 frames. At 30 fps and a 128-frame stride, consecutive windows are separated by approximately 4.3 seconds, a span sufficient for meaningful object state changes to emerge in manipulation tasks while keeping the index size tractable for placement stage work. As work within the broader M.AI.ESTRO research programme progressed, this fixed stride scheme was subsequently superseded by an adaptive segmentation strategy that divides each episode into a configurable number of relative position segments and applies a variable frame skip proportional to episode length, achieving substantially better coverage across variable duration demonstrations. The transition from the initial fixed stride approach to this adaptive extraction method therefore reflects a direct line of development that originated in the placement prototype.




Figure 1: Fixed stride window sampling used for placement stage preprocessing. Each indexed sample contains 16 observation frames and 16 prediction frames.

2.1 Semantic Alignment

The first semantic candidate used during development was Qwen3-VL-8B-Instruct, selected because the early project goal required a model capable of generating scene level descriptions and structured semantic signals from a small set of representative frames [2]. The initial implementation was tightly coupled to that model, with prompt templates, output parsing, and quality checks written specifically for its response format. Once the semantic stage was shown to function reliably end-to-end, the code was refactored into a model agnostic benchmark contract with a shared prompt interface, standardised output schema, and interchangeable embedding backend. This refactoring was essential because the objective was not simply to make the pipeline run once, but to produce a reusable and comparatively inspectable semantic stage.

In this context, embeddings are the vector representations used to place text and visual content into a common comparison space [3]. The semantic generator produced captions and structured tags, while the embedding backend was used to assess how well those generated descriptions still corresponded to the sampled window frames. The embedding model was therefore not merely an implementation detail, but part of the semantic quality check loop itself. SigLIP became the main operational backend because its pairwise sigmoid training objective produces per-pair similarity scores that do not depend on batch composition, making it well suited to the role of an independent semantic quality check on individual window caption pairs [4]. The project later also explored OpenCLIP as an alternative similarity model in order to test whether the semantic ranking remained stable under a different comparison space [5].

The candidate generator set then widened. The Qwen family was compared with LLaVA-v1.6 [6], SmolVLM2-2.2B-Instruct [7], and Idefics3-8B [8]. In practical terms, the Qwen family showed the strongest overall balance for this task. During this benchmarking stage, the release of the Qwen 3.5 family provided an opportunity to extend the comparison further and examine whether the new variants changed the quality versus latency trade off enough to justify a change in the operational pipeline [9]. These later checks became part of the

placement process rather than a separate line of work.

The semantic outputs themselves also evolved during development. Initially, the stage emphasised raw human readable artefacts so that the output could be inspected directly and used to judge whether the model was describing the task window sensibly. As the placement moved towards a multi stage pipeline, human readable artefacts alone were no longer sufficient: the grounding stage required a consistent, machine parseable object list rather than free text captions to function without manual intervention. The outputs therefore had to be written in a structured schema so that the next stage could consume them reliably. This is why the semantic stage now writes an aligned manifest, captions, tags, embeddings, and especially the object list information that acts as the prompt source for downstream grounding. The remaining manifest fields are also useful later as richer conditioning signals, particularly when contact relations or state changes are detected in the window.

2.2 Bounding Box Grounding

Once the semantic stage had stabilised, its outputs were used to drive the next stage of the pipeline. Rather than entering text prompts manually, prompt candidates were derived from the semantic results themselves, with particular attention to the contact object and the deduplicated object list. This design choice preserved the pipeline logic established in the semantic stage: the model first interpreted the task window and then supplied the vocabulary required for localisation.

Those prompts were passed into Grounding DINO, which produced text conditioned bounding boxes [10]. This step was important because the placement was not limited to scene description and had to move from semantic interpretation towards explicit image space localisation. The resulting boxes therefore became the first direct spatial outputs produced by the implemented pipeline. Grounding DINO was selected on the basis of published benchmark performance rather than through an internal benchmark as extensive as the semantic VLM bakeoff. At the time of selection it represented the state of the art on zero-shot open-set detection, with published results reporting 52.5 AP on COCO without any in distribution training images and a mean 26.1 AP on the ODinW zero-shot benchmark, outperforming prior alternatives including GLIP and OWL-ViT [10]. In use, the stage showed the expected weaknesses of zero-shot grounding: small or visually similar objects could be confused, and category labels were not always assigned consistently. Nevertheless, the

detector generally placed boxes in the correct image regions, which was the main requirement at this stage because those regions were then passed forward to mask generation. The implementation also enabled prompt label variants for difficult categories such as lids; in qualitative inspection, these variants recovered detections that the primary label missed in visually ambiguous frames, though no formal recall metric was computed across the evaluation sample.




Figure 2: Excerpt from a structured semantic manifest showing the object list field used as the automated prompt source for Grounding DINO. Each entry in the list is passed as a text query to the detector without manual intervention.

2.3 Object Mask Generation

The final implemented stage of the placement used SAM 2 after the bounding box step so that the detected objects also yielded binary masks [11]. This gave the pipeline denser spatial outputs than boxes alone and made the results more useful for later conditioning and inspection. The purpose of this stage was not to recover geometry or stable 3D identity. Its role was to provide a practical 2D segmentation signal aligned to the same windows already used by the semantic and grounding stages.

This distinction is central to how the placement should be interpreted. The mask stage belongs to an object pose estimation path only in a limited sense. What was implemented here was the bounding box and binary mask portion of that path, not the final geometry aware stage. That narrower reading is methodologically important because it explains both the value of the delivered work and the reason the 6 DoF stage was not forced into the placement before it was justified by the results.

In practical terms, integrating SAM 2 required implementing three concrete components: a box expansion step that enlarged detector outputs by a configurable pixel margin before passing them to the segmenter, a frame level batching loop that applied SAM 2 to each observation frame in the window, and a run length encoding serialiser

that stored the resulting masks compactly together with their spatial shape metadata. Just as importantly, the mask rows remained aligned with the original semantic window order, which preserved a stable correspondence between captions, tags, detections, and masks for downstream reuse. SAM 2, which extends the original Segment Anything Model [12] to video and streaming contexts, was chosen largely from the literature and from its strong practical reputation as a general purpose segmentation model, rather than from a dedicated in house benchmark across many maskers.

This stage also clarified the practical interpretation of the outputs. A named limitation of the grounding and mask stages is category confusion between visually similar objects: qualitative inspection of notebook outputs showed cup-versus-lid misclassification in frames where both objects were present and similarly sized, consistent with the known sensitivity of zero-shot grounding to fine grained visual distinctions. In this placement context, that ambiguity did not invalidate the stage, because the value of the masks lay in providing spatially reasonable support regions for later conditioning rather than exact class level identification. This limitation becomes consequential only when the downstream objective requires precise object identity rather than approximate spatial localisation.

3 Results

3.1 Semantic Benchmark Results

The semantic benchmark expanded the model comparison first introduced in the one page progress report. Six candidate generators were evaluated: Qwen3-VL-2B, Qwen3-VL-8B, LLaVA-v1.6-7B, Idefics3-8B, Qwen3.5-4B, and SmolVLM2-2.2B. The evaluation sample consisted of 64 windows drawn deterministically from the EgoDex index using a fixed random seed, stratified to include a range of object categories and task phases. The benchmark slice was held fixed across all runs, with only the embedding backend varied, so that the contribution of the generator could be separated cleanly from the contribution of the similarity model.

Table 1 summarises the main outcome across all evaluated generators. Both SigLIP and OpenCLIP scores are reported for each model; the OpenCLIP column is included as a cross family stability check rather than as a competing quality metric.

Table 1: Semantic benchmark results over the shared 64-window sample. SigLIP score and OpenCLIP score are cosine similarity proxies computed with each embedding family independently; they are not directly comparable across families. Lower seconds per window indicates lower runtime cost.

Model	SigLIP score	OpenCLIP score	Sec./win
Qwen3-VL-2B	0.365	0.409	1.432
Qwen3-VL-8B	0.431	0.473	2.435
Qwen3.5-4B	0.434	0.478	5.763
LLaVA-v1.6-7B	0.422	0.463	3.317
Idefics3-8B	0.405	0.450	4.186
SmolVLM2-2.2B	0.319	0.361	7.046

The comparison led to two practical results. First, OpenCLIP consistently produced higher similarity style scores than SigLIP on matched runs. Across all six table rows, the mean OpenCLIP score exceeded the corresponding SigLIP score by approximately 0.043. However, this uplift did not justify replacing SigLIP as the main operational backend. The score divergence has a principled mechanistic basis: SigLIP is trained with a pairwise sigmoid loss that evaluates each image text pair independently, whereas OpenCLIP follows the original CLIP softmax contrastive objective in which each score is normalised against the full batch [4, 3]. The two families therefore produce scores at different scales and with different calibration geometry, making their values not directly comparable. This scale incomparability between embedding families is a recognised phenomenon in applied vision research: work on downstream robotic tasks has independently observed that CLIP based representations produce a characteristically narrower cosine similarity range than sigmoid trained alternatives, which necessitates family specific thresholding when both are used in the same evaluation context. The higher OpenCLIP figures in the placement benchmark therefore reflect the scale properties of the softmax normalisation space rather than a genuine quality advantage in the generated descriptions. Crucially, the cross backend comparison demonstrated that the relative ordering of candidate generators remained stable across both scoring spaces, confirming that the choice of embedding backend is less critical to the ranking outcome than the choice of generator. SigLIP was therefore retained as the primary operational choice, with OpenCLIP serving as a secondary comparative reference that confirmed the robustness of the generator ranking rather than motivating a change in the operational pipeline.

Second, the Qwen family remained the strongest practical choice once quality and latency were considered together. Although larger models could improve raw semantic scores, they did so at a much higher runtime cost. For the placement, the most useful production choice remained Qwen3-VL-2B-

Instruct together with SigLIP because it produced strong semantic outputs while keeping inference practical for repeated runs and later downstream stages.

3.2 Grounding and Mask Results

The grounding and mask stages produced the first explicit spatial outputs from the placement pipeline. Using the semantic object list as the prompt source, Grounding DINO generated text conditioned bounding boxes, and SAM 2 then produced the corresponding binary masks. These results confirmed that the semantic stage was not only interpretable in isolation, but also useful as a prompt generator for downstream perception modules.

The strongest positive result here was that the semantic-to-grounding pipeline could be automated without freezing on manual prompt entry. Across the qualitatively inspected sample, the upstream object list triggered at least one bounding box detection in the majority of windows where manipulable objects were visible; failures were predominantly restricted to frames with heavy occlusion or where the object appeared only partially in view. No formal precision or recall metric was computed against ground truth annotations, which represents a limitation of the current evaluation. Nevertheless, the results confirmed that semantic object lists are a practically viable automated prompt source, giving the placement a coherent multi stage workflow rather than a set of disconnected experiments.




Figure 4: Example placement output from the grounding and mask stages. Semantic object prompts become bounding boxes and binary masks for the sampled window frames.




Figure 5: Additional qualitative Object Pose Estimation (OPE) pipeline example from a ring handling episode, included to show a second object configuration beyond the cup dominated case.

3.3 Limitations

Three principal limitations bound the scope of these results. First, the semantic benchmark was a bounded comparison over a 64-window shared sample; it is not a universal ranking and may not generalise to object categories or task types not represented in the sample. Second, the grounding and mask stages were evaluated through qualitative inspection and runtime measurement rather than against a formal annotated validation set with ground truth boxes or masks; no precision, recall, or IoU figures are reported. Third, the placement pipeline was developed and tested exclusively on EgoDex; cross dataset generalisation to other egocentric manipulation datasets was not evaluated.

3.4 Interpretation

This bounded reading also explains the 6 DoF decision. By the time the semantic, bounding box, and mask stages had matured into a coherent prototype, the next step towards canonical 6 DoF pose recovery would have required a substantially different level of geometric machinery, stronger CAD dependence, and a more specialised evaluation setup [13, 14]. The placement therefore reached its best value by consolidating the 2D perception and conditioning pipeline rather than pushing prematurely into a heavier geometric stage.

4 Placement Effort and Supervision

The placement unit corresponds to approximately 100 to 140 hours of work. The hours worked over the placement period can be summarised as shown in Table 2, which presents a reconstructed allocation of 110 hours across the completed stages of the implementation. The allocation is estimated rather than logged: a conservative basis of 5 hours per day across 22 placement equivalent days was used, cross referenced against calendar records and




Figure 3: Semantic quality versus runtime trade off used to support the final production choice.

implementation commit history. These activities were not always sequential, since model runtime, reruns, and waiting periods were interleaved with implementation, inspection, and refinement across multiple stages of the pipeline.

Table 2: Retrospective placement effort allocation reconstructed from the completed implementation stages.

Placement activity	Days	Hours
Environment preparation, Docker checks, and kr2280 setup	1	5
Semantic pipeline implementation and early validation	6	30
Object grounding implementation and prompt construction	5	25
Mask generation integration and refinement	4	20
Model rebenchmarking and later candidate checks	2	10
Notebook walkthroughs, result consolidation, and visual inspection	2	10
Report preparation and final packaging	2	10
Total	22	110

This placement was also shaped by continuous supervision. The implementation state, benchmark outcomes, and main design decisions were communicated regularly to Prof. Alexiei Dingli throughout the placement period. This supervisory loop proved particularly important when new candidate models became available and when benchmark deltas required careful interpretation before any change to the operational pipeline was warranted.

5 Conclusion and Future Work

The strongest evidence for the value of this placement comes from its alignment with current re-

search practice. As work within M.A.I.ESTRO progressed and an academic paper currently in preparation undertook a systematic comparison of state-of-the-art world models, it became clear that the enrichment strategy central to this placement recurs independently across leading architectures. Within the M.A.I.ESTRO programme itself, DexWM [15] encodes scene states using a frozen DINOv2 encoder whose semantically rich patch-level features serve the same function as the structured semantic outputs produced in this placement: both provide the enriched representation layer that a downstream world model predictor consumes. More broadly, PointWorld [16] enriches raw scene point clouds with frozen DINOv3 ViT-L visual features before world model conditioning, confirming that the same principle has been independently adopted at the frontier of open world robotic manipulation research. Together these two examples demonstrate that the enrichment strategy developed here is not an incidental design choice but a practice convergently validated by leading work in the field.

This placement successfully moved from raw EgoDex windows to structured semantic outputs, image space localisation, and binary masks, while benchmarking the model choices required to make that pipeline practical and reproducible. The contribution lies not only in the implemented stages but in the progression from a tightly coupled prototype to a modular, benchmark driven pipeline that can be extended or replaced at the component level. The work was carried out as supervised and iteratively reviewed development within the agreed project framework. Returning to the central research question posed in the Introduction, the evidence confirms that vision language model outputs can serve as reliable automated prompt sources for downstream open-set object grounding in egocentric manipulation data: the semantic object lists generated by Qwen3-VL-2B-Instruct triggered consistent detections across the inspected

sample without any manual prompt entry, satisfying the core condition the question set out to test.

The principal limitations of the delivered work are: the absence of a ground truth annotated validation set for formal grounding and mask evaluation; the sensitivity of zero-shot grounding to visually similar object categories; the restriction of all testing to EgoDex with no cross dataset evaluation; and the fixed stride windowing approach, which has since been superseded by an adaptive strategy in the broader research programme. These limitations define the boundaries of the current contribution and motivate the future work below.

Six directions for future work follow from this placement. First, the enriched semantic outputs, bounding boxes, and masks can be used directly as structured conditioning signals for world model and policy training, the two immediate next stages in the M.AI.ESTRO pipeline. Second, a formally annotated subset of EgoDex windows should be created to enable precision and recall evaluation of the grounding and mask stages. Third, the 6 DoF pose estimation stage, explicitly deferred in this placement, represents the natural geometric extension once the 2D perception layer has been validated against ground truth. Fourth, the adaptive windowing strategy developed in subsequent WACT work should be integrated with the semantic pipeline end-to-end to replace the fixed stride approach. Fifth, cross dataset evaluation on non EgoDex egocentric manipulation data would establish the generalisability of the semantic and grounding stages beyond the training distribution of the current source material. Sixth, the implemented pipeline also sits within a wider Malta–Poland research collaboration being developed to address Digital Product Passport requirements through industrial AI methods, in view of the Ecodesign for Sustainable Products Regulation [17]. The placement thereby delivered a reusable, validated perception layer that connects raw egocentric demonstration data to the structured representations required by downstream learning and policy systems, establishing a foundation on which both the immediate M.AI.ESTRO objectives and the longer term research programme can build.

References

- [1] R. Hoque, P. Huang, D. J. Yoon, M. Sivapurapu, and J. Zhang, “Egodex: Learning dexterous manipulation from large-scale egocentric video,” 2026.
- [2] S. B. et al., “Qwen3-vl technical report,” 2025.
- [3] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.
- [4] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, “Sigmoid loss for language image pre-training,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 11975–11986, 2023.
- [5] M. Cherti, R. Beaumont, R. Wightman, M. Wortsman, G. Ilharco, C. Gordon, C. Schuhmann, L. Schmidt, and J. Jitsev, “Reproducible scaling laws for contrastive language-image learning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2818–2829, 2023.
- [6] H. Liu, C. Li, Y. Li, and Y. J. Lee, “Improved baselines with visual instruction tuning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [7] A. Marafioti, O. Zohar, M. Farré, M. Noyan, E. Bakouch, P. Cuenca, C. Zakka, L. Ben Alal, A. Lozhkov, N. Tazi, V. Srivastav, J. Lochner, H. Larcher, M. Morlon, L. Tunstall, L. von Werra, and T. Wolf, “Smolvlm: Redefining small and efficient multimodal models,” 2025.
- [8] H. Laurençon, A. Marafioti, V. Sanh, and L. Tronchon, “Building and better understanding vision-language models: Insights and future directions,” 2024.
- [9] Qwen Team, “Qwen3.5: Towards native multimodal agents.” <https://qwen.ai/blog?id=qwen3.5>, 2026. Qwen Team blog post.
- [10] T. Ren, Q. Jiang, S. Liu, Z. Zeng, W. Liu, H. Gao, H. Huang, Z. Ma, X. Jiang, Y. Chen, Y. Xiong, H. Zhang, F. Li, P. Tang, K. Yu, and L. Zhang, “Grounding dino 1.5: Advance the “edge” of open-set object detection,” 2024.
- [11] N. Ravi, V. Gabeur, Y.-T. Hu, R. Hu, C. Ryali, T. Ma, H. Khedr, R. Rädle, C. Rolland, L. Gustafson, E. Mintun, J. Pan, K. V. Alwala, N. Carion, C.-Y. Wu, R. Girshick, P. Dollár, and C. Feichtenhofer, “Sam 2: Segment anything in images and videos,” 2024.
- [12] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollár, and R. Girshick, “Segment anything,” 2023.
- [13] B. Wen, W. Yang, J. Kautz, and S. Birchfield,

- “Foundationpose: Unified 6d pose estimation and tracking of novel objects,” 2024.
- [14] J. Lin, L. Liu, D. Lu, and K. Jia, “Sam-6d: Segment anything model meets zero-shot 6d object pose estimation,” in Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024.
 - [15] R. G. Goswami, A. Bar, D. Fan, T.-Y. Yang, G. Zhou, P. Krishnamurthy, M. Rabat, F. Khorrami, and Y. LeCun, “World models can leverage human videos for dexterous manipulation,” 2025.
 - [16] W. Huang, Y.-W. Chao, A. Mousavian, M.-Y. Liu, D. Fox, K. Mo, and L. Fei-Fei, “Point-world: Scaling 3d world models for in-the-wild robotic manipulation,” 2026.
 - [17] European Parliament and Council of the European Union, “Regulation (EU) 2024/1781 of the european parliament and of the council on ecodesign requirements for sustainable products,” EU Regulation 2024/1781, Official Journal of the European Union, 2024.